

Get the best out of Oracle Partitioning

A practical guide and reference

Version 5.0



Before we start ..

- Oracle wants to hear from you!
 - There's still lots of ideas and things to do
 - Input steers the direction
- Let us know about
 - Interesting use cases and implementations
 - Enhancement requests
 - Complaints
 - Contact dw-pm_us@oracle.com or hermann.baer@oracle.com



Table of Contents

- [What's New in 12.2?](#)
- [Partitioning Summary](#)
- [Partitioning Benefits](#)
- [Partitioning Concepts](#)
- [Partitioning Methods](#)
 - [Range Partitioning](#)
 - [Hash Partitioning](#)
 - [List Partitioning](#)
 - [Interval Partitioning](#)
 - [Range versus Interval](#)
 - [Reference Partitioning](#)
 - [Interval Reference Partitioning](#)
 - [Virtual Column Based Partitioning](#)
- [Partitioning Methods](#)
 - [Composite Partitioning](#)
 - [Multi-Column Range Partitioning](#)
 - [Range Partitioned Hash Cluster](#)
- [Indexing of Partitioned Tables](#)
 - [Local Indexing](#)
 - [Global Non-Partitioned Indexing](#)
 - [Global Partitioned Indexing](#)
 - [Indexing for unique constraints and primary keys](#)
 - [Partial Indexing](#)
 - [Unusable versus partial Indexes](#)
- [Partitioning for Performance](#)
- [Partition Maintenance](#)
 - [Maintenance on multiple Partitions](#)
 - [Cascading Truncate and Exchange for Reference Partitioning](#)
 - [Online Move Partition](#)
 - [Asynchronous Global Index Maintenance](#)
- [Stats Management for Partitioning](#)
- [Attribute Clustering/ Zone Maps](#)
- [Tips and tricks](#)
 - [Think about partitioning strategy](#)
 - [Physical and logical attributes](#)
 - [Eliminate hot spots](#)
 - [Smart partial exchange](#)
 - [Exchange with PK and UK](#)



What's New in Oracle Database 12c Release 2 ?



What's New in 12.2

- New Core Functionality Features

- Auto-list partitioning
- Multi-column list partitioning
- Partitioned external tables

- New Performance Features

- Online partition maintenance operations
- Online table conversion to partitioned table
- Reduced cursor invalidations for DDL's

- New Manageability Features

- Filtered partition maintenance operations
- Read only partitions
- Create table for exchange

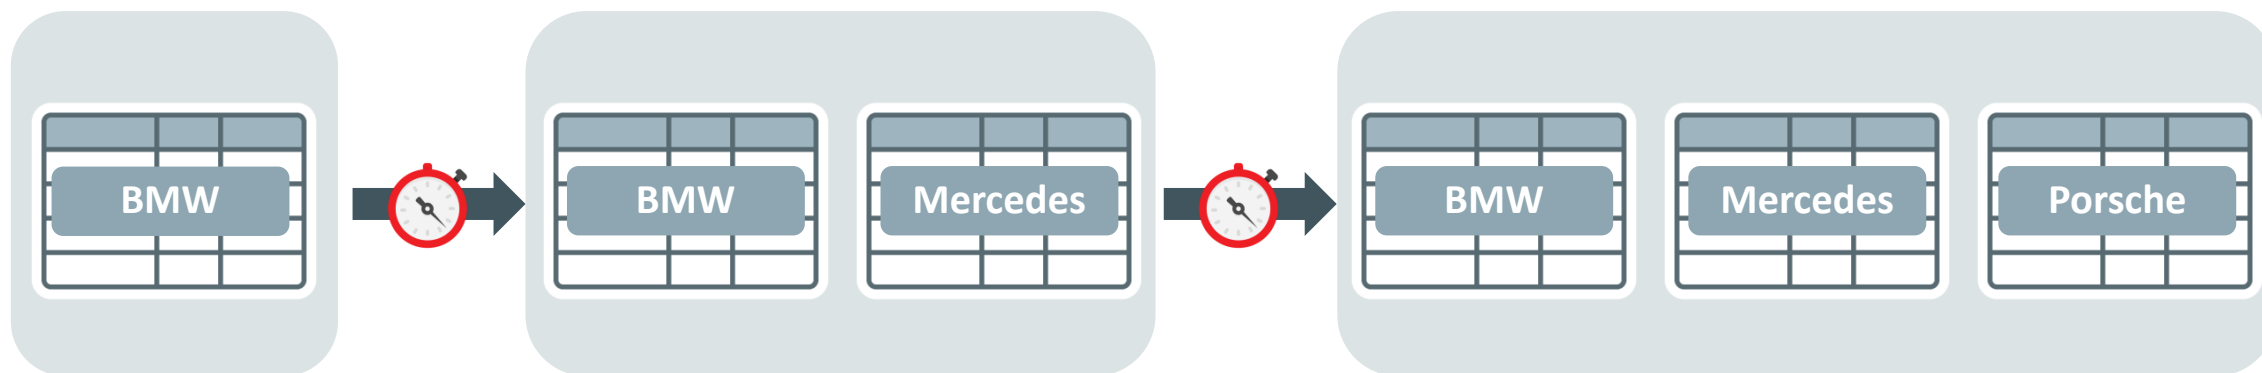
New Core Functionality Features



Auto-list Partitioning



Auto-List Partitioning



- Partitions are created automatically as data arrives
 - Extension to LIST partitioning
 - Every distinct partition key value will be stored in separate partition

Details of Auto-List Strategy

- Automatically creates new list partitions that contain one value per partition
 - Only available as top-level partitioning strategy in 12.2.0.1
- No notion of default partition
- System generated partition names for auto-created partitions
 - Use FOR VALUES clause for deterministic [sub]partition identification
- Can evolve list partitioning into auto-list partitioning
 - Only requirement is having no DEFAULT partition
 - Protection of customer investment



Auto-List Partitioned Table

Syntax example

```
CREATE TABLE sales ( brand VARCHAR2(50),  
                      model VARCHAR2(50), ...)  
PARTITION BY LIST (brand) AUTOMATIC  
( partition p1 values ('BMW'));
```

Auto-List is not equivalent to List + DEFAULT

- Different use case scenarios
- List with DEFAULT partitioning
 - Targeted towards multiple large distinct list values plus “noise”
- Auto-list partitioning
 - Expects ‘critical mass of records’ per partition key value
 - Could be used as pre-cursor state for using List + DEFAULT

Auto-List is not equivalent to List + DEFAULT

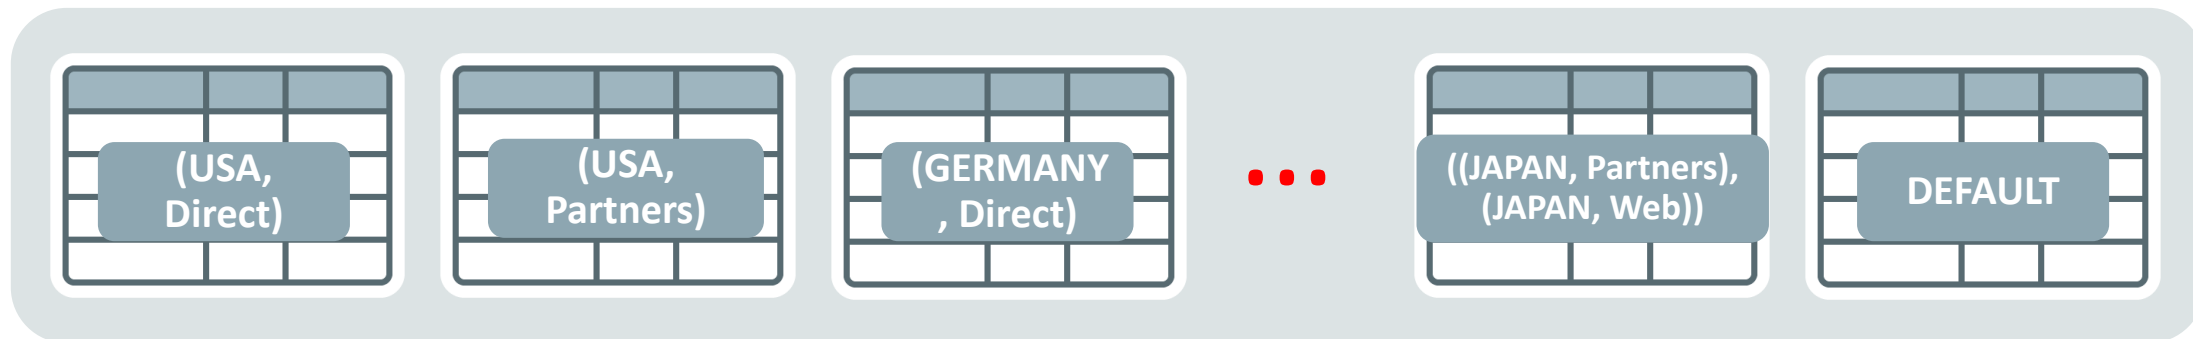
- Different use case scenarios
- List with DEFAULT partitioning
 - Targeted towards multiple large distinct list values plus “noise”
- Auto-list partitioning
 - Expects ‘critical mass of records’ per value
 - Could be used as pre-cursor state for using List + DEFAULT
- .. Plus they are functionally conflicting and cannot be used together
 - Either you get a new partition for a new partition key value
 - .. Or “dump” it in the catch-it-all bucket



Multi-Column List Partitioning



Multi-Column List Partitioning



- Data is organized in lists of multiple values (multiple columns)
 - Individual partitions can contain sets of multiple values
 - Functionality of DEFAULT partition (catch-it-all for unspecified values)
- Ideal for segmentation of distinct value pairs, e.g. (region, channel)

Details of Multi-Column List Strategy

- Allow specification of more than one column as partitioning key
 - Up to 16 partition key columns
 - Each set of partitioning keys must be unique
- Notation of one DEFAULT partition
- Functional support
 - Supported as both partition and subpartition strategy
 - Support for heap tables
 - Support for external tables
 - Supported with Reference Partitioning and Auto-List

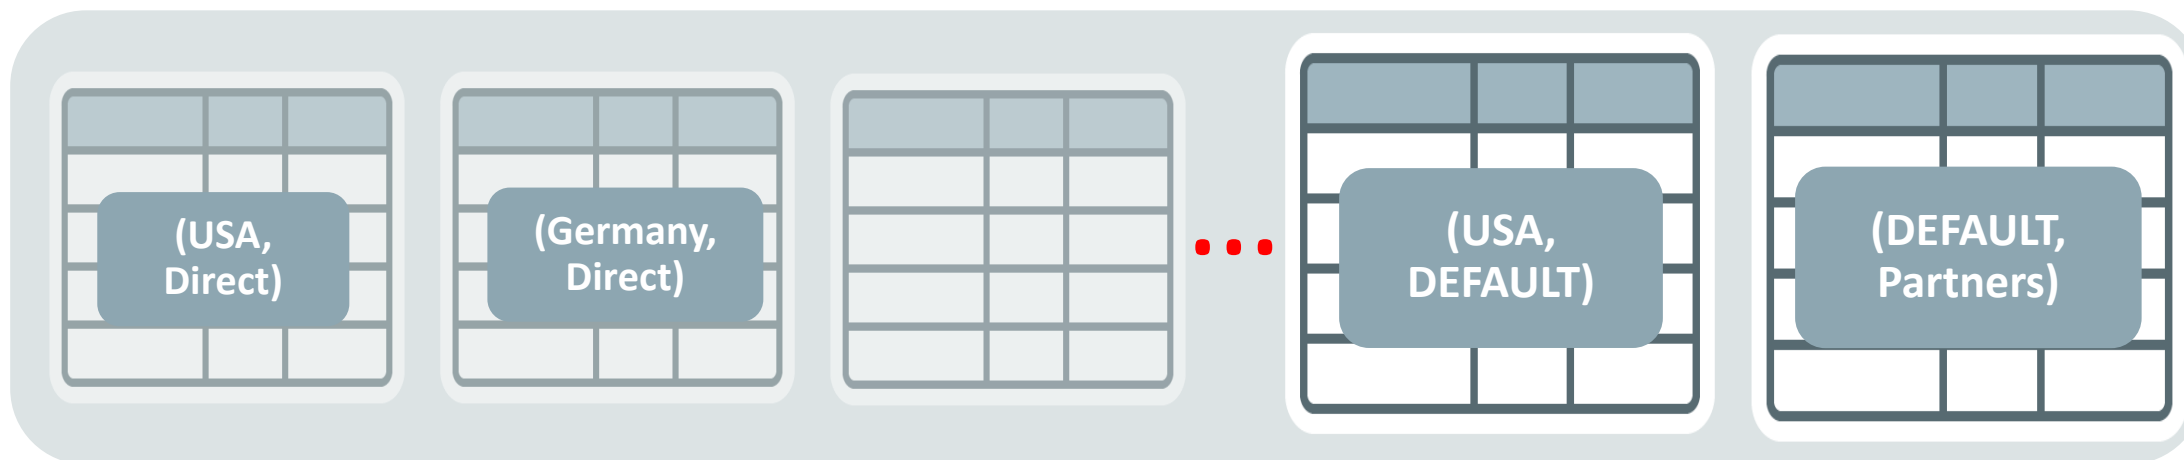
Multi-Column List Partitioned Table

Syntax example

```
CREATE TABLE sales ( region VARCHAR2(50),  
                      channel VARCHAR2(50), ...)  
PARTITION BY LIST (region, channel)  
( partition p1 values ('USA', 'Direct'),  
  partition p2 values ('USA', 'Partners'),  
  partition p3 values ('GERMANY', 'Direct'),  
  ...  
  partition p44 values (('JAPAN', 'Partners'),  
                       ('JAPAN', 'Web')),  
  partition p45 values (DEFAULT)  
);
```

Multi-Column List Partitioning

What if there was a DEFAULT per column?



- Where do we store (USA, Partners) ????

Multi-Column List Partitioning

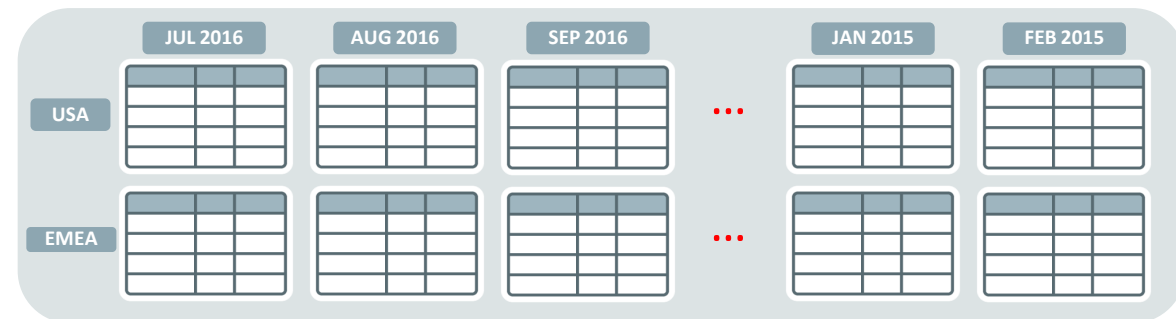
What if there was a DEFAULT per column?



- Where do we store (USA, Partners) ????
 - In the one-and-only DEFAULT partition

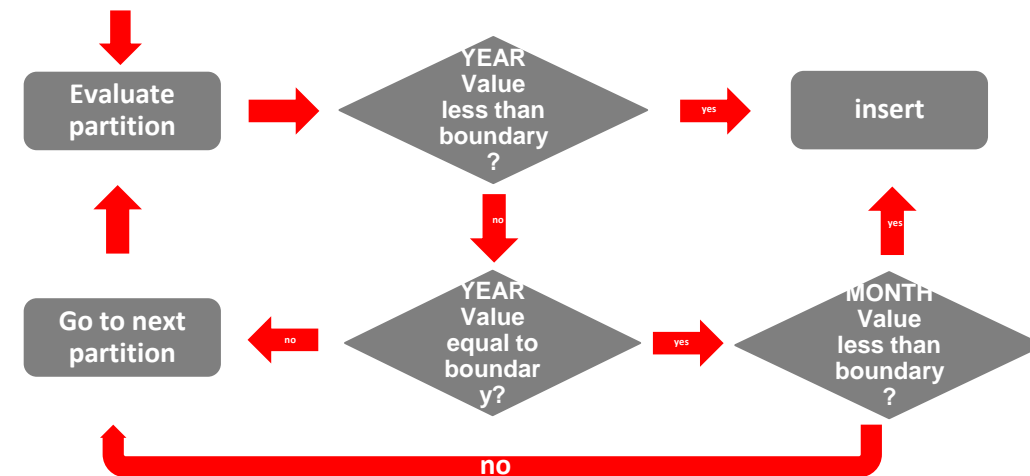
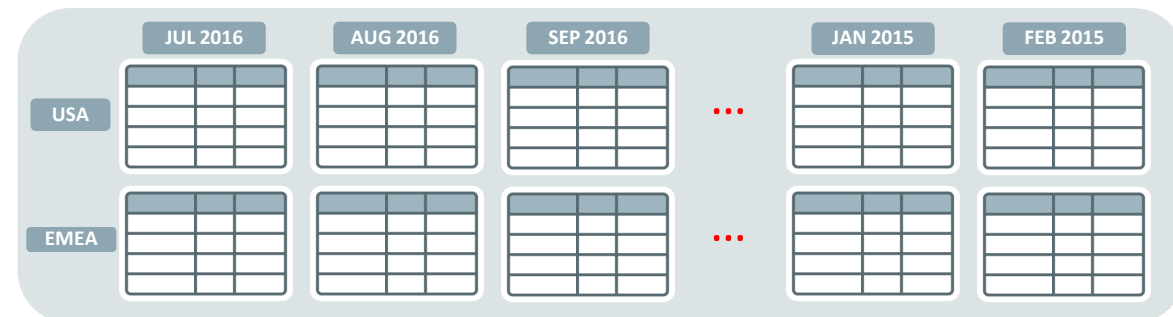
Multi-column List Partitioning prior to 12.2

- List – List partitioning
 - Almost equivalent
 - Only two columns as keys (two levels)
 - Conceptual symmetrical



Multi-column List Partitioning prior to 12.2

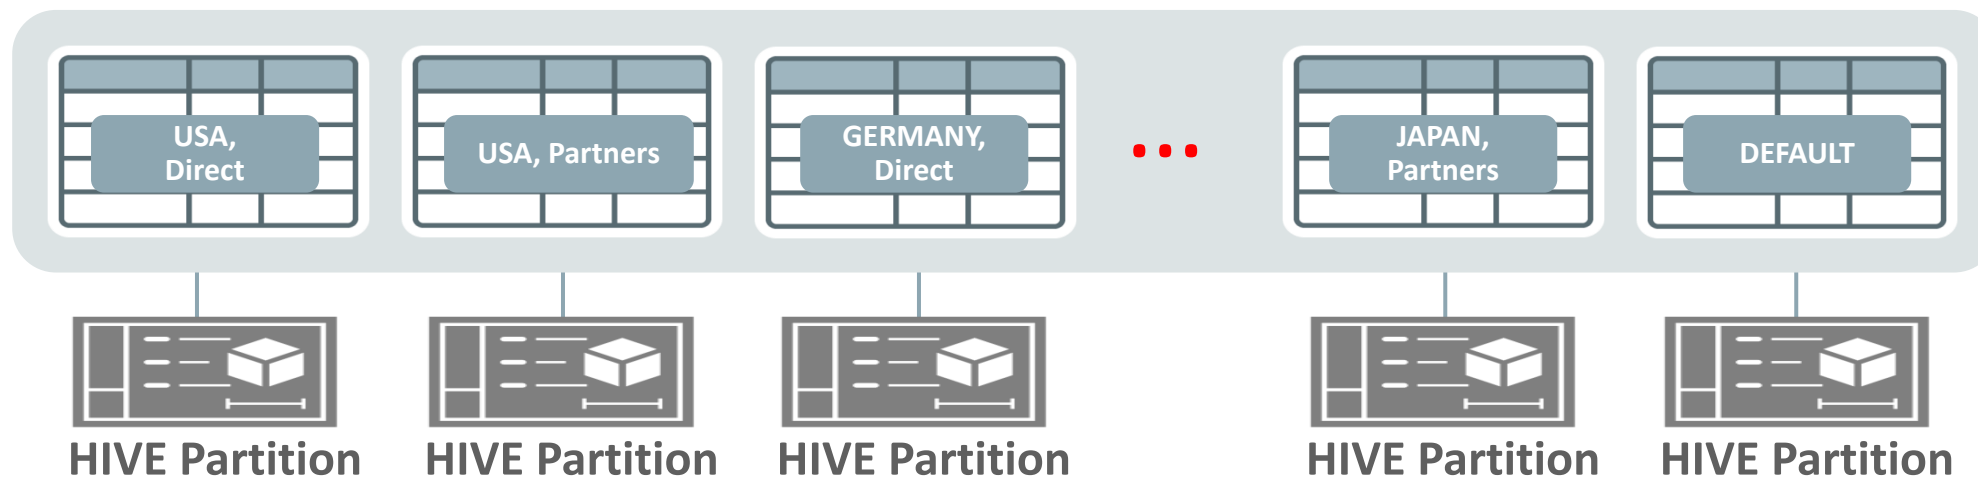
- List – List partitioning
 - Almost equivalent
 - Only two columns as key (two levels)
 - Conceptual symmetrical
- Multi-column range partitioning
 - NOT equivalent
 - Hierarchical evaluation of predicates only in case of disambiguity



Partitioned External Tables



Partitioned External Tables



- External tables can be partitioned, using all partitioning techniques
 - Multi-column partitioning optimally suited for partitioned HIVE tables
- Partition pruning and limited partition maintenance
 - Support of add partition, drop partition, exchange partition

New Performance Features



Online Partition Maintenance Operations



More Online DDL functionality

- Alter table modify non-partitioned table to partitioned table
- Alter table move online for heap tables
- Alter table split partition online

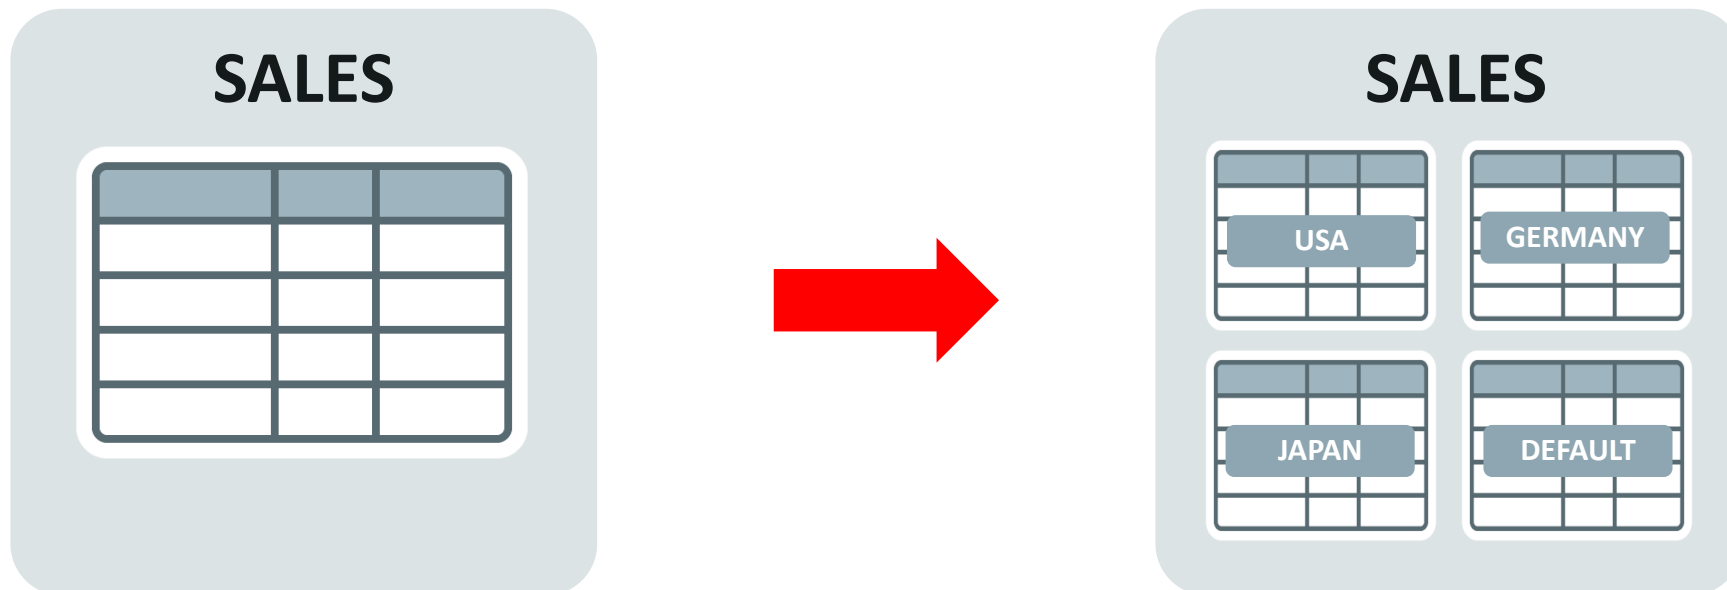


12.2

Online Table Conversion to Partitioned Table



Online Table Conversion



- Completely non-blocking (online) DDL

Online Table Conversion

Syntax Example

```
CREATE TABLE sales ( order_num NUMBER,  
                      region VARCHAR2 (10), ... );
```

```
ALTER TABLE sales MODIFY  
PARTITION BY LIST (region)  
  (partition p1 values ('USA'),  
   partition p2 values ('Germany'),  
   partition p3 values ('Japan'),  
   partition p4 values (DEFAULT))
```

```
UPDATE INDEXES ONLINE;
```


Online Table Conversion

Indexing

- Indexes are converted and kept online throughout the conversion process
- Full flexibility for indexes, following today's rules
- Default indexing rules to provide minimal to no access change behavior
 - Global partitioned indexes will retain the original partitioning shape.
 - Non-prefixed indexes will become global nonpartitioned indexes.
 - Prefixed indexes will be converted to local partitioned indexes.
 - Bitmap indexes will become local partitioned indexes



Reduced Cursor Invalidations for DDL's



Reduced Cursor Invalidations for DDL's

- Reduces the number of hard parses caused by DDL's
 - If hard parses are unavoidable, workload is spread over time
- New optional clause “[DEFERRED | IMMEDIATE] INVALIDATION” for several DDL's
 - If DEFERRED, Oracle will avoid invalidating dependent cursors when possible
 - If IMMEDIATE, Oracle will immediately invalidate dependent cursors
 - If neither, CURSOR_INVALIDATION parameter controls default behavior
- Supported DDL's:
 - Create, drop, alter index
 - Alter table column operations
 - Alter table segment operations
 - Truncate table



Reduced Cursor Invalidations for DDL's

Syntax Example

```
DROP INDEX emp_index DEFERRED INVALIDATION;
```

New Manageability Features

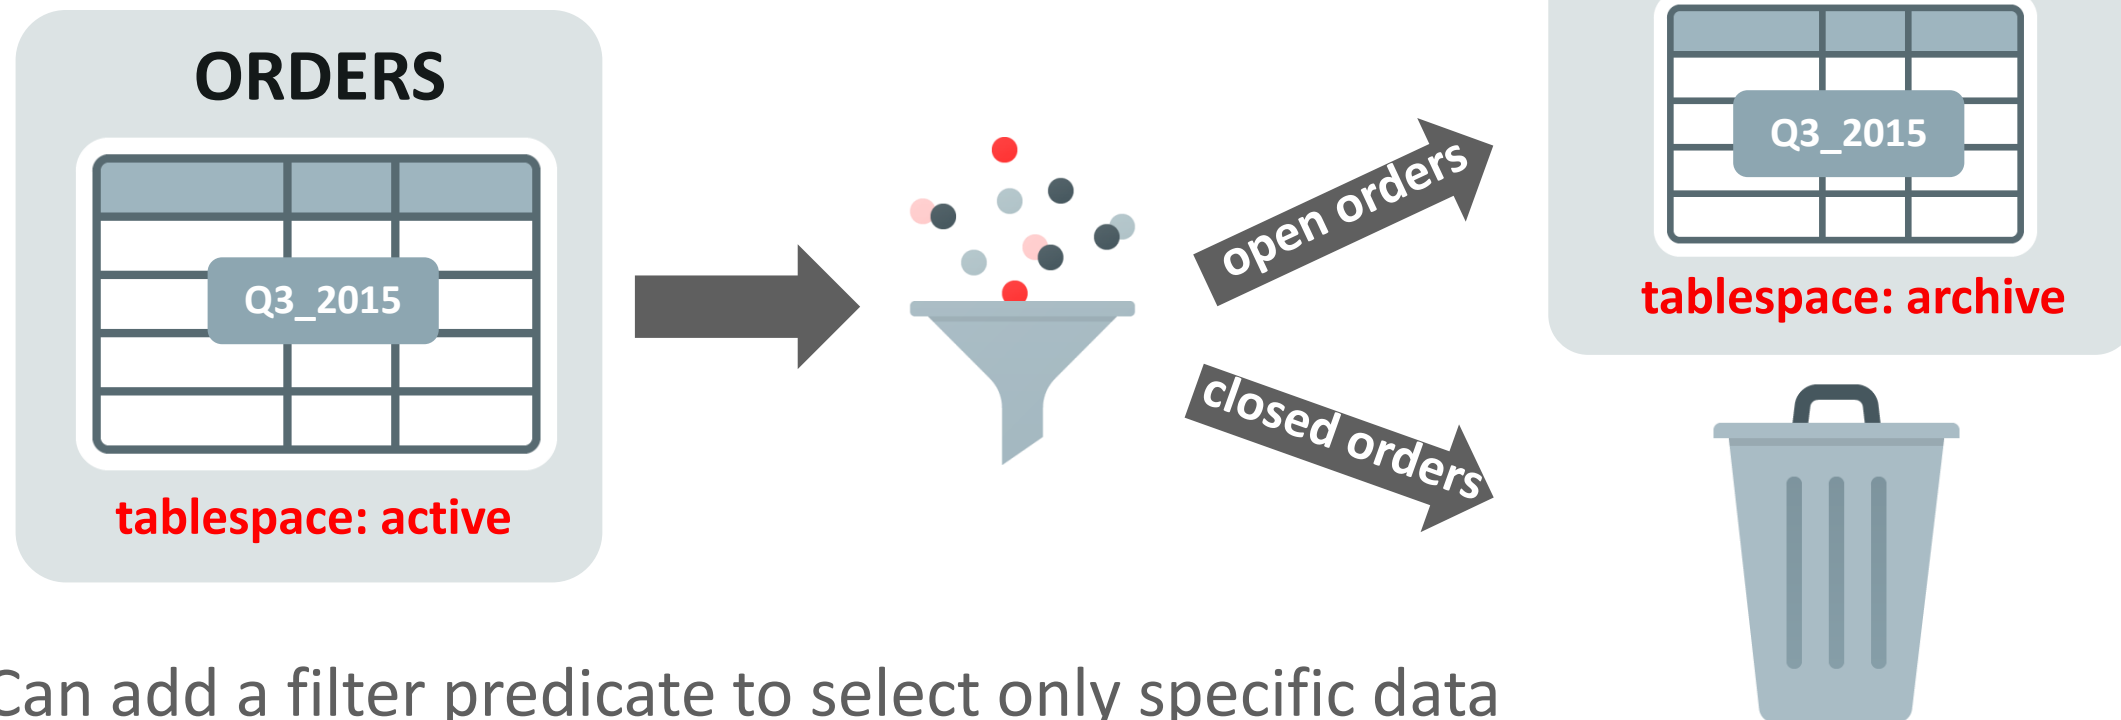


Filtered Partition Maintenance Operations



Filtered Partition Maintenance Operations

Move Partition Example



- Can add a filter predicate to select only specific data
- Combines data maintenance with partition maintenance

Details of Filtered Partition Maintenance Operations

- Can specify a single table filter predicate to MOVE, SPLIT and MERGE operations
 - Specification must be consistent across all partition maintenance
 - Specification needs to clearly specify the data of interest
- Specification will be added to the recursively generated CTAS command for the creation of the various new partition or subpartitions segments
- Filter predicates work for both offline and new online PMOP's

Filtered Partition Maintenance Operations

Move Partition Syntax Example

```
ALTER TABLE orders MOVE PARTITION q3_2015  
TABLESPACE archive  
INCLUDING ROWS WHERE order_state = 'open';
```

Filtered Partition Maintenance Operations

Move Partition Syntax Example

```
ALTER TABLE orders MOVE PARTITION q3_2015  
TABLESPACE archive  
INCLUDING ROWS WHERE order_state = 'open';
```

.. And what about online?

Filtered Partition Maintenance Operation

DML Behavior for online operations

- Filter condition is NOT applied to ongoing concurrent DML

```
INCLUDING ROWS WHERE order_state = 'open'
```

Filtered Partition Maintenance Operation

DML Behavior for online operations

- Filter condition is NOT applied to ongoing concurrent DML

```
INCLUDING ROWS WHERE order_state = 'open'
```

- Inserts will always go through

```
INSERT VALUES (order_state = 'closed')
```

Filtered Partition Maintenance Operation

DML Behavior for online operations

- Filter condition is NOT applied to ongoing concurrent DML

```
INCLUDING ROWS WHERE order_state = 'open'
```

- Inserts will always go through

```
INSERT VALUES (order_state = 'closed')
```

- Deletes on included data will always go through

```
DELETE WHERE order_state = 'open'
```

- Deletes on deleted data are void

```
DELETE WHERE order_state = 'closed'
```



Filtered Partition Maintenance Operation

DML Behavior for online operations

- Filter condition is NOT applied to ongoing concurrent DML

```
INCLUDING ROWS WHERE order_state = 'open'
```

- Inserts will always go through

```
INSERT VALUES (order_state = 'closed')
```

- Deletes on included data will always go through

```
DELETE WHERE order_state = 'open'
```

- Deletes on deleted data are void

```
DELETE WHERE order_state = 'closed'
```

- Updates on included data always goes through

```
UPDATE set order_status = 'closed'  
WHERE order_state = 'open'
```

- Updates on excluded data are void

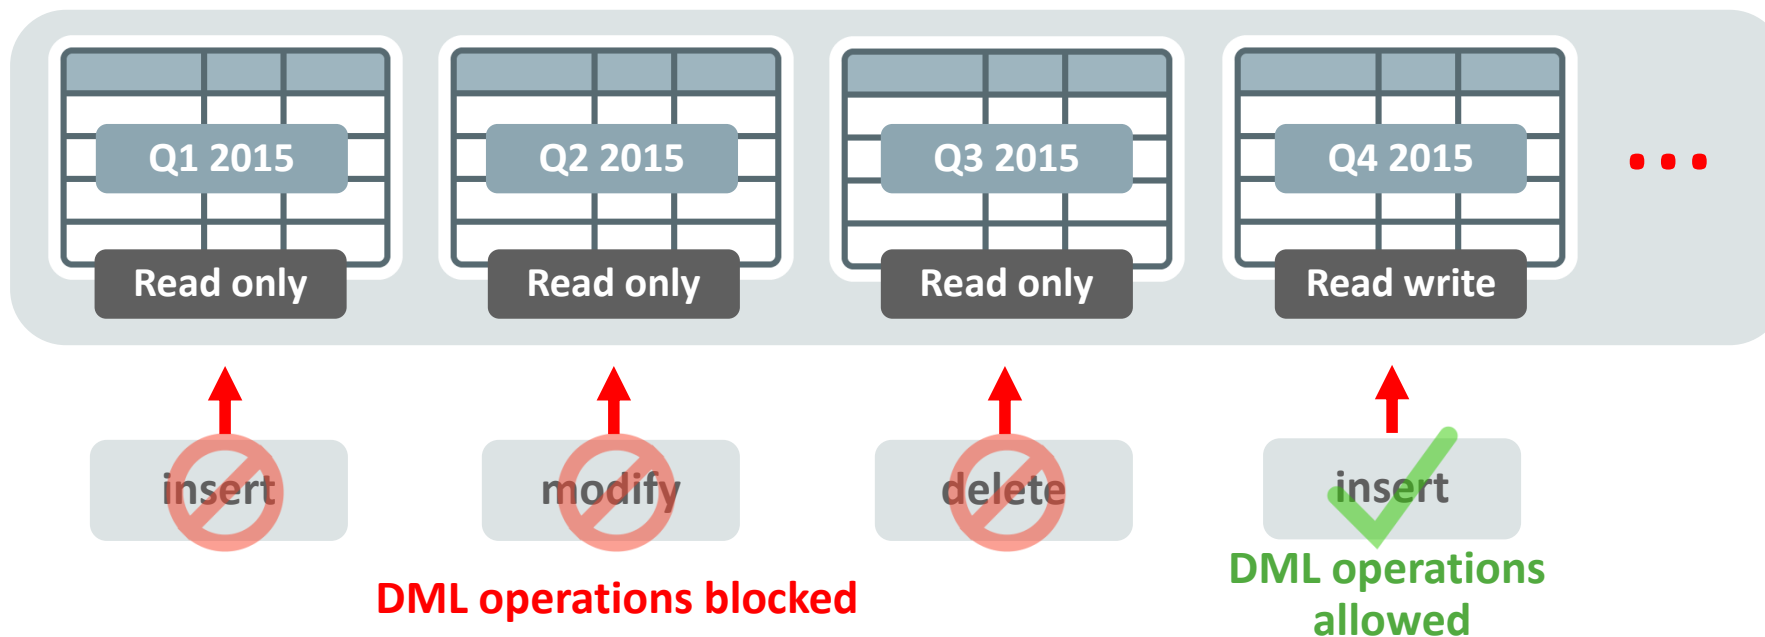
```
UPDATE set order_status = 'open'  
WHERE order_state = 'closed'
```



Read Only Partitions



Read Only Partitions



- Partitions and subpartitions can be set to read only or read write
- Any attempt to alter data in a read only partition will result in an error
- Ideal for protecting data from unintentional DML by any user or trigger

Details of Read Only Partitions

- Read only attribute guarantees data immutability
 - “SELECT <column_list> FROM <table>” will always return the same data set after a table or [sub]partition is set to read only
- If not specified, each partition and subpartition will inherit read only property from top level parent
 - Modifying lower level read only property will override higher level property
 - Alter tablespace has highest priority and cannot be overwritten
- Data immutability does not prevent all structural DDL for a table
 - ADD and MODIFY COLUMN are allowed and do not violate data immutability of existing data
 - Others like DROP/RENAME/SET UNUSED COLUMN are forbidden
 - DROP [read only] PARTITION forbidden, too - - violates data immutability of the table



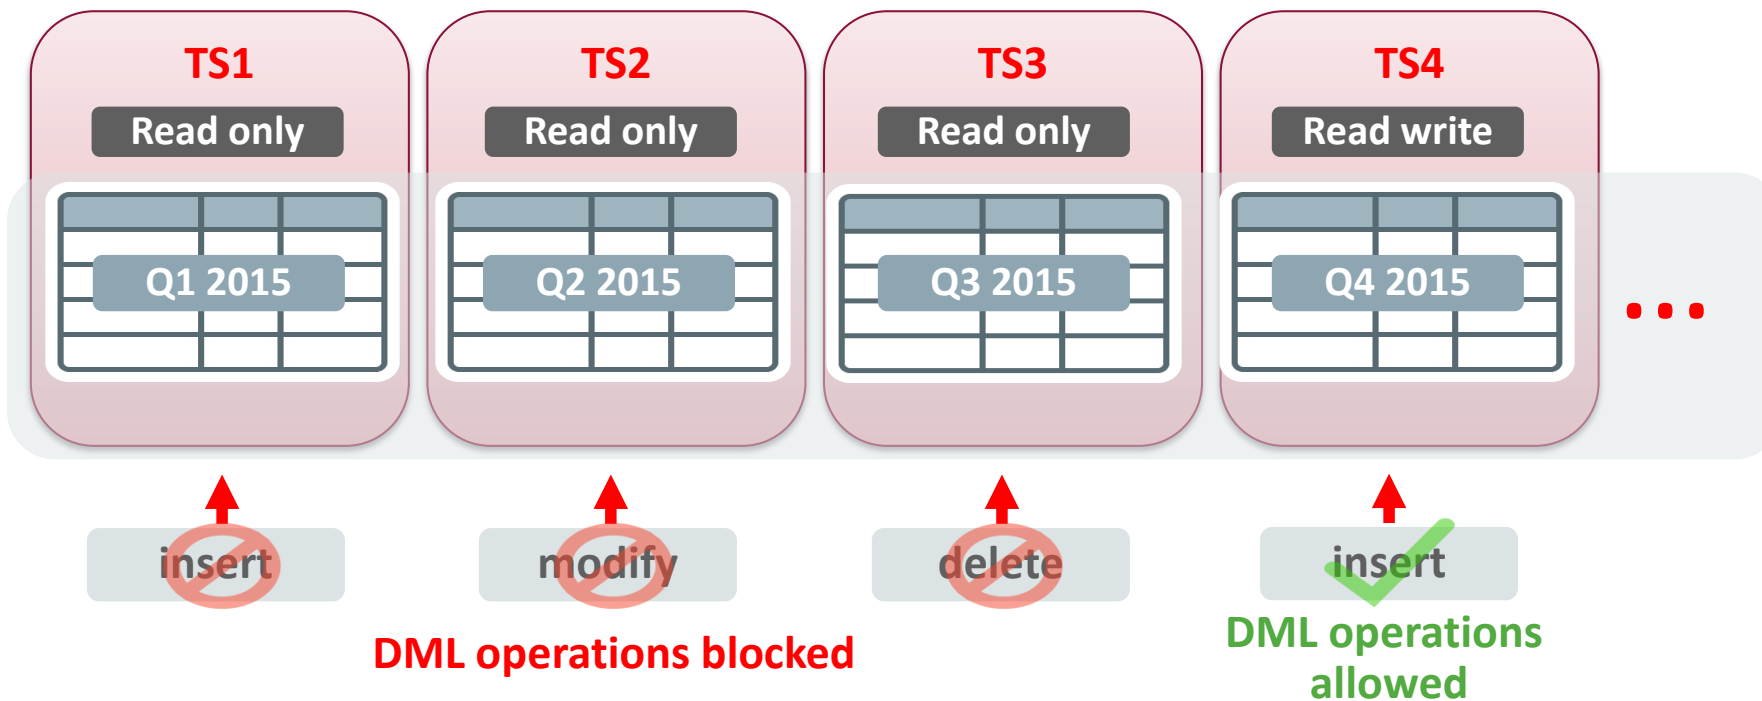
Read Only Partitions

```
CREATE TABLE orders ( order_id number,  
                        order_date DATE, ... ) read only  
  
PARTITION BY RANGE(order_date)  
( partition q1_2015 values less than ('2016-10-01'),  
  partition q2_2015 values less than ('2015-01-01'),  
  partition q3_2015 values less than ('2015-04-01'),  
  partition q4_2015 values less than ('2015-07-01') read write  
);
```

Read only tablespace versus read only partitions

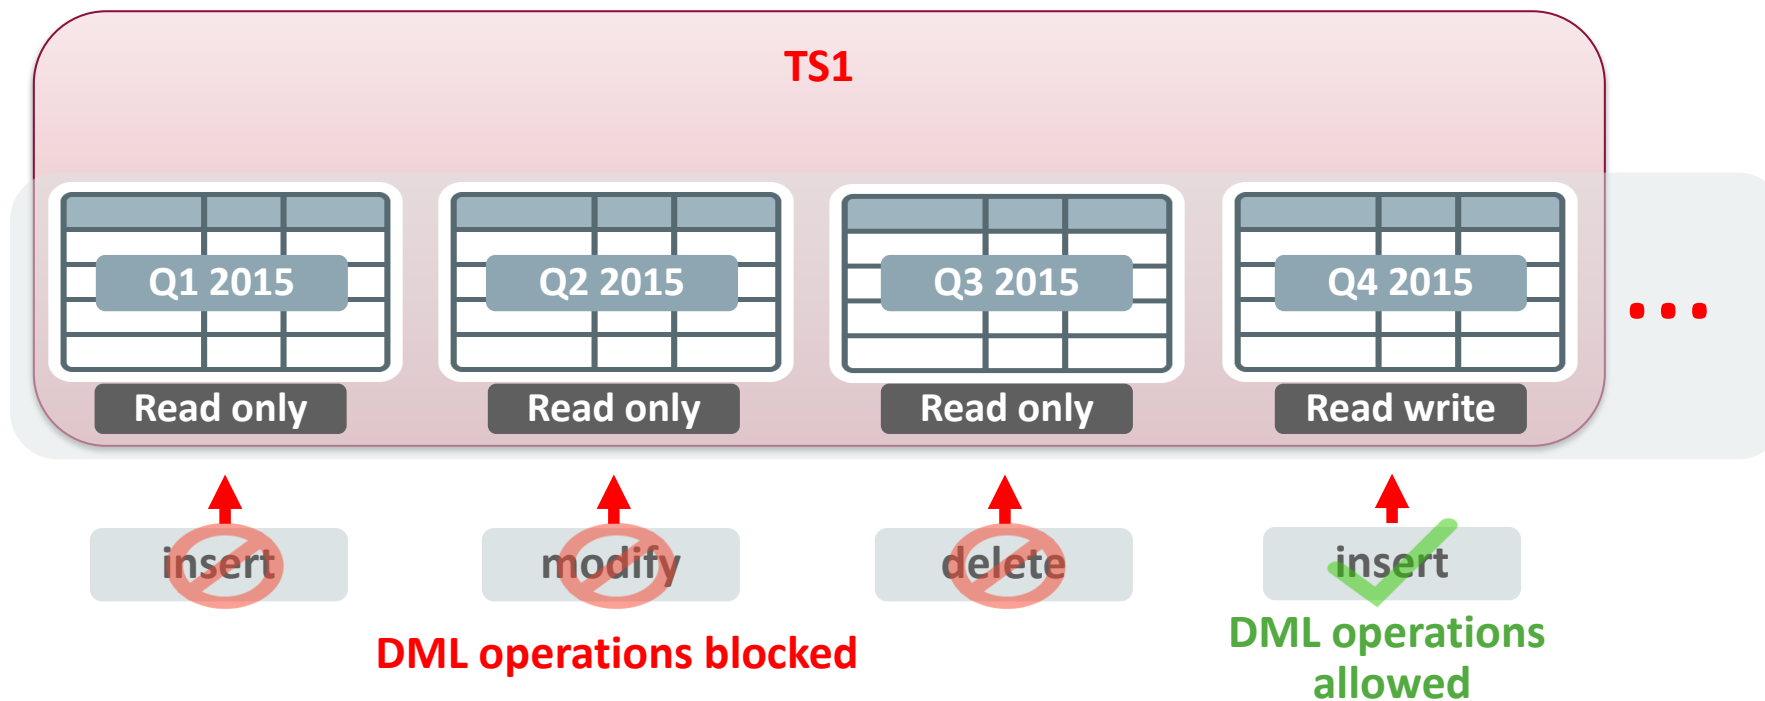
Read only partitions introduced in Oracle Database 12c Release 2

Read Only Tablespaces and Partitions



- Partitions and subpartitions can be placed in read only tablespaces
- Any attempt to alter data in a read only tablespace will result in an error

Read Only Partitions



- **Partitions and subpartitions can be set to read only or read write**
- Any attempt to alter data in a read only partition will result in an error

Read Only Object versus Read Only Tablespace

- Read Only Tablespaces protect physical storage from updates
 - DDL operations that are not touching the storage are allowed
 - E.g. ALTER TABLE SET UNUSED, DROP TABLE
 - No guaranteed data immutability
- Read Only Objects protect data from updates
 - ‘Data immutability’
 - Does not prevent changes on storage
 - E.g. ALTER TABLE MOVE COMPRESS, ALTER TABLE MERGE PARTITIONS

Read Only Partitions

- Read only attribute guarantees data immutability
 - “SELECT <column_list> FROM <table>” will always return the same data set after a table or [sub]partition is set to read only
- Data immutability does not prevent all structural DDL for a table
 - ADD and MODIFY COLUMN are allowed and do not violate data immutability of existing data
 - Others like DROP/RENAME/SET UNUSED COLUMN are forbidden
 - DROP [read only] PARTITION forbidden, too - - violates data immutability of the table



Create Table for Exchange



Create Table for Exchange

- Simple DDL command
- Ensures both the semantic and internal table shape are identical so partition exchange command will always succeed
- Operates like a special CREATE TABLE AS SELECT operation
- Always creates an empty table

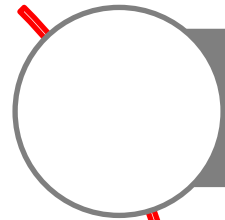
Create Table for Exchange

Syntax Example

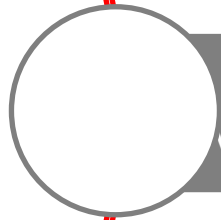
```
CREATE TABLE emp_ex TABLESPACE sysaux  
FOR EXCHANGE WITH emp;
```

Partitioning Summary

What is Oracle Partitioning?

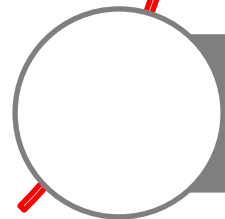


Powerful functionality to logically divide objects into smaller pieces



Key requirement for large databases needing high performance and availability

I high

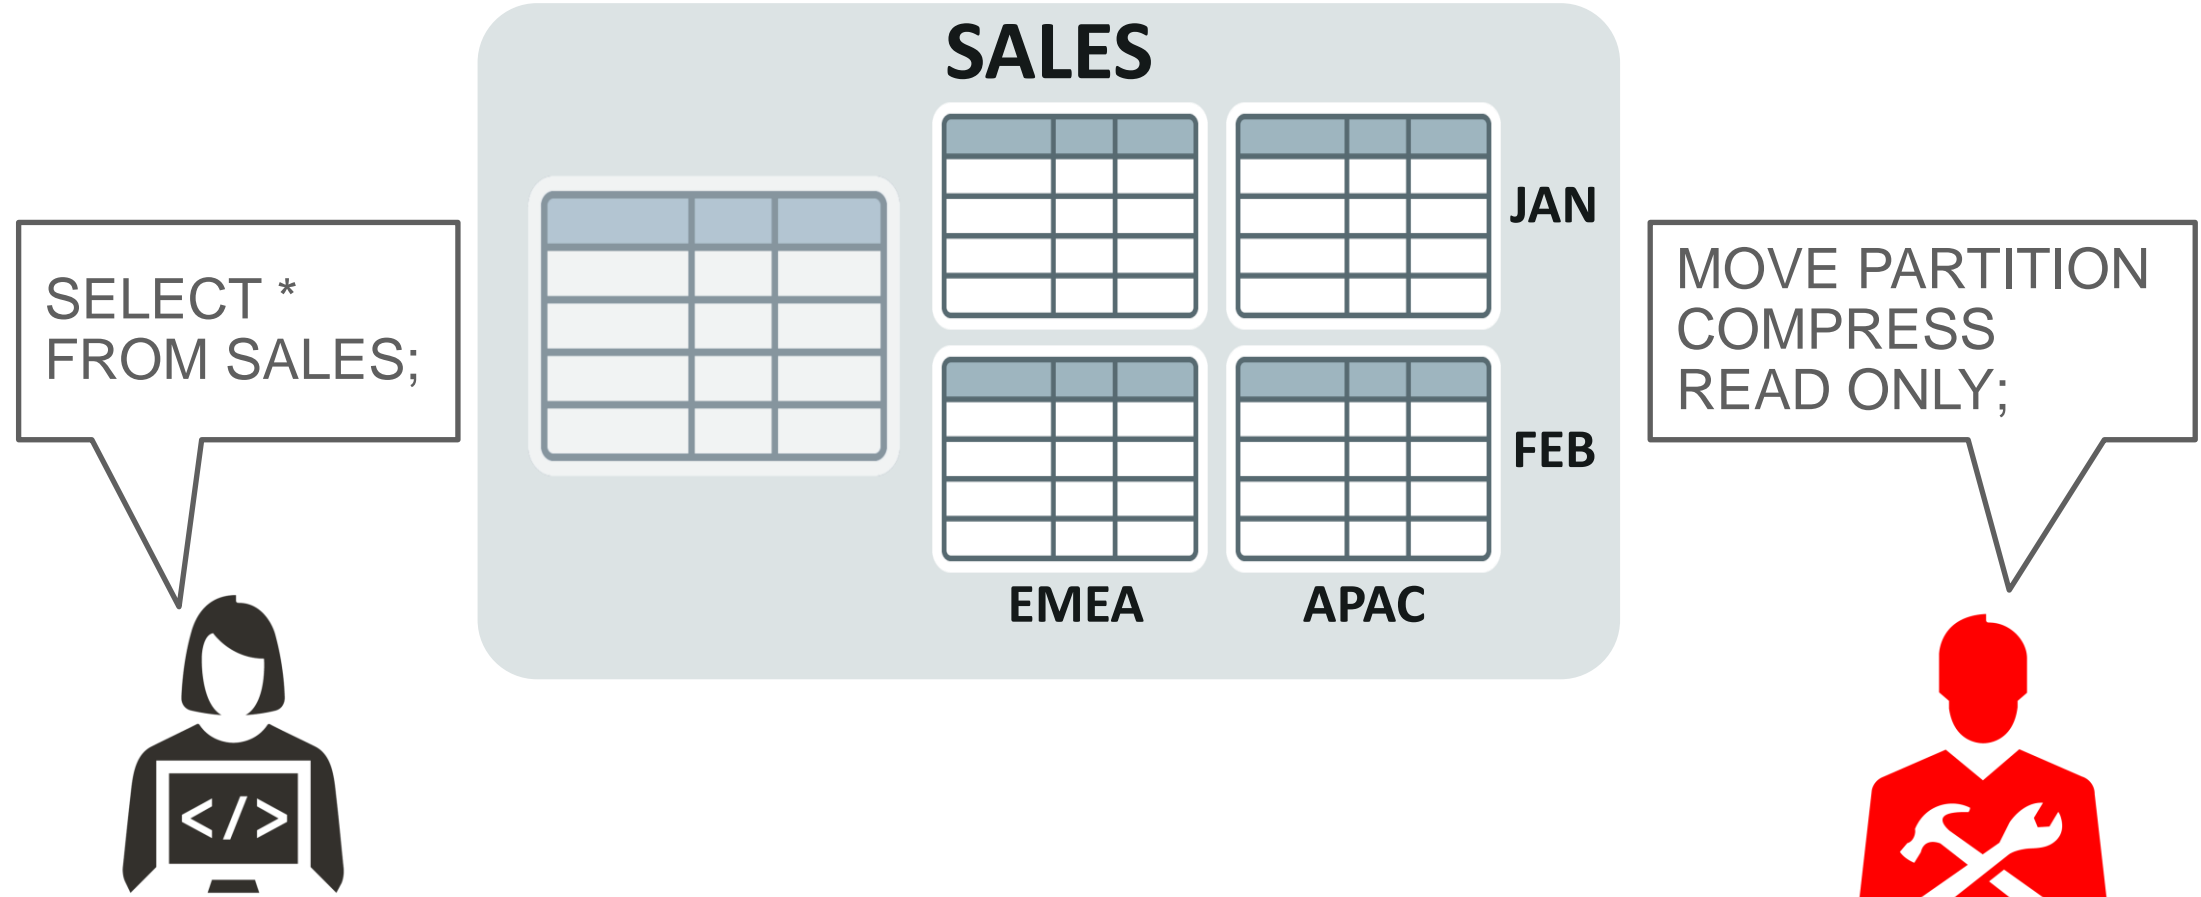


Driven by business requirements

Why use Oracle Partitioning?

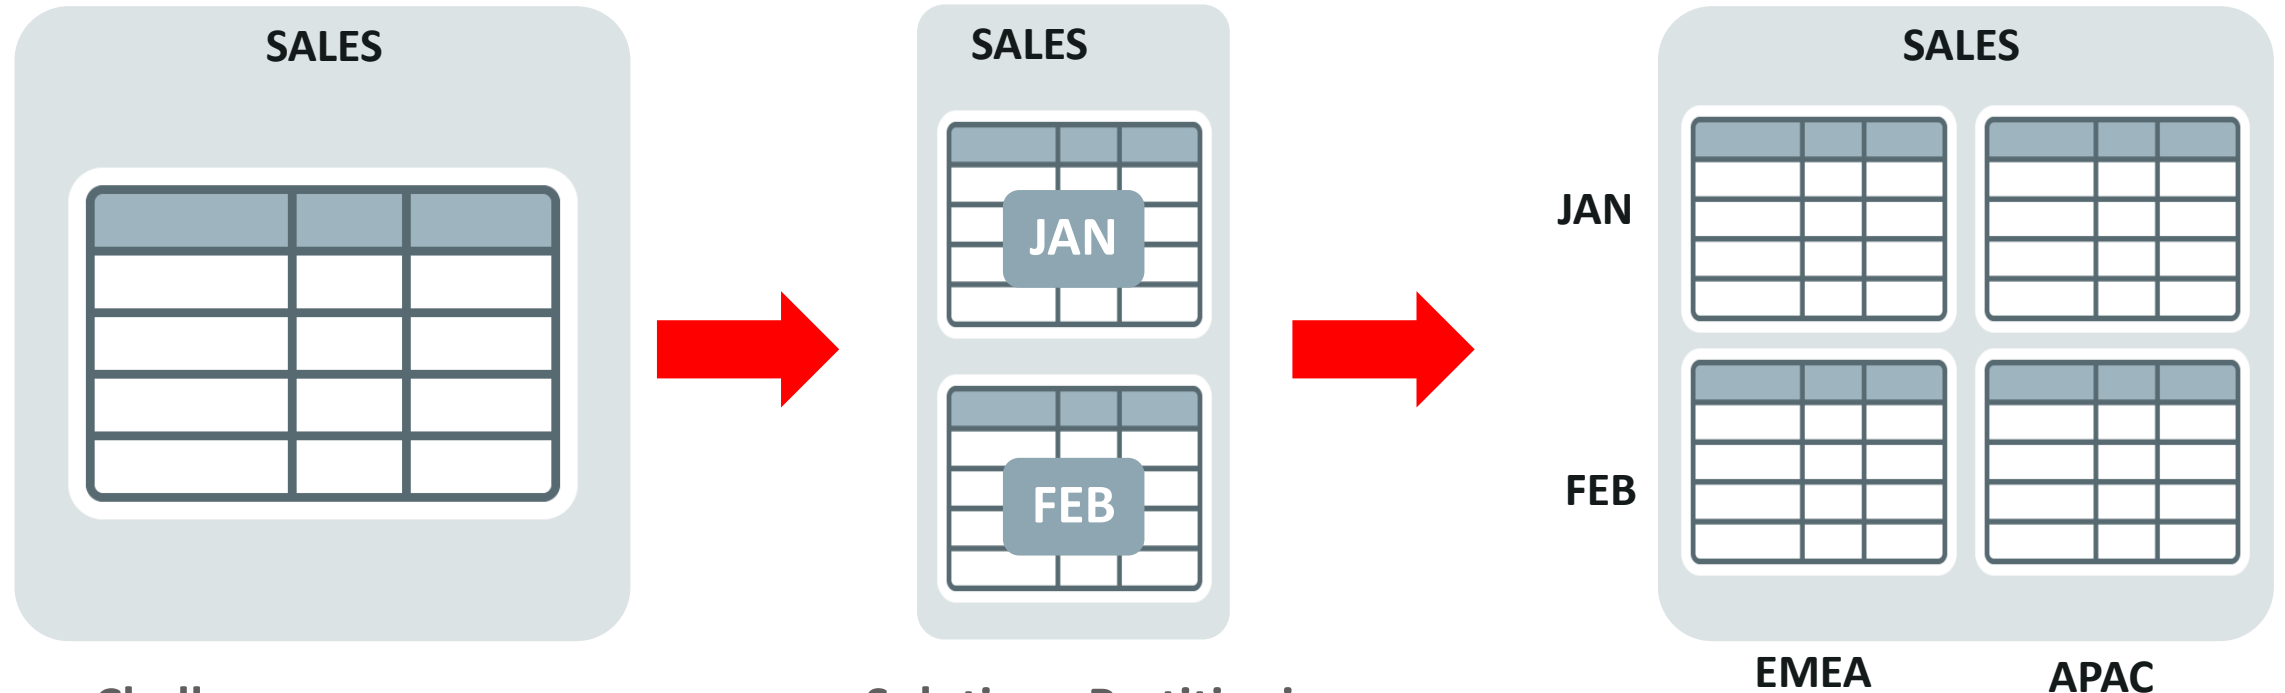
- ↑ Performance – lowers data access times
- ↑ Availability – improves access to critical information
- ↓ Costs – leverages multiple storage tiers
- ✓ Easy Implementation – requires no changes to applications and queries
- ✓ Mature Feature – supports a wide array of partitioning methods
- ✓ Well Proven – used by thousands of Oracle customers

The two Personalities of Partitioning



How does Partitioning work?

Enables large databases and indexes to be split into smaller, more manageable pieces



Challenges:

Large tables are difficult to manage

Solution: Partitioning

- Divide and conquer
- Easier data management
- Improve performance

Partitioning Benefits



Increased Performance

Only work on the data that is relevant

Partitioning enables data management operations such as...

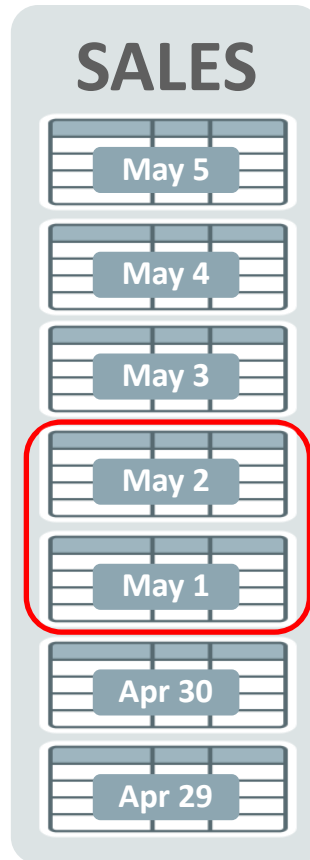
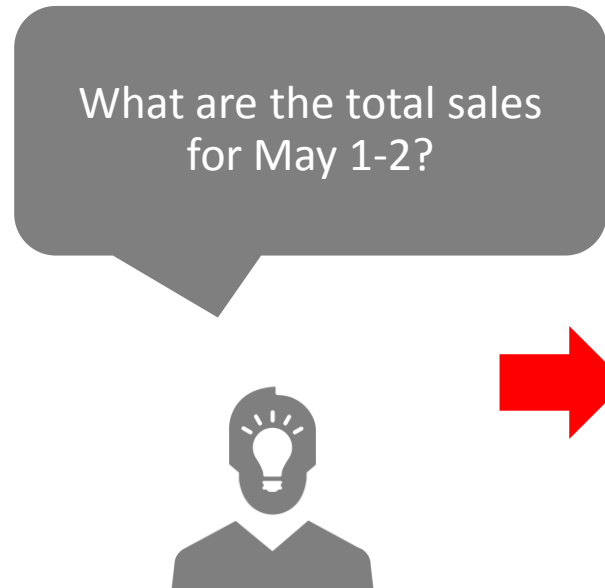
- Data loads, joins and pruning,
- Index creation and rebuilding,
- Backup and recovery

...At partition level instead of on the entire table

Result: Order of magnitude gains on performance

Increased Performance - Example

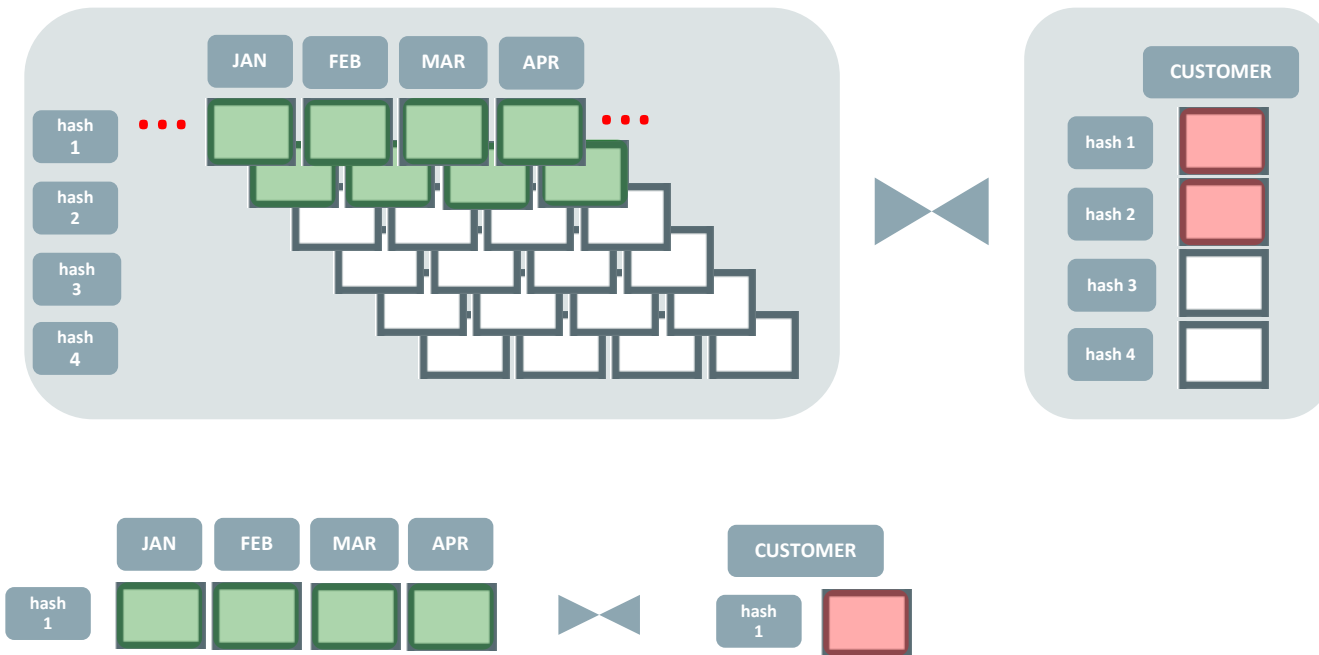
Partition Pruning



- Partition elimination
 - Dramatically reduces amount of data retrieved from storage
 - Performs operations only on relevant partitions
 - Transparently improves query performance and optimizes resource utilization

Increased Performance - Example

Partition-wise joins



- A large join is divided into multiple smaller joins, executed in parallel
 - # of partitions to join must be a multiple of DOP
 - Both tables must be partitioned the same way on the join column

Decreased Costs

Store data in the most appropriate manner

Partitioning finds the balance between...

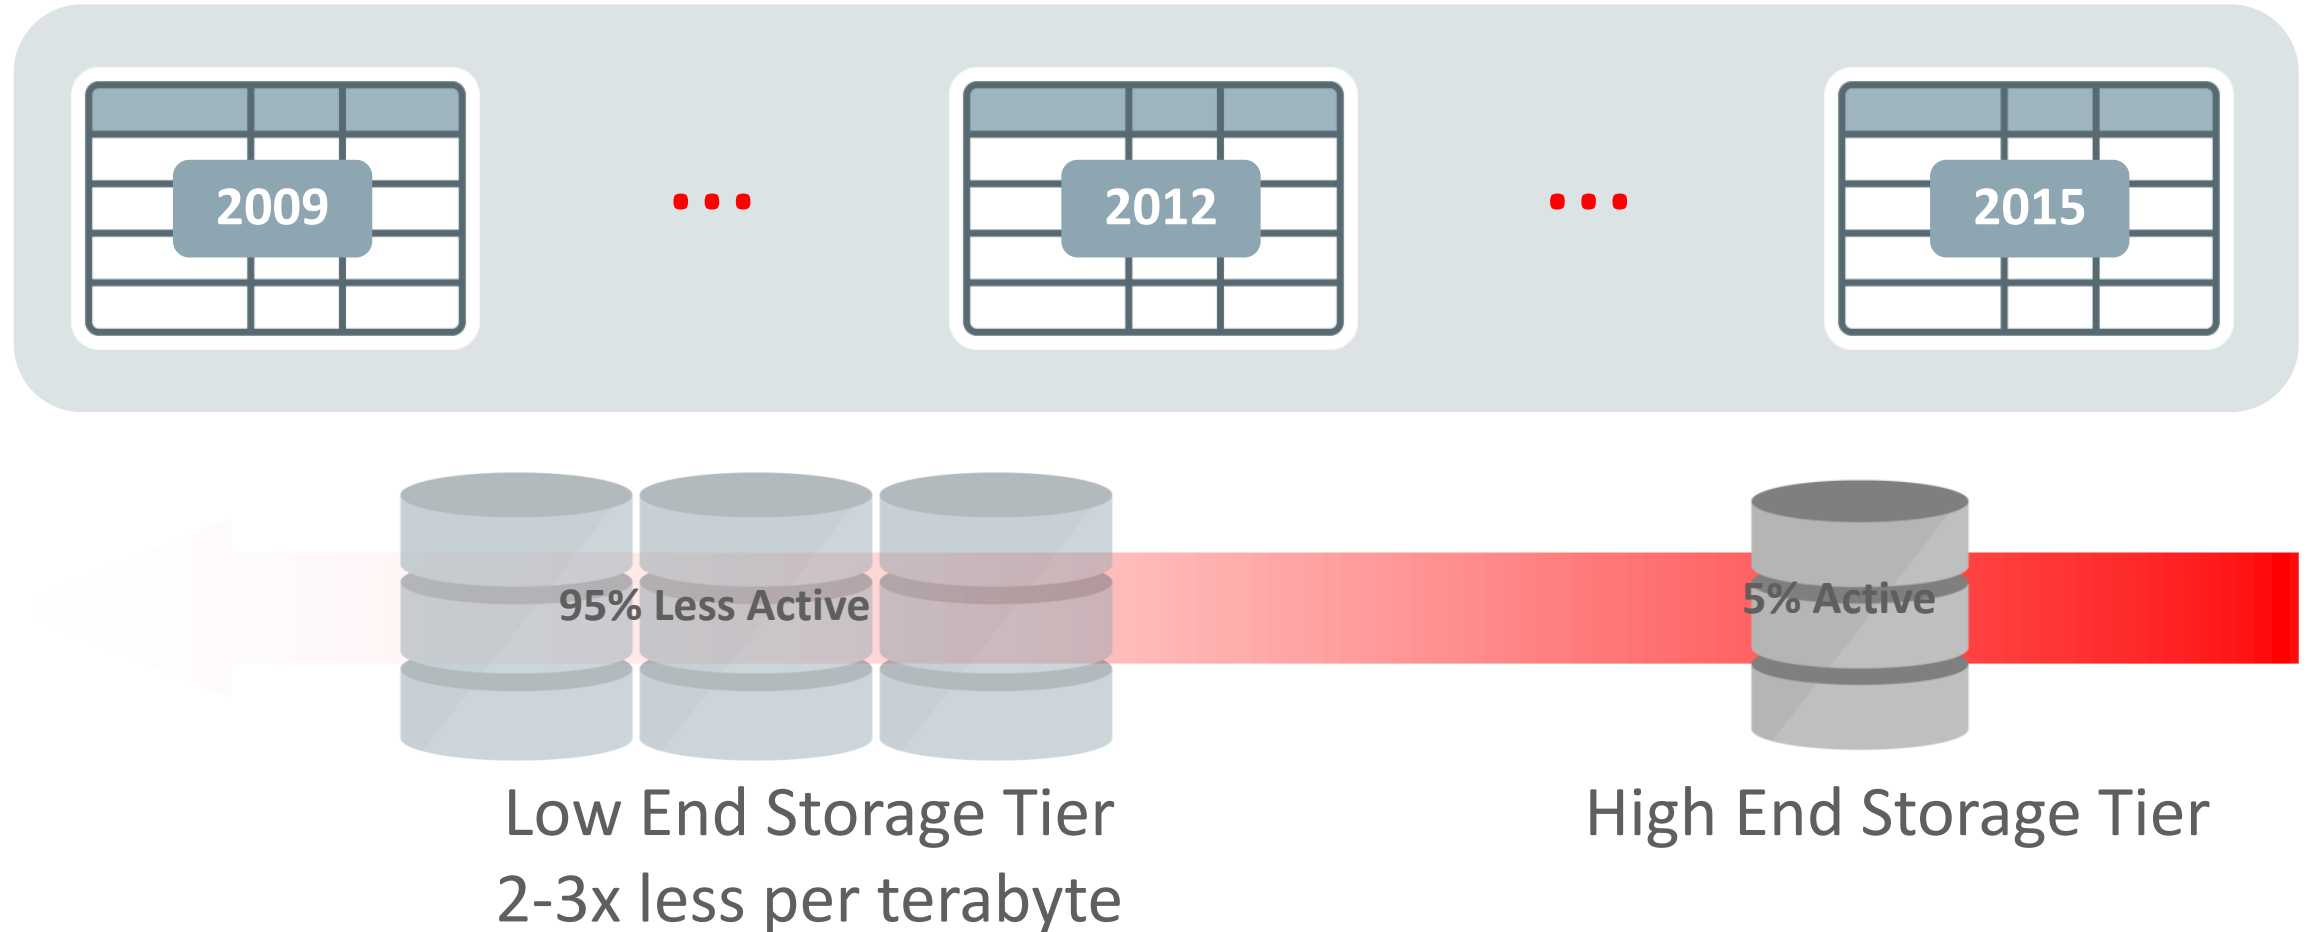
- data importance,
- storage performance,
- storage reliability,
- storage form

... allowing you to leverage multiple storage tiers

Result: Reduce storage costs by 2x or more

Decreased Costs - Example

Partition for Tiered Storage



Increased Availability

Individual partition manageability

Partitioning reduces...

- Maintenance windows
- Impact of scheduled downtime and failures,
- Recovery times

... if critical tables and indexes are partitioned

Result: Improves access to critical information

Increased Availability - Example

Partition for Manageability/Availability



Other partitions visible and usable

Easy Implementation

Transparent to applications

- Partitioning requires NO changes to applications and queries
 - Adjustments might be necessary to fully exploit the benefits of Partitioning

Mature, Well Proven Functionality

Over a decade of development

- Used by tens of thousands of Oracle customers
- Supports a wide array of partitioning methods

Oracle Partitioning

Over a decade of development **and better than ever before**

	Core functionality	Performance	Manageability
Oracle 8.0	Range partitioning Local and global Range indexing	Static partition pruning	Basic maintenance: ADD, DROP, EXCHANGE
Oracle 8i	Hash partitioning Range-Hash partitioning	Partition-wise joins Dynamic partition pruning	Expanded maintenance: MERGE
Oracle 9i	List partitioning		Global index maintenance
Oracle 9i R2	Range-List partitioning	Fast partition SPLIT	
Oracle 10g	Global Hash indexing		Local Index maintenance
Oracle 10g R2	1M partitions per table	Multi-dimensional pruning	Fast DROP TABLE
Oracle 11g	Virtual column based partitioning More composite choices Reference partitioning		Interval partitioning Partition Advisor Incremental stats mgmt
Oracle 11g R2	Hash-* partitioning Expanded Reference partitioning	“AND” pruning	Multi-branch execution (aka table or-expansion)
Oracle 12c R1	Interval-Reference partitioning	Partition Maintenance on multiple partitions Asynchronous global index maintenance Zone maps	Online partition MOVE Cascading TRUNCATE Partial indexing
Oracle 12c R2	Auto-list partitioning Multi-column list partitioning Partitioned external tables	Online partition maintenance operations Online table conversion to partitioned table Reduced cursor invalidations for DDL's	Filtered partition maintenance Read only partitions Create table for exchange

Partitioning Concepts

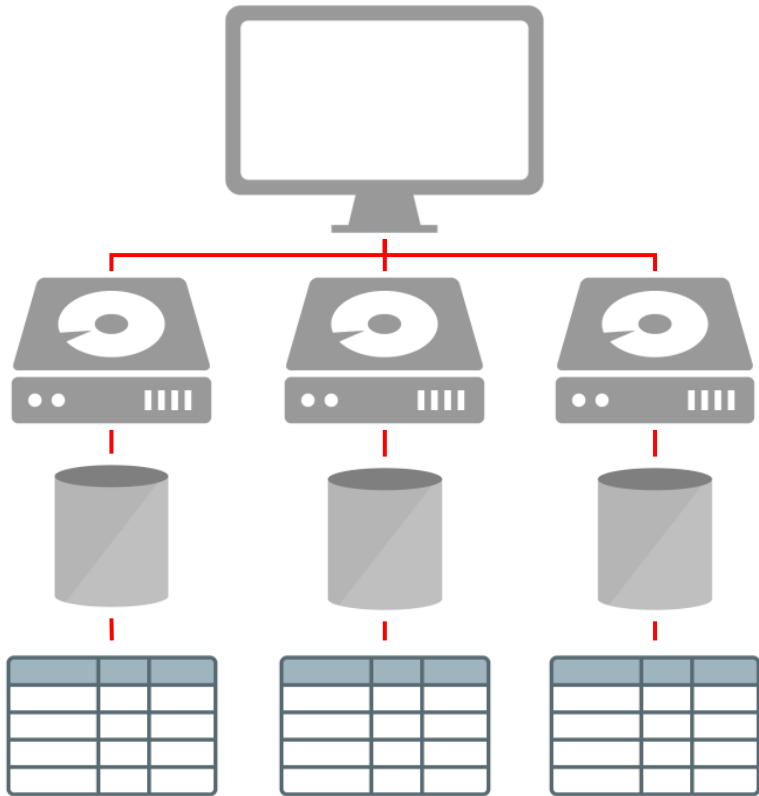
def **Par•ti•tion**

To divide (something) into parts

– “Miriam Webster Dictionary”

Physical Partitioning

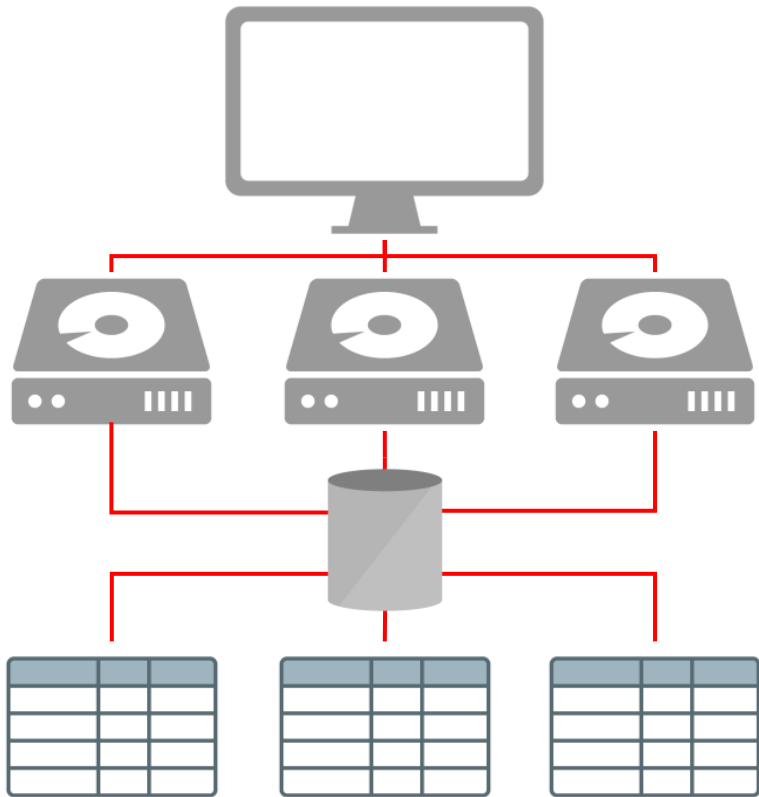
Shared Nothing Architecture



- Fundamental system setup requirement
 - Node owns piece of DB
 - Enables parallelism
- Number of partitions is equivalent to minimum required parallelism
 - Always needs HASH or random distribution
- Equally sized partitions per node required for proper load balancing

Logical Partitioning

Shared Everything Architecture - Oracle

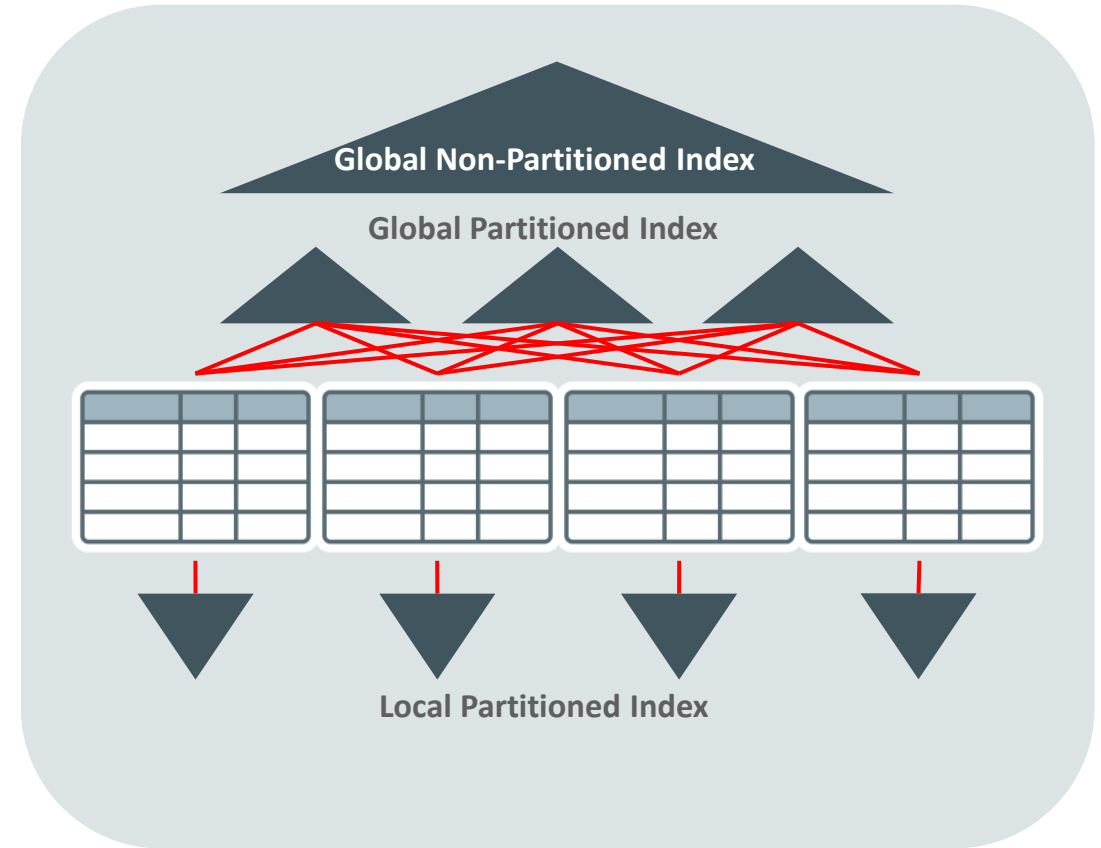


- Does not underlie any constraints
 - SMP, MPP, Cluster, Grid does not matter
- Purely based on the business requirement
 - Availability, Manageability, Performance
- Beneficial for every environment
 - Provides the most comprehensive functionality

Partitioning Methods

What can be partitioned?

- Tables
 - Heap tables
 - Index-organized tables
- Indexes
 - Global Indexes
 - Local Indexes
- Materialized Views
- Hash Clusters



Partitioning Methods

Single-level partitioning

- Range
- List
- Hash

Composite-level partitioning

- [Range | List | Hash | Interval] –
[Range | List | Hash]

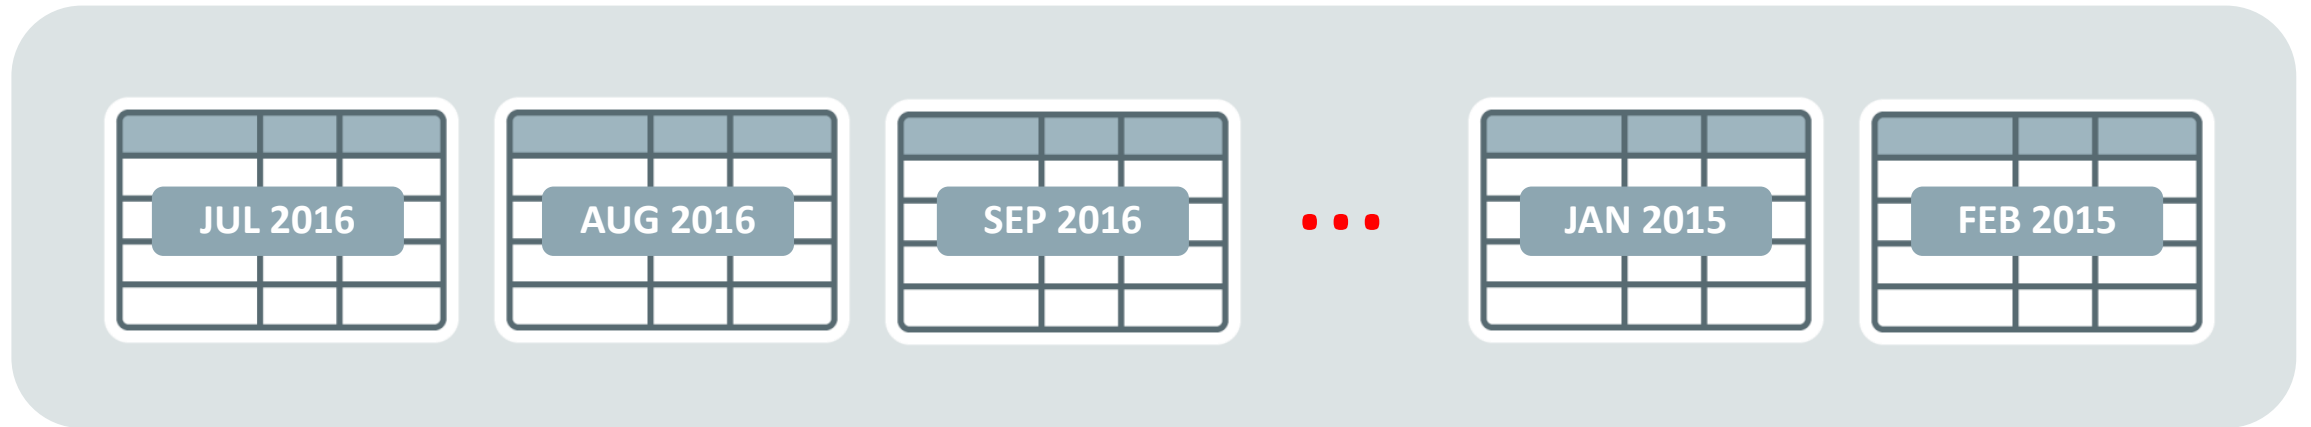
Partitioning extensions

- Interval
- Reference
- Interval Reference
- Virtual Column Based

Range Partitioning

Introduced in Oracle 8.0

Range Partitioning

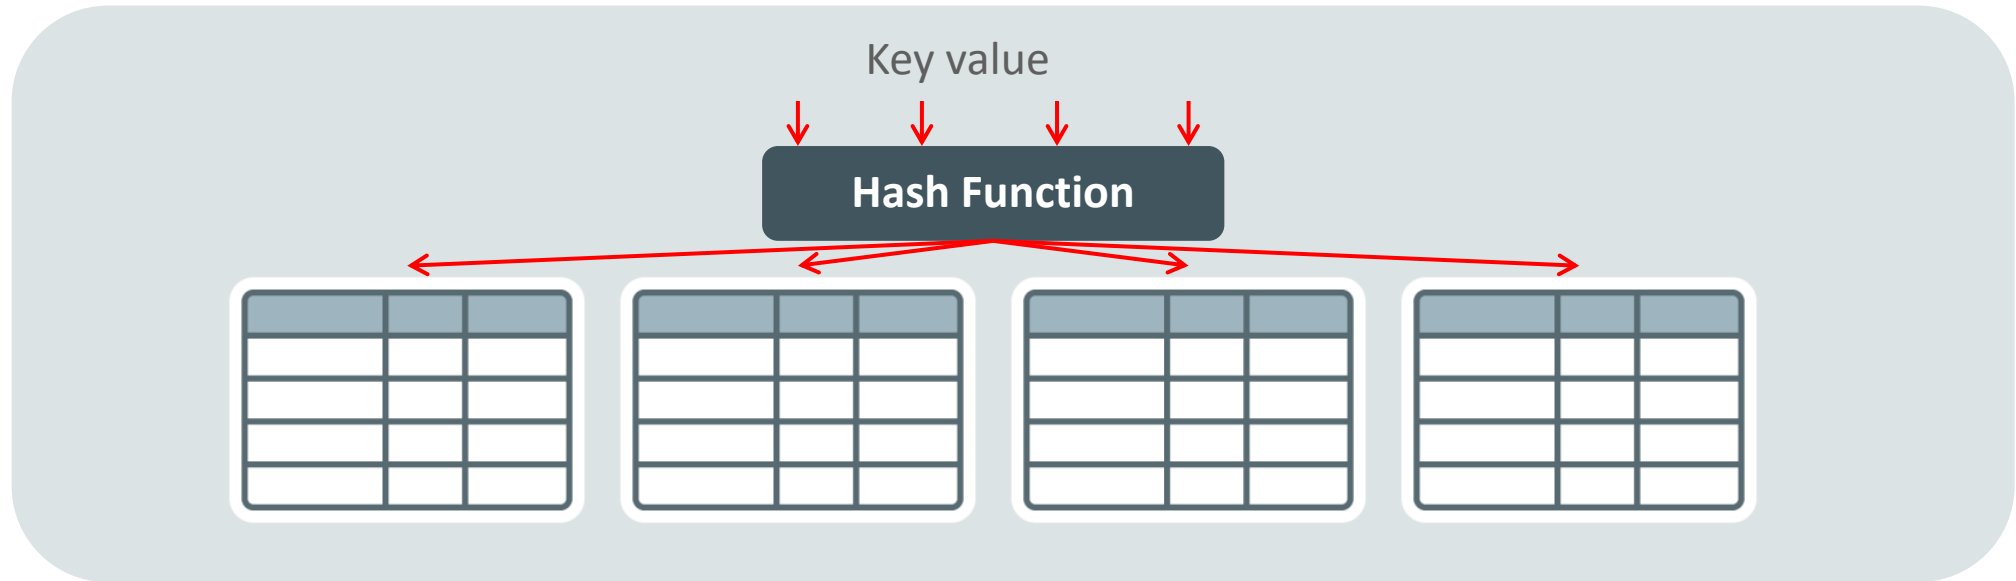


- Data is organized in ranges
 - Lower boundary derived by upper boundary of preceding partition
 - No gaps
- Ideal for chronological data

Hash Partitioning

Introduced in Oracle 8i (8.1)

Hash Partitioning

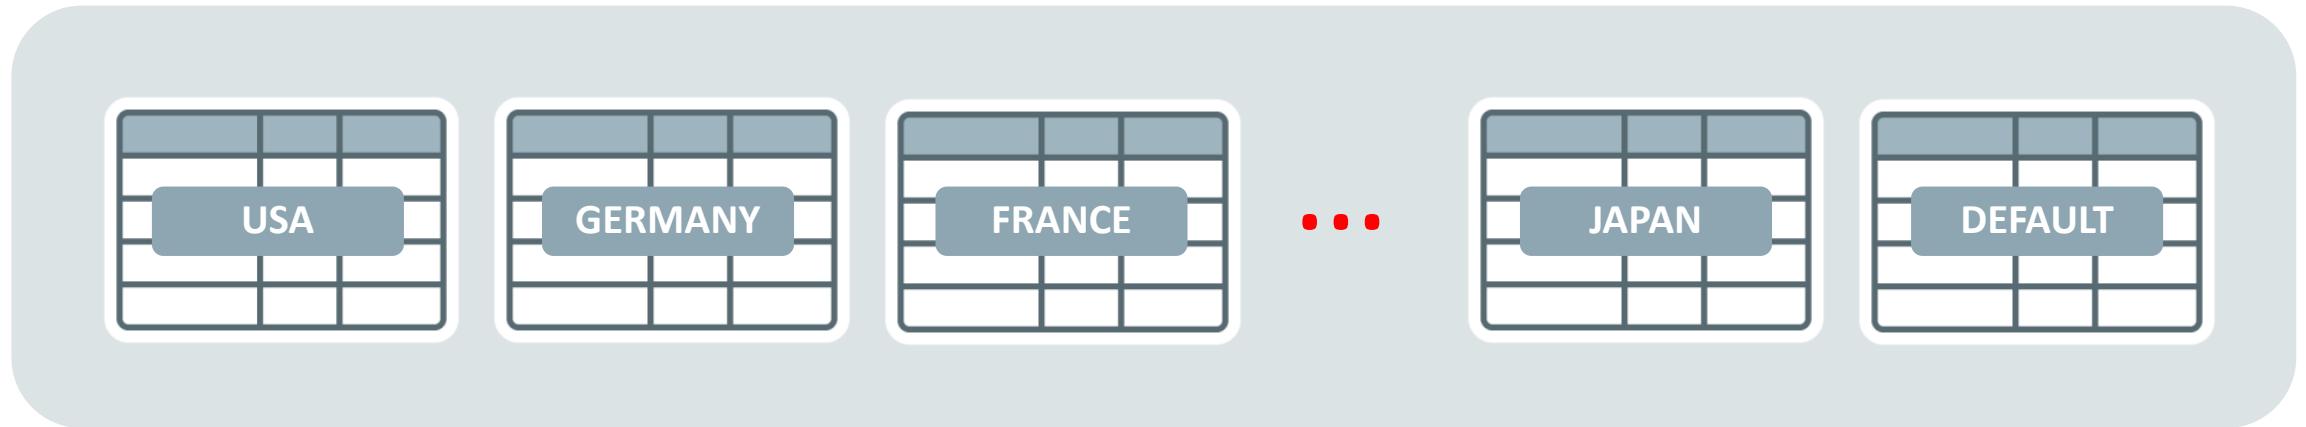


- Data is placed based on hash value of partition key
 - Number of hash buckets equals number of partitions
- Ideal for equal data distribution
 - Number of partitions should be a power of 2 for equal data distribution

List Partitioning

Introduced in Oracle 9i (9.0)

List Partitioning



- Data is organized in lists of values
 - One or more unordered distinct values per list
 - Functionality of DEFAULT partition (Catch-it-all for all unspecified values)
- Ideal for segmentation of distinct values, e.g. region

Interval Partitioning

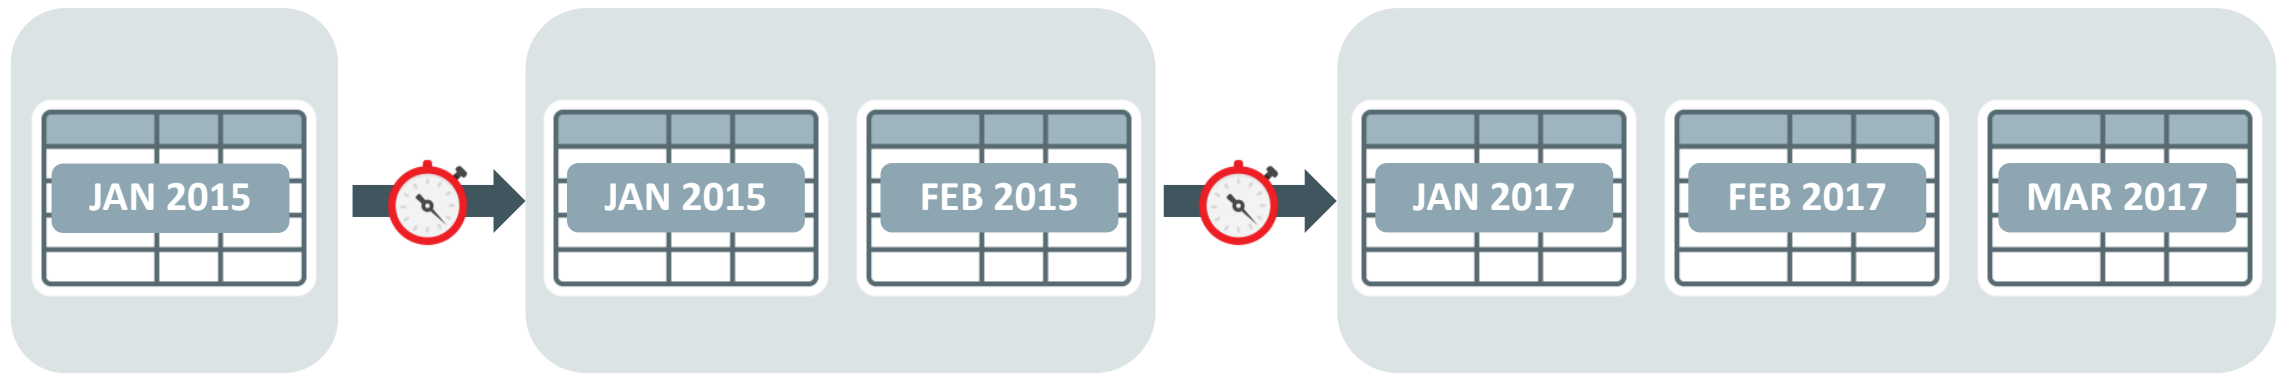
Introduced in Oracle 11g Release 1 (11.1)

Interval Partitioning

- Extension to Range Partitioning
- Full automation for equi-sized range partitions
- Partitions are created as metadata information only
 - Start Partition is made persistent
- Segments are allocated as soon as new data arrives
 - No need to create new partitions
 - Local indexes are created and maintained as well

No need for any partition management

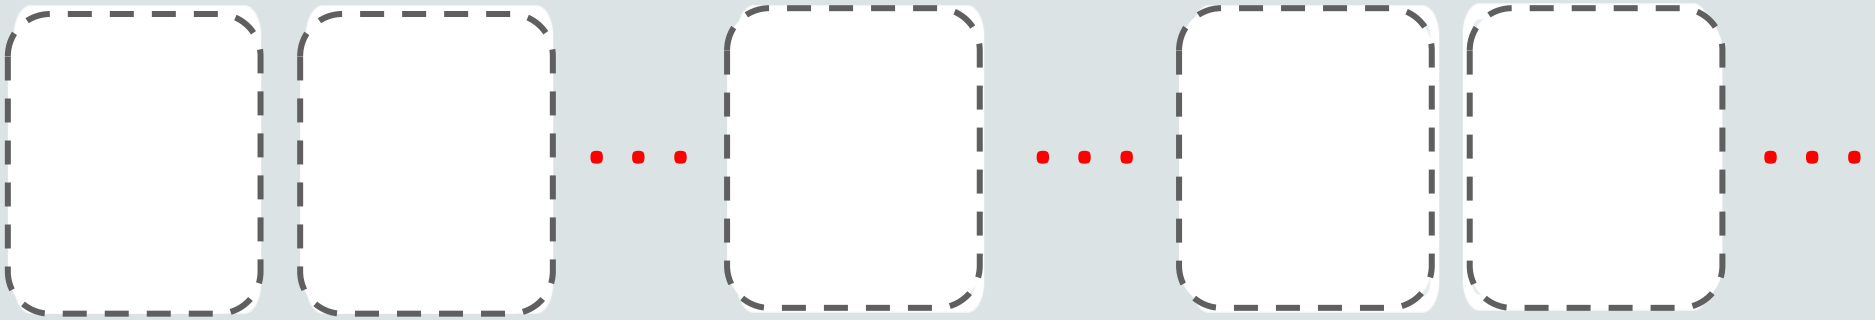
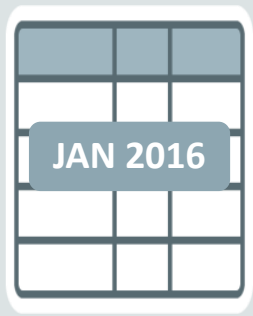
Interval Partitioning



- Partitions are created automatically as data arrives
 - Extension to RANGE partitioning

Interval Partitioning

As easy as One, Two, Three...

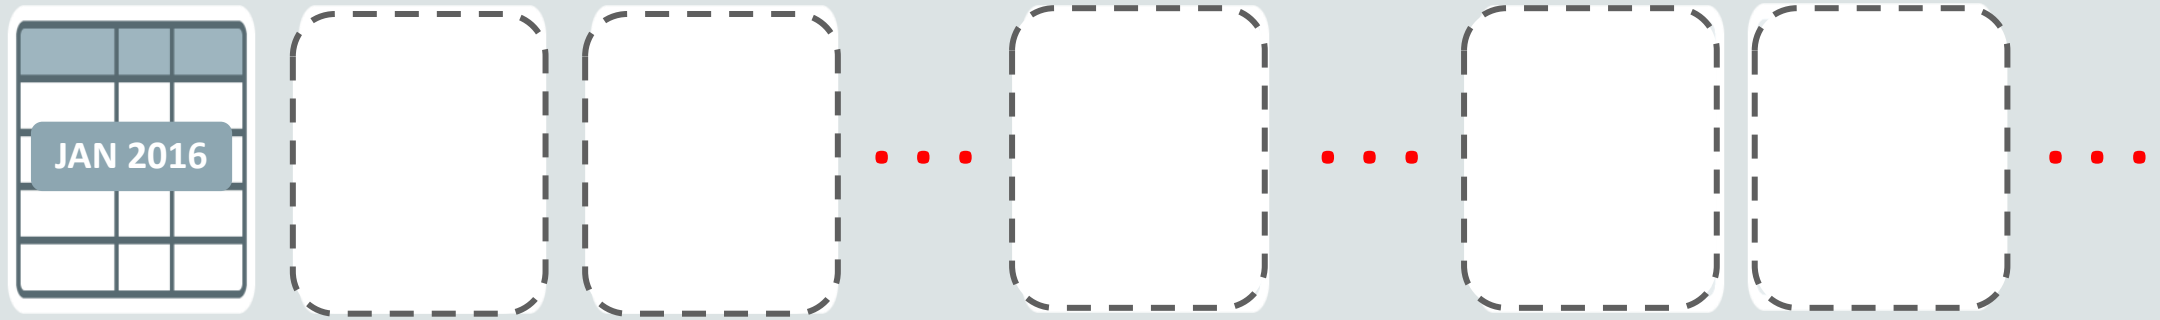


**First
partition
is created**

```
CREATE TABLE sales (order_date DATE, ...)
PARTITION BY RANGE (order_date)
INTERVAL (NUMTOYMINTERVAL(1, 'month'))
(PARTITION p_first VALUES LESS THAN ('01-JAN-2013'));
```

Interval Partitioning

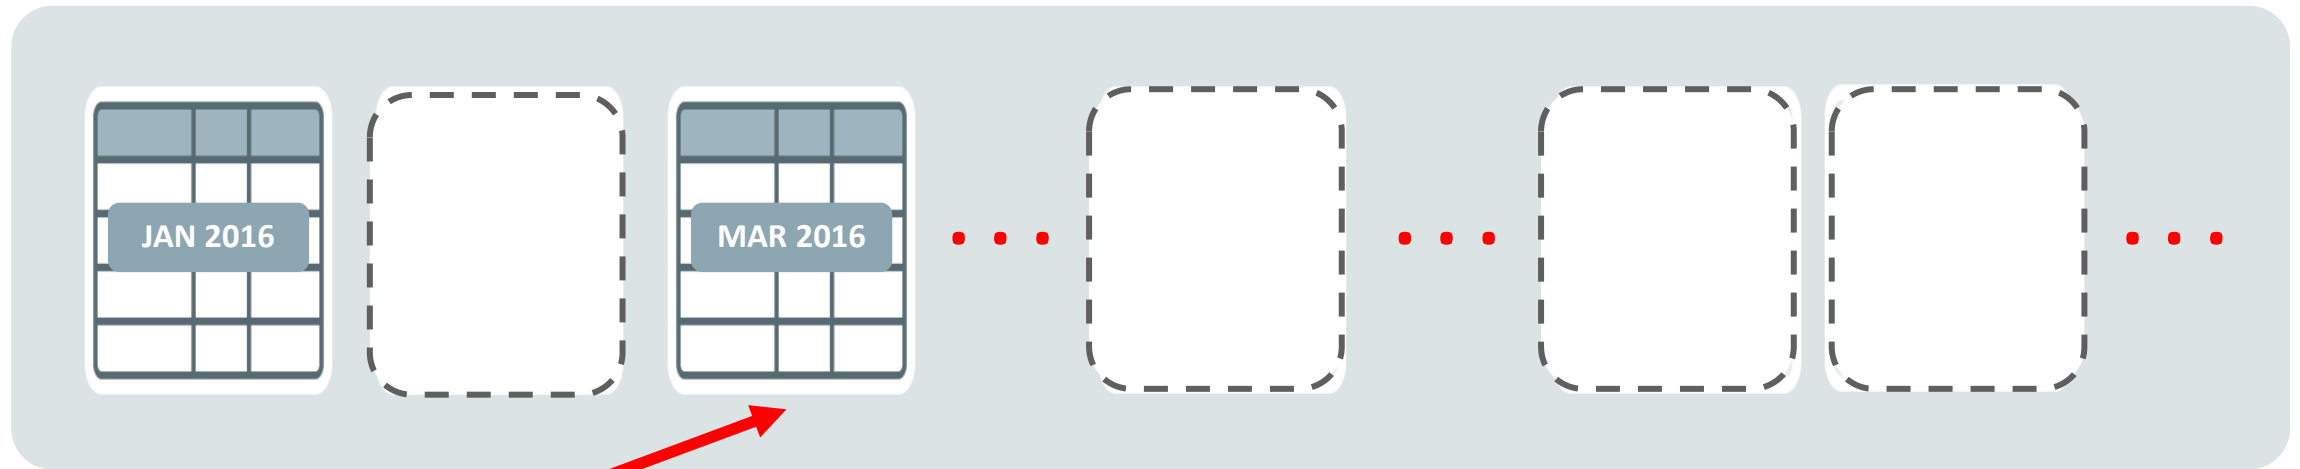
As easy as One, Two, Three...



Other partitions only exist in table metadata

Interval Partitioning

As easy as One, Two, Three...

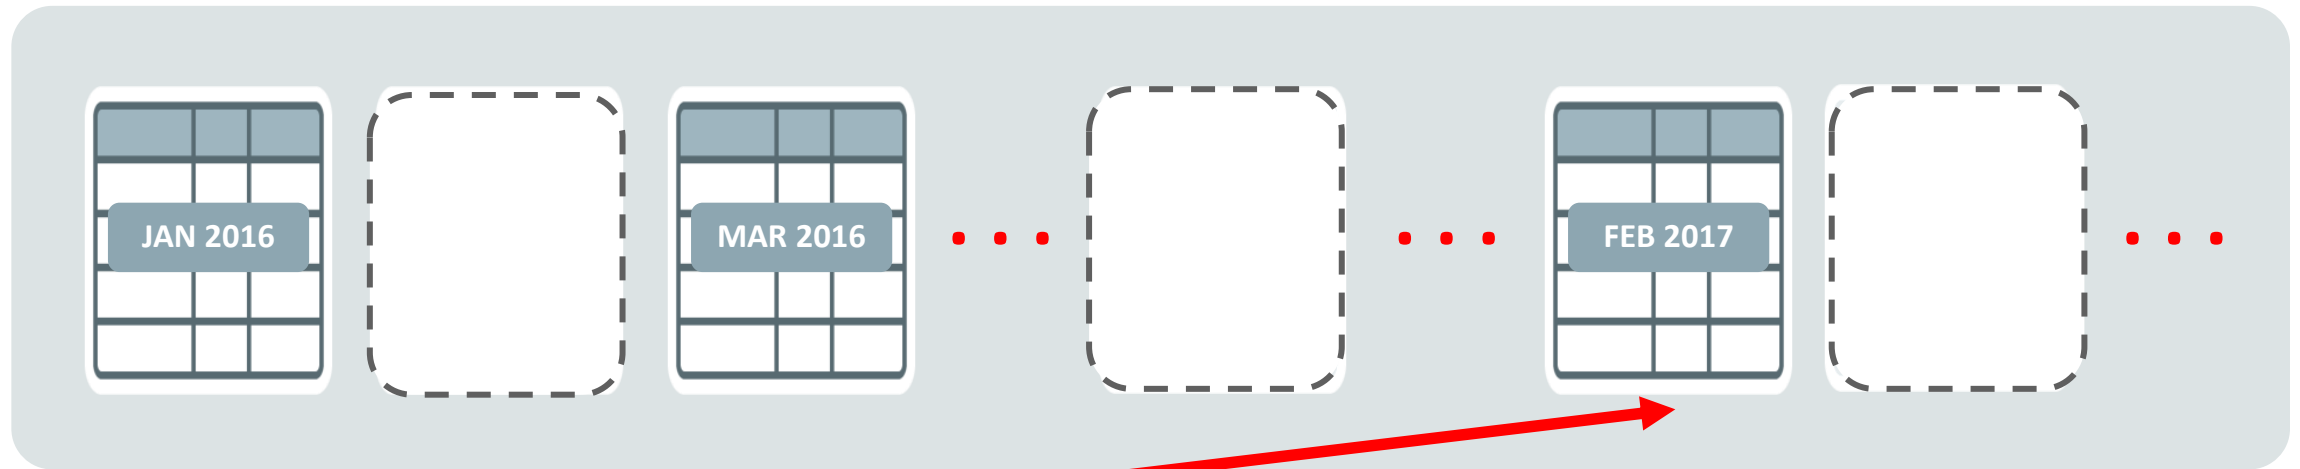


**New partition is
automatically instantiated**

```
INSERT INTO sales (order_date DATE, ...)
VALUES ('30-MAR-2013', ...);
```

Interval Partitioning

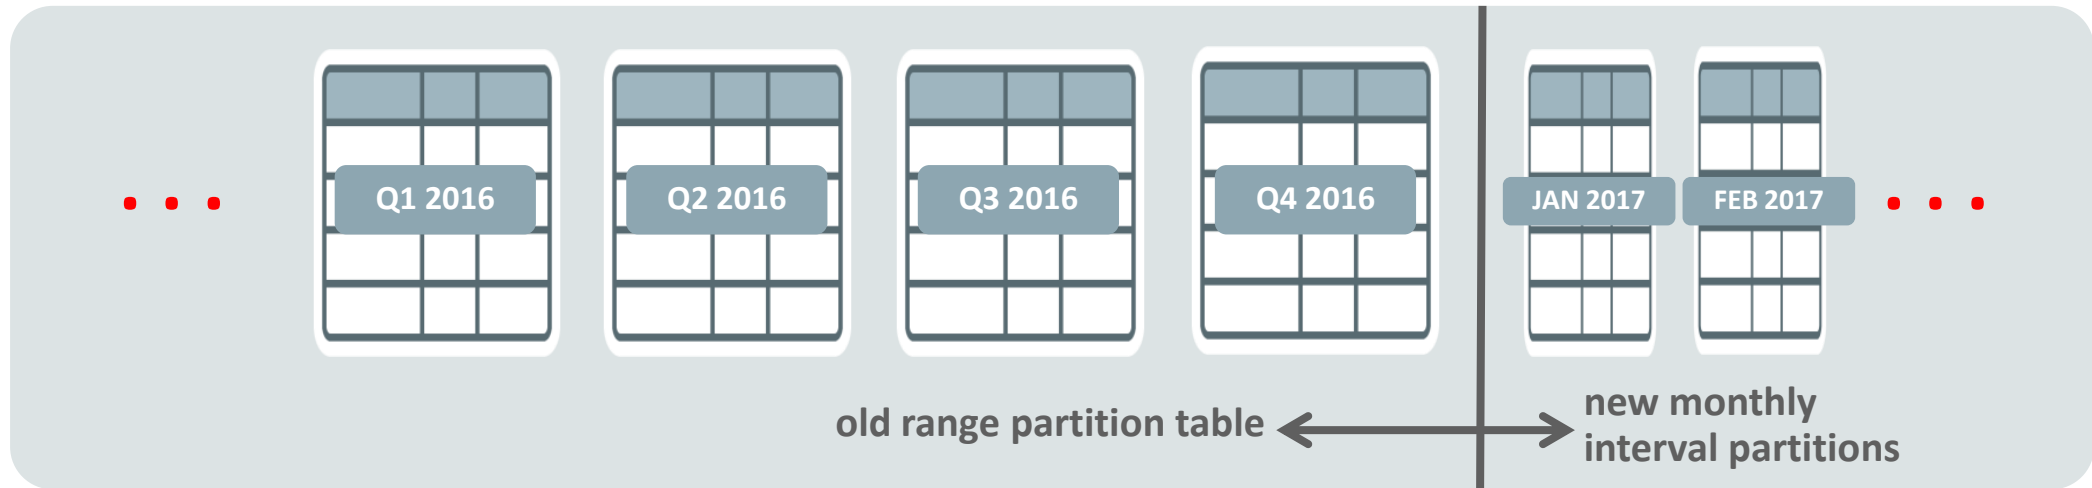
As easy as One, Two, Three...



**Whenever data for
a new partition arrives**

```
INSERT INTO sales (order_date DATE, ...)
VALUES ('04-FEB-2017',...);
```

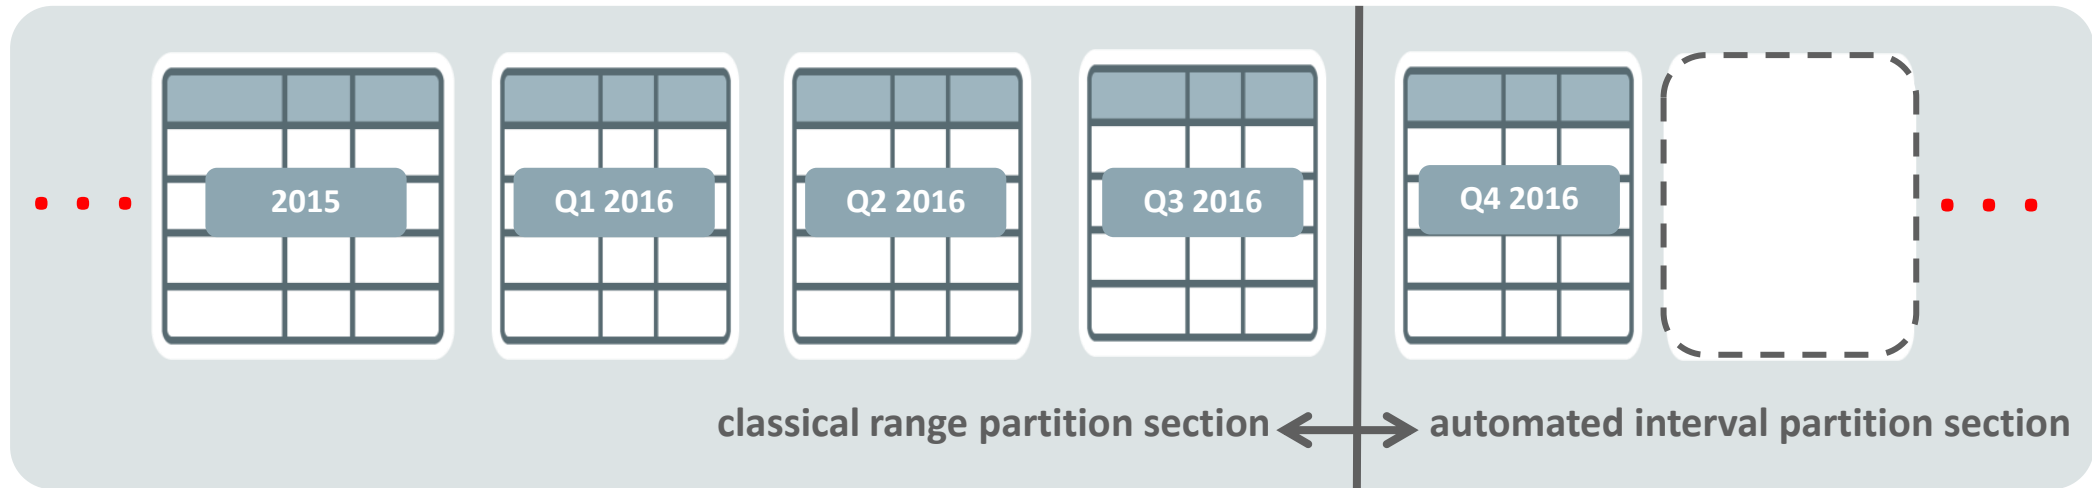
Interval Partitioning



- Range partitioned tables can be extended into interval partitioned tables
 - Simple metadata command
 - Investment protection

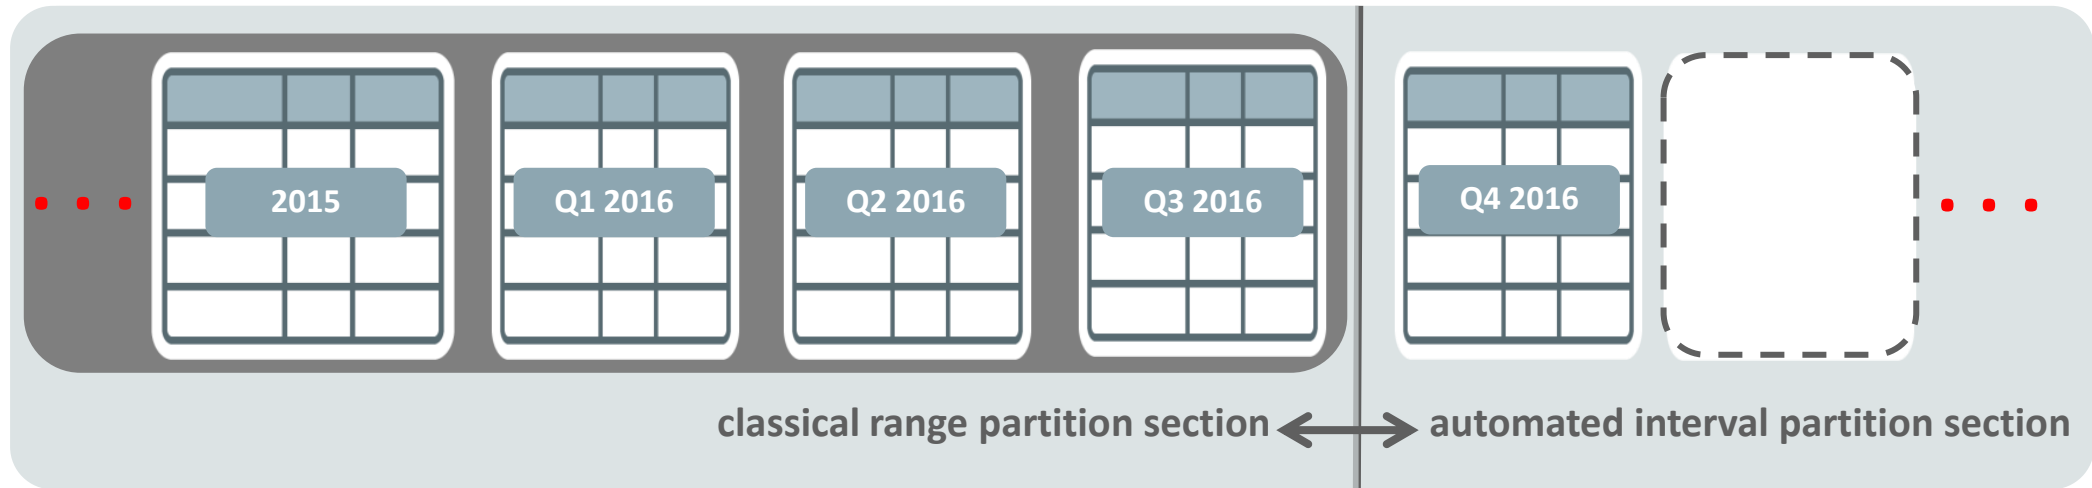
```
ALTER TABLE sales  
SET INTERVAL (NUMTOYMINTERVAL(1, 'month'));
```

Interval Partitioning



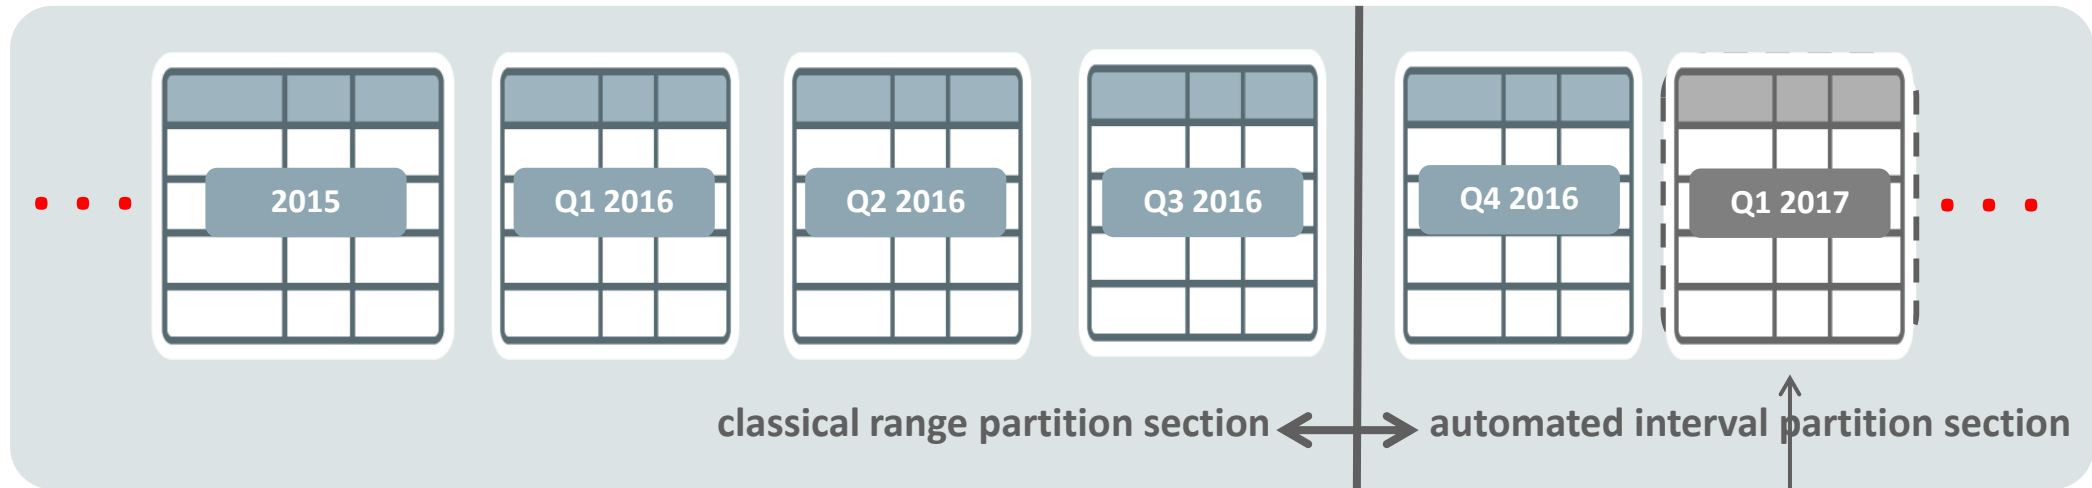
- Interval partitioned table has classical range and automated interval section
 - Automated new partition management plus full partition maintenance capabilities: ***“Best of both worlds”***

Interval Partitioning



1. Merge and move old partitions for ILM

Interval Partitioning



1. Merge and move old partitions for ILM
2. Insert new data
 - Automatic partition instantiation

Deferred Segment Creation vs Interval Partitioning

Interval Partitioning

- Maximum number of one million partitions are pre-defined
 - Explicitly defined plus interval-based partitions
- No segments are allocated for partitions without data
 - New record insertion triggers segment creation
- Ideal for “ever-growing” tables

“Standard” Partitioning with deferred segment creation

- Only explicitly defined partitions are existent
 - New partitions added via DDL
- No segments are allocated for partitions without data
 - New record insertion triggers segment creation when data matches pre-defined partitions
- Ideal for sparsely populated pre-defined tables

Difference Between Range and Interval

Interval Partitioning

- Full automation for equi-sized range partitions
- Partitions are created as metadata information only
 - Start Partition is made persistent
- Segments are allocated as soon as new data arrives
 - No need to create new partitions
 - Local indexes are created and maintained as well
- Interval Partitioning is almost a transparent extension to range partitioning
 - .. But interval implementation introduces some subtle differences

Interval versus Range Partitioning

Partition bounds

- Interval partitions have lower and upper bound
 - No infinite upper bound (MAXVALUES)
- Range partitions only have upper bounds
 - Lower bound derived by previous partition
 - Upper bound infinite (MAXVALUES)

Partition naming

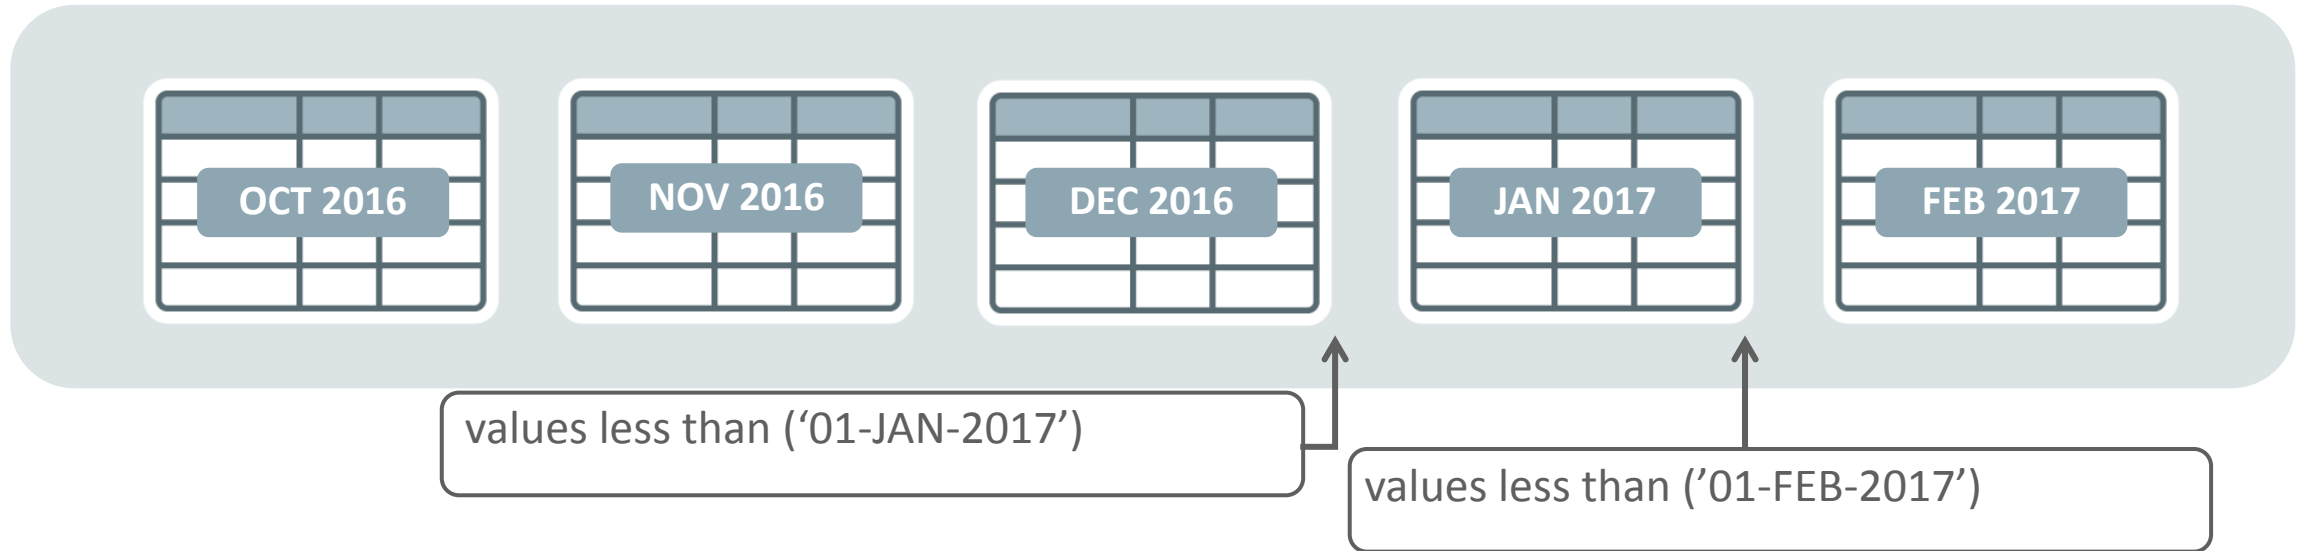
- Interval partitions cannot be named in advance
 - Use the PARTITION FOR (<value>) clause
- Range partitions must be named

Interval versus Range Partitioning, cont.

- Partition merge
 - Multiple non-existent interval partitions are silently merged
 - Only two adjacent range partitions can be merged at any point in time
- Number of partitions
 - Interval partitioned tables have always one million partitions
 - Non-existent partitions “exist” through INTERVAL clause
 - No MAXVALUE clause for interval partitioning
 - Maximum value defined through number of partitions and INTERVAL clause
 - Range partitioning can have up to one million partitions
 - MAXVALUE clause defines most upper partition

Interval Versus Range Partitioning

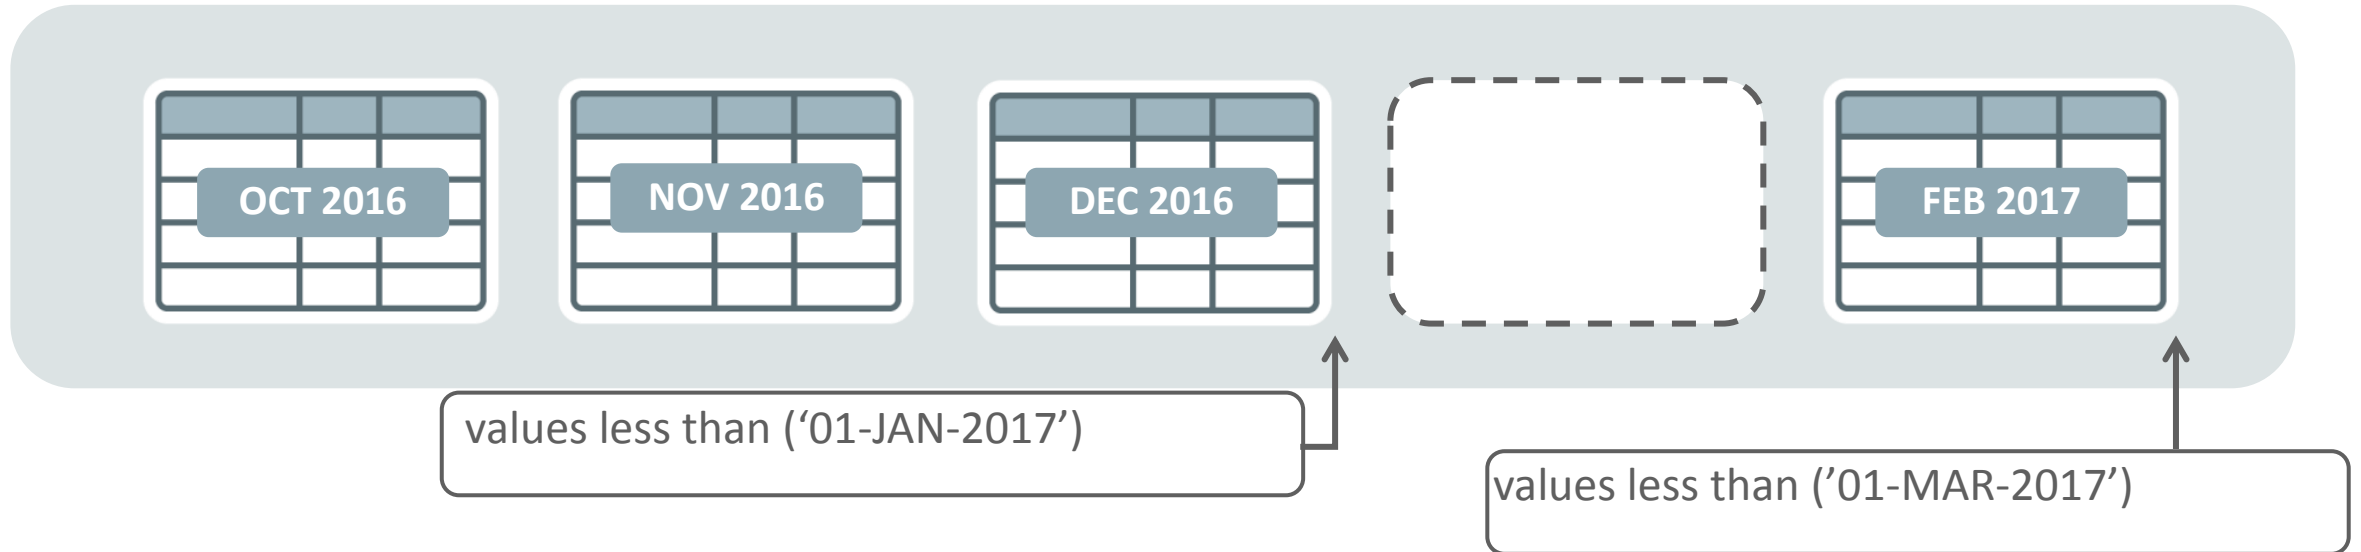
Partition Bounds for **Range Partitioning**



- Partitions only have upper bounds
 - Lower bound derived through upper bound of previous partition

Interval Versus Range Partitioning

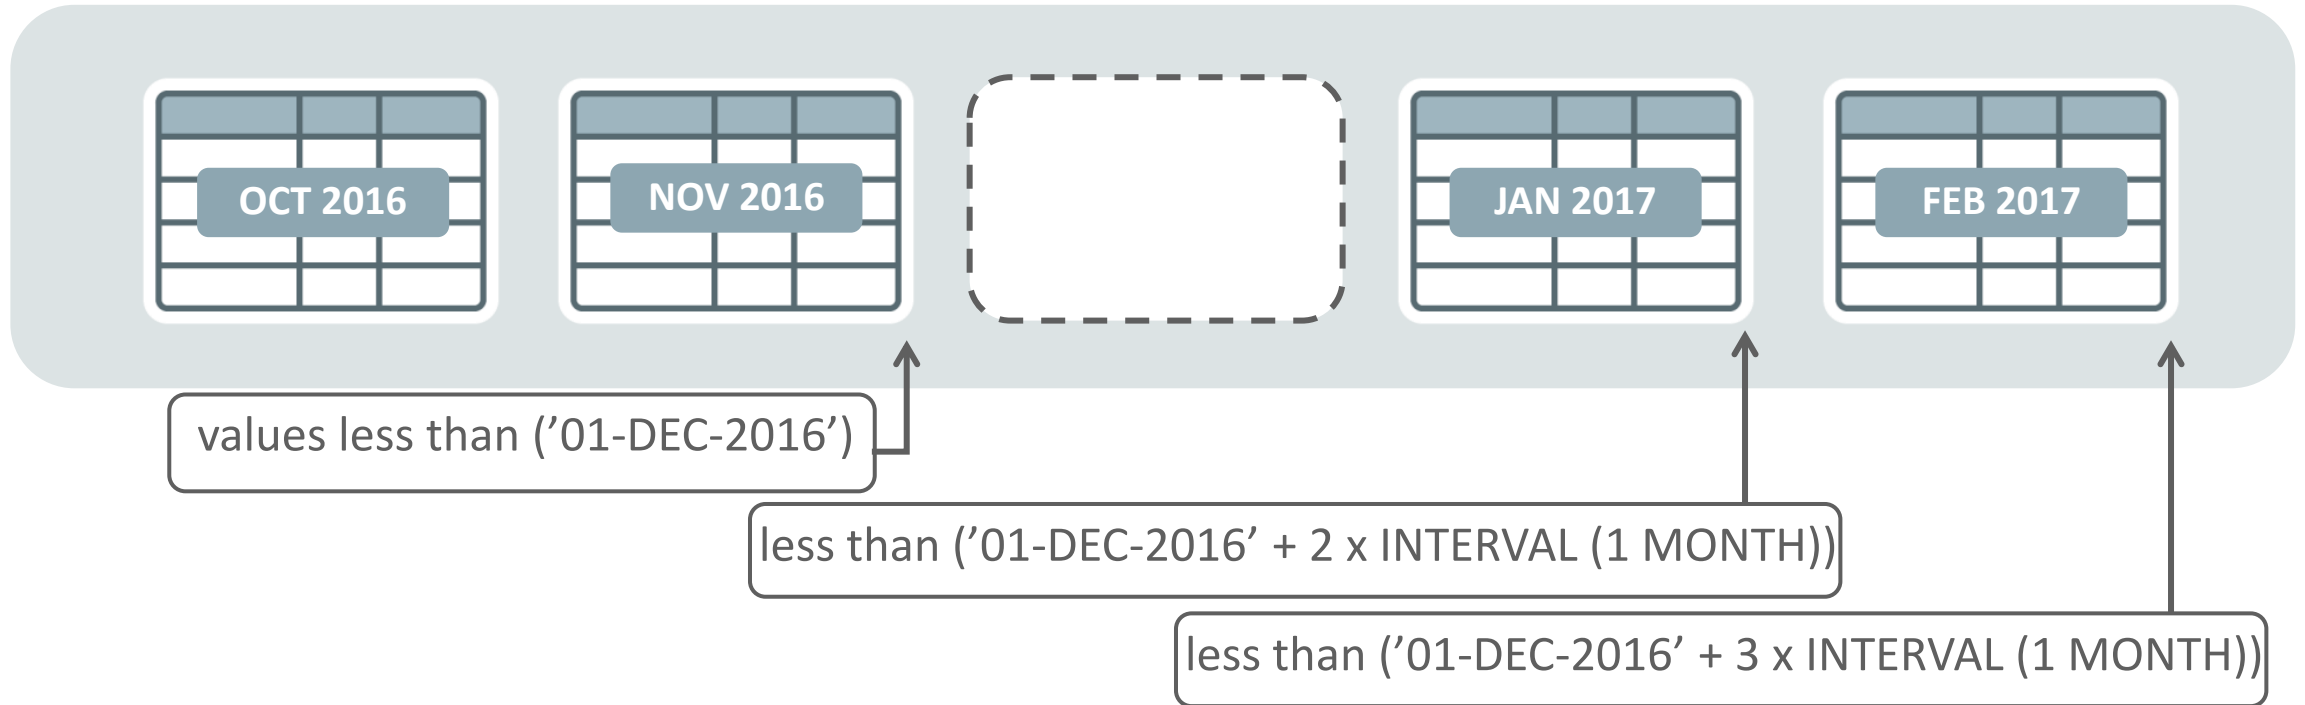
Partition Bounds for **Range Partitioning**



- Drop of previous partition moves lower boundary
 - “Feb 2017” now spawns 01-JAN-2017 to 28-FEB-2017

Interval Versus Range Partitioning

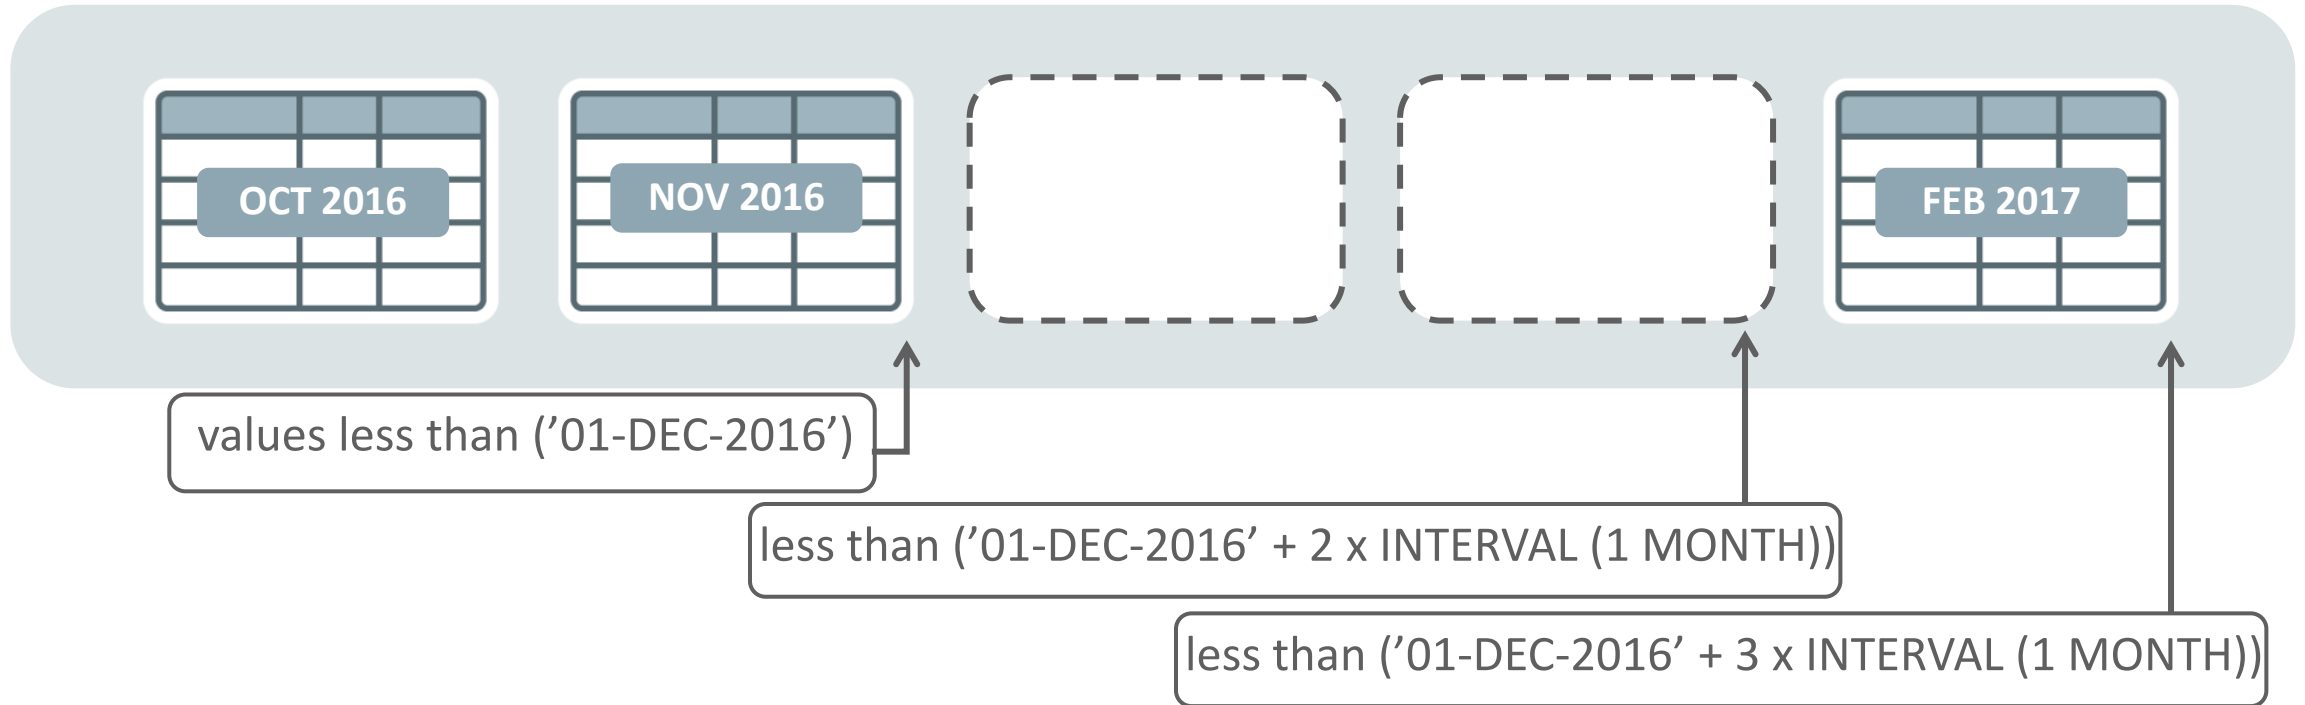
Partition Bounds for **Interval Partitioning**



- Partitions have upper and lower bounds
 - Derived by INTERVAL function and last range partition

Interval Versus Range Partitioning

Partition Bounds for **Interval Partitioning**



- Drop does not impact partition boundaries
 - “Feb 2017” still spawns 01-FEB-2017 to 28-FEB-2017

Interval versus Range Partitioning

Partition Naming

- Range partitions **can** be named
 - System generated name if not specified

```
SQL> alter table t add partition values less than(20);  
Table altered.  
SQL> alter table t add partition P30 values less than(30);  
Table altered.
```

- Interval partitions **cannot** be named
 - Always system generated name

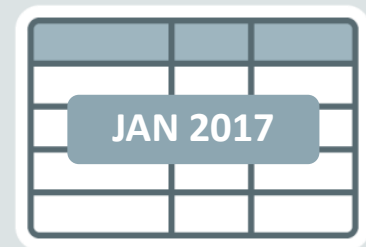
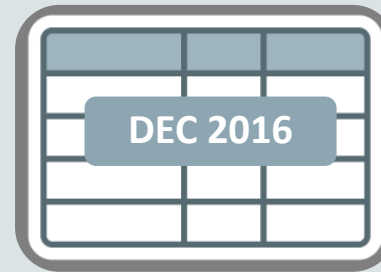
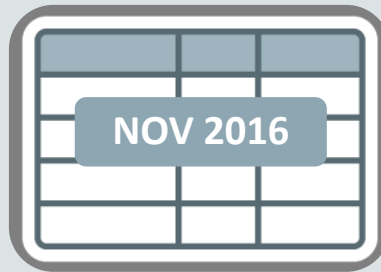
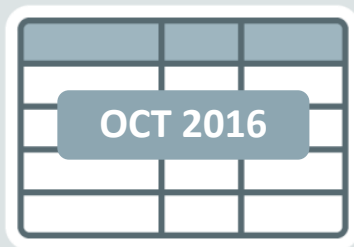
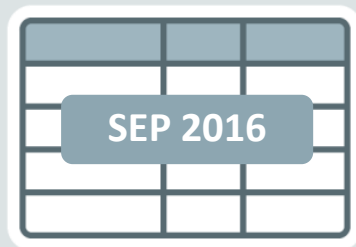
```
SQL> alter table t add partition values less than(20);  
*  
ERROR at line 1: ORA-14760: ADD PARTITION is not permitted  
on Interval partitioned objects
```

- Use new deterministic PARTITION FOR () extension

```
SQL> alter table t1 rename partition for (9) to p_10;  
Table altered.
```

Interval Versus Range Partitioning

Partition Merge – **Range Partitioning**

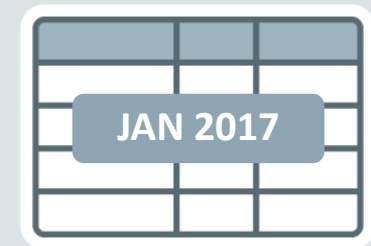
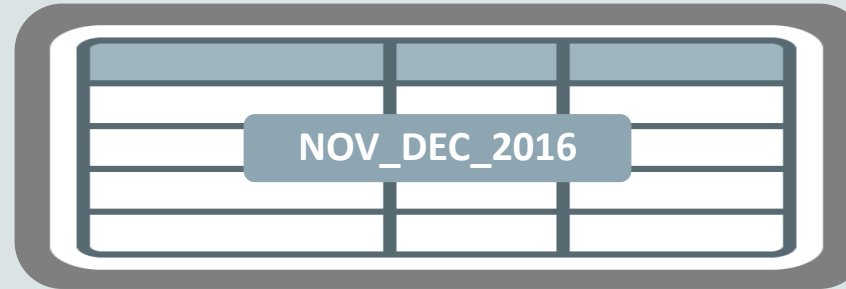
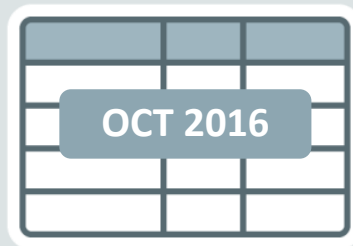
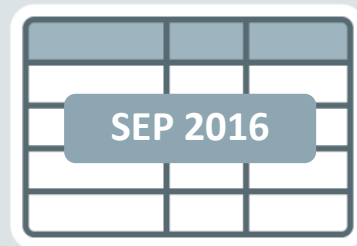


```
MERGE PARTITIONS NOV_2016, DEC_2016 INTO PARTITION NOV_DEC_2016
```

- Merge two adjacent partitions for range partitioning
 - Upper bound of higher partition is new upper bound
 - Lower bound derived through upper bound of previous partition

Interval Versus Range Partitioning

Partition Merge – **Range Partitioning**

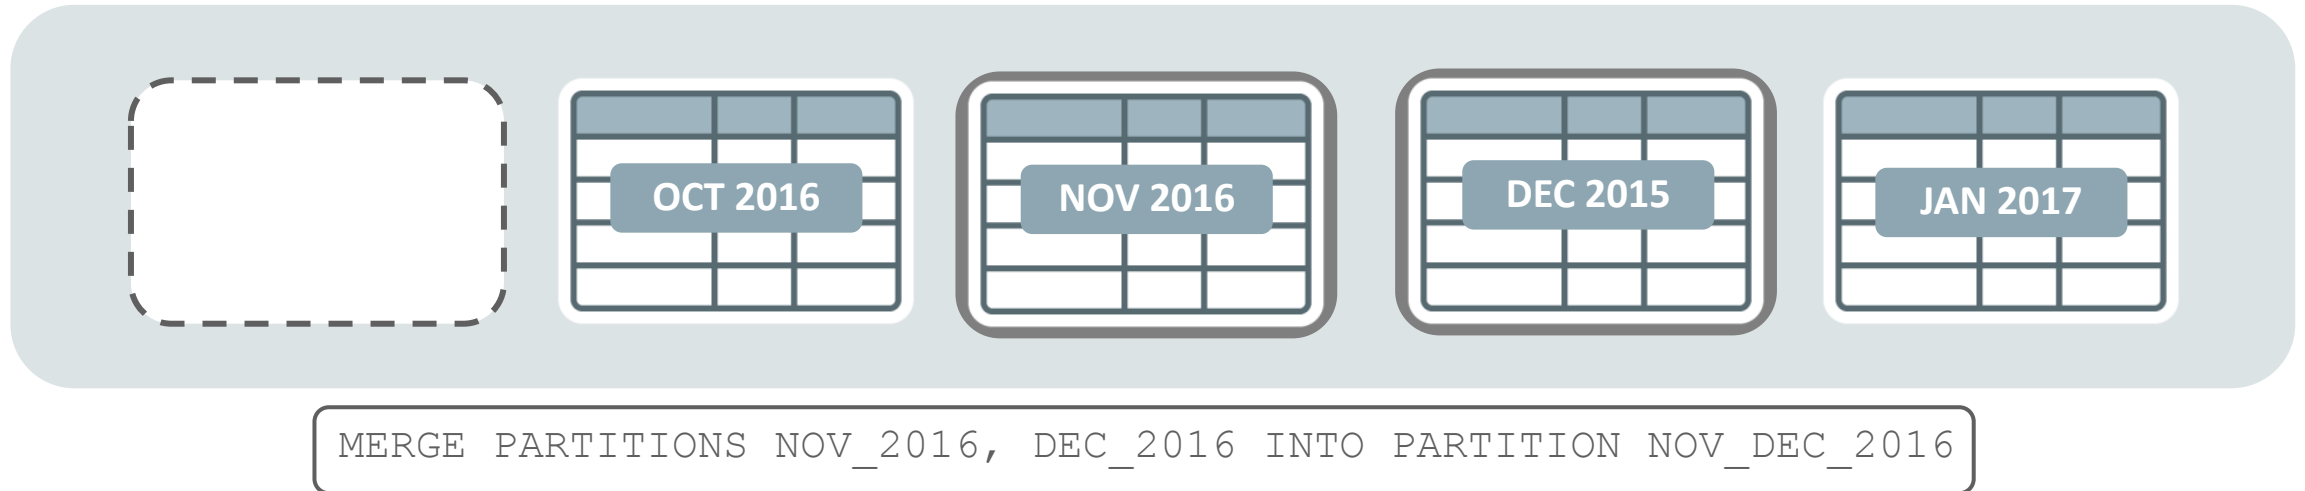


```
MERGE PARTITIONS NOV_2016, DEC_2016 INTO PARTITION NOV_DEC_2016
```

- New segment for merged partition is created
 - Rest of the table is unaffected

Interval Versus Range Partitioning

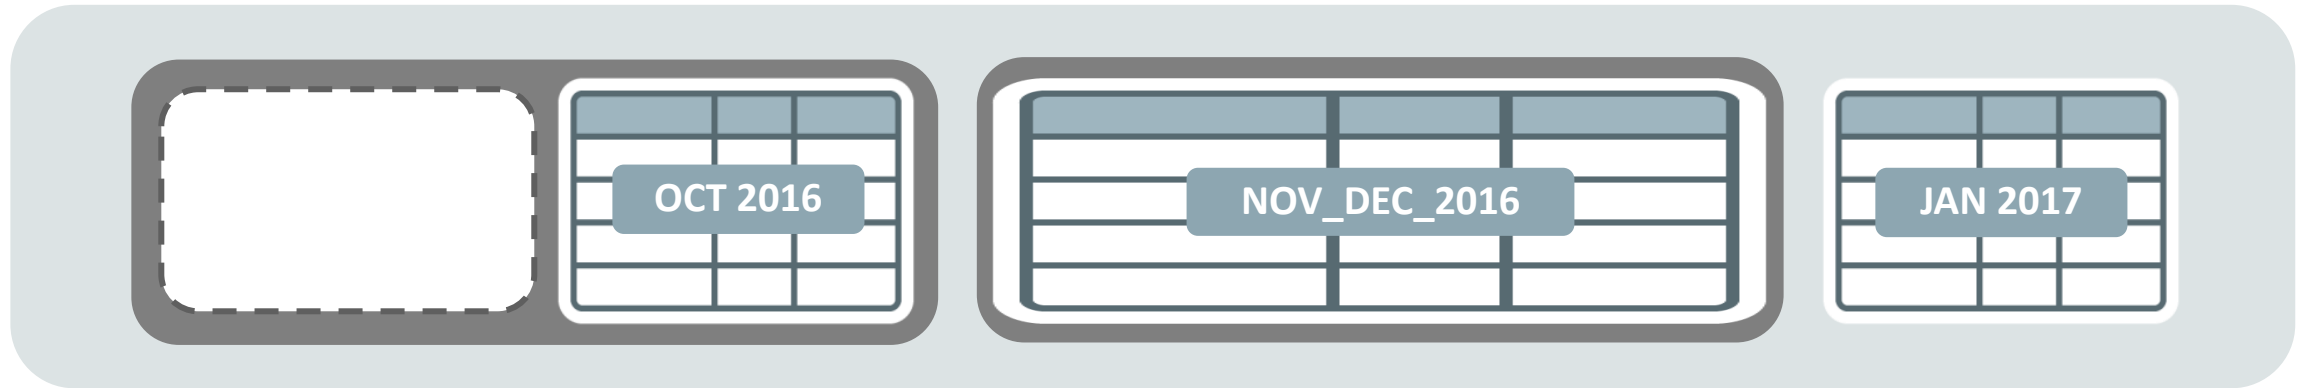
Partition Merge – **Interval Partitioning**



- Merge two adjacent partitions for interval partitioning
 - Upper bound of higher partition is new upper bound
 - Lower bound derived through lower bound of first partition

Interval Versus Range Partitioning

Partition Merge – **Interval Partitioning**



```
MERGE PARTITIONS NOV_2016, DEC_2016 INTO PARTITION NOV_DEC_2016
```

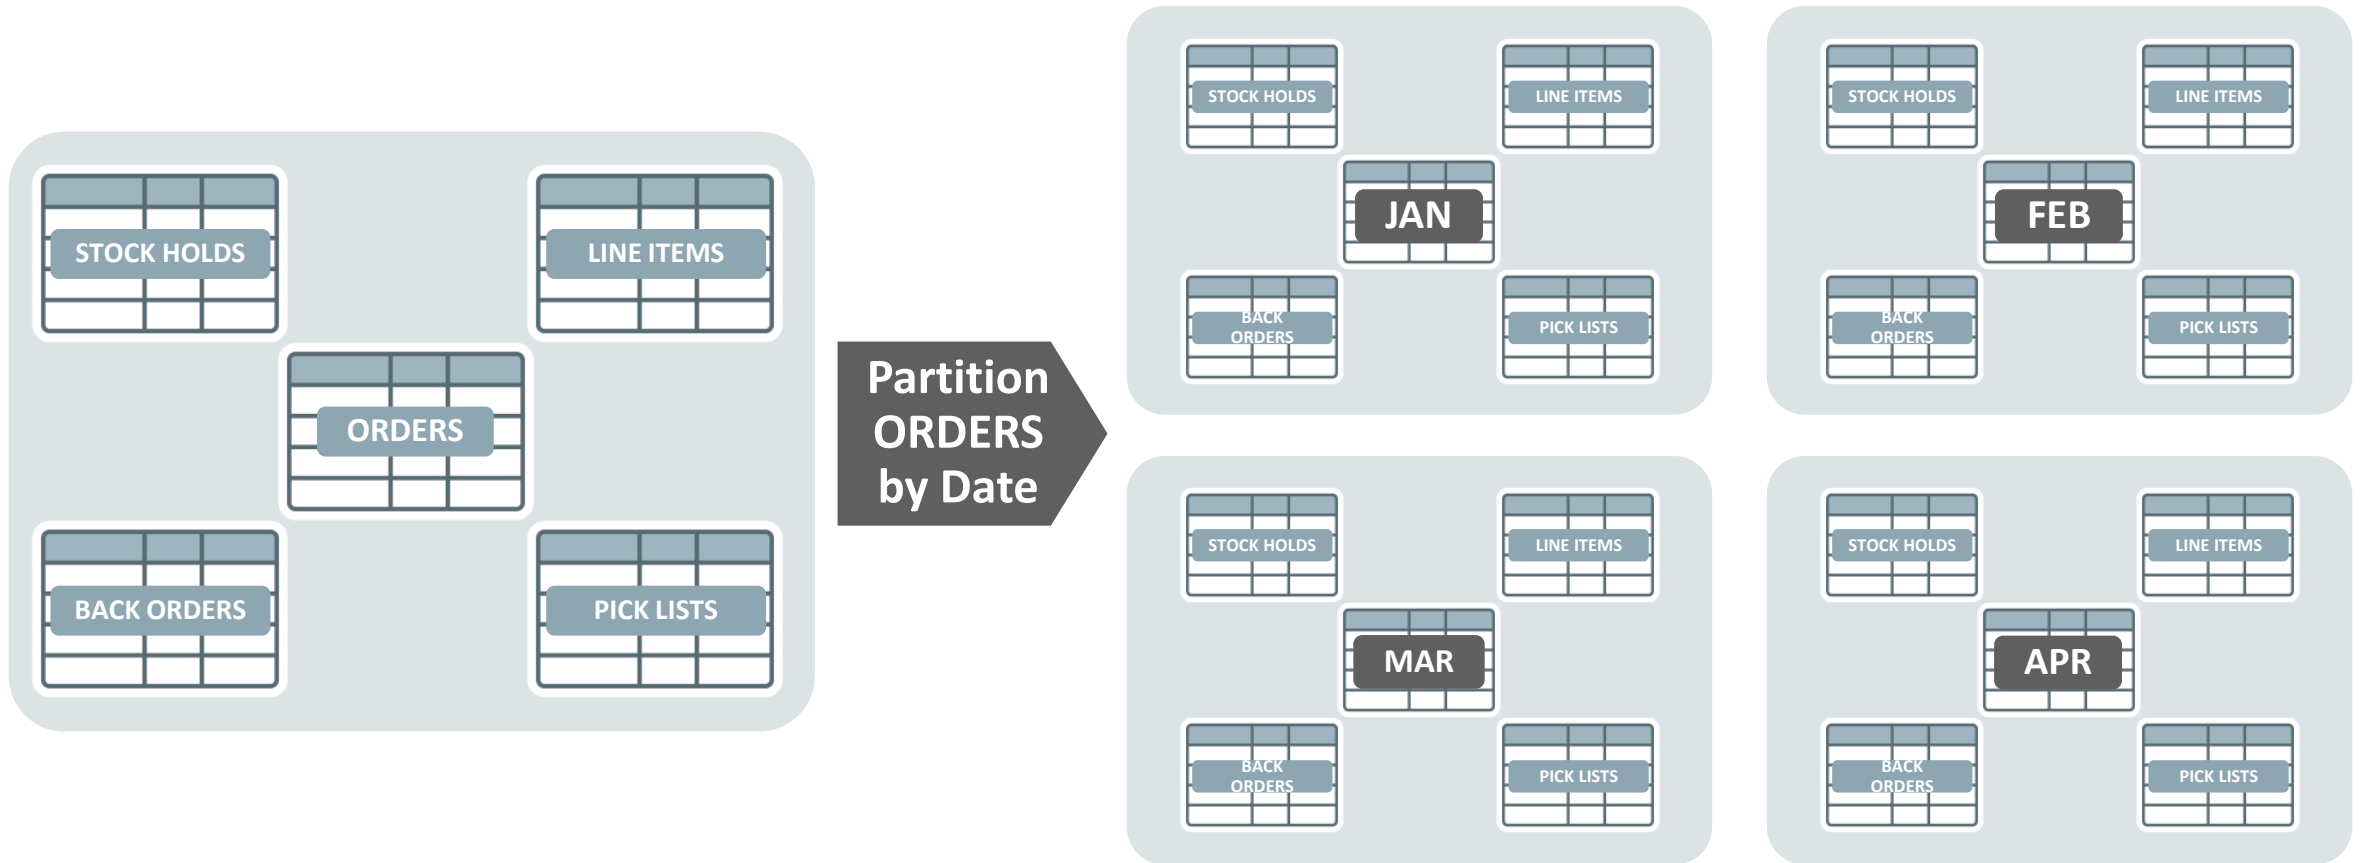
- New segment for merged partition is created
 - Holes before highest non-interval partition will be silently “merged” as well
 - Interval only valid beyond the highest non-interval partition

Reference Partitioning

Introduced in Oracle 11g Release 1 (11.1)

Reference Partitioning

Inherit partitioning strategy



Reference Partitioning

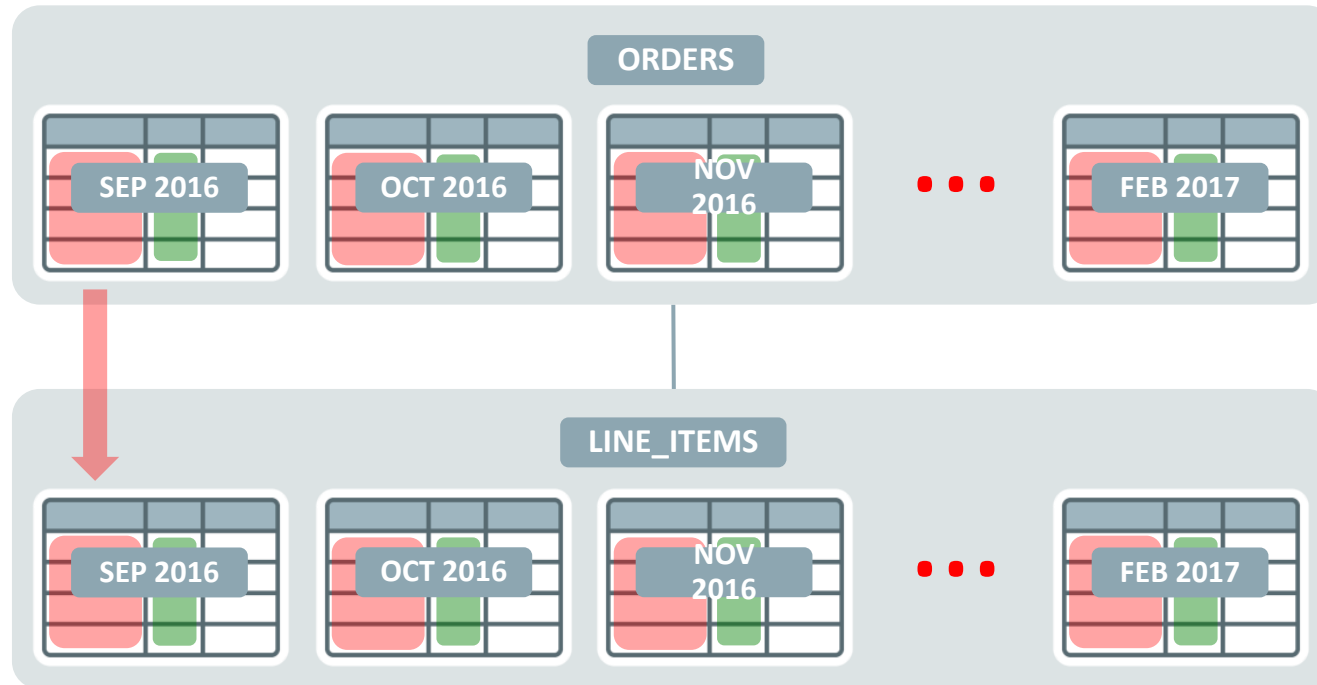
Business Problem

- Related tables benefit from same partitioning strategy
 - Sample 3NF order entry data model
- Redundant storage of same information solves problem
 - Data and maintenance overhead

Solution

- Oracle Database 11g introduces Reference Partitioning
 - Child table inherits the partitioning strategy of parent table through PK-FK
 - Intuitive modelling
- Enhanced Performance and Manageability

Without Reference Partitioning

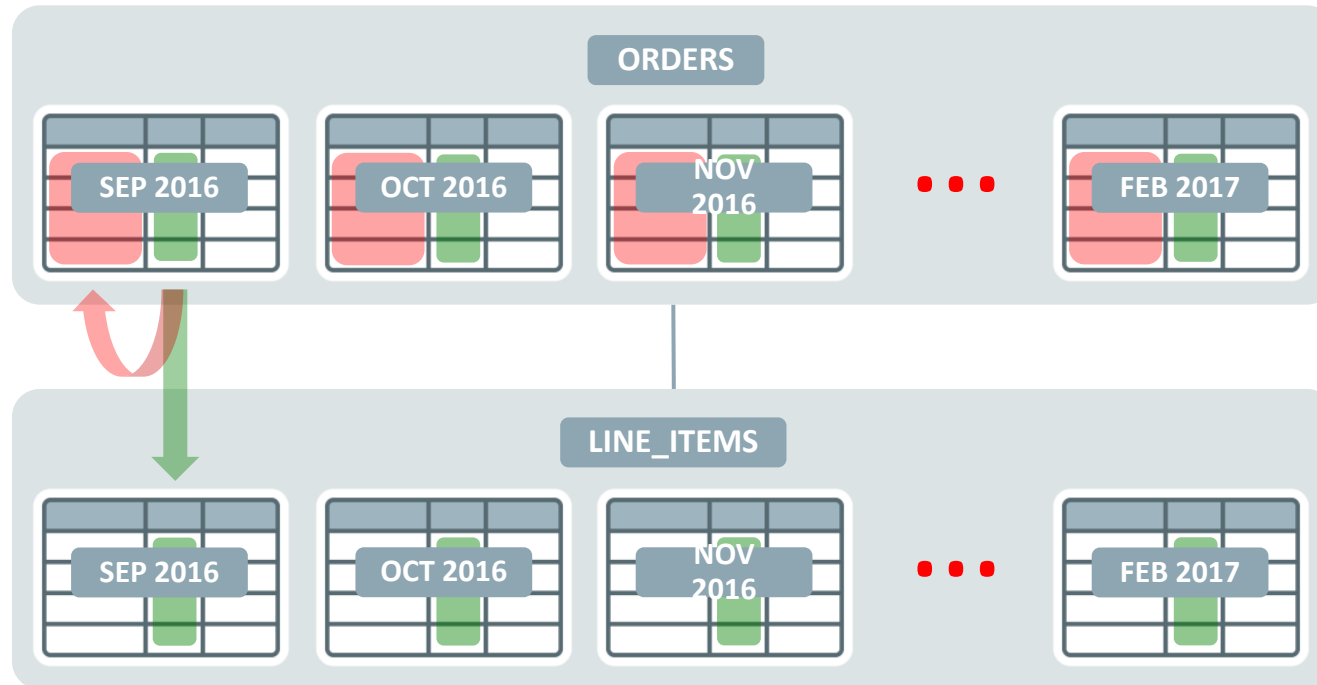


```
RANGE (order_date)
Primary key order_id
```

- Redundant storage
- Redundant maintenance

```
RANGE (order_date)
Foreign key order_id
```

With Reference Partitioning



```
RANGE (order_date)
Primary key order_id
```

- Partitioning key inherited through PK-FK relationship

```
RANGE (order_date)
Foreign key order_id
```

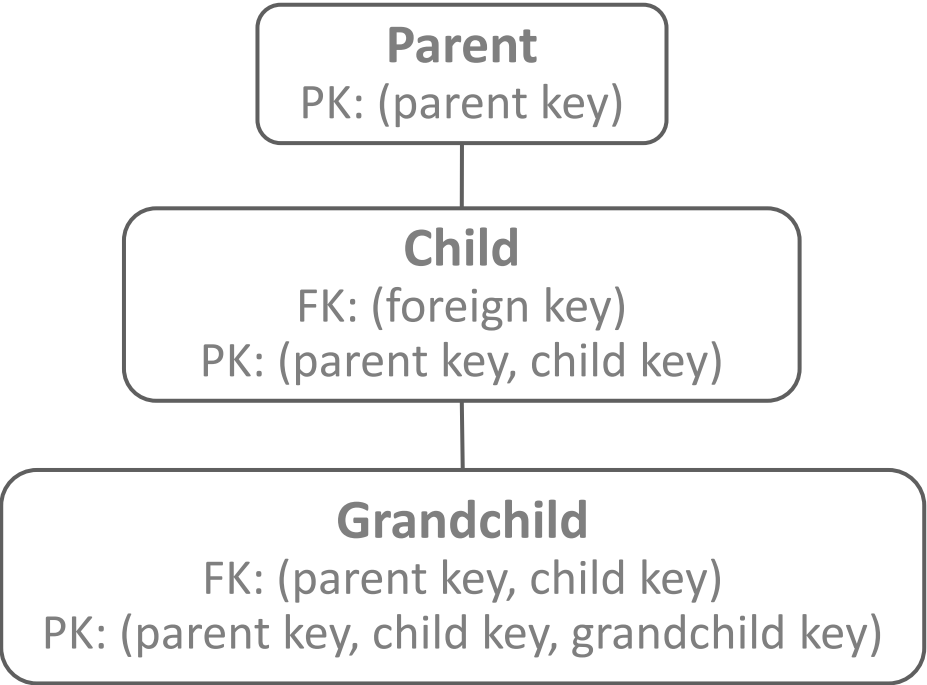
Reference Partitioning

Use Cases

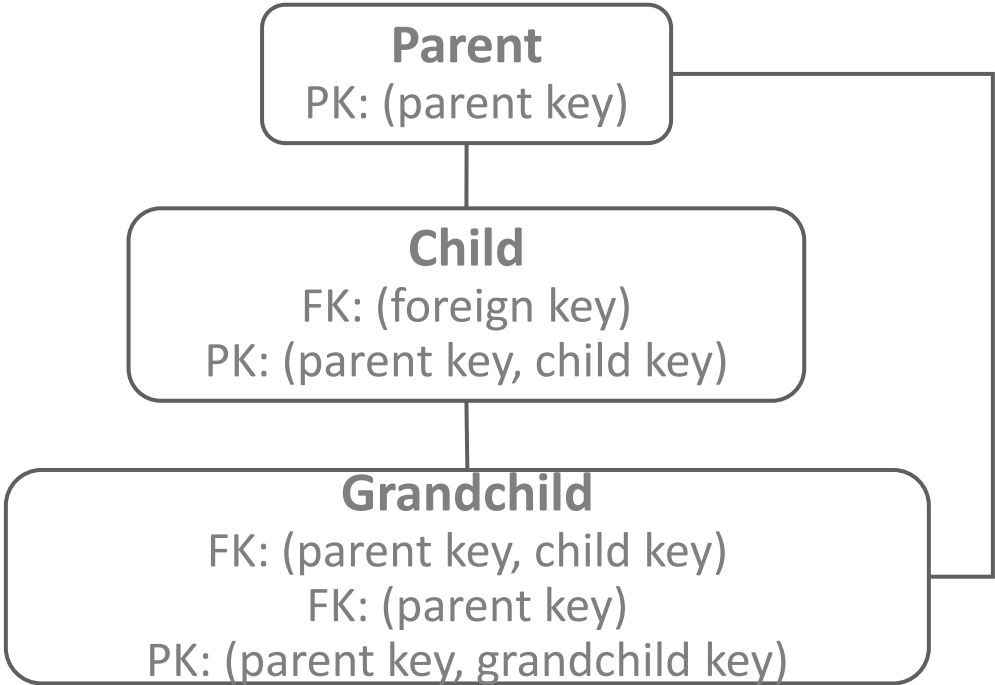
- Traditional relational model
 - Primary key inherits down to all levels of children and becomes part of an (elongated) primary key definition
- Object oriented-like model
 - Several levels of primary-foreign key relationship
 - Primary key on each level is primary key + “object ID”
- Reference Partitioning optimally suited to address both modeling techniques

Reference Partitioning

Relational Model



“Object-like” model



Reference Partitioning

Example

```
create table project (project_id number not null,  
                      project_number varchar2(30),  
                      project_name varchar2(30), ...  
                      constraint proj_pk primary key (project_id))  
partition by list (project_id)  
(partition p1 values (1),  
 partition p2 values (2),  
 partition pd values (DEFAULT));
```

```
create table project_customer (project_cust_id number not null,  
                               project_id number not null,  
                               cust_name varchar2(30),  
                               constraint pk_proj_cust primary key  
                                 (project_id, project_cust_id),  
                               constraint proj_cust_proj_fk foreign key  
                                 (project_id) references project(project_id))  
partition by reference (proj_cust_proj_fk);
```

Reference Partitioning

Example, cont.

```
create table proj_cust_address (project_cust_addr_id number not null,  
                                project_cust_id number not null,  
                                project_id number not null,  
                                cust_address varchar2(30),  
                                constraint pk_proj_cust_addr primary key  
                                    (project_id, project_cust_addr_id),  
                                constraint proj_c_addr_proj_cust_fk foreign key  
                                    (project_id, project_cust_id)  
                                    references project_customer  
                                        (project_id, project_cust_id))  
partition by reference (proj_c_addr_proj_cust_fk);
```

Reference Partitioning

Some metadata

Table information

```
SQL> SELECT table_name, partitioning_type, ref_ptn_constraint_name
       FROM   user_part_tables
       WHERE  table_name IN ('PROJECT','PROJECT_CUSTOMER','PROJ_CUST_ADDRESS');
```

TABLE_NAME	PARTITION	REF_PTN_CONSTRAINT_NAME
PROJECT	LIST	
PROJECT_CUSTOMER	REFERENCE	PROJ_CUST_PROJ_FK
PROJ_CUST_ADDRESS	REFERENCE	PROJ_C_ADDR_PROJ_FK

Partition information

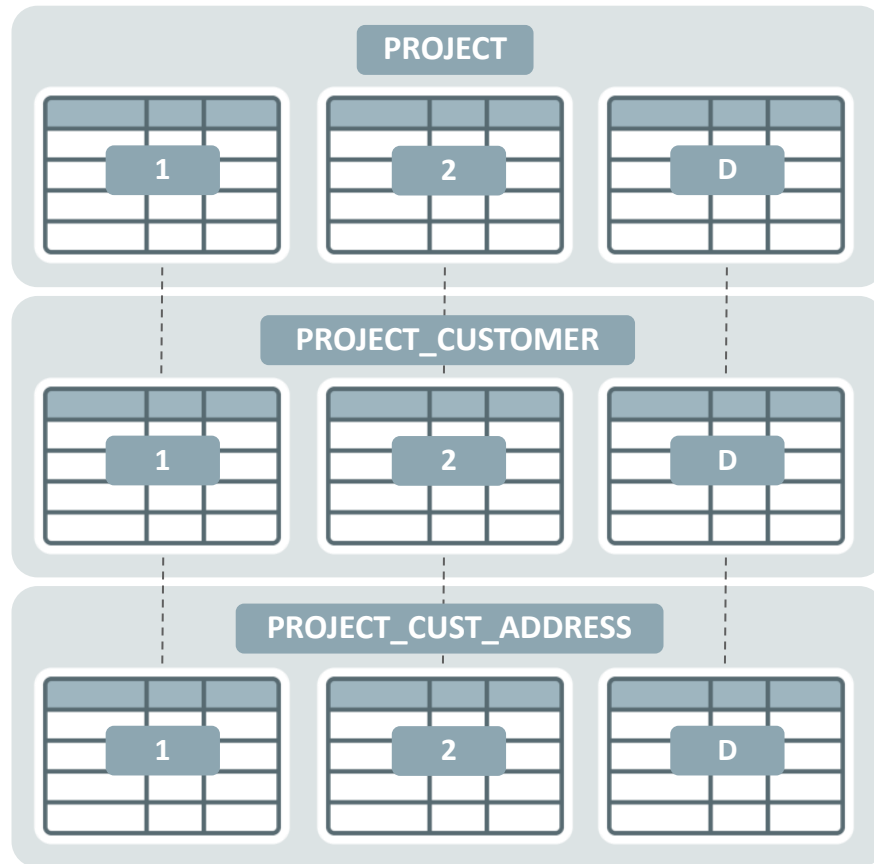
```
SQL> SELECT table_name, partition_name, high_value
       FROM   user_tab_partitions
       WHERE  table_name in ('PROJECT','PROJECT_CUSTOMER')
       ORDER BY table_name, partition_position;
```

TABLE_NAME	PARTITION_NAME	HIGH_VALUE
PROJECT	P1	1
PROJECT	P2	2
PROJECT	PD	DEFAULT
PROJECT_CUSTOMER	P1	
PROJECT_CUSTOMER	P2	
PROJECT_CUSTOMER	PD	



Reference Partitioning

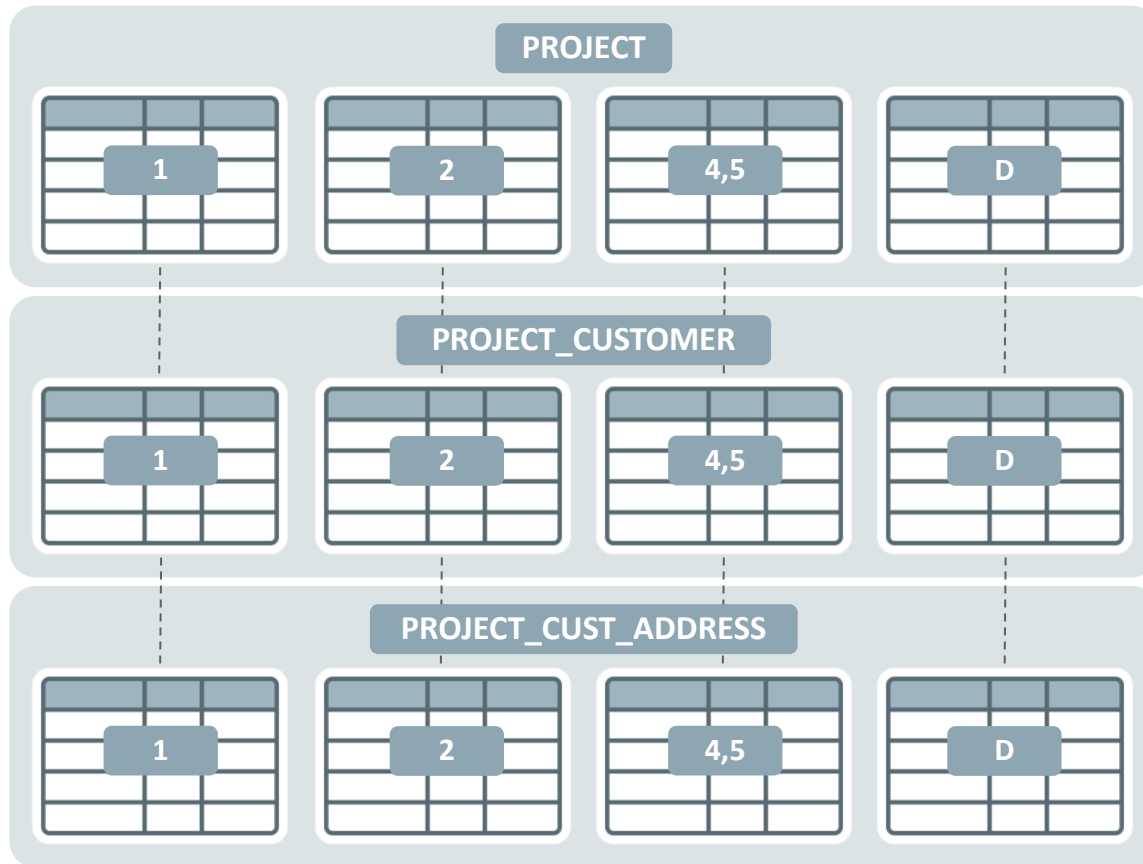
Partition Maintenance



```
ALTER TABLE project  
SPLIT PARTITION pd VALUES (4,5) INTO  
(PARTITION pd, PARTITION p45);
```

Reference Partitioning

Partition Maintenance

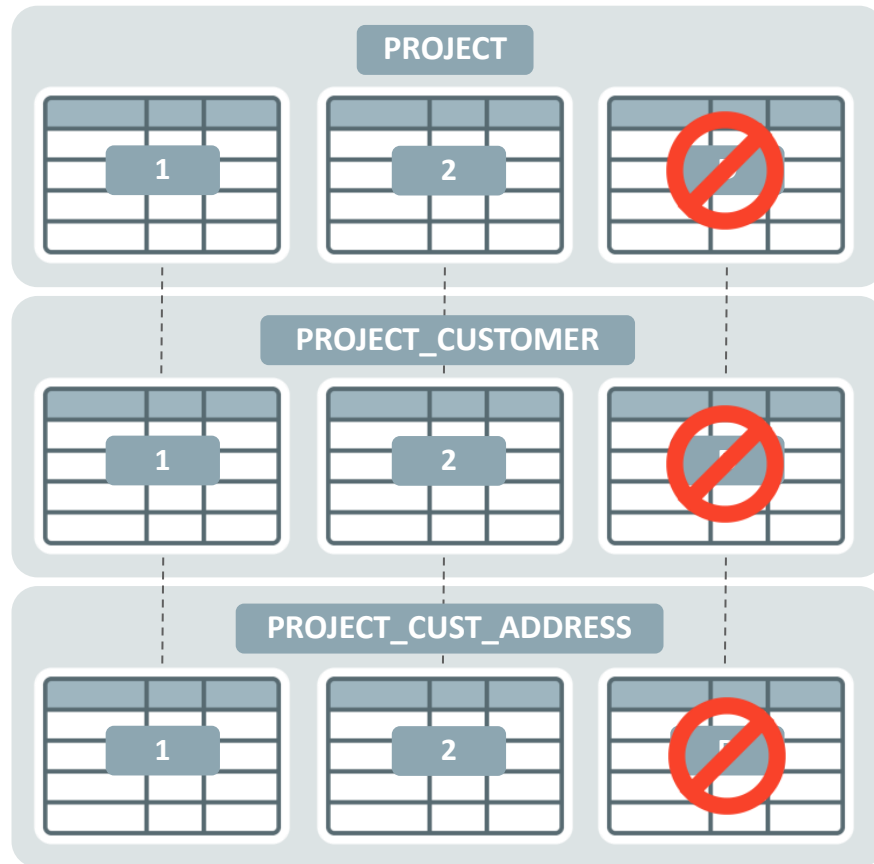


```
ALTER TABLE project  
SPLIT PARTITION pd VALUES (4,5) INTO  
(PARTITION pd, PARTITION p45);
```

- PROJECT partition PD will be split
 - “Default” and (4,5)
- PROJECT_CUSTOMER will split its dependent partition
 - Co-location with equivalent parent record of PROJECT
 - Parent record in (4,5) means child record in (4.5)
- PROJECT_CUST_ADDRESS will split its dependent partition
 - Co-location with equivalent parent record of PROJECT_CUSTOMER
- One-level lookup required for both placements

Reference Partitioning

Partition Maintenance



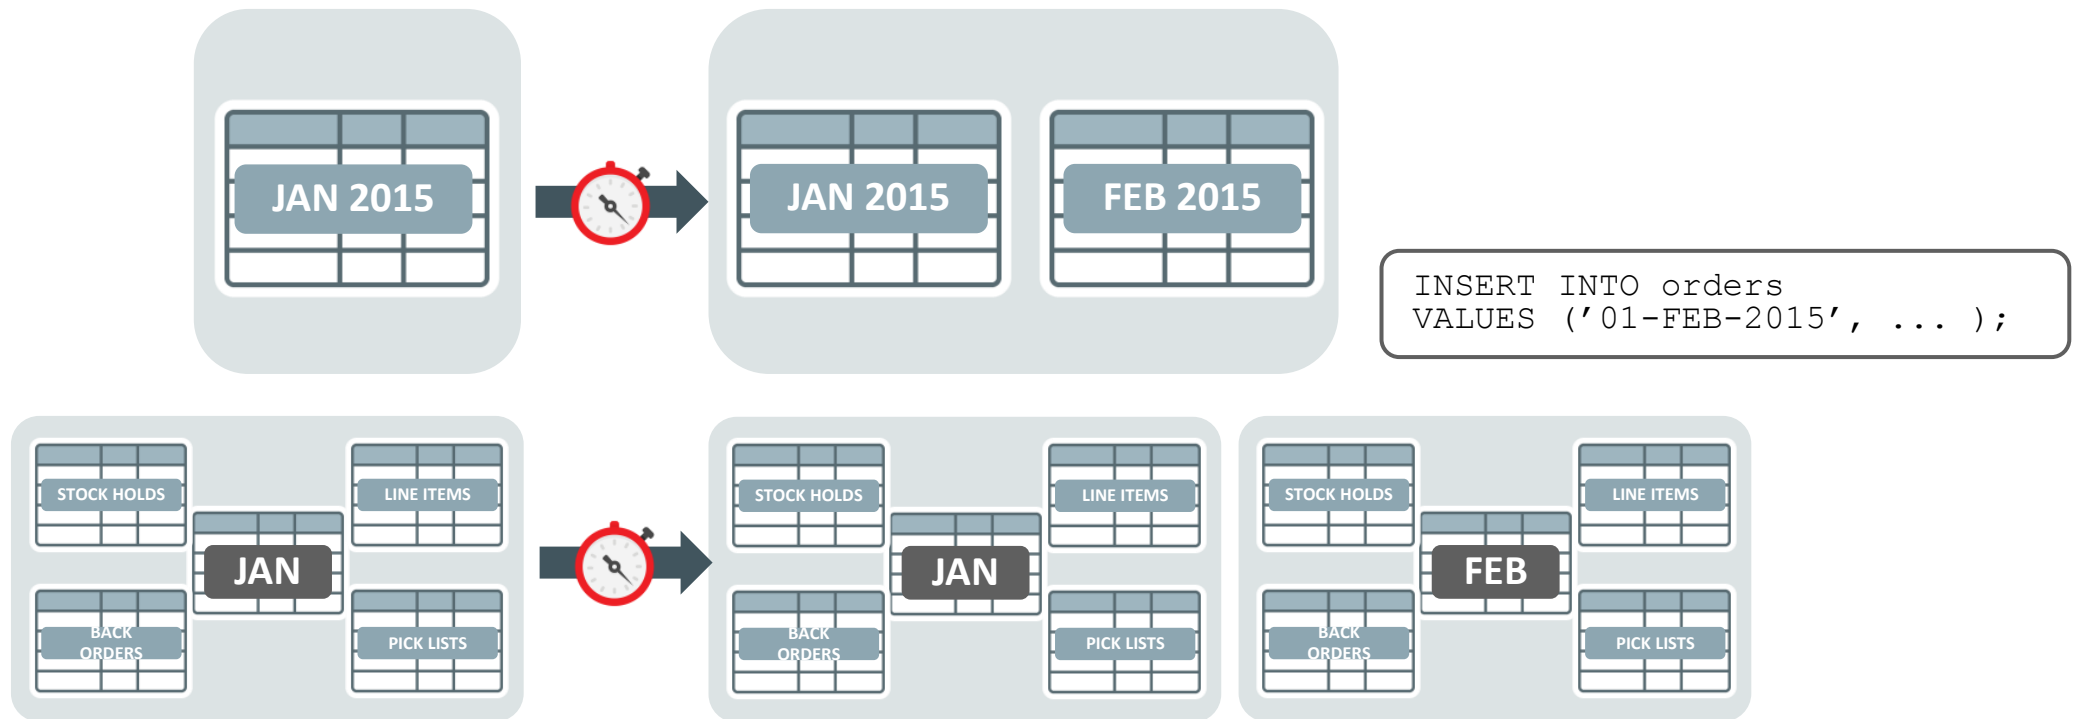
```
ALTER TABLE project_cust_address  
DROP PARTITION pd;
```

- PROJECT partition PD will be dropped
 - PK-FK is guaranteed not to be violated
- PROJECT_CUSTOMER will drop its dependent partition
- PROJECT_CUST_ADDRESS will drop its dependent partition
- Unlike “normal” partitioned tables, PK-FK relationship stays enabled
 - You cannot arbitrarily drop or truncate a partition with the PK of a PK-FK relationship
- Same is true for TRUNCATE
 - Bottom-up operation

Interval Reference Partitioning

Introduced in Oracle 12c Release 1 (12.1)

Interval-Reference Partitioning



- New partitions are automatically created when new data arrives
- All child tables will be automatically maintained
- Combination of two successful partitioning strategies for better business modeling

Interval-Reference Partitioning

```
SQL> REM create some interval-referenced tables ..
```

```
SQL> create table intRef_p (pkcol number not null, col2 varchar2(200),  
2                          constraint pk_intref primary key (pkcol))  
3  partition by range (pkcol) interval (10)  
4  (partition p1 values less than (10));
```

Table created.

```
SQL>
```

```
SQL> create table intRef_c1 (pkcol number not null, col2 varchar2(200), fkcol number not null,  
2                          constraint pk_c1 primary key (pkcol),  
3                          constraint fk_c1 foreign key (fkcol) references intRef_p(pkcol) ON DELETE CASCADE)  
4  partition by reference (fk_c1);
```

Table created.

```
SQL>
```

```
SQL> create table intRef_c2 (pkcol number primary key not null, col2 varchar2(200), fkcol number not null,  
2                          constraint fk_c2 foreign key (fkcol) references intRef_p(pkcol) ON DELETE CASCADE)  
3  partition by reference (fk_c2);
```

Table created.

Interval-Reference Partitioning

- New partitions only created when data arrives
 - No automatic partition instantiation for complete reference tree
 - Optimized for sparsely populated reference partitioned tables
- Partition names inherited from already existent partitions
 - Name inheritance from direct relative
 - Parent partition p100 will result in child partition p100
 - Parent partition p100 and child partition c100 will result in grandchild partition c100

Virtual Column Based Partitioning

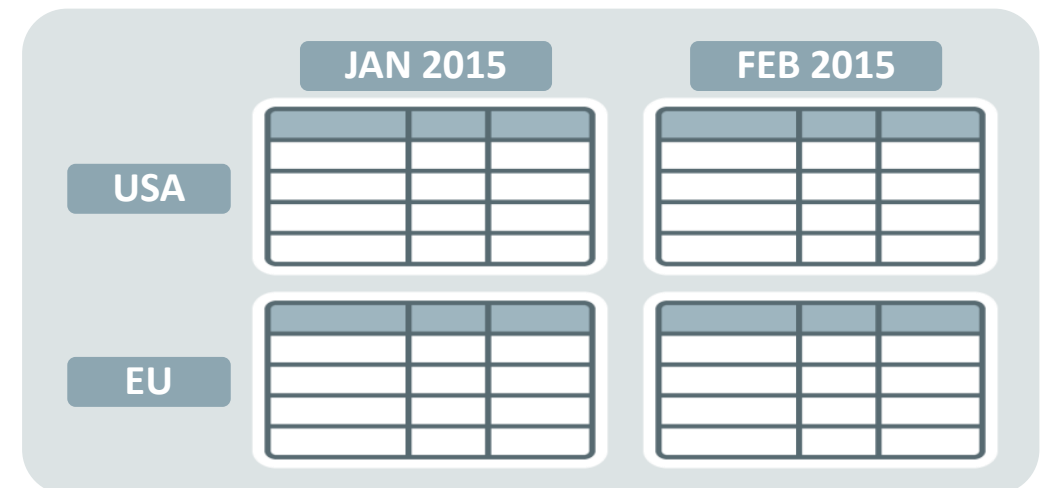
Introduced in Oracle 11g Release 1 (11.1)

Virtual Column Based Partitioning

ORDERS

ORDER_ID	ORDER_DATE	CUSTOMER_ID...	REGION AS (SUBSTR(ORDER_ID,6,2))
9834-US-14	12-JAN-2015	65920	US
8300-EU-97	14-FEB-2015	39654	EU
3886-EU-02	16-JAN-2015	4529	EU
2566-US-94	19-JAN-2015	15327	US
3699-US-63	02-FEB-2015	18733	US

- REGION requires no storage
- Partition by ORDER_DATE, REGION



Virtual Columns

Example

- Base table with all attributes ...

```
CREATE TABLE accounts  
(acc_no      number(10)    not null,  
 acc_name    varchar2(50) not null, ...
```

12500	Adams	
12507	Blake	
12666	King	
12875	Smith	

Virtual Columns

Example

- Base table with all attributes ...
... is extended with the virtual (derived) column

```
CREATE TABLE accounts  
(acc_no      number(10)    not null,  
 acc_name    varchar2(50) not null, ...  
 acc_branch  number(2)     generated always as  
             (to_number(substr(to_char(acc_no),1,2)))
```

12500	Adams	12
12507	Blake	12
12666	King	12
12875	Smith	12

Virtual Columns

Example

- Base table with all attributes ...

... is extended with the virtual (derived) column
... and the virtual column is used as partitioning key

```
CREATE TABLE accounts  
(acc_no      number(10)    not null,  
 acc_name    varchar2(50) not null, ...  
 acc_branch  number(2)     generated always as  
             (to_number(substr(to_char(acc_no),1,2)))  
 partition by list (acc_branch) ...
```

12500	Adams	12
12507	Blake	12
12666	King	12
12875	Smith	12

...

32320	Jones	32
32407	Clark	32
32758	Hurd	32
32980	Kelly	32

Virtual Columns

Partition Pruning

- Conceptual model considers virtual columns as **visible** and **used** attributes
- Partition pruning currently only works with predicates on the virtual column (partition key) itself
 - No transitive predicates
 - Enhancement planned for the future

Composite Partitioning

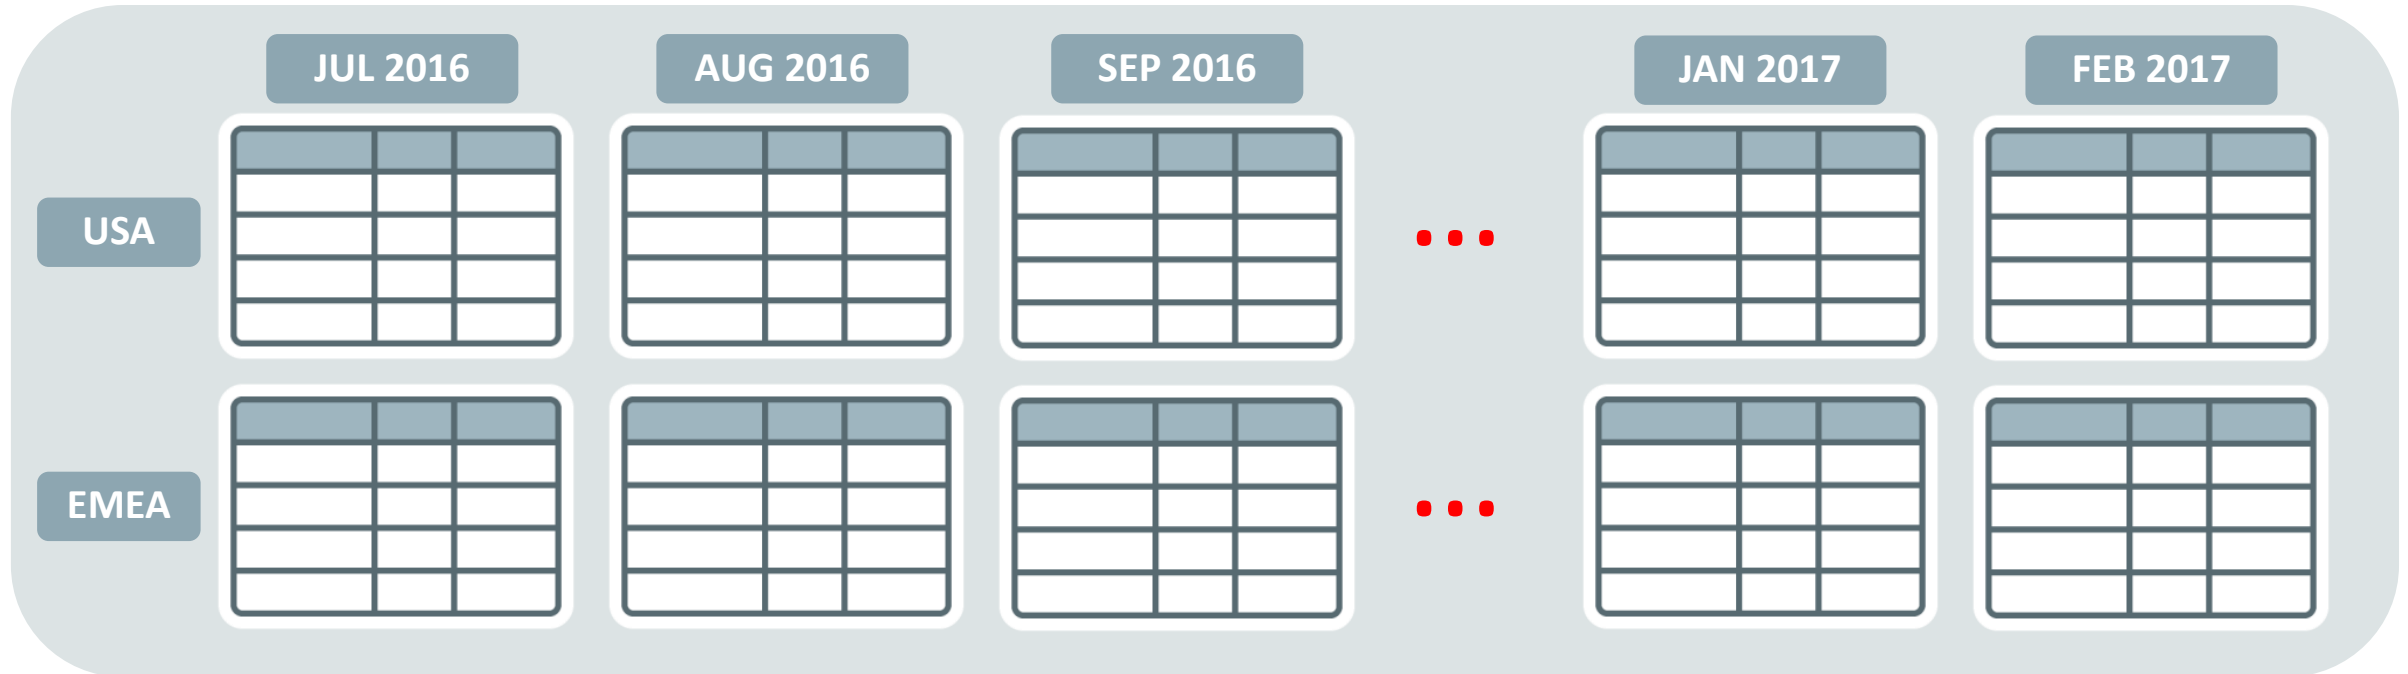
Range-Hash introduced in Oracle 8i

Range-List introduced in Oracle 9i Release 2

[Interval | Range | List | Hash]-[Range | List | Hash] introduced in Oracle 11g
Release 1 | 2

*Hash-Hash in 11.2

Composite Partitioning



- Data is organized along two dimensions
 - Record placement is deterministically identified by dimensions
 - Example RANGE-LIST

Composite Partitioning

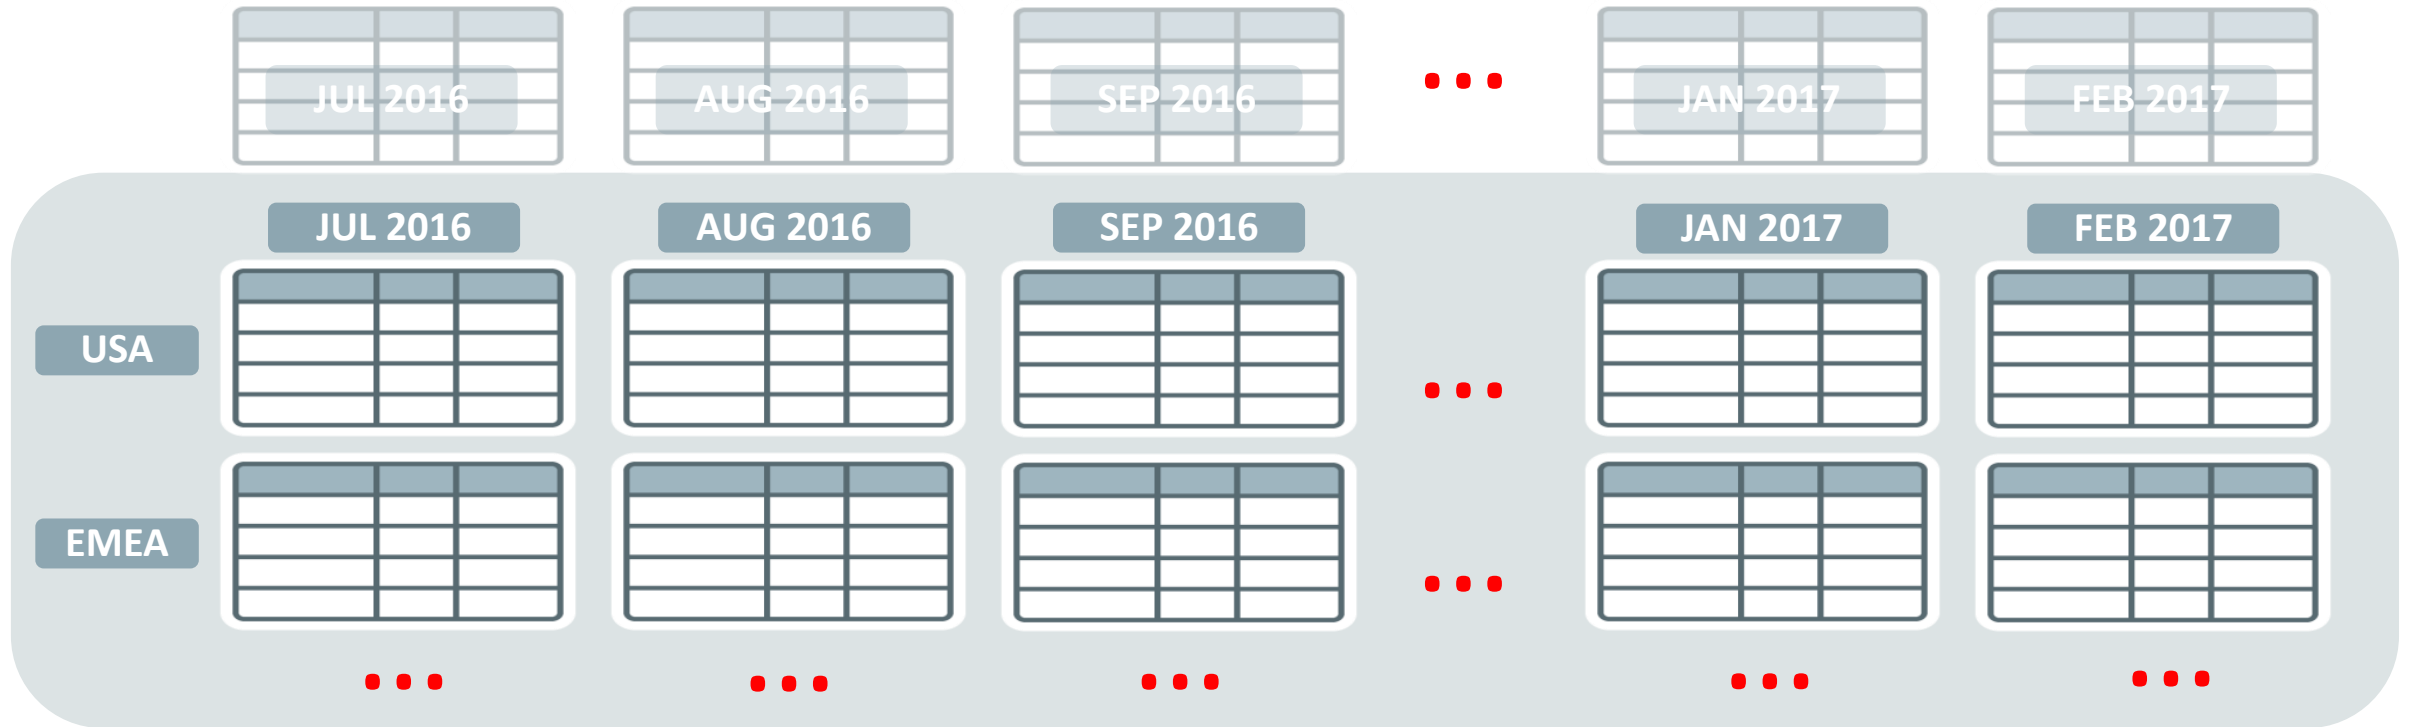
Concept



```
CREATE TABLE SALES ..PARTITION BY RANGE (time_id)
```


Composite Partitioning

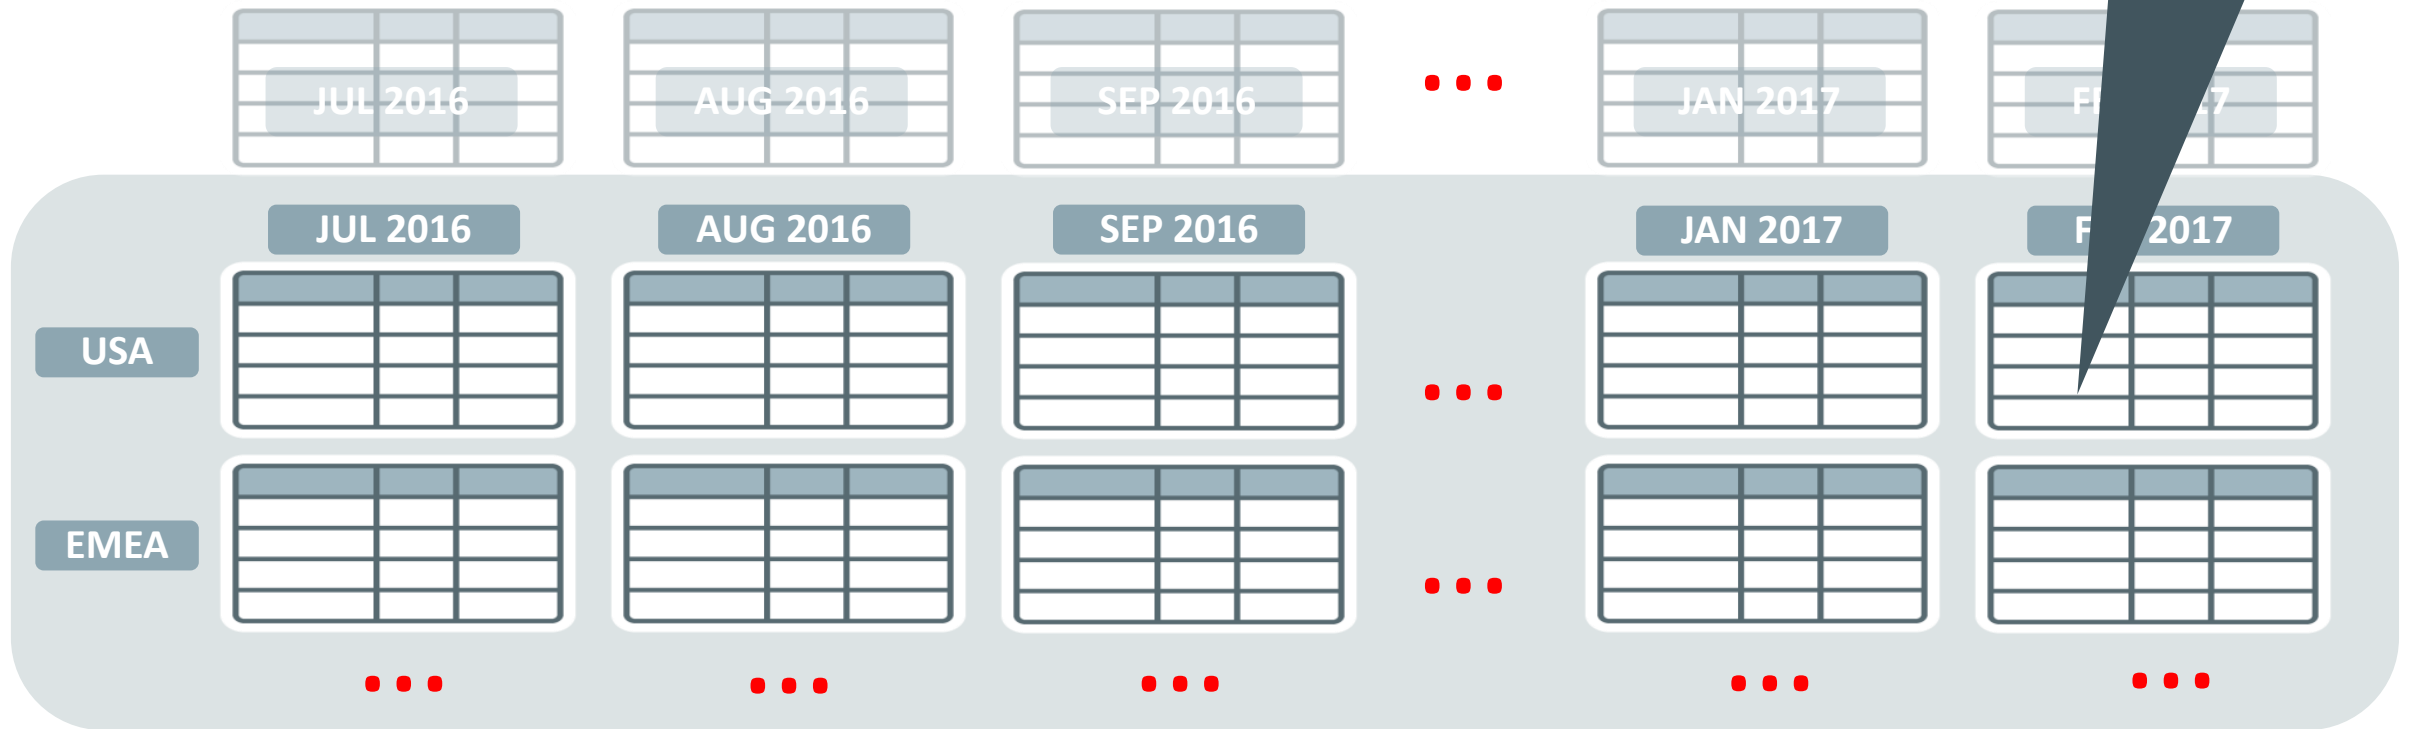
Concept



```
CREATE TABLE SALES ..PARTITION BY RANGE (time_id)
                        SUPARTITION BY LIST (region)
```

Composite Partitioning

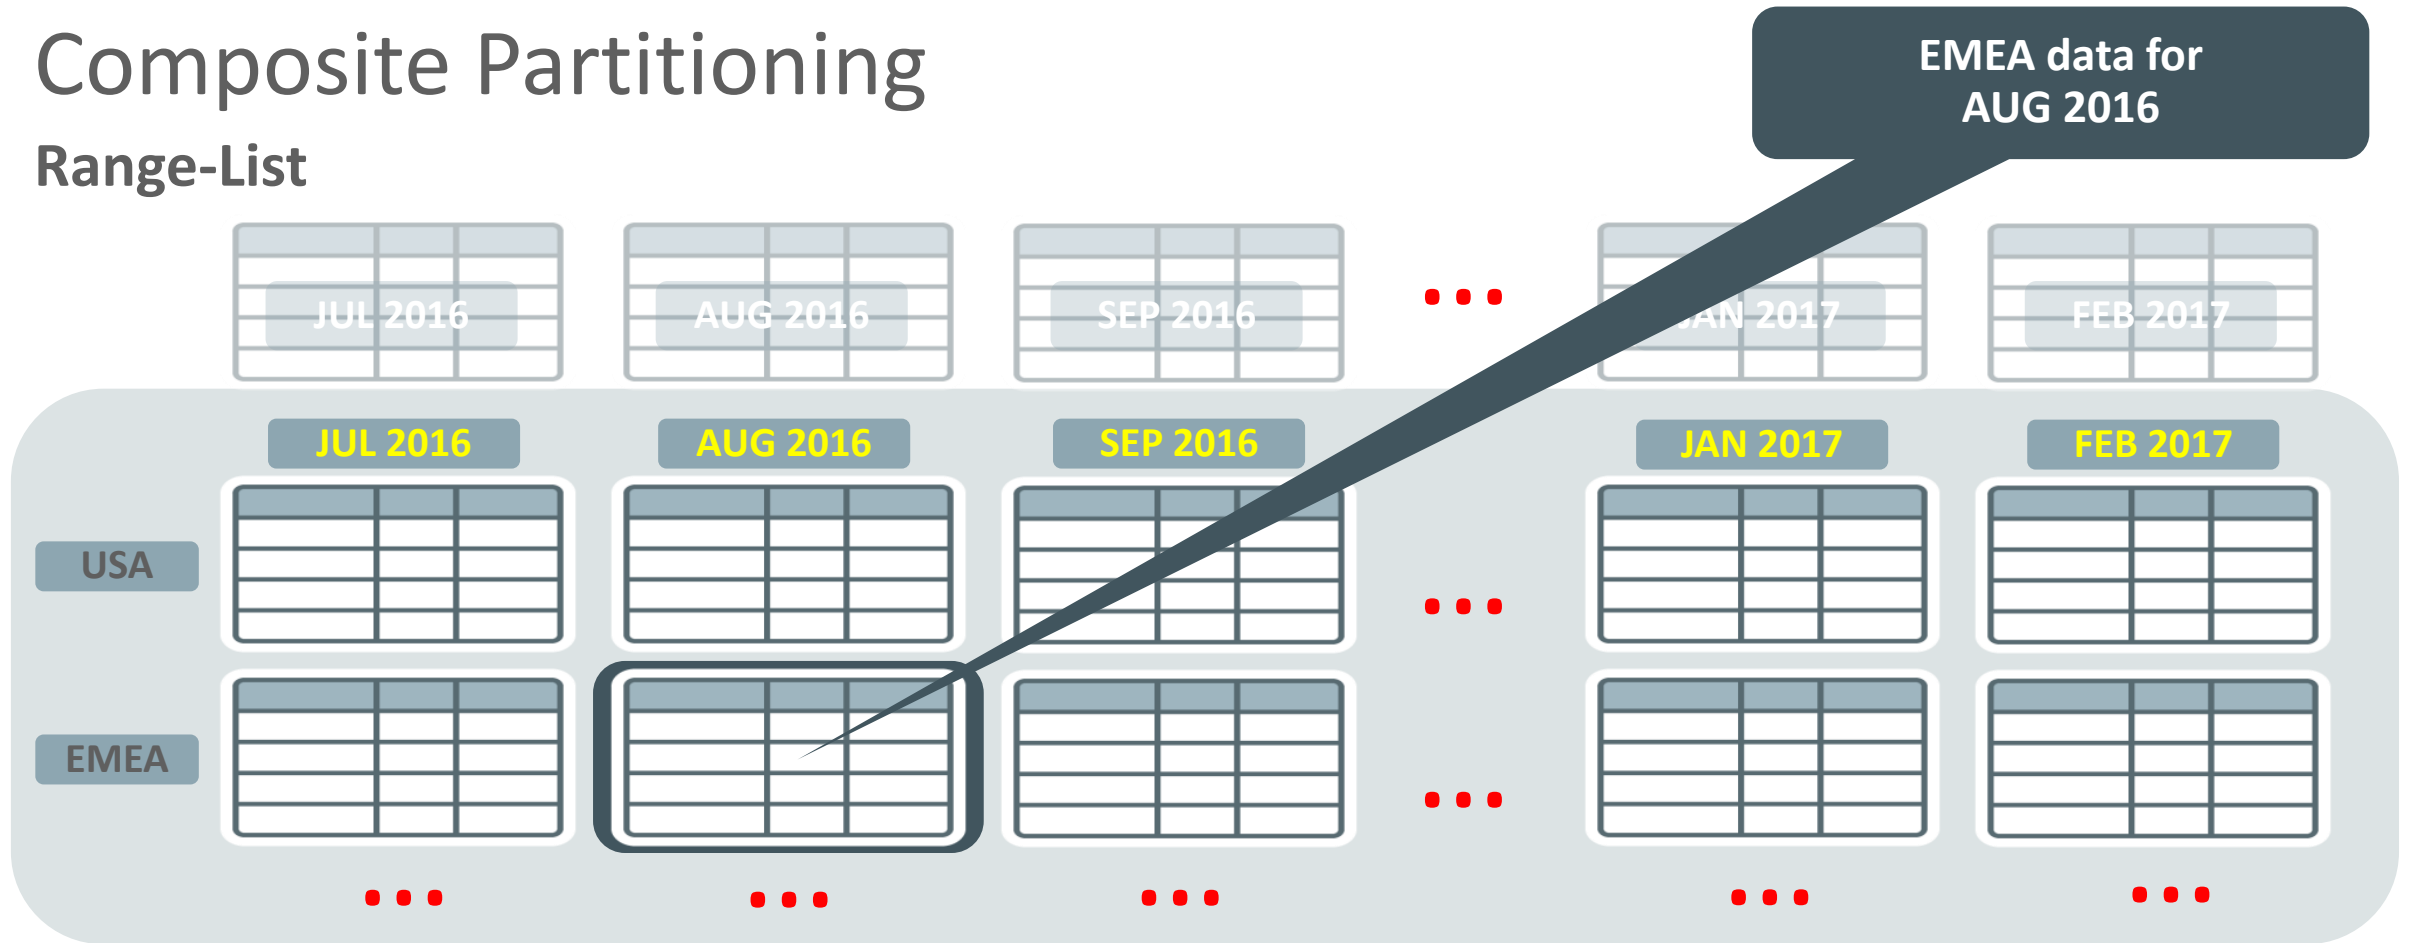
Concept



```
CREATE TABLE SALES ..PARTITION BY RANGE (time_id)
                        SUPARTITION BY LIST (region)
```

Composite Partitioning

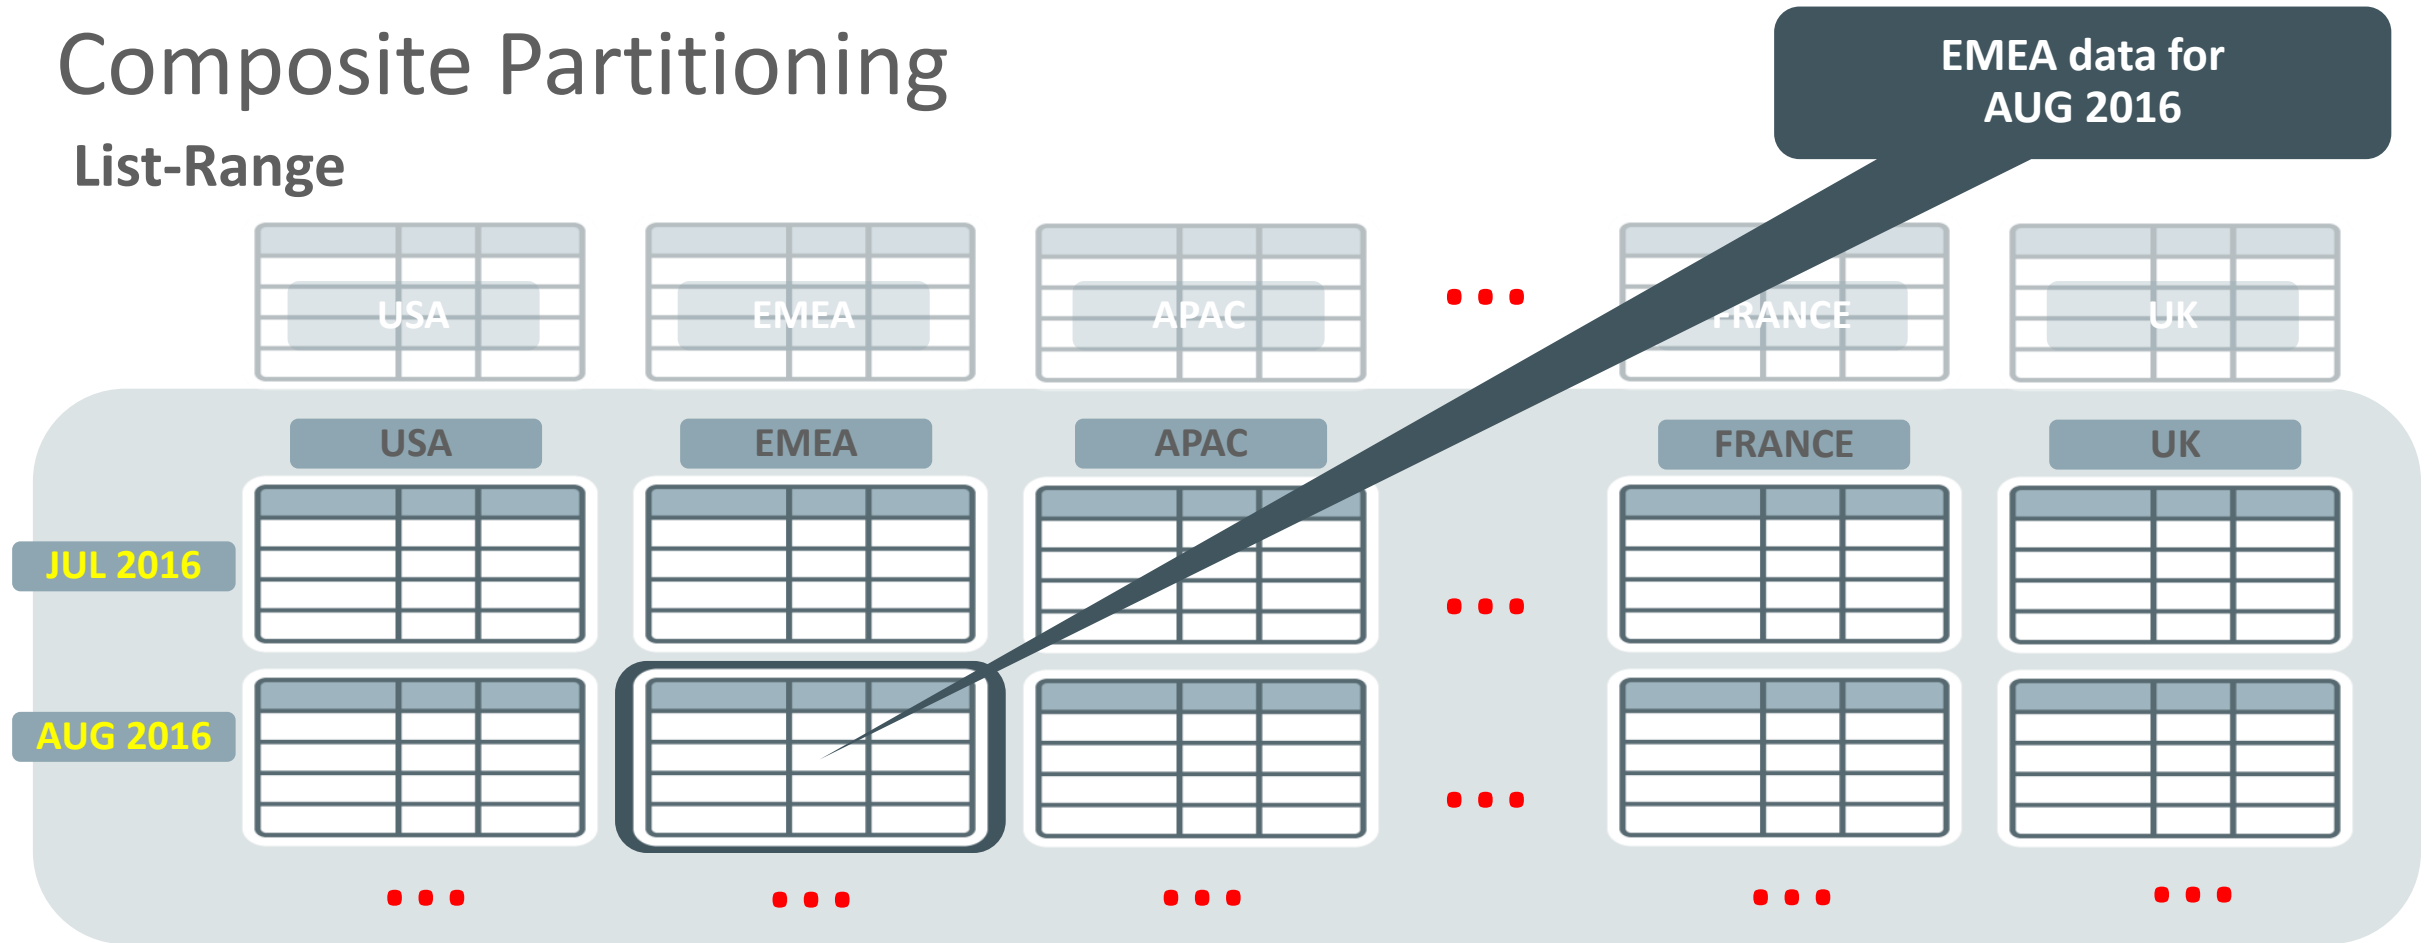
Range-List



```
CREATE TABLE SALES ..PARTITION BY RANGE (time_id)
SUPARTITION BY LIST (region)
```

Composite Partitioning

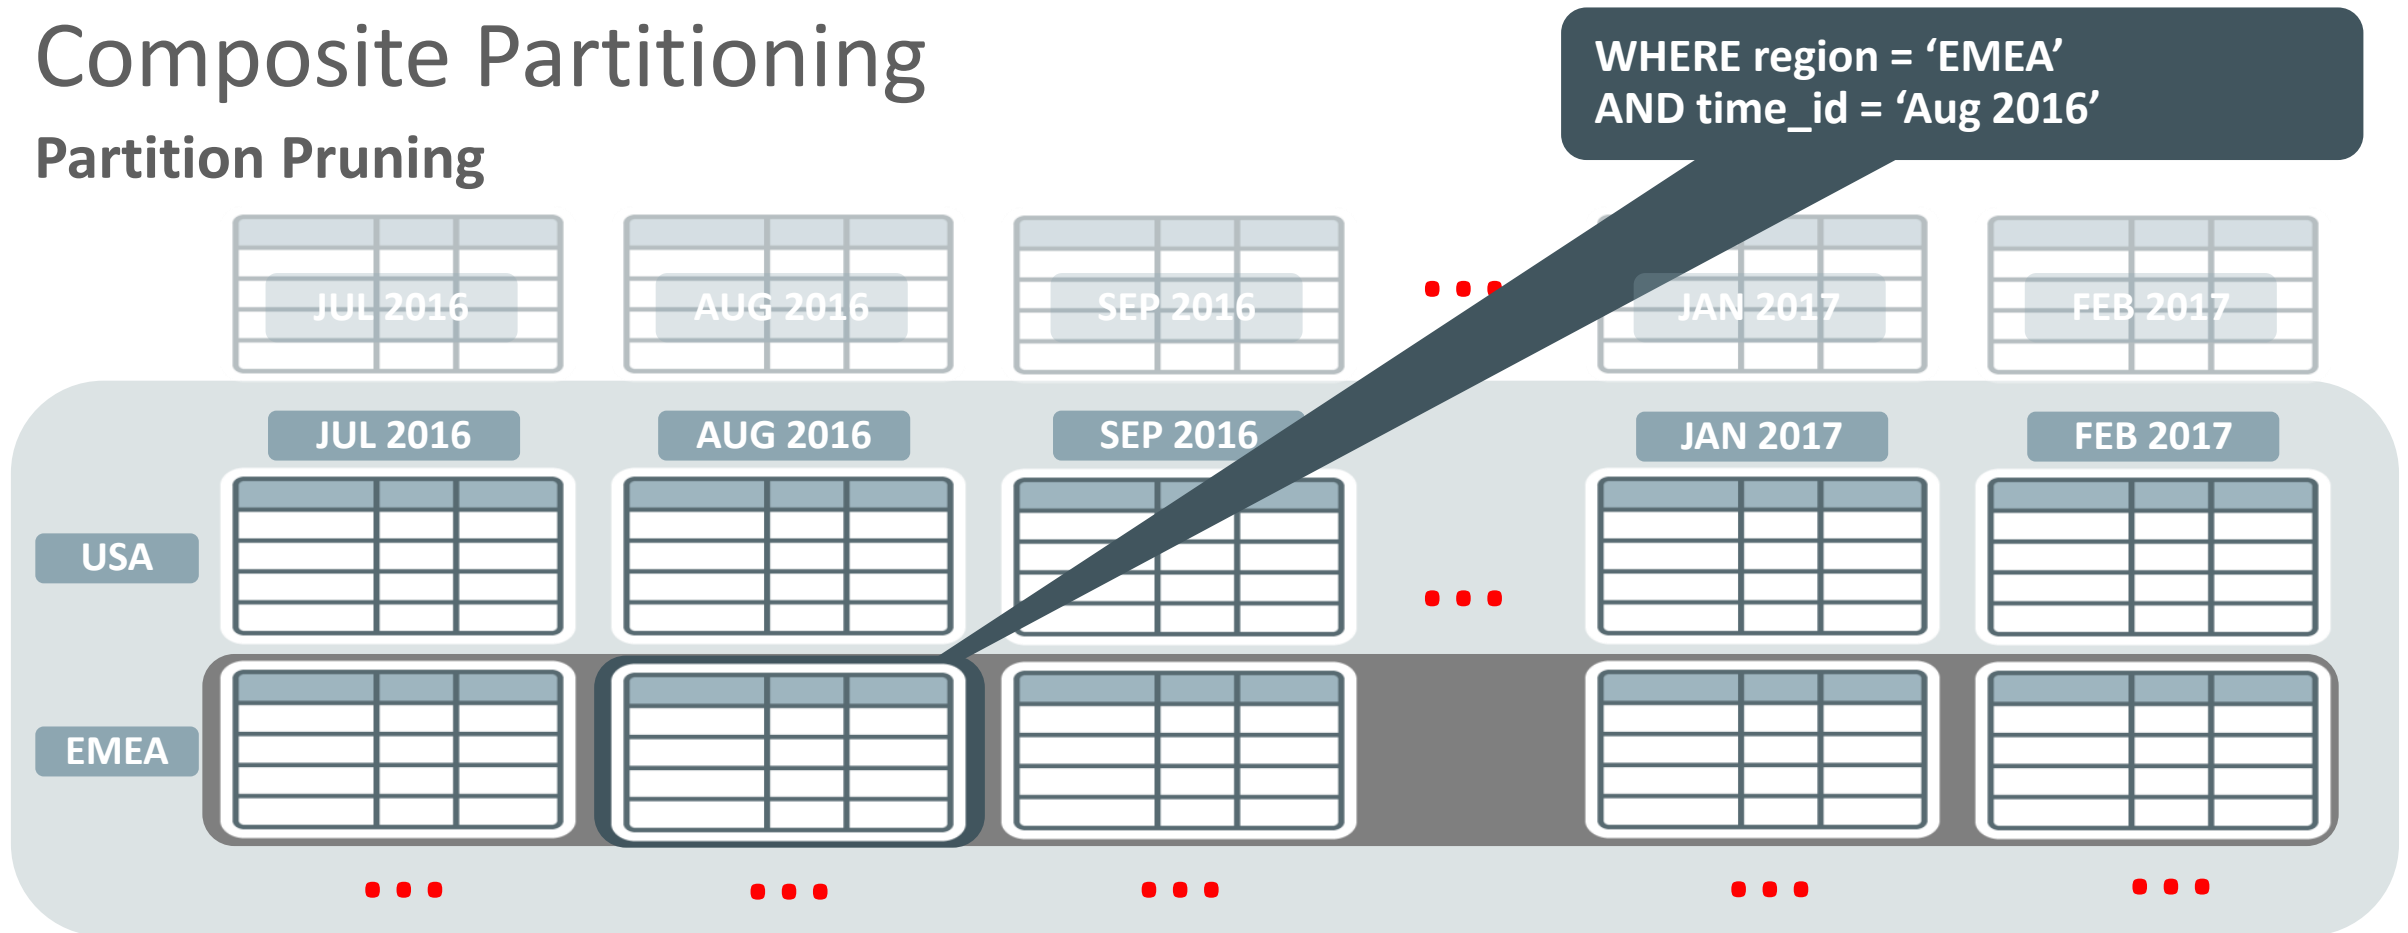
List-Range



```
CREATE TABLE SALES ..PARTITION BY LIST (region)
SUPARTITION BY RANGE (time_id)
```

Composite Partitioning

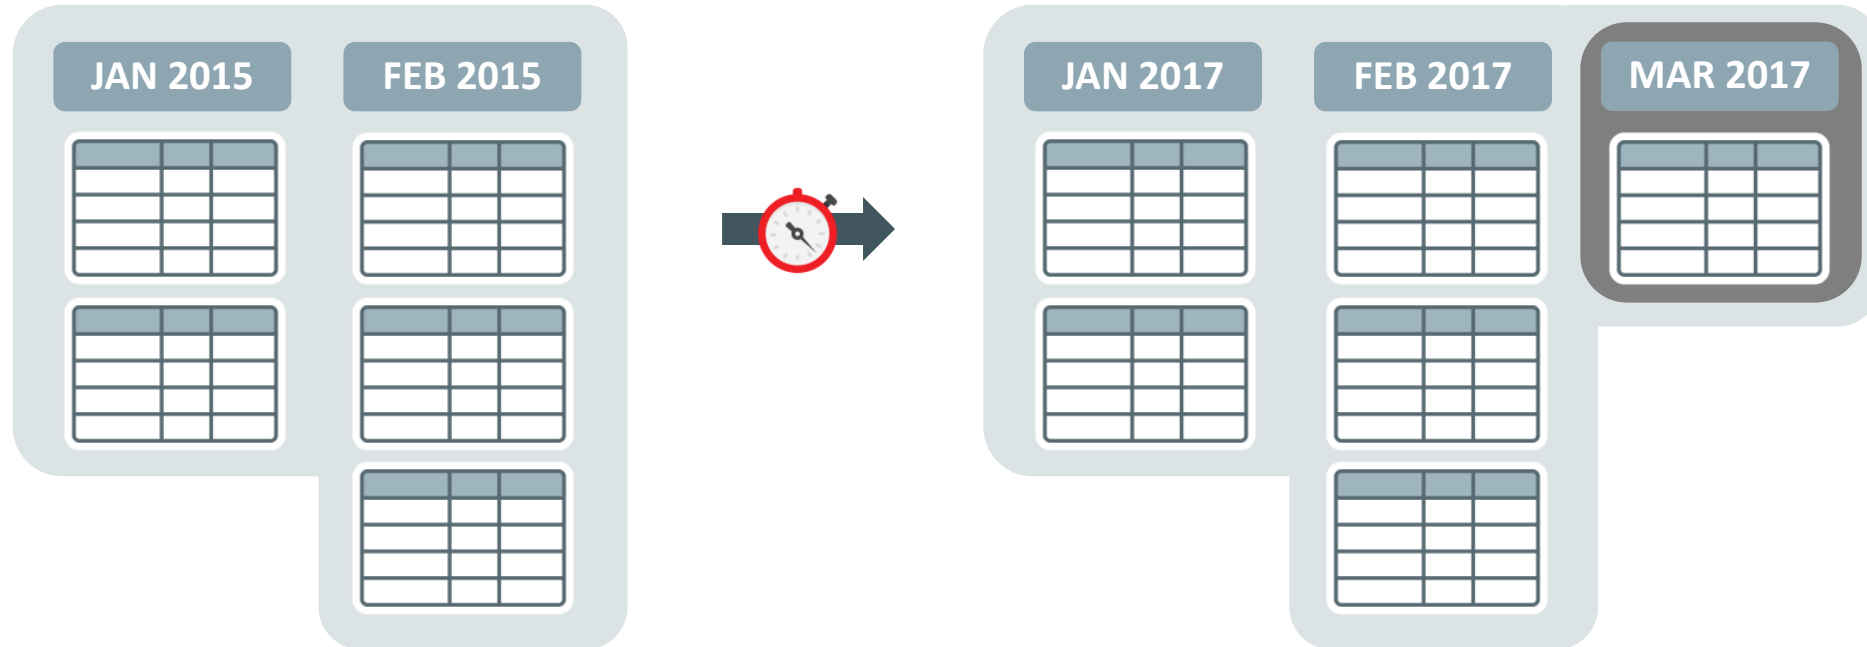
Partition Pruning



- Partition pruning is independent of composite order
 - Pruning along one or both dimensions
 - Same pruning for RANGE-LIST and LIST_RANGE

Composite Interval Partitioning

Add Partition



- Without subpartition template, only **one** subpartition will be created
 - Range: MAXVALUE
 - List: DEFAULT
 - Hash: one hash bucket

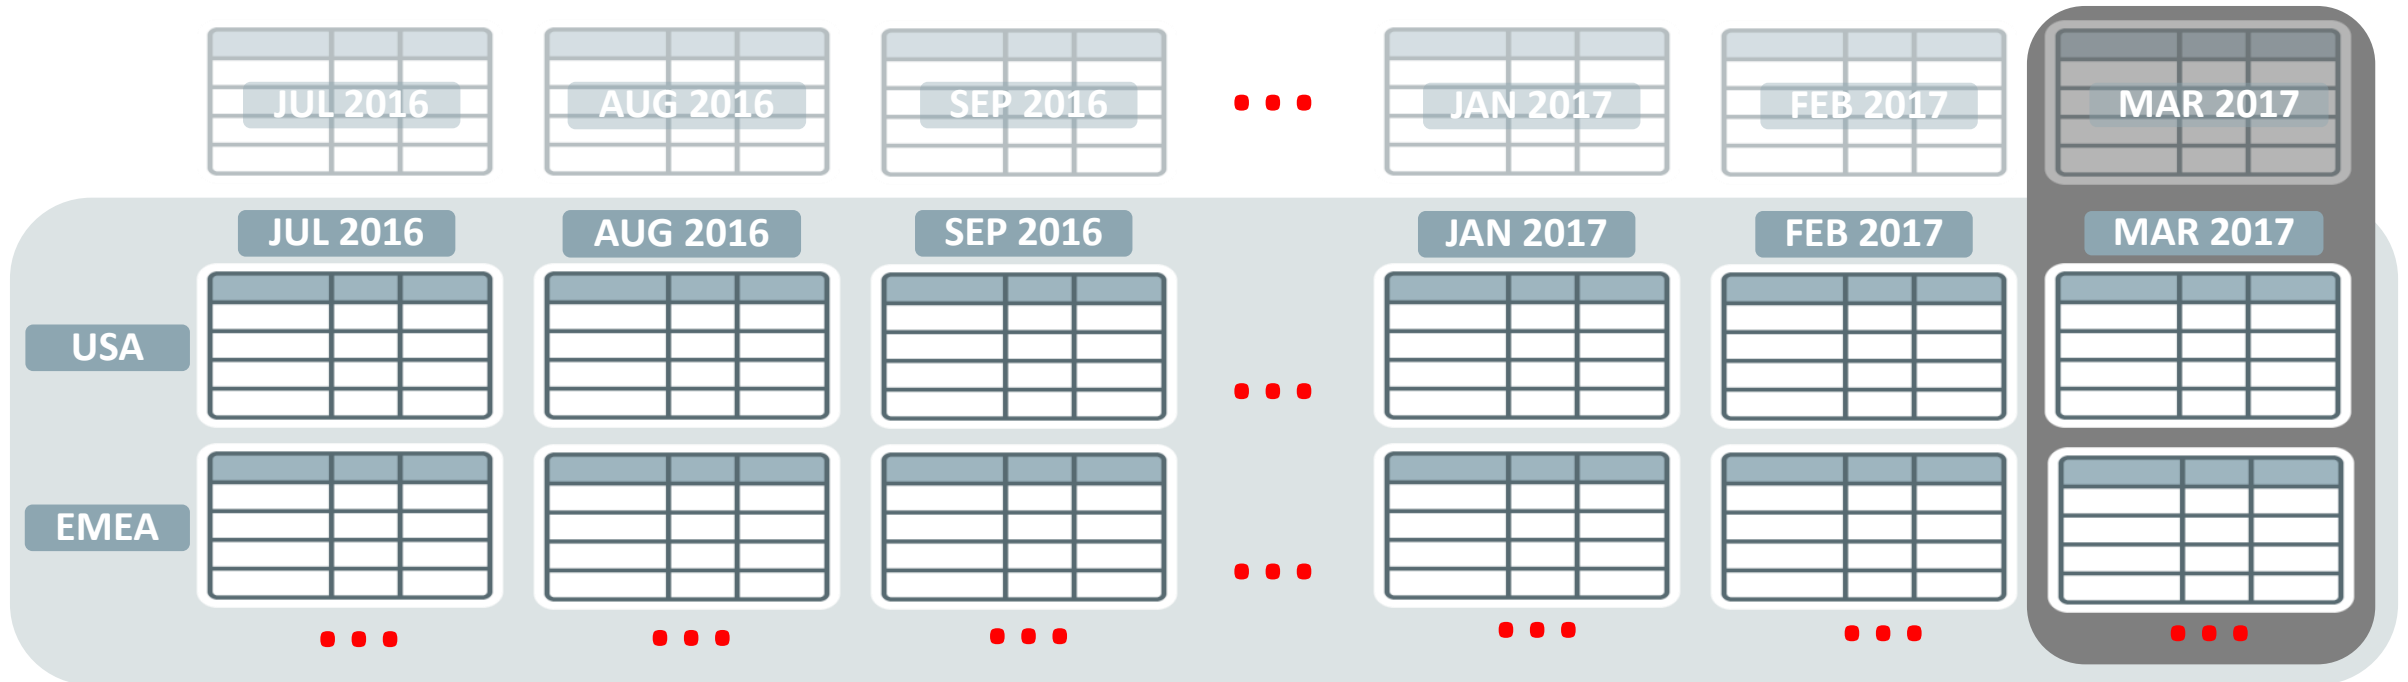
Composite Interval Partitioning

Subpartition template

- Subpartition template defines shape of **future** subpartitions
 - Can be added and/or modified at any point in time
 - No impact on existing [sub]partitions
- Controls physical attributes for subpartitions as well
 - Just like the default settings for a partitioned table does for partitions
- Difference Interval and Range Partitioning
 - Naming template only for Range
 - System-generated names for Interval

Composite Partitioning

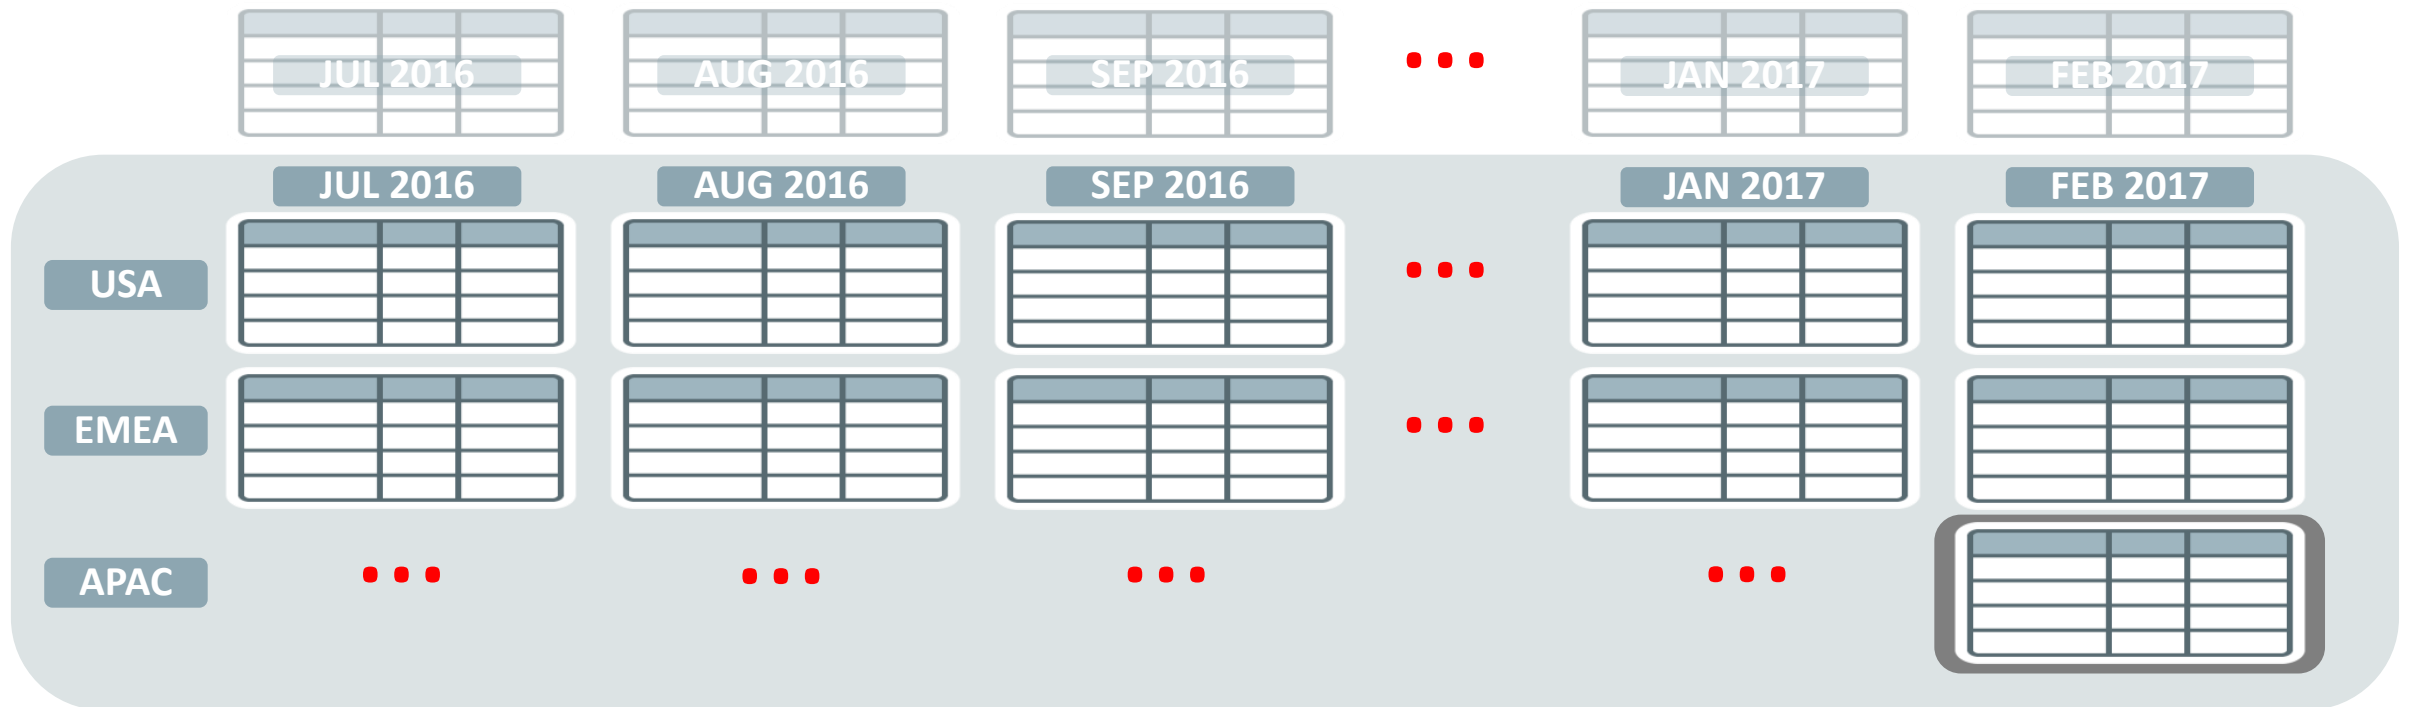
Add Partition



- ADD PARTITION always on top-level dimension
 - Identical for all newly added subpartitions
 - RANGE-LIST: new time_id range
 - LIST-RANGE: new list of region values

Composite Partitioning

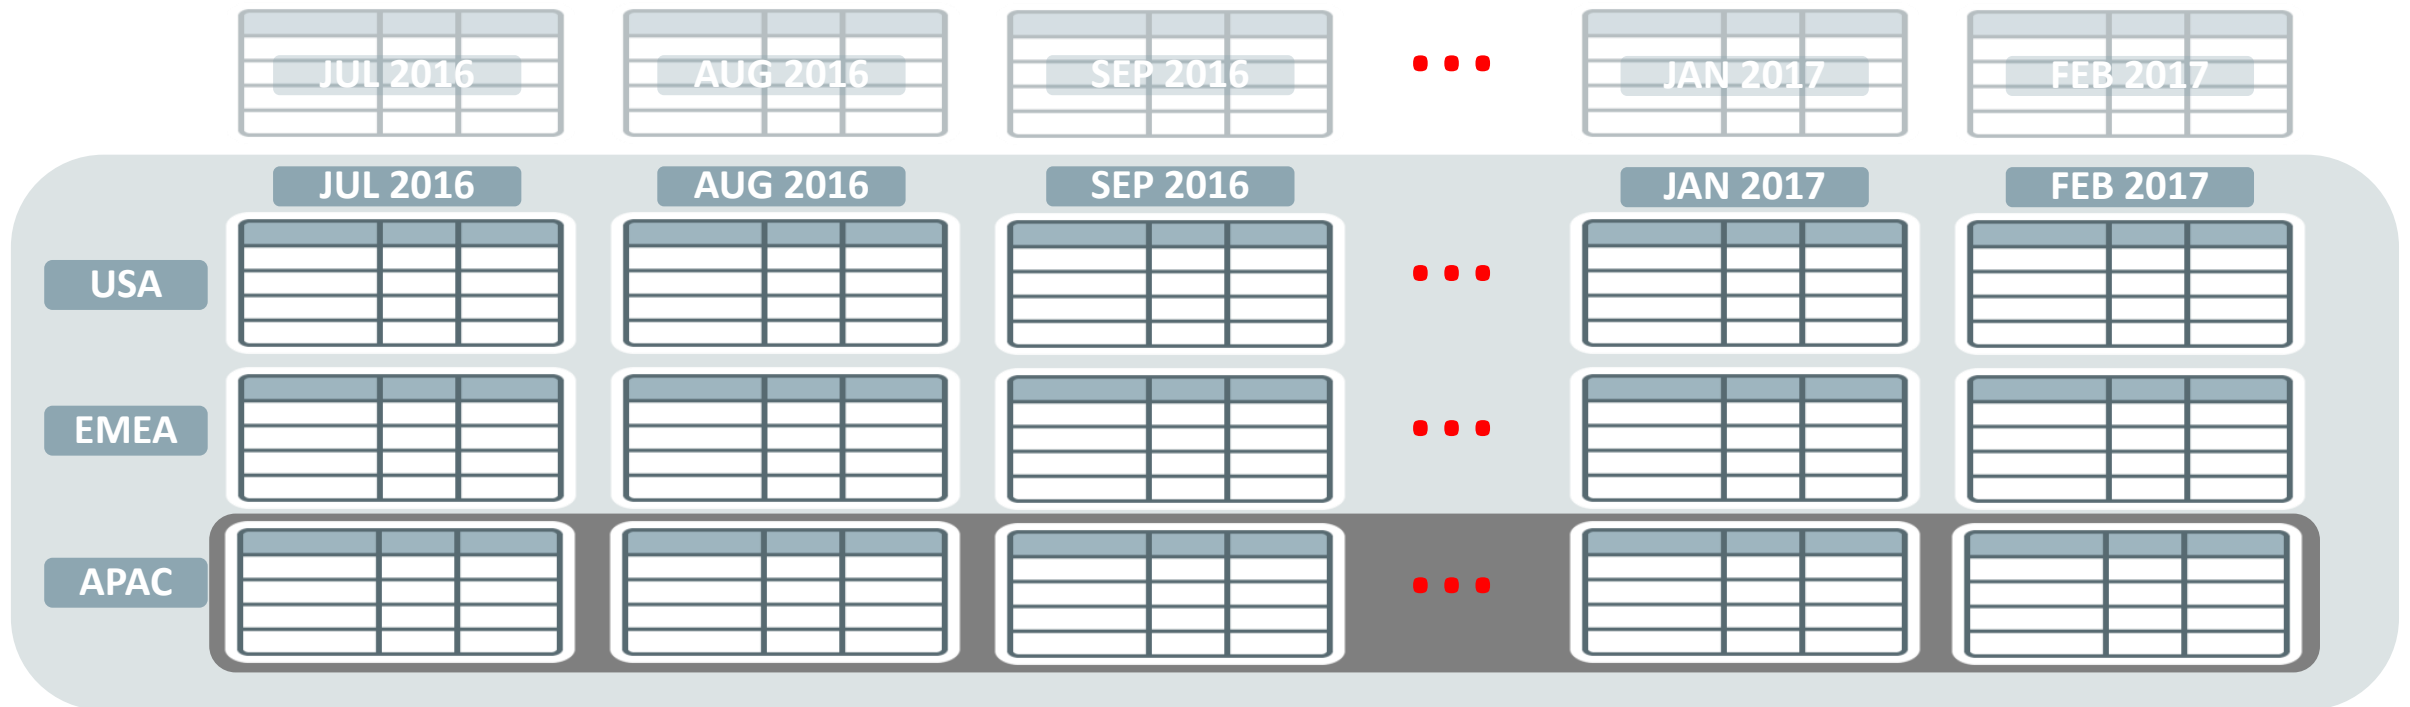
Add Subpartition



- **ADD SUBPARTITION** only for one partition
 - Asymmetric, only possible on subpartition level
 - Impact on partition-wise joins

Composite Partitioning

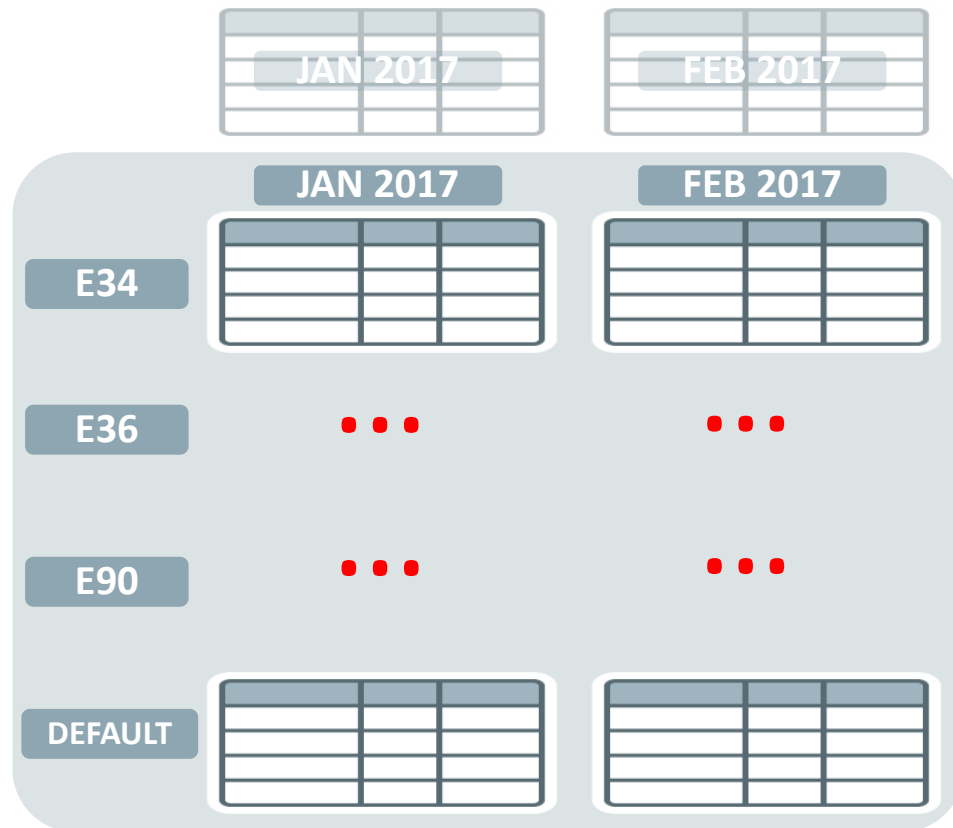
Add Subpartition



- ADD SUBPARTITION for all partitions
 - N operations necessary (for each existing partition)
 - Adjust subpartition template for future partitions

Composite Partitioning

Asymmetric subpartitions

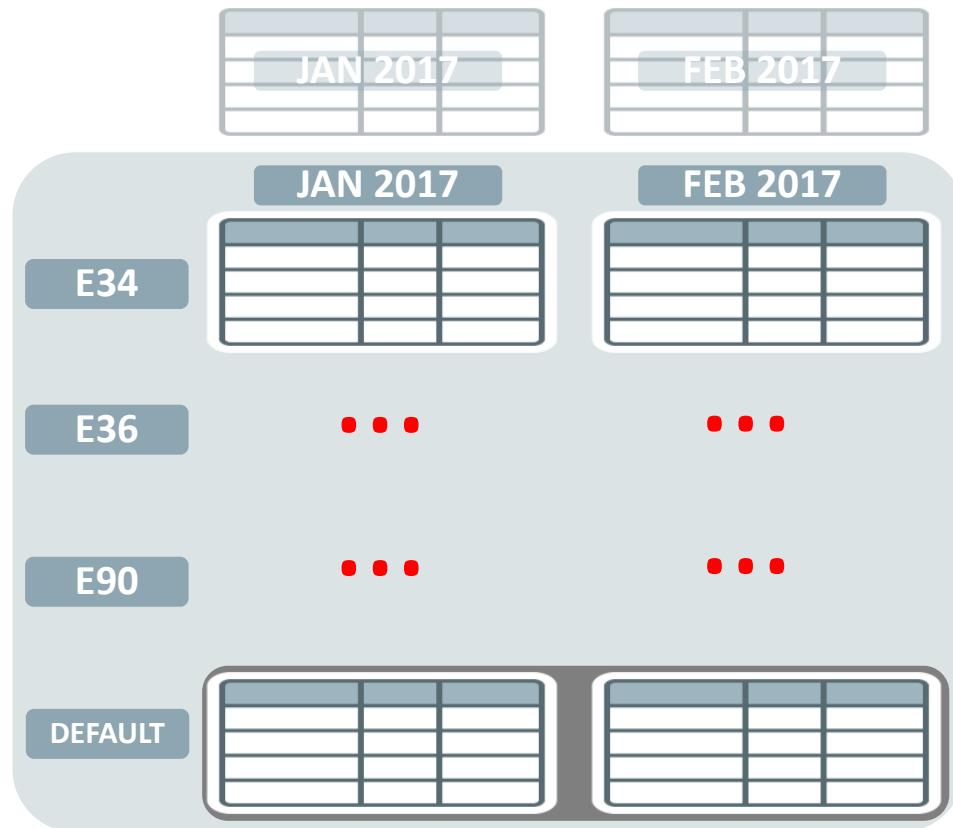


- Number of subpartitions varies for individual partitions
 - Most common for LIST subpartition strategies

```
CREATE TABLE CARS..  
PARTITION BY RANGE (time_id)  
SUBPARTITION BY LIST (model)
```

Composite Partitioning

Asymmetric subpartitions

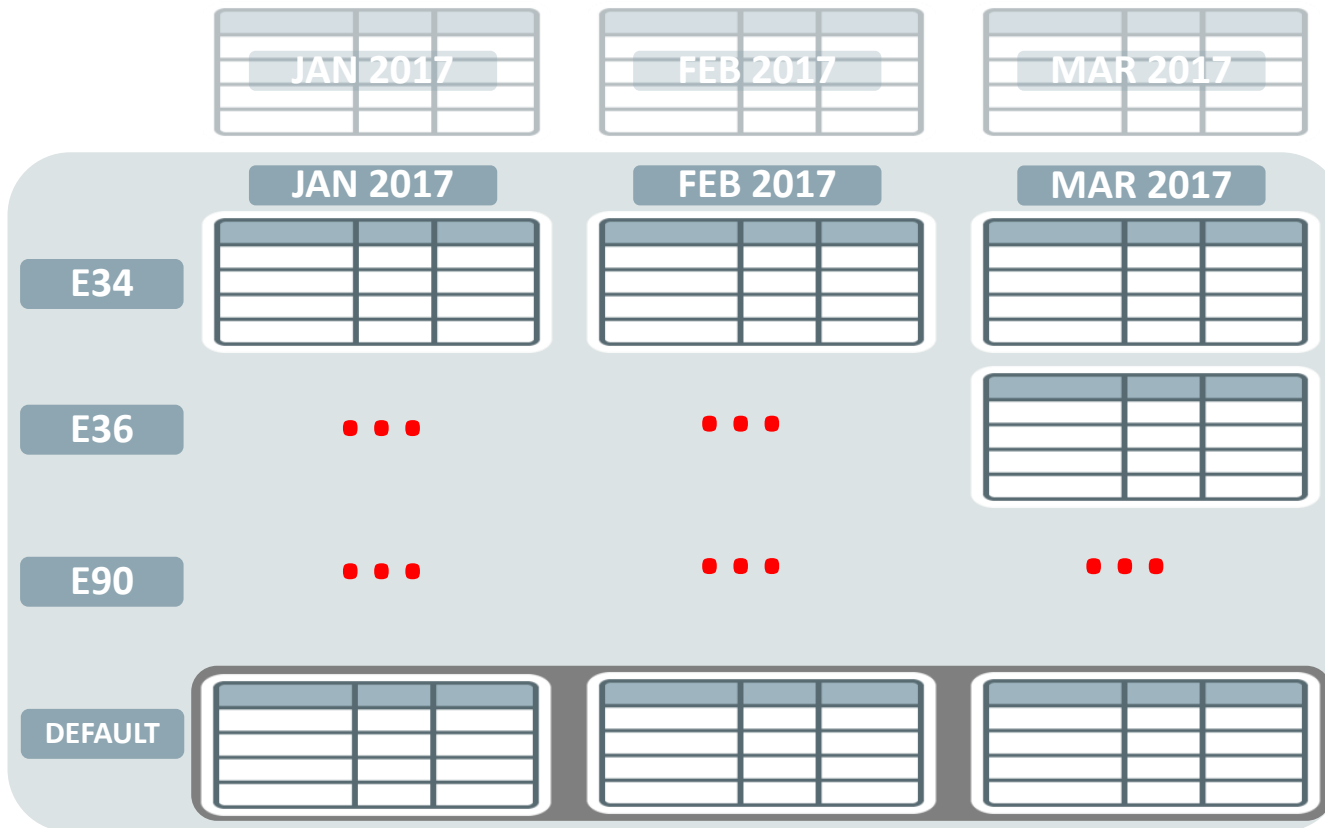


- Number of subpartitions varies for individual partitions
 - Most common for LIST subpartition strategies
- Zero impact on partition pruning capabilities

```
SELECT .. FROM cars  
WHERE model = 'E90';
```

Composite Partitioning

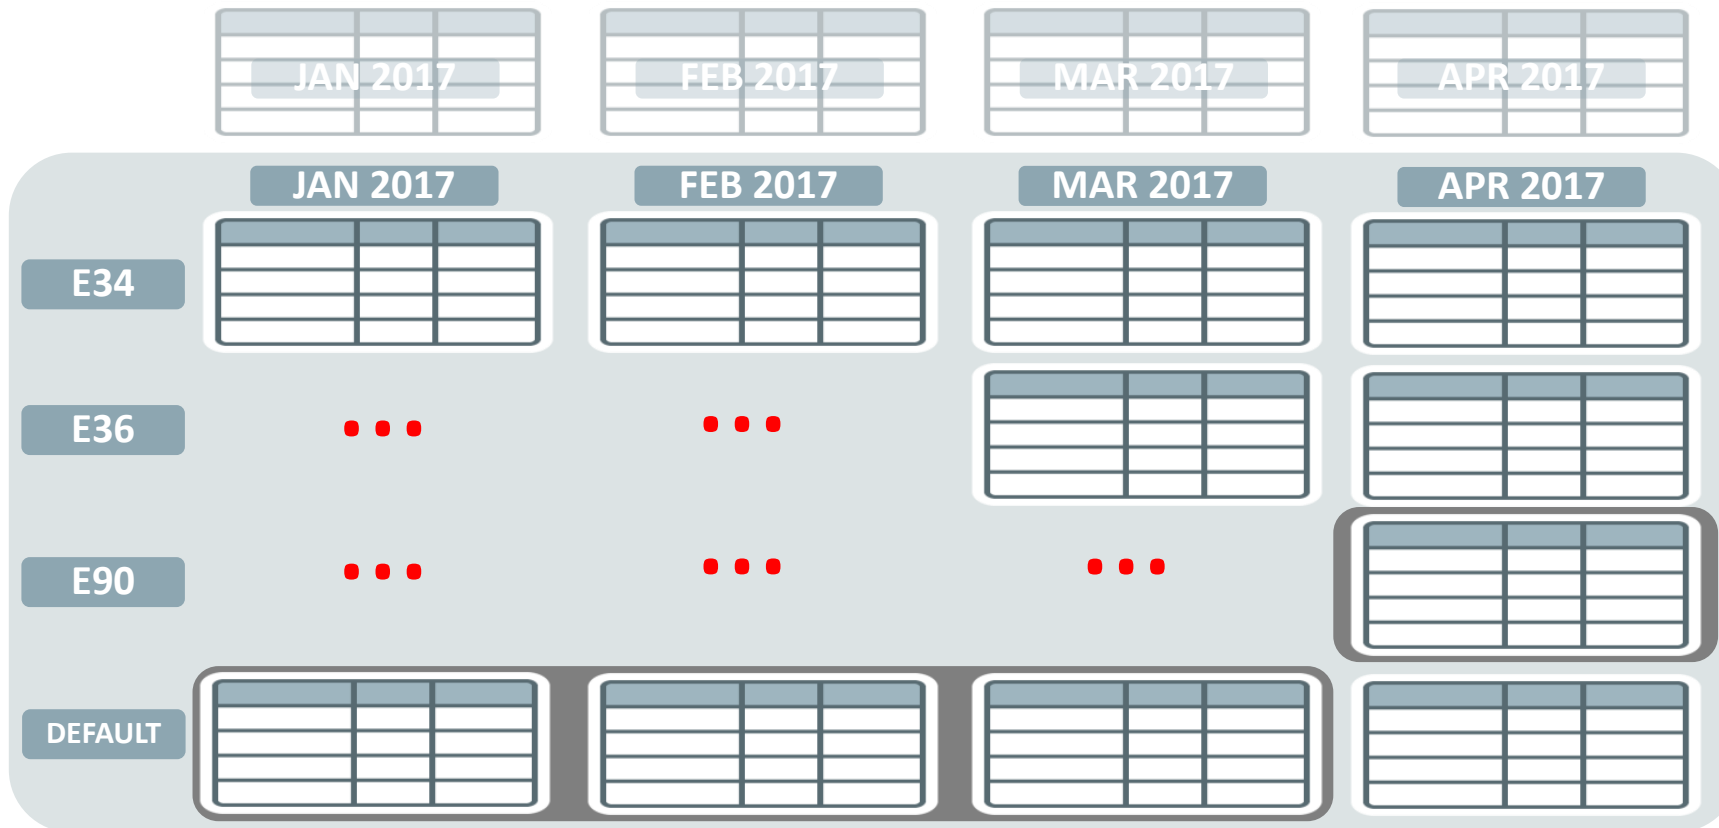
Asymmetric subpartitions



```
SELECT .. FROM cars  
WHERE model = 'E90';
```

Composite Partitioning

Asymmetric subpartitions



```
SELECT .. FROM cars  
WHERE model = 'E90';
```

Composite Partitioning

- Always use appropriate composite strategy
 - Top-level dimension mainly chosen for Manageability
 - E.g. add and drop time ranges
 - Sub-level dimension chosen for performance or manageability
 - E.g. load_id, customer_id
 - Asymmetry has advantages but should be thought through
 - E.g. different time granularity for different regions
 - Remember the impact of asymmetric composite partitioning

Multi-Column Range Partitioning

Introduced in Oracle 8i (8.1)

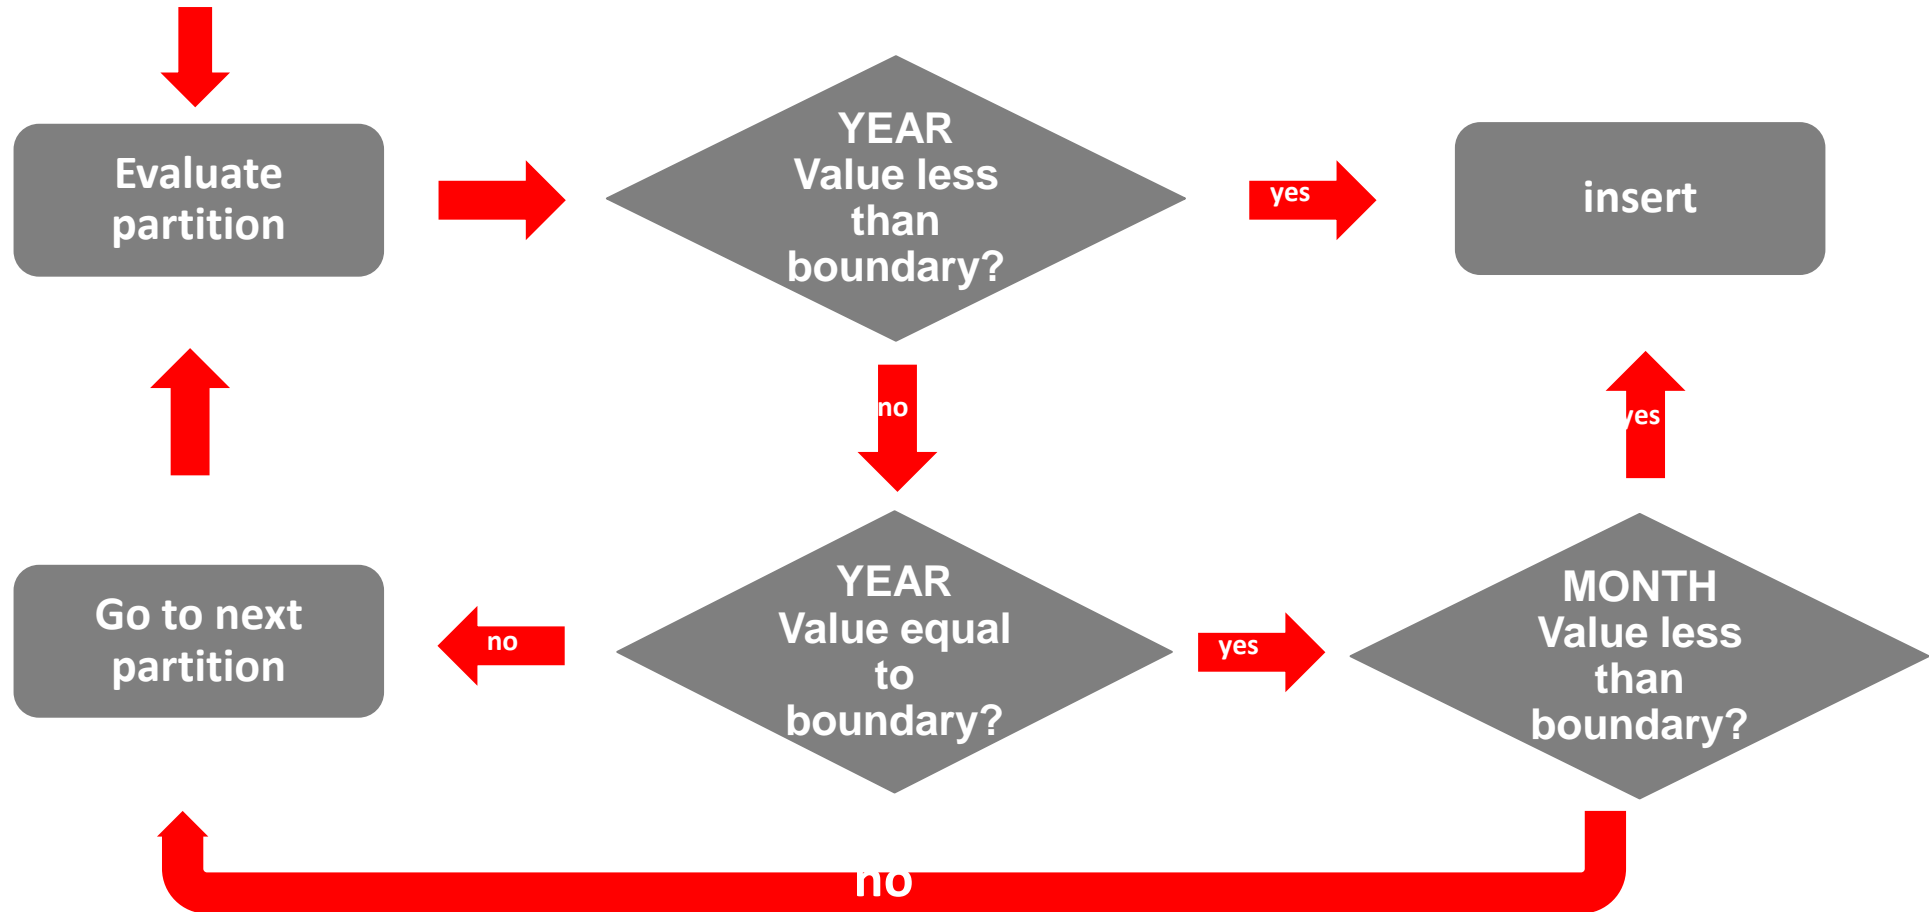
Multi-column Range Partitioning

Concept

- Partitioning key is composed of several columns and subsequent columns define a higher granularity than the preceding one
 - E.g. (YEAR, MONTH, DAY)
 - It is NOT an n-dimensional partitioning
- Major watch-out is difference of how partition boundaries are evaluated
 - For simple RANGE, the boundaries are **less than** (exclusive)
 - Multi-column RANGE boundaries are **less than or equal**
 - The n^{th} column is investigated only when all previous (n-1) values of the multicolumn key exactly match the (n-1) bounds of a partition

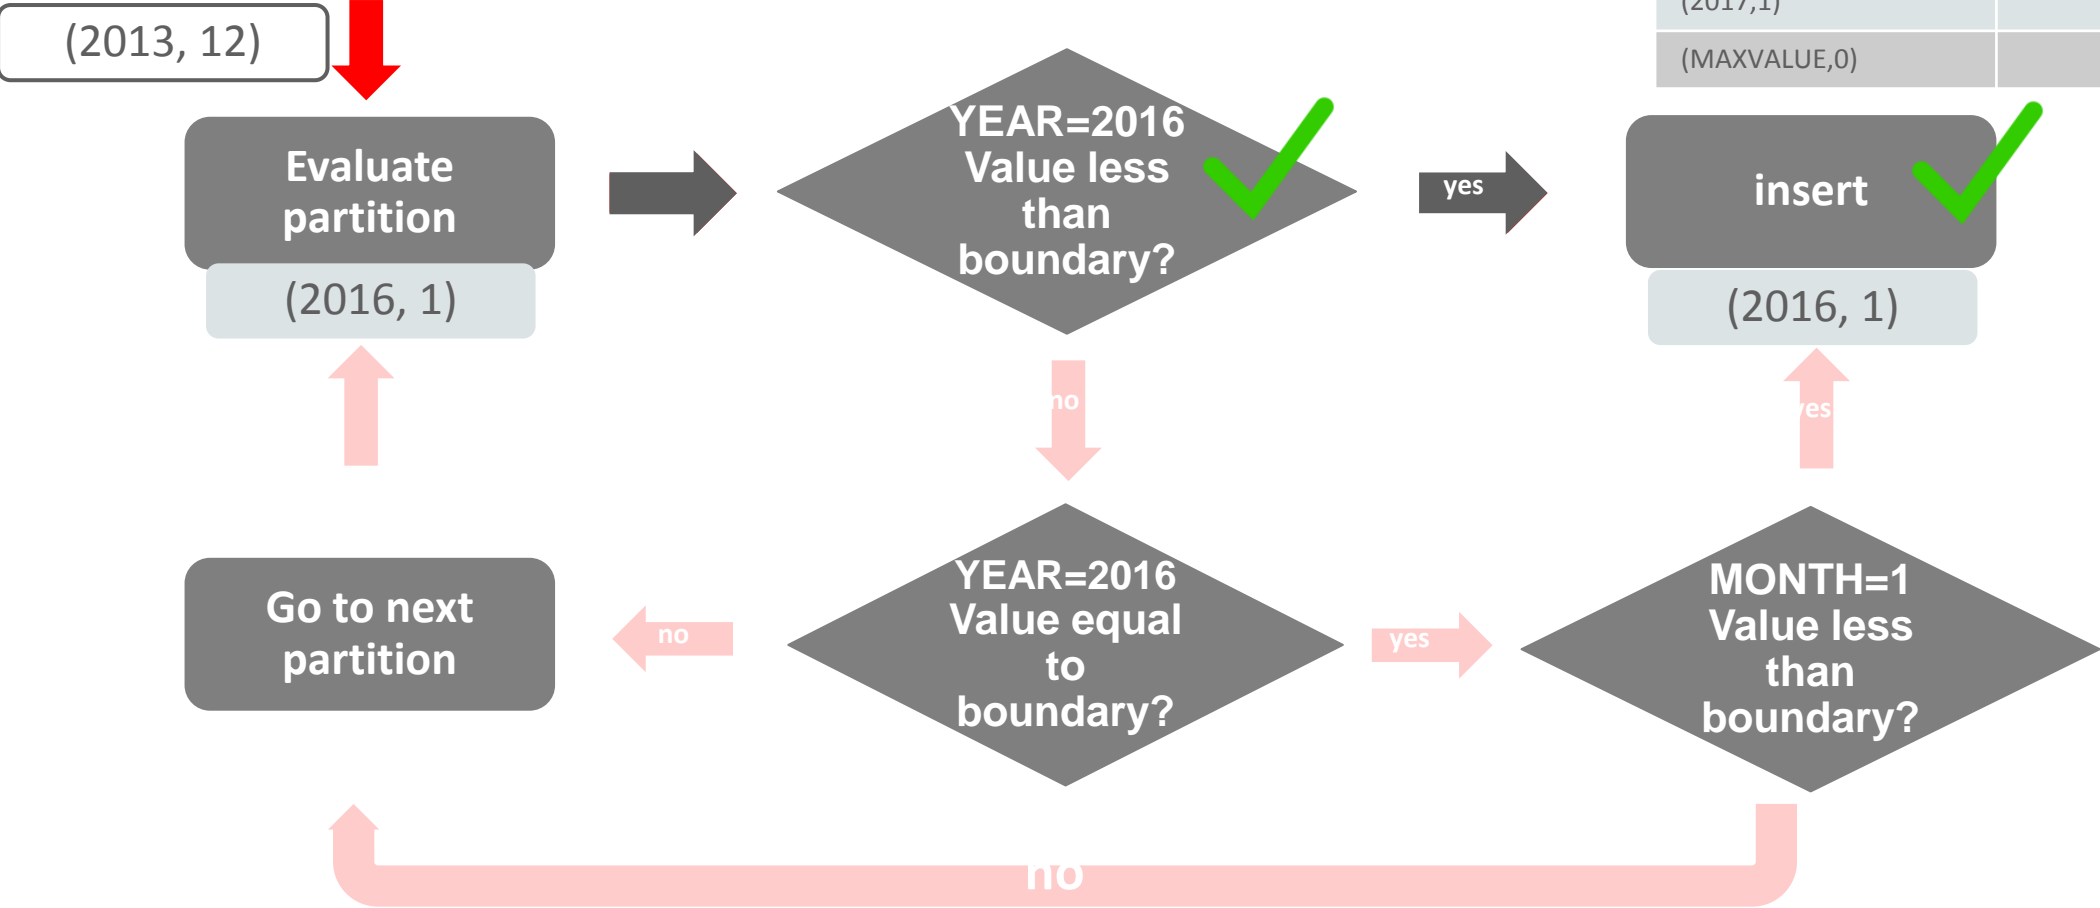
Multi-Column Range Partition

Sample Decision Tree (YEAR, MONTH)



Multi-Column Range Partition

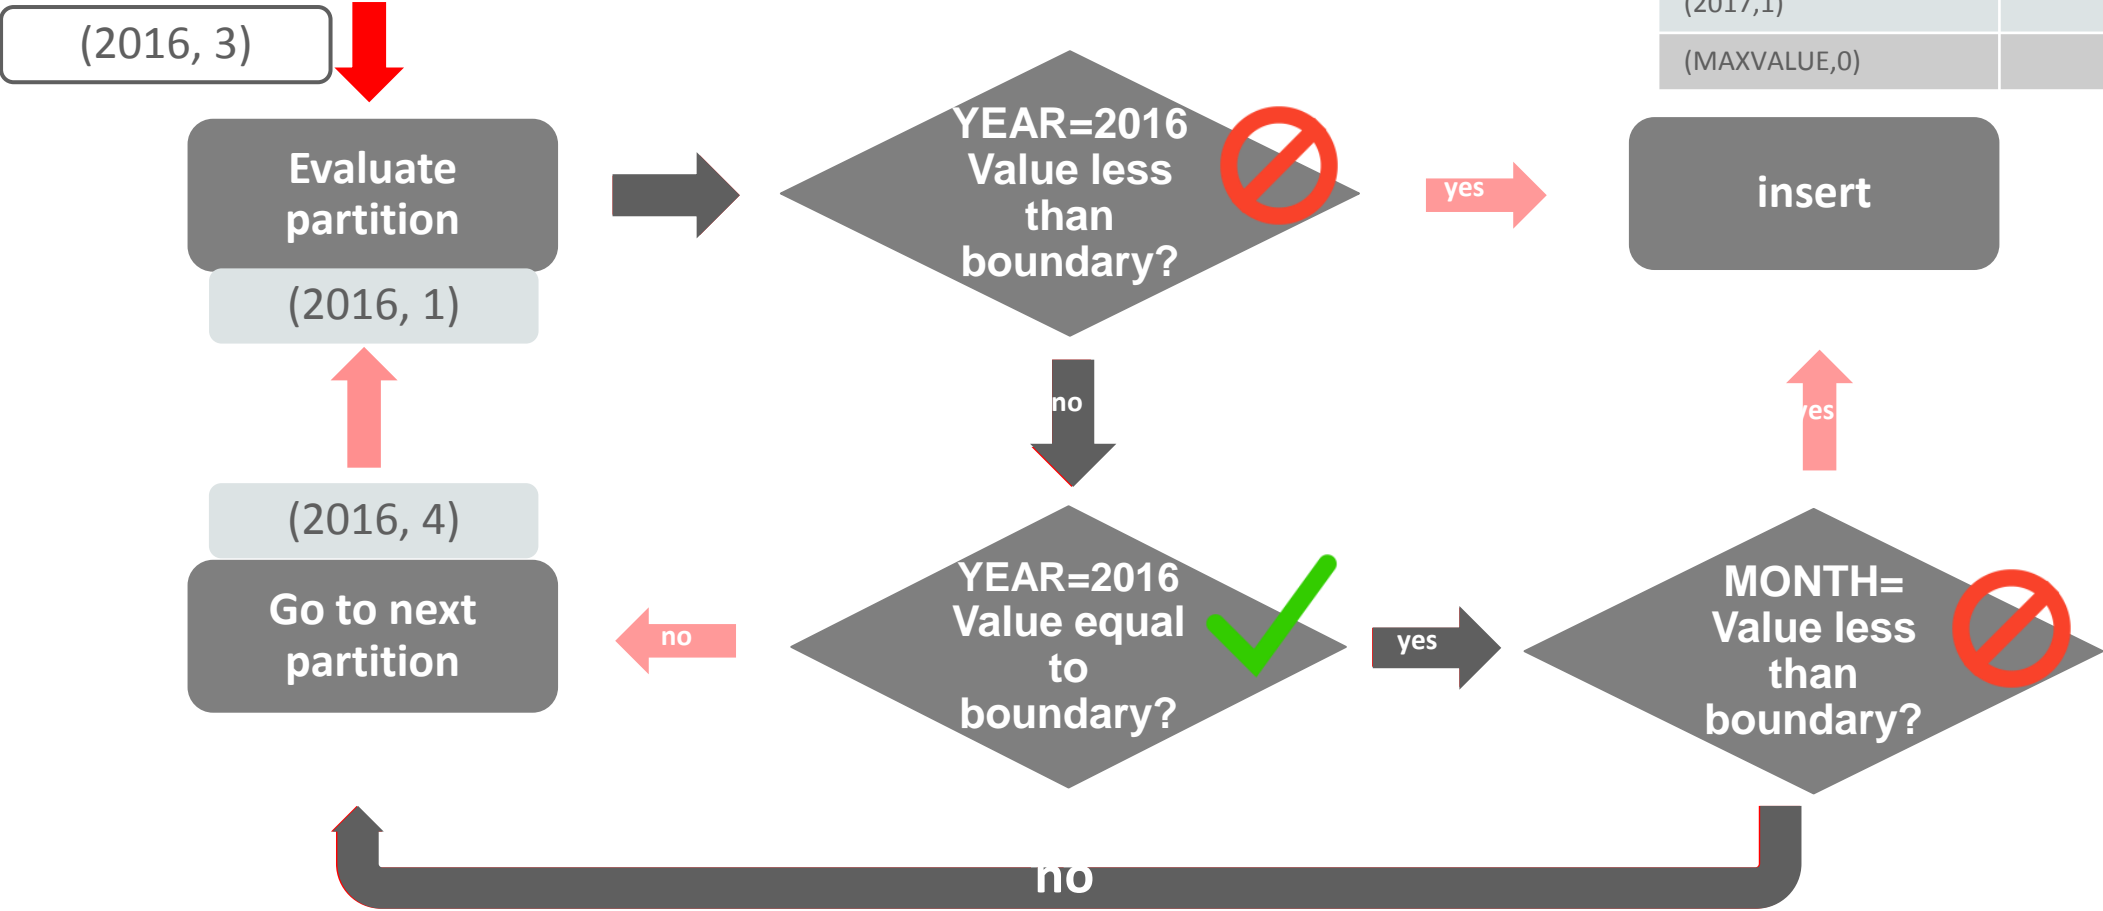
Example



(YEAR,MONTH) Boundaries	Values
(2016,1)	(2013, 12)
(2016,4)	
(2016,7)	
(2016,10)	
(2017,1)	
(MAXVALUE,0)	

Multi-Column Range Partition

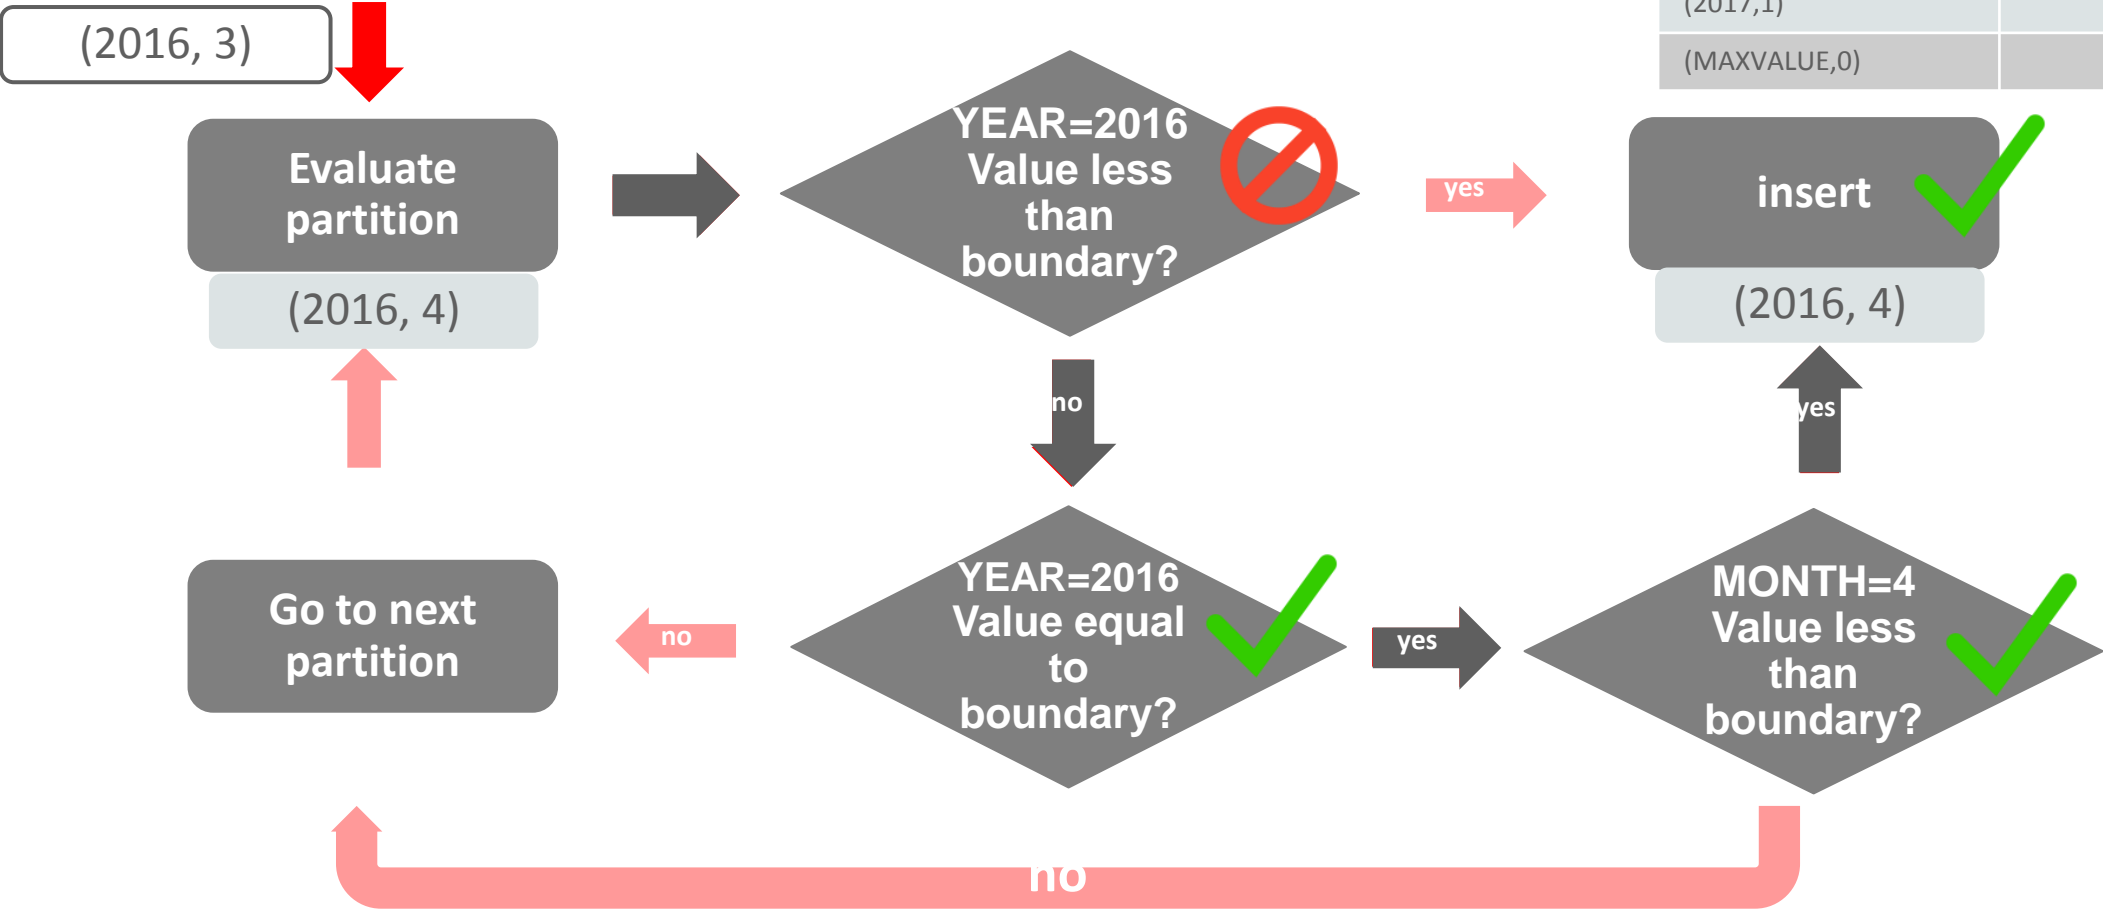
Example Cont'd



(YEAR,MONTH) Boundaries	Values
(2016,1)	(2013, 12)
(2016,4)	
(2016,7)	
(2016,10)	
(2017,1)	
(MAXVALUE,0)	

Multi-Column Range Partition

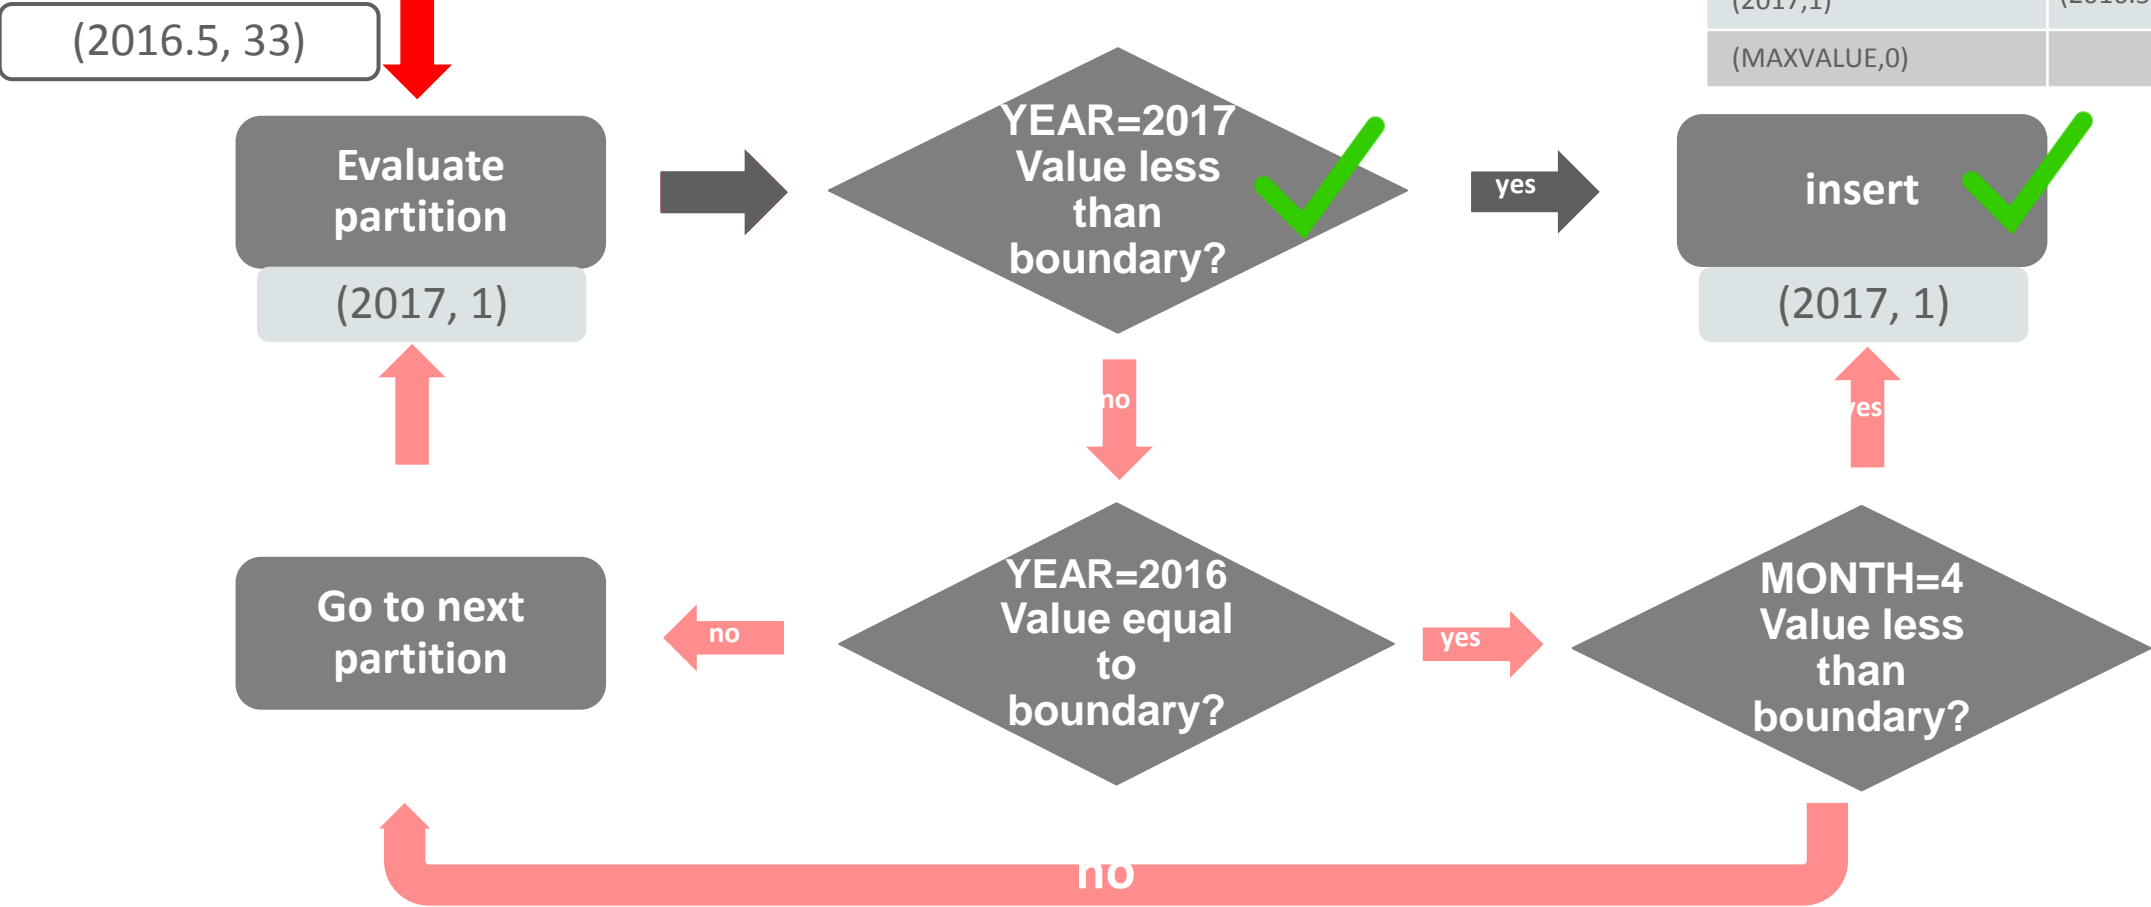
Example Cont'd



(YEAR,MONTH) Boundaries	Values
(2016,1)	(2013, 12)
(2016,4)	(2016, 3)
(2016,7)	
(2016,10)	
(2017,1)	
(MAXVALUE,0)	

Multi-Column Range Partition

Example Cont'd



(YEAR,MONTH) Boundaries	Values
(2016,1)	(2013, 12)
(2016,4)	(2016, 3)
(2016,7)	
(2016,10)	
(2017,1)	(2016.5, 33)
(MAXVALUE,0)	

Multi-Column Range Partitioning

Some things to bear in mind

- ✓ • Powerful partitioning mechanism to add a third (or more) dimensions
 - Smaller data partitions
- Pruning works also for trailing column predicates without filtering the leading column(s)
- △ • Boundaries are not enforced by the partition definition
 - Ranges are consecutive
- Logical ADD partition can mean SPLIT partition in the middle of the table

Multi-Column Range Partition

A slightly different real-world scenario

- Multi-column range used to introduce a third (non-numerical) dimension


```
CREATE TABLE product_sales (prod_id number, site_id CHAR(2), start_date date)
PARTITION BY RANGE (site_id, start_date)
SUBPARTITION BY HASH (prod_id) SUBPARTITIONS 16
(PARTITION de_2013 VALUES LESS THAN ('DE', to_date('01-JAN-2014', 'dd-mon-yyyy')),
 PARTITION de_2016 VALUES LESS THAN ('DE', to_date('01-JAN-2015', 'dd-mon-yyyy')),
 PARTITION us_2013 VALUES LESS THAN ('US', to_date('01-JAN-2014', 'dd-mon-yyyy')),
 PARTITION us_2016 VALUES LESS THAN ('US', to_date('01-JAN-2015', 'dd-mon-yyyy')),
 PARTITION za_2013 VALUES LESS THAN ('ZA', to_date('01-JAN-2014', 'dd-mon-yyyy')),
 PARTITION za_2016 VALUES LESS THAN ('ZA', to_date('01-JAN-2015', 'dd-mon-yyyy'))
);
```

Character SITE_ID has to be defined in an ordered fashion

Multi-Column Range Partition

A slightly different real-world scenario

- Multi-column range used to introduce a third (non-numerical) dimension



```
CREATE TABLE product_sales (prod_id number, site_id CHAR(2), start_date date)
PARTITION BY RANGE (site_id, start_date)
SUBPARTITION BY HASH (prod_id) SUBPARTITIONS 16
(PARTITION de_2013 VALUES LESS THAN ('DE',to_date('01-JAN-2014','dd-mon-yyyy')),
PARTITION de_2016 VALUES LESS THAN ('DE',to_date('01-JAN-2015','dd-mon-yyyy')),
PARTITION us_2013 VALUES LESS THAN ('US',to_date('01-JAN-2014','dd-mon-yyyy')),
PARTITION us_2016 VALUES LESS THAN ('US',to_date('01-JAN-2015','dd-mon-yyyy')),
PARTITION za_2013 VALUES LESS THAN ('ZA',to_date('01-JAN-2014','dd-mon-yyyy')),
PARTITION za_2016 VALUES LESS THAN ('ZA',to_date('01-JAN-2015','dd-mon-yyyy'))
);
```

Non-defined SITE_ID will follow the LESS THAN probe and always end in the lowest partition of a defined SITE_ID

Multi-Column Range Partition

A slightly different real-world scenario

- Multi-column range used to introduce a third (non-numerical) dimension

(DE, 2015) (US, 2015) (ZA, 2015) ❌

```
CREATE TABLE product_sales (prod_id number, site_id CHAR(2), start_date date)
PARTITION BY RANGE (site_id, start_date)
SUBPARTITION BY HASH (prod_id) SUBPARTITIONS 16
(PARTITION de_2013 VALUES LESS THAN ('DE', to_date('01-JAN-2014', 'dd-mon-yyyy')),
 PARTITION de_2016 VALUES LESS THAN ('DE', to_date('01-JAN-2015', 'dd-mon-yyyy')),
 PARTITION us_2013 VALUES LESS THAN ('US', to_date('01-JAN-2014', 'dd-mon-yyyy')),
 PARTITION us_2016 VALUES LESS THAN ('US', to_date('01-JAN-2015', 'dd-mon-yyyy')),
 PARTITION za_2013 VALUES LESS THAN ('ZA', to_date('01-JAN-2014', 'dd-mon-yyyy')),
 PARTITION za_2016 VALUES LESS THAN ('ZA', to_date('01-JAN-2015', 'dd-mon-yyyy'))
);
```

?

Future dates will always go in the lowest partition of the next higher SITE_ID or being rejected

Multi-Column Range Partition

A slightly different real-world scenario

- Multi-column range used to introduce a third (non-numerical) dimension

AC, CN

EE, ES, UK

```
create table product_sales (prod_id number, site_id CHAR(2), start_date date)
partition by range (site_id, start_date)
subpartition by hash (prod_id) subpartitions 16
(partition below_de values less than ('DE',to_date('01-JAN-1492','dd-mon-yyyy')),
partition de_2013 values less than ('DE',to_date('01-JAN-2014','dd-mon-yyyy')),
partition de_2016 values less than ('DE',to_date('01-JAN-2015','dd-mon-yyyy')),
partition de_max values less than ('DE',MAXVALUE),
partition below_us values less than ('US',to_date('01-JAN-1492','dd-mon-yyyy')),
...
partition za_max values less than ('ZA',MAXVALUE),
partition pmax values less than (MAXVALUE,MAXVALUE));
```

Introduce a dummy 'BELOW_X' partition to catch "lower" nondefined SITE_ID

Multi-Column Range Partition

A slightly different real-world scenario

- Multi-column range used to introduce a third (non-numerical) dimension

```
create table product_sales (prod_id number, site_id CHAR(2), start_date date)
partition by range (site_id, start_date)
subpartition by hash (prod_id) subpartitions 16
(partition below_de values less than ('DE',to_date('01-JAN-1492','dd-mon-yyyy')),
partition de_2013 values less than ('DE',to_date('01-JAN-2014','dd-mon-yyyy')),
partition de_2016 values less than ('DE',to_date('01-JAN-2015','dd-mon-yyyy')),
partition de_max values less than ('DE',MAXVALUE),
partition below_us values less than ('US',to_date('01-JAN-1492','dd-mon-yyyy')),
...
partition za_max values less than ('ZA',MAXVALUE),
partition pmax values less than (MAXVALUE,MAXVALUE));
```

(DE, 2015)

(ZA, 2015)

Introduce a MAXVALUE 'X_FUTURE' partition to catch future dates

Multi-Column Range Partition

A slightly different real-world scenario

- Multi-column range used to introduce a third (non-numerical) dimension

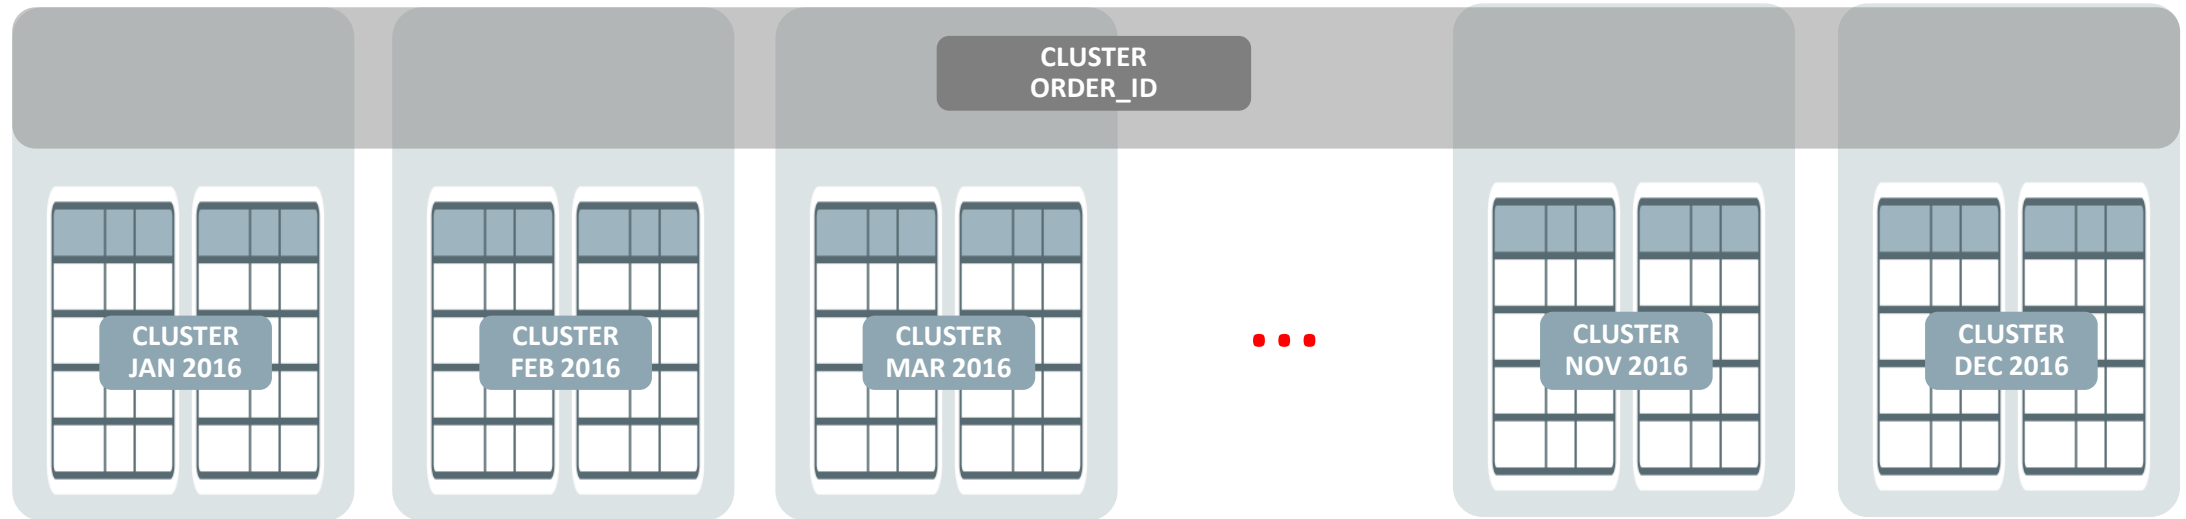
```
create table product_sales (prod_id number, site_id CHAR(2), start_date date)
partition by range (site_id, start_date)
subpartition by hash (prod_id) subpartitions 16
(partition below_de values less than ('DE',to_date('01-JAN-1492','dd-mon-yyyy')),
partition de_2013 values less than ('DE',to_date('01-JAN-2014','dd-mon-yyyy')),
partition de_2016 values less than ('DE',to_date('01-JAN-2015','dd-mon-yyyy')),
partition de_max values less than ('DE',MAXVALUE),
partition below_us values less than ('US',to_date('01-JAN-1492','dd-mon-yyyy')),
...
partition za_max values less than ('ZA',MAXVALUE),
partition pmax values less than (MAXVALUE,MAXVALUE));
```

If necessary, catch the open-ended SITE_ID (leading key column)

Range-Partitioned Hash Cluster

Introduced in Oracle 12c Release 1 (12.1.0.2)

Range-Partitioned Hash Cluster

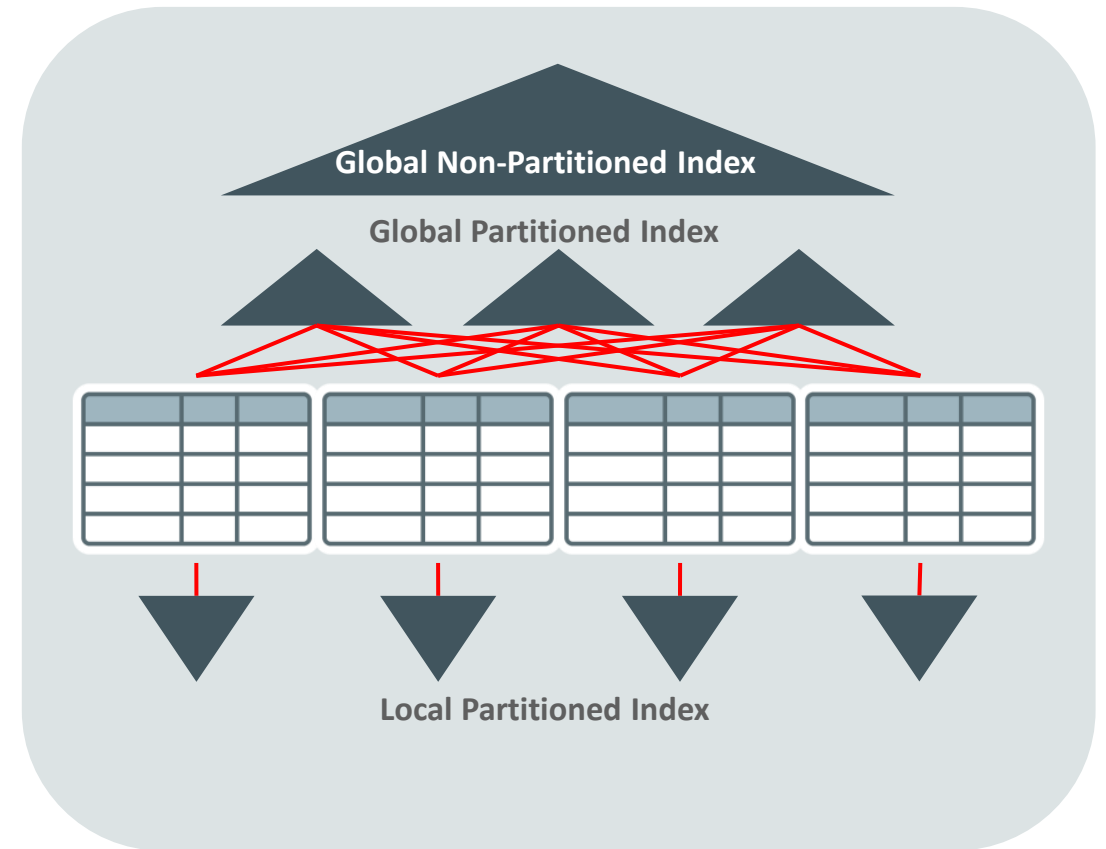


- Single-level range partitioning
 - No composite partitioning
 - No index clusters

Indexing of Partitioned Tables

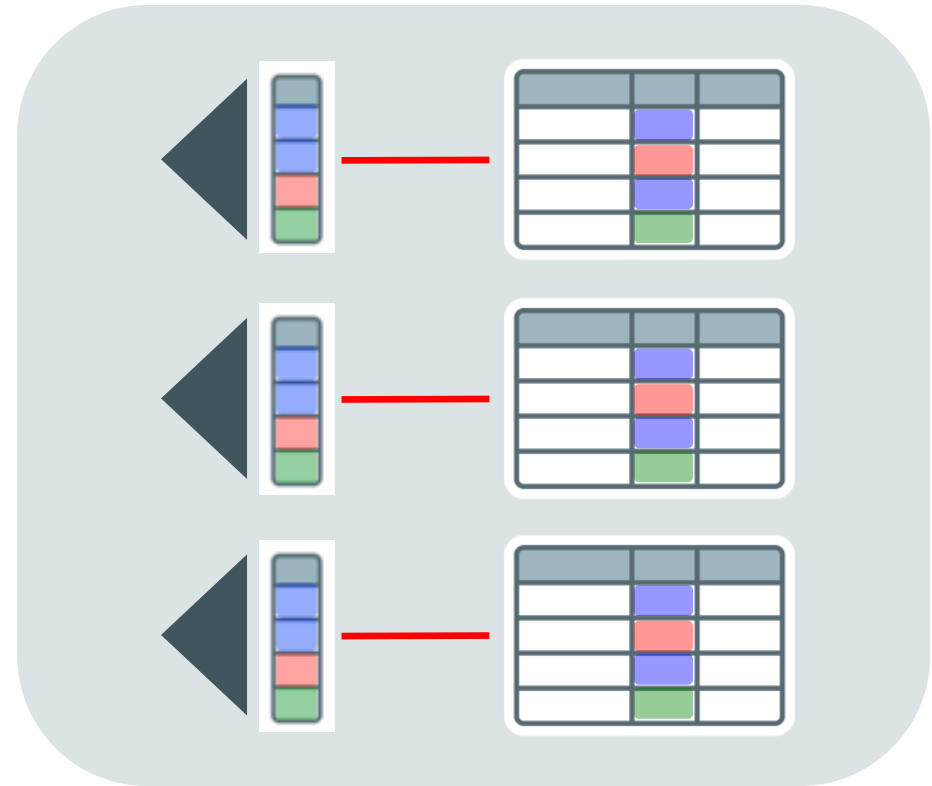
Index Partitioning

- GLOBAL index points to rows in any partition
 - Index can be partitioned or not
 - Partition maintenance affects entire index
- LOCAL index is partitioned same as table
 - Index partitioning key can be different from index key
 - Index partitions can be maintained separately



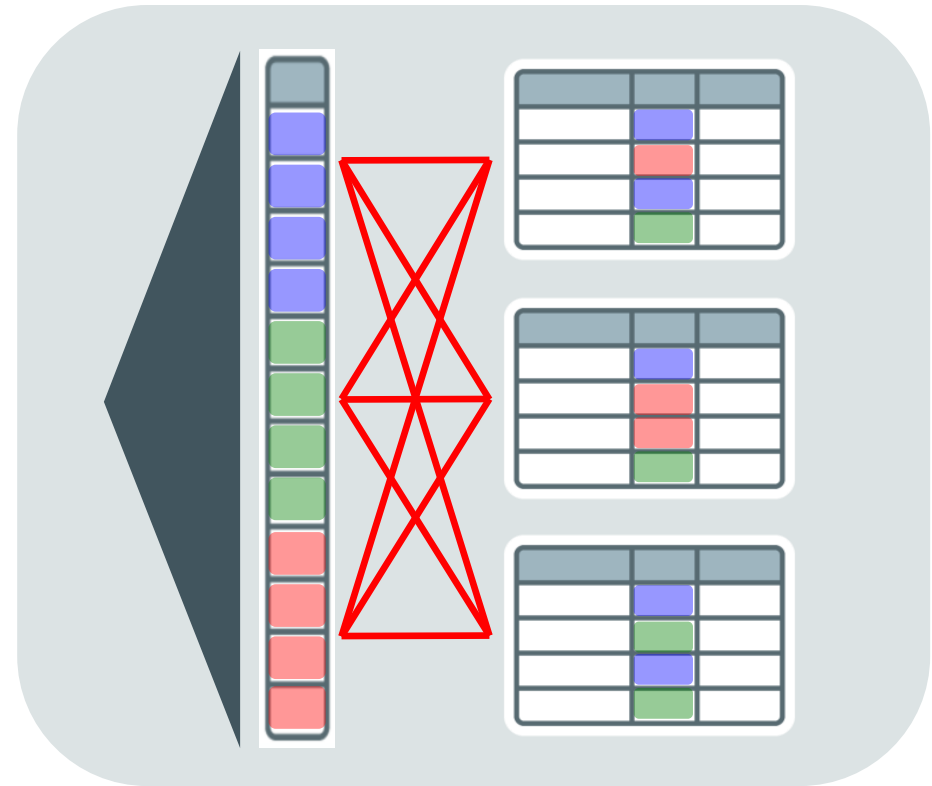
Local Index

- Index is partitioned along same boundaries as data
 - B-tree or bitmap
- Pros
 - Easy to manage
 - Parallel index scans
- Cons
 - Less efficient for retrieving small amounts of data



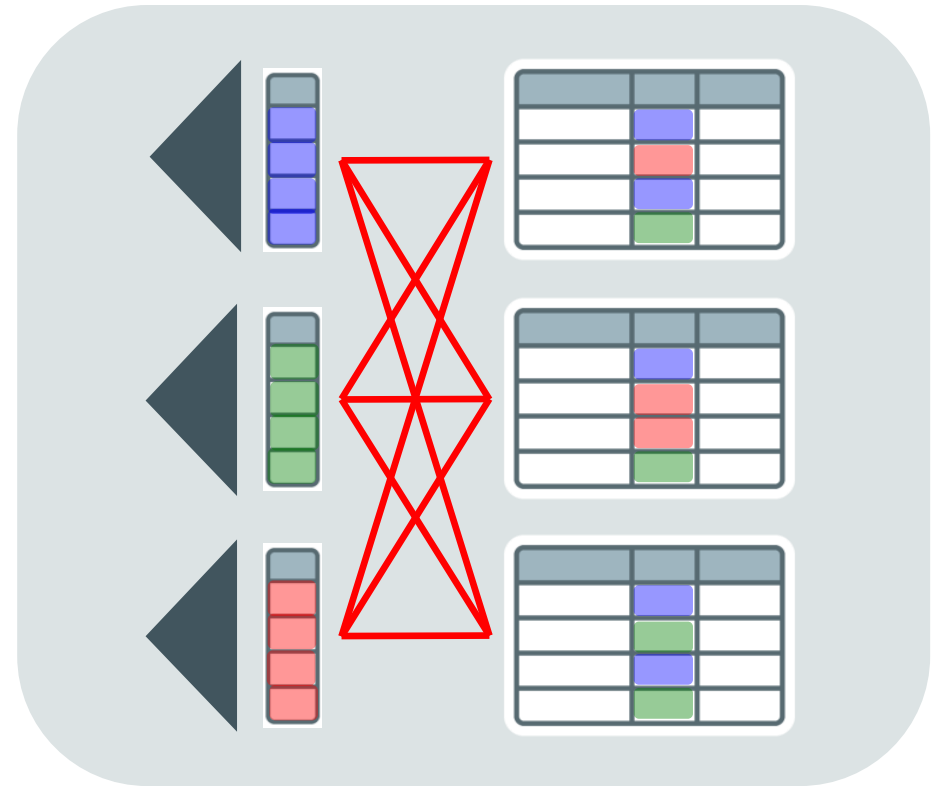
Global Non-Partitioned Index

- One index b-tree structure that spans all partitions
- Pros
 - Efficient access to any individual record
- Cons
 - Partition maintenance always involves index maintenance



Global Partitioned Index

- Index is partitioned independently of data
 - Each index structure may reference any and all partitions.
- Pros
 - Availability and manageability
- Cons
 - Partition maintenance always involves index maintenance



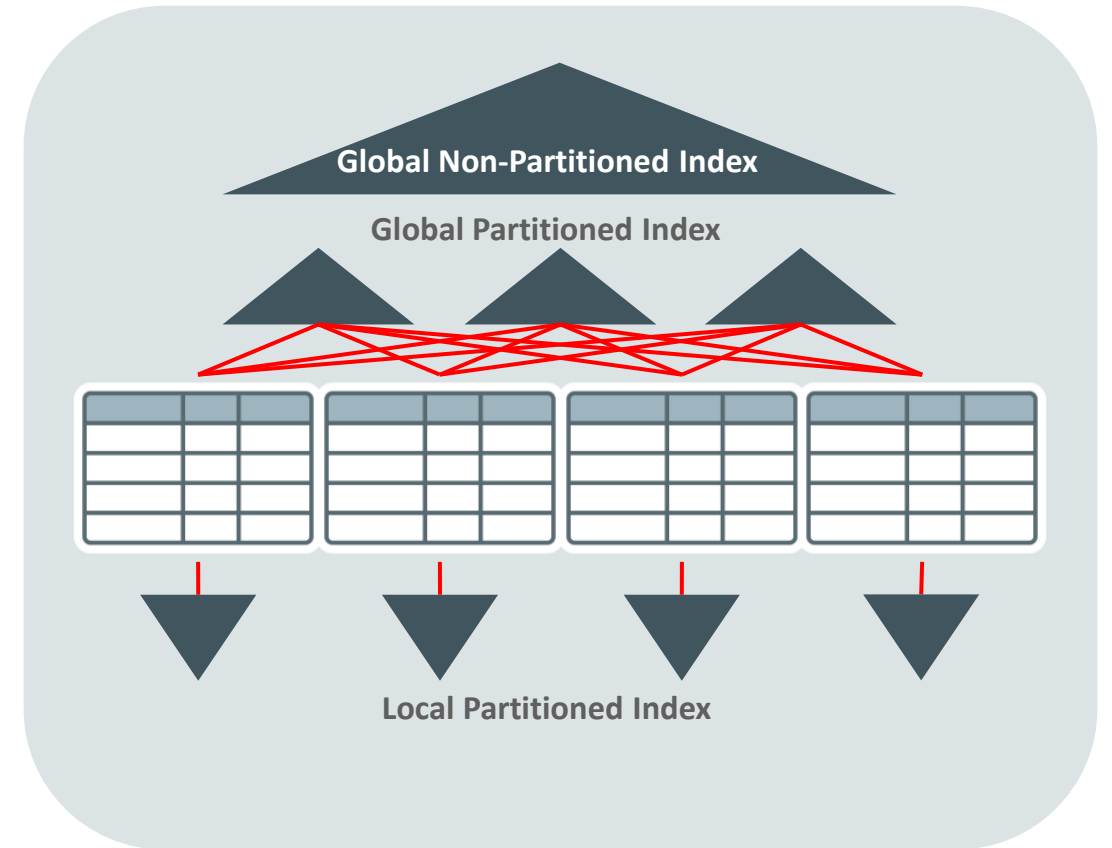
Indexing for unique constraints and primary keys

Unique Constraints/Primary Keys

- Unique constraints are enforced with unique indexes
 - Primary key constraint adds NOT NULL to column
 - Table can have only one primary key (“unique identifier”)
- Partitioned tables offer two types of indexes
 - Local indexes
 - Global index, both partitioned and non-partitioned
- Which one to pick?
 - Do I even have a choice?

Index Partitioning

- GLOBAL index points to rows in any partition
 - Index can be partitioned or not
 - Partition maintenance affects entire index
- LOCAL index is partitioned same as table
 - Index partitioning key can be different from index key
 - Index partitions can be maintained separately



Unique Constraints/Primary Keys

Applicability of Local Indexes

- Local indexes are equi-partitioned with the table
 - Follow autonomy concept of a table partition
 - “I only care about myself”
- Requirement for local indexes to enforce uniqueness
 - Partition key column(s) to be a subset of the unique key

Unique Constraints/Primary Keys, cont.

Applicability of Local Indexes

- Local indexes are equi-partitioned with the table
 - Follow autonomy concept of a table partition
 - “I only care about myself”
- Requirement for local indexes to enforce uniqueness
 - Partition key column(s) to be a subset of the unique key



The diagram shows three tables, each with 4 columns and 4 rows. A green checkmark is placed to the right of the third table, indicating that this configuration is valid for local indexes.

PARTITION BY (col1), PK(col1)



The diagram shows three tables, each with 4 columns and 4 rows. A red prohibition sign (a circle with a diagonal line) is placed to the right of the third table, indicating that this configuration is invalid for local indexes.

PARTITION BY (col1), PK(col2)

Unique Constraints/Primary Keys, cont.

Applicability of Global Indexes

- Global indexes do not have any relation to the partitions of a table
 - By definition, a global index contains data from all partitions
 - True for both partitioned and non-partitioned global indexes
- Global index can always be used to enforce uniqueness



PARTITION BY (col1), PK(col1)



PARTITION BY (col1), PK(col2)

Partial Indexing

Introduced in Oracle 12c Release 1 (12.1)

Enhanced Indexing with Oracle Partitioning

Indexing prior to Oracle Database 12c

- Local indexes
- Non-partitioned or partitioned global indexes
- Usable or unusable index segments
 - Non-persistent status of index, no relation to table

Enhanced Indexing with Oracle Partitioning

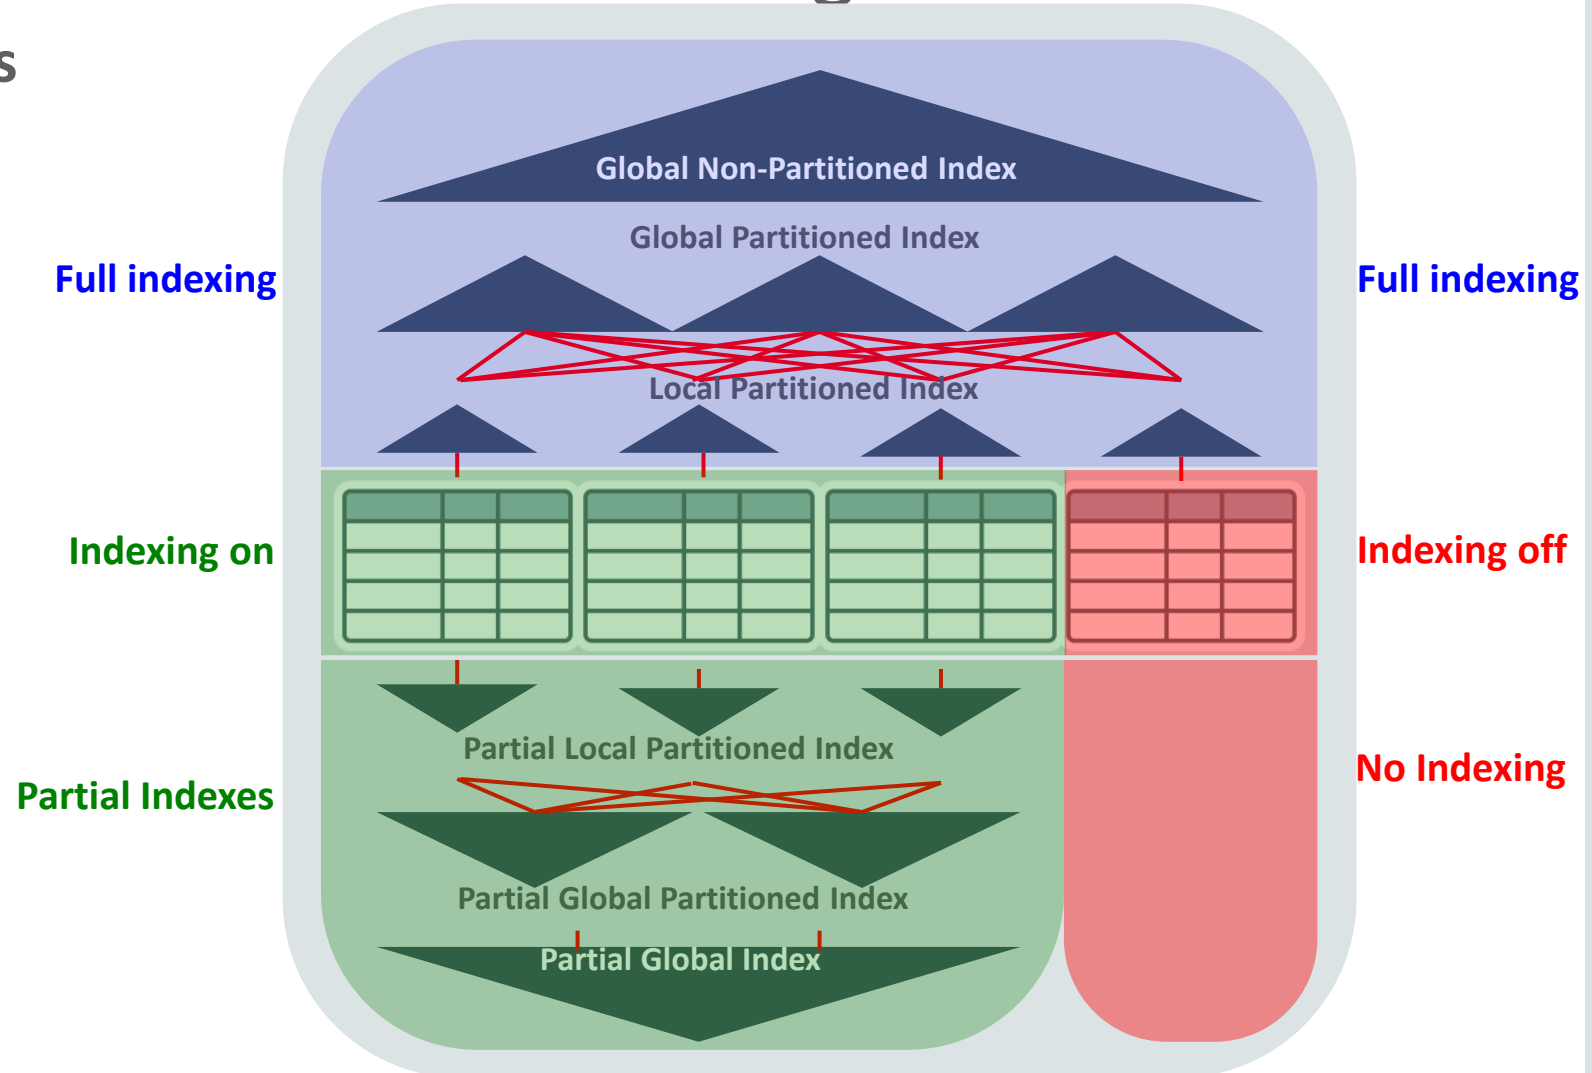
Indexing with Oracle Database 12c

- Local indexes
- Non-partitioned or partitioned global indexes
- Usable or unusable index segments
 - Non-persistent status of index, no relation to table
- Partial local and global indexes
 - Partial indexing introduces table and [sub]partition level metadata
 - Leverages usable/unusable state for local partitioned indexes
 - Policy for partial indexing can be overwritten

Enhanced Indexing with Oracle Partitioning

Partial Local and Global Indexes

- Partial indexes span only some partitions
- Applicable to local and global indexes
- Complementary to full indexing
- Enhanced business modeling



Enhanced Indexing with Oracle Partitioning

Partial Local and Global Indexes

Before

```
SQL> create table pt (col1, col2, col3, col4)
2 indexing off
3 partition by range (col1)
4 interval (1000)
5 (partition p100 values less than (101) indexing on,
6 partition p200 values less than (201) indexing on,
7 partition p300 values less than (301) indexing on);
```

Table created.

```
SQL> REM partitions and its indexing status
SQL> select partition_name, high_value, indexing
2 from user_tab_partitions where table_name='PT';
```

PARTITION_NAME	HIGH_VALUE	INDEXING
P100	101	ON
P200	201	ON
P300	301	ON
SYS_P1256	1301	OFF

After

```
SQL> REM local indexes
SQL> create index i_l_partpt on pt(col1) local indexing partial;
SQL> create index i_l_pt on pt(col4) local;

SQL> REM global indexes
SQL> create index i_g_partpt on pt(col2) indexing partial;
SQL> create index i_g_pt on pt(col3);
```

```
SQL> REM index status
SQL> select index_name, partition_name, status, null
2 from user_ind_partitions where index_name in ('I_L_PARTPT','I_L_PT')
3 union all
4 select index_name, indexing, status, orphaned_entries
5 from user_indexes where index_name in ('I_G_PARTPT','I_G_PT');
```

INDEX_NAME	PARTITION_NAME	STATUS	ORPHAN
I_L_PARTPT	P100	USABLE	
I_L_PARTPT	P200	USABLE	
I_L_PARTPT	P300	USABLE	
I_L_PARTPT	SYS_P1257	UNUSABLE	
I_L_PT	P200	USABLE	
I_L_PT	P300	USABLE	
I_L_PT	SYS_P1258	USABLE	
I_L_PT	P100	USABLE	
I_G_PT	FULL	VALID	NO
I_G_PARTPT	PARTIAL	VALID	NO

10 rows selected.



Enhanced Indexing with Oracle Partitioning

Partial Local and Global Indexes

- Partial global index excluding partition 4

```
SQL> explain plan for select count(*) from pt where col2 = 3;
```

Explained.

```
SQL> select * from table(dbms_xplan.display);
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT		1	22	54 (12)	00:00:01		
1	SORT AGGREGATE		1	22				
2	VIEW	VW_TE_2	2		54 (12)	00:00:01		
3	UNION-ALL							
* 4	TABLE ACCESS BY GLOBAL INDEX ROWID BATCHED	PT	1	26	2 (0)	00:00:01	ROWID	ROWID
* 5	INDEX RANGE SCAN	I_G_PARTPT	1		1 (0)	00:00:01		
6	PARTITION RANGE SINGLE		1	26	52 (12)	00:00:01	4	4
* 7	TABLE ACCESS FULL	PT	1	26	52 (12)	00:00:01	4	4

Predicate Information (identified by operation id):

```
4 - filter("PT"."COL1"<301)
5 - access("COL2"=3)
7 - filter("COL2"=3)
```


Unusable versus Partial Indexes

Unusable Indexes

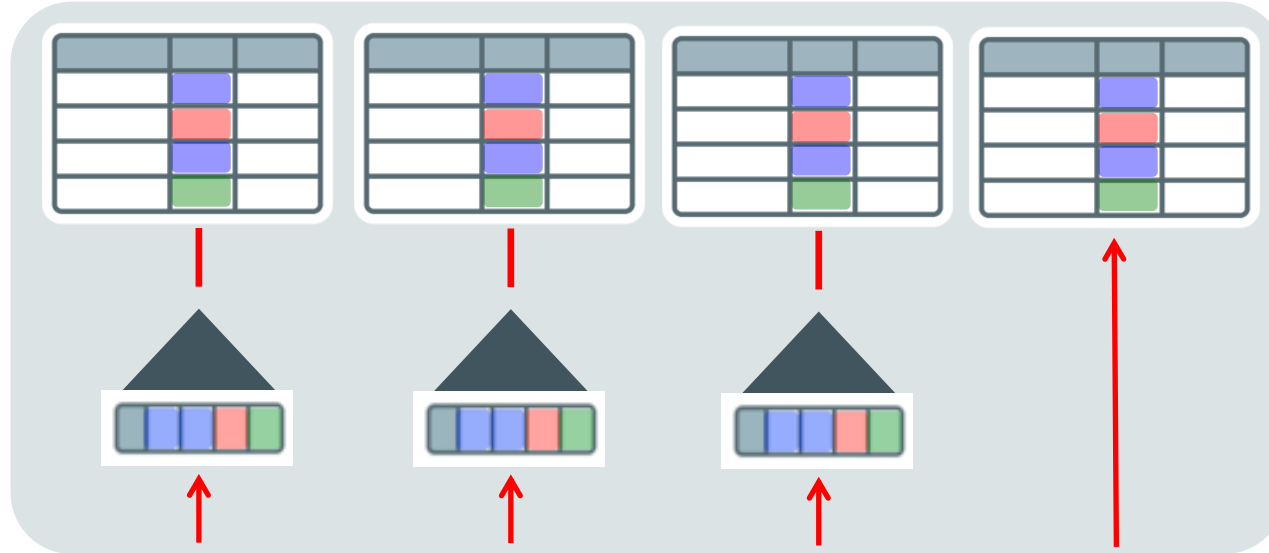
- Unusable index partitions are commonly used in environments with fast load requirements
 - “Save” the time for index maintenance at data insertion
 - Unusable index segments do not consume any space (11.2)
- Unusable indexes are ignored by the optimizer

```
SKIP_UNUSABLE_INDEXES = [TRUE | FALSE ]
```

- Partitioned indexes can be used by the optimizer even if some partitions are unusable
 - Prior to 11.2, static pruning and only access of usable index partitions mandatory
 - With 11.2, intelligent rewrite of queries using UNION ALL

Table-OR-Expansion

Multiple SQL branches are generated and executed



- Intelligent UNION ALL expansion in the presence of partially unusable indexes
 - Transparent internal rewrite
 - Usable index partitions will be used
 - Full partition access for unusable index partitions

Table-OR-Expansion

Sample Plan - Multiple SQL branches are generated and executed

```
select count(*) from toto where name = 'FOO' and rn between 1300 and 1400
```

Plan hash value: 2830852558

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT				27M(100)			
1	SORT AGGREGATE		1	21				
2	VIEW	VW_TE_2	2		27M (3)	92:15:22		
3	UNION-ALL							
4	PARTITION RANGE SINGLE		1	20	2 (0)	00:00:01	14	14
5	TABLE ACCESS BY LOCAL INDEX ROWID	TOTO	1	20	2 (0)	00:00:01	14	14
* 6	INDEX RANGE SCAN	I_TOTO	1		1 (0)	00:00:01	14	14
7	PARTITION RANGE SINGLE		1	22	27M (3)	92:15:22	15	15
* 8	TABLE ACCESS FULL	TOTO	1	22	27M (3)	92:15:22	15	15

Predicate Information (identified by operation id):

```
6 - access("NAME"='FOO')
8 - filter(("NAME"='FOO' AND "TOTO"."RN">=1400))
```

27 rows selected.

Partitioning for Performance

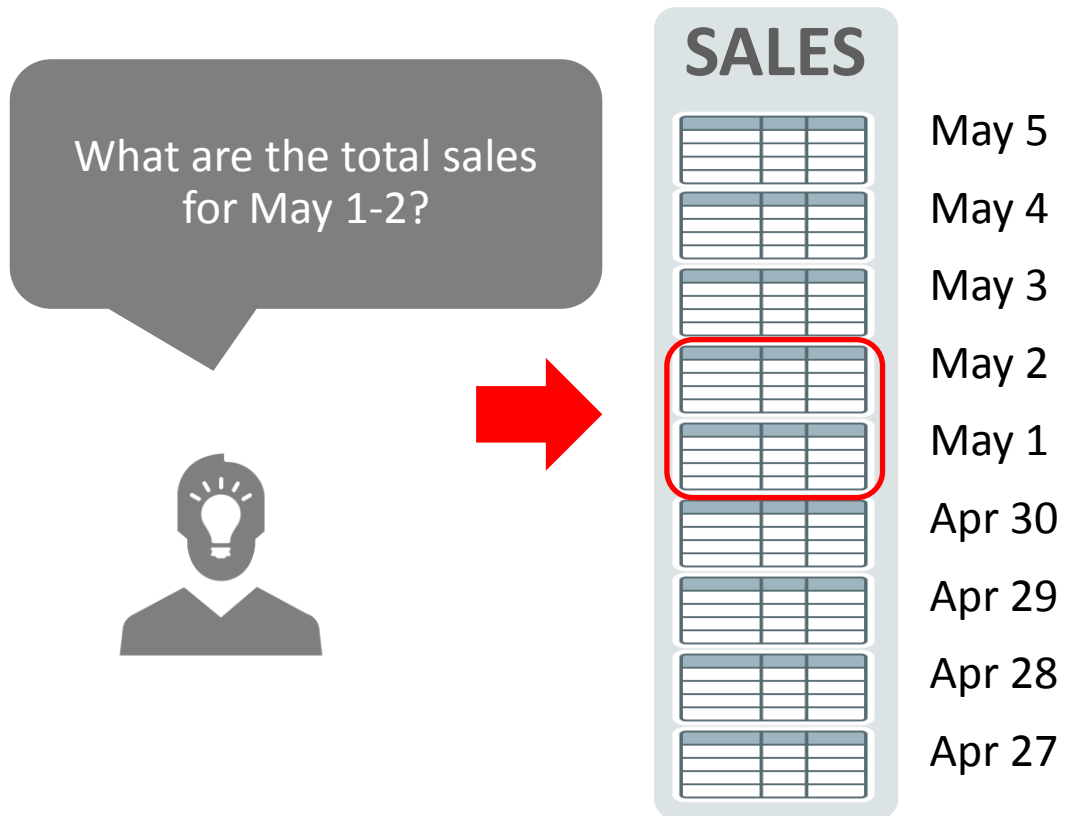


Partitioning for Performance

- Partitioning is transparently leveraged to improve performance
- Partition pruning
 - Using partitioning metadata to access only partitions of interest
- Partition-wise joins
 - Join equi-partitioned tables with minimal resource consumption
 - Process co-location capabilities for RAC environments
- Partition-Exchange loading
 - “Load” new data through metadata operation

Partitioning for Performance

Partition Pruning

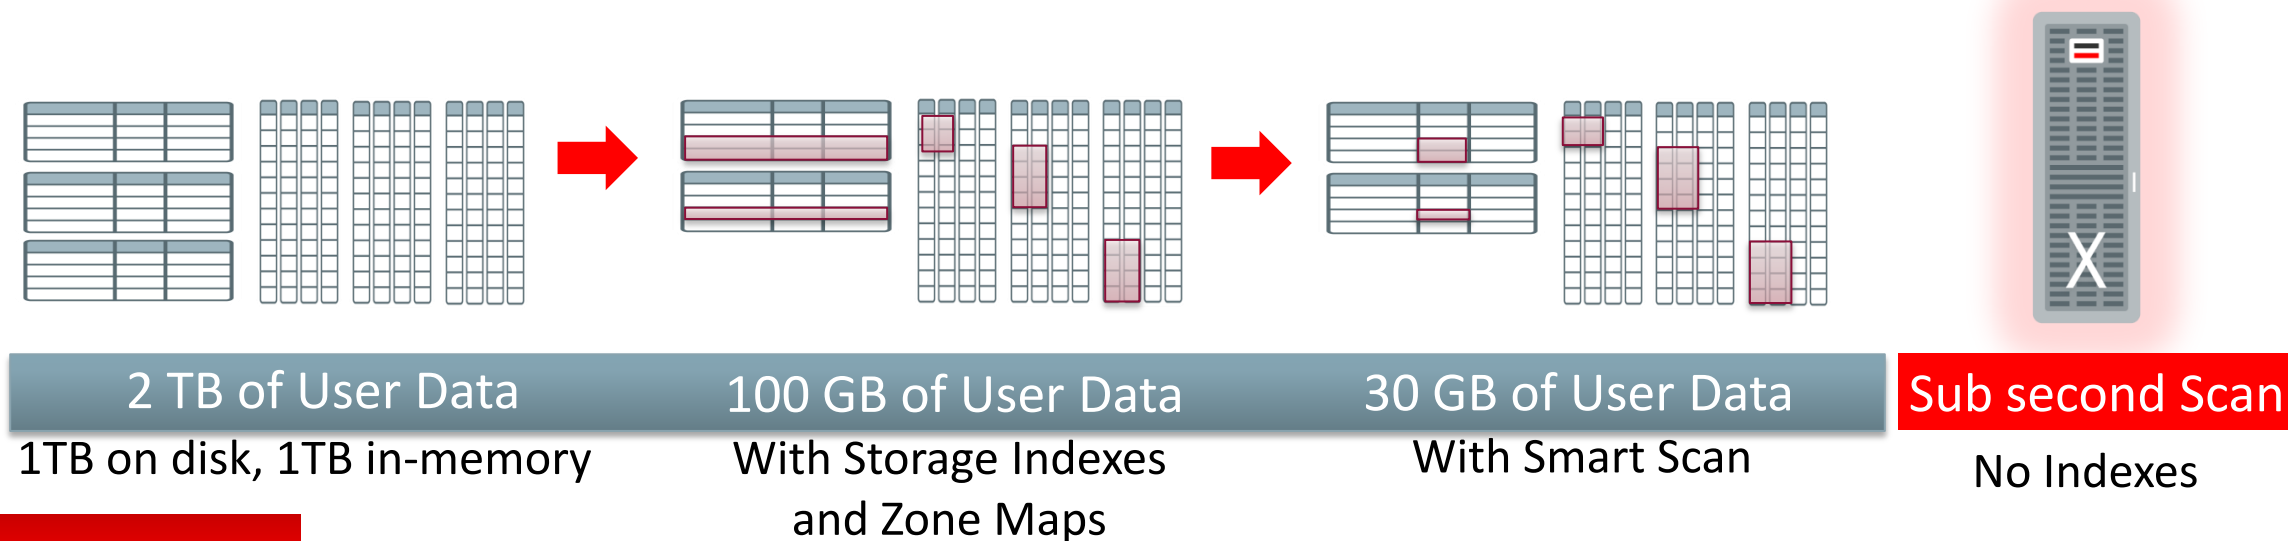
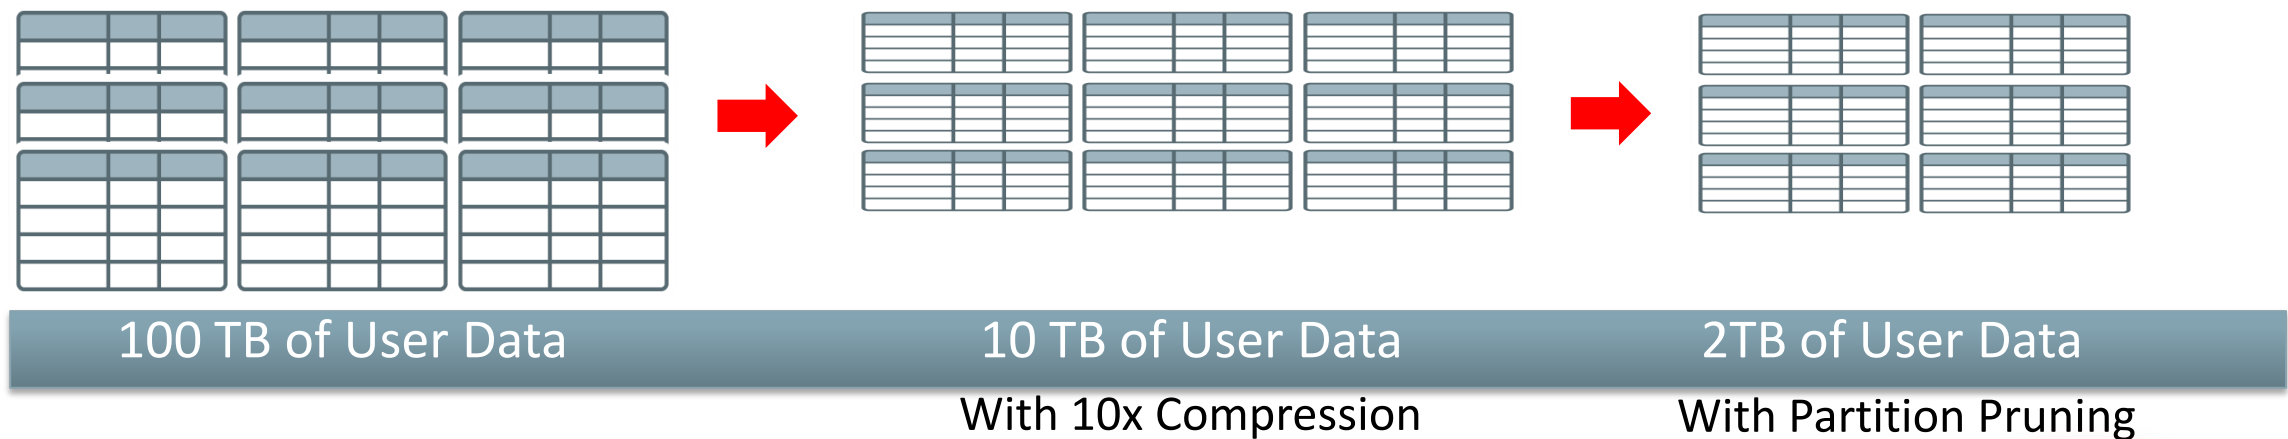


- Partition elimination
 - Dramatically reduces amount of data retrieved from storage
 - Performs operations only on relevant partitions
 - Transparently improves query performance and optimizes resource utilization

Partition Pruning

- Works for simple and complex SQL statements
- Transparent to any application
- Two flavors of pruning
 - Static pruning at compile time
 - Dynamic pruning at runtime
- Complementary to Exadata Storage Server
 - Partitioning prunes logically through partition elimination
 - Exadata prunes physically through storage indexes
 - Further data reduction through filtering and projection

Performance Features Multiply the Benefits



Static Partition Pruning

```
SELECT sum(amount_sold) FROM sales  
WHERE times_id  
BETWEEN '01-MAR-2016' and '31-MAY-2016';
```

14-JAN

14-FEB

14-MAR

14-APR

14-MAY

14-JUN

- Relevant Partitions are known at compile time
 - Look for actual values in PSTART/PSTOP columns in the plan
- Optimizer has most accurate information for the SQL statement

Static Pruning

Sample Plan

```
SELECT sum(amount_sold)
FROM sh.sales s, sh.times t
WHERE s.time_id = t.time_id
AND s.time_id between TO_DATE('01-JAN-2016', 'DD-MON-YYYY')
    and TO_DATE('01-JAN-2015', 'DD-MON-YYYY')
```

Plan hash value: 2025449199

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT				3 (100)			
1	SORT AGGREGATE		1	12				
2	PARTITION RANGE ITERATOR		313	3756	3 (0)	00:00:01	9	13
* 3	TABLE ACCESS FULL	SALES	313	3756	3 (0)	00:00:01	9	13

Predicate Information (identified by operation id):

```
3 - filter("S"."TIME_ID"<=TO_DATE(' 2015-01-01 00:00:00', 'syyy-mm-dd hh24:mi:ss'))
22 rows selected.
```



Static Pruning

Sample Plan

```
SELECT sum(amount_sold)
FROM sh.sales s, sh.times t
WHERE s.time_id = t.time_id
AND s.time_id between TO_DATE('01-JAN-2016', 'DD-MON-YYYY')
    and TO_DATE('01-JAN-2015', 'DD-MON-YYYY')
```

Plan hash value: 2025449199

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT				3 (100)			
1	SORT AGGREGATE		1	12				
2	PARTITION RANGE ITERATOR		313	3756	3 (0)	00:00:01	9	13
* 3	TABLE ACCESS FULL	SALES	313	3756	3 (0)	00:00:01	9	13

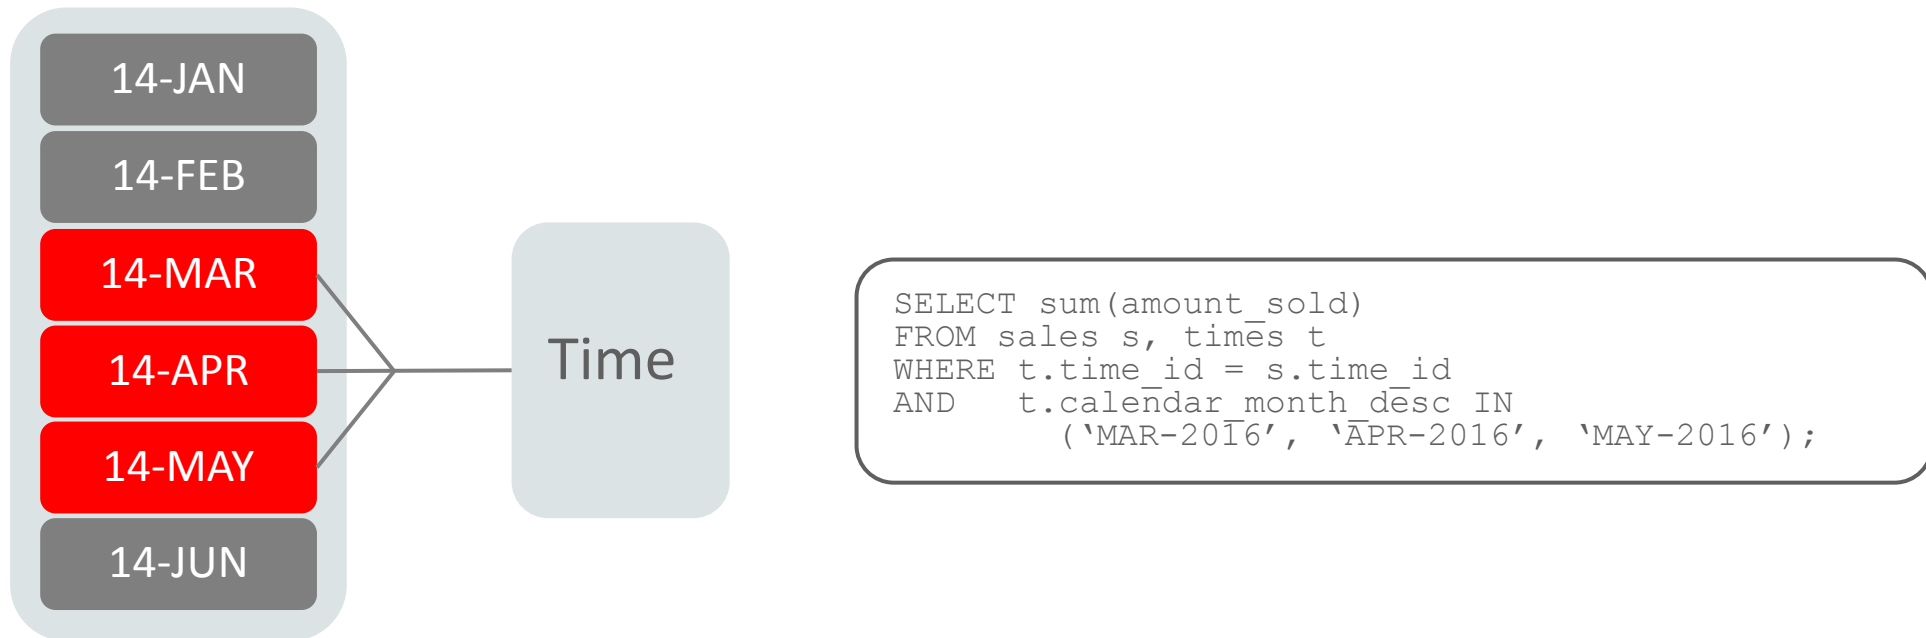
Predicate Information (identified by operation id):

```
3 - filter("S"."TIME_ID"<=TO_DATE(' 2015-01-01 00:00:00', 'syyyy-mm-dd hh24:mi:ss'))
```

22 rows selected.



Dynamic Partition Pruning



- Advanced Pruning mechanism for complex queries
- Relevant partitions determined at runtime
 - Look for the word 'KEY' in PSTART/PSTOP columns in the Plan

Dynamic Partition Pruning

Sample Plan – Nested Loop

```
SELECT sum(amount_sold)
FROM sales s, times t
WHERE s.time_id = t.time_id
AND t.calendar_month_desc in ('MAR-2016', 'APR-2016', 'MAY-2016')
```

Plan hash value: 1350851517

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT				13 (100)			
1	SORT AGGREGATE		1	28				
2	NESTED LOOP		2	56	13 (0)	00:00:01		
* 3	TABLE ACCESS FULL	TIMES	2	32	13 (8)	00:00:01		
4	PARTITION RANGE ITERATOR		2	24	0 (0)		KEY	KEY
* 5	TABLE ACCESS FULL	SALES	2	24	0 (0)		KEY	KEY

Predicate Information (identified by operation id):

- 3 - filter(("T"."CALENDAR_MONTH_DESC"='MAR-2016' OR "T"."CALENDAR_MONTH_DESC"='APR-2016' OR "T"."CALENDAR_MONTH_DESC"='MAY-2016'))
- 5 - filter("T"."TIME_ID"="S"."TIME_ID")

26 rows selected.

Dynamic Partition Pruning

Sample Plan – Nested Loop

```
SELECT sum(amount_sold)
FROM sales s, times t
WHERE s.time_id = t.time_id
AND t.calendar_month_desc in ('MAR-2016', 'APR-2016', 'MAY-2016')
```

Plan hash value: 1350851517

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT				13 (100)			
1	SORT AGGREGATE		1	28				
2	NESTED LOOP		2	56	13 (0)	00:00:01		
* 3	TABLE ACCESS FULL	TIMES	2	32	13 (8)	00:00:01		
4	PARTITION RANGE ITERATOR		2	24	0 (0)		KEY	KEY
* 5	TABLE ACCESS FULL	SALES	2	24	0 (0)		KEY	KEY

Predicate Information (identified by operation id):

```
3 - filter(("T"."CALENDAR_MONTH_DESC"='MAR-2016' OR "T"."CALENDAR_MONTH_DESC"='APR-2016'
          OR "T"."CALENDAR_MONTH_DESC"='MAY-2016'))
5 - filter("T"."TIME_ID"="S"."TIME_ID")
```

26 rows selected.

Dynamic Partition Pruning

Sample Plan - Subquery pruning

```
SELECT sum(amount_sold)
FROM sales s, times t
WHERE s.time_id = t.time_id
AND t.calendar_month_desc in ('MAR-2016', 'APR-2016', 'MAY-2016')
```

Plan hash value: 2475767165

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT				2000K(100)			
1	SORT AGGREGATE		1	28				
* 2	HASH JOIN		24M	646M	2000K(100)	06:40:01		
* 3	TABLE ACCESS FULL	TIMES	2	32	43 (8)	00:00:01		
4	PARTITION RANGE SUBQUERY		10G	111G	1166K(100)	03:53:21	KEY(SQ)	KEY(SQ)
5	TABLE ACCESS FULL	SALES	10G	111G	1166K(100)	03:53:21	KEY(SQ)	KEY(SQ)

Predicate Information (identified by operation id):

- 2 - access("S"."TIME_ID"="T"."TIME_ID")
- 3 - filter(("T"."CALENDAR_MONTH_DESC"='MAR-2016' OR "T"."CALENDAR_MONTH_DESC"='APR-2016' OR "T"."CALENDAR_MONTH_DESC"='MAY-2016'))

26 rows selected.



Dynamic Partition Pruning

Sample Plan - Bloom filter pruning

```
SELECT sum(amount_sold)
FROM sales s, times t
WHERE s.time_id = t.time_id
AND t.calendar_month_desc in ('MAR-2016', 'APR-2016', 'MAY-2016')
```

Plan hash value: 365741303

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT				19 (100)			
1	SORT AGGREGATE		1	28				
* 2	HASH JOIN		2	56	19 (100)	00:00:01		
3	PART JOIN FILTER CREATE	:BF0000	2	32	13 (8)	00:00:01		
* 4	TABLE ACCESS FULL	TIMES	2	32	13 (8)	00:00:01		
5	PARTITION RANGE JOIN-FILTER		960	11520	5 (0)	00:00:01	:BF0000	:BF0000
6	TABLE ACCESS FULL	SALES	960	11520	5 (0)	00:00:01	:BF0000	:BF0000

Predicate Information (identified by operation id):

```
2 - access("S"."TIME_ID"="T"."TIME_ID")
4 - filter(("T"."CALENDAR_MONTH_DESC"='MAR-2016' OR "T"."CALENDAR_MONTH_DESC"='APR-2016'
        OR "T"."CALENDAR_MONTH_DESC"='MAY-2016'))
```

27 rows selected.



“AND” Pruning

```
FROM sales s, times t ...  
WHERE s.time_id = t.time_id ..  
AND t.fiscal_year in (2016,2015)  
AND s.time_id  
    between TO_DATE('01-JAN-2016','DD-MON-YYYY')  
    and TO_DATE('01-JAN-2017','DD-MON-YYYY')
```

Dynamic pruning

Static pruning

- All predicates on partition key will be used for pruning
 - Dynamic and static predicates will now be used combined
- Example:
 - Star transformation with pruning predicate on both the FACT table and a dimension

“AND” Pruning

Sample Plan

Plan hash value: 552669211

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT		1	24	17 (12)	00:00:01		
1	SORT AGGREGATE		1	24				
* 2	HASH JOIN		204	4896	17 (12)	00:00:01		
3	PART JOIN FILTER CREATE	:BF0000	185	2220	13 (8)	00:00:01		
* 4	TABLE ACCESS FULL	TIMES	185	2220	13 (8)	00:00:01		
5	PARTITION RANGE AND		313	3756	3 (0)	00:00:01	KEY (AP)	KEY (AP)
* 6	TABLE ACCESS FULL	SALES	313	3756	3 (0)	00:00:01	KEY (AP)	KEY (AP)

Predicate Information (identified by operation id):

- 2 - access("S"."TIME_ID"="T"."TIME_ID")
- 4 - filter("T"."TIME_ID"<=TO_DATE(' 2015-01-01 00:00:00', 'syyy-mm-dd hh24:mi:ss') AND ("T"."FISCAL_YEAR"=2016 OR "T"."FISCAL_YEAR"=2015) AND "T"."TIME_ID">=TO_DATE(' 2016-01-01 00:00:00', 'syyy-mm-dd hh24:mi:ss'))
- 6 - filter("S"."TIME_ID"<=TO_DATE(' 2015-01-01 00:00:00', 'syyy-mm-dd hh24:mi:ss'))

22 rows selected.



Ensuring Partition Pruning

Don't use functions on partition key filter predicates

```
SELECT sum(amount_sold)
FROM sh.sales s, sh.times t
WHERE s.time_id = t.time_id
AND TO_CHAR(s.time_id, 'YYYYMMDD') between '20160101' and '20170101'
```

Plan hash value: 672559287

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT				6 (100)			
1	SORT AGGREGATE		1	12				
2	PARTITION RANGE ALL		2	24	6 (17)	00:00:01	1	16
* 3	TABLE ACCESS FULL	SALES	2	24	6 (17)	00:00:01	1	16

Predicate Information (identified by operation id):

```
3 - filter((TO_CHAR(INTERNAL_FUNCTION("S"."TIME_ID"), 'YYYYMMDD') >= '20160101' AND
            TO_CHAR(INTERNAL_FUNCTION("S"."TIME_ID"), 'YYYYMMDD') <= '20170101'))
```

23 rows selected.

Ensuring Partition Pruning

Don't use functions on partition key filter predicates

```
SELECT sum(amount_sold)
FROM sh.sales s, sh.times t
WHERE s.time_id = t.time_id
AND TO_CHAR(s.time_id, 'YYYYMMDD') between '20160101' and '20170101'

Plan hash value: 672559287
```

```
SELECT sum(amount_sold)
FROM sh.sales s, sh.times t
WHERE s.time_id = t.time_id
AND s.time_id between TO_DATE('20160101','YYYYMMDD') and TO_DATE('20170101','YYYYMMDD')

Plan hash value: 2025449199
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	Pstart	Pstop
0	SELECT STATEMENT				3 (100)			
1	SORT AGGREGATE		1	12				
2	PARTITION RANGE ITERATOR		313	3756	3 (0)	00:00:01	9	13
* 3	TABLE ACCESS FULL	SALES	313	3756	3 (0)	00:00:01	9	13

```
Predicate Information (identified by operation id):

3 - filter("S"."TIME_ID"<=TO_DATE(' 2015-01-01 00:00:00', 'syyyy-mm-dd hh24:mi:ss'))

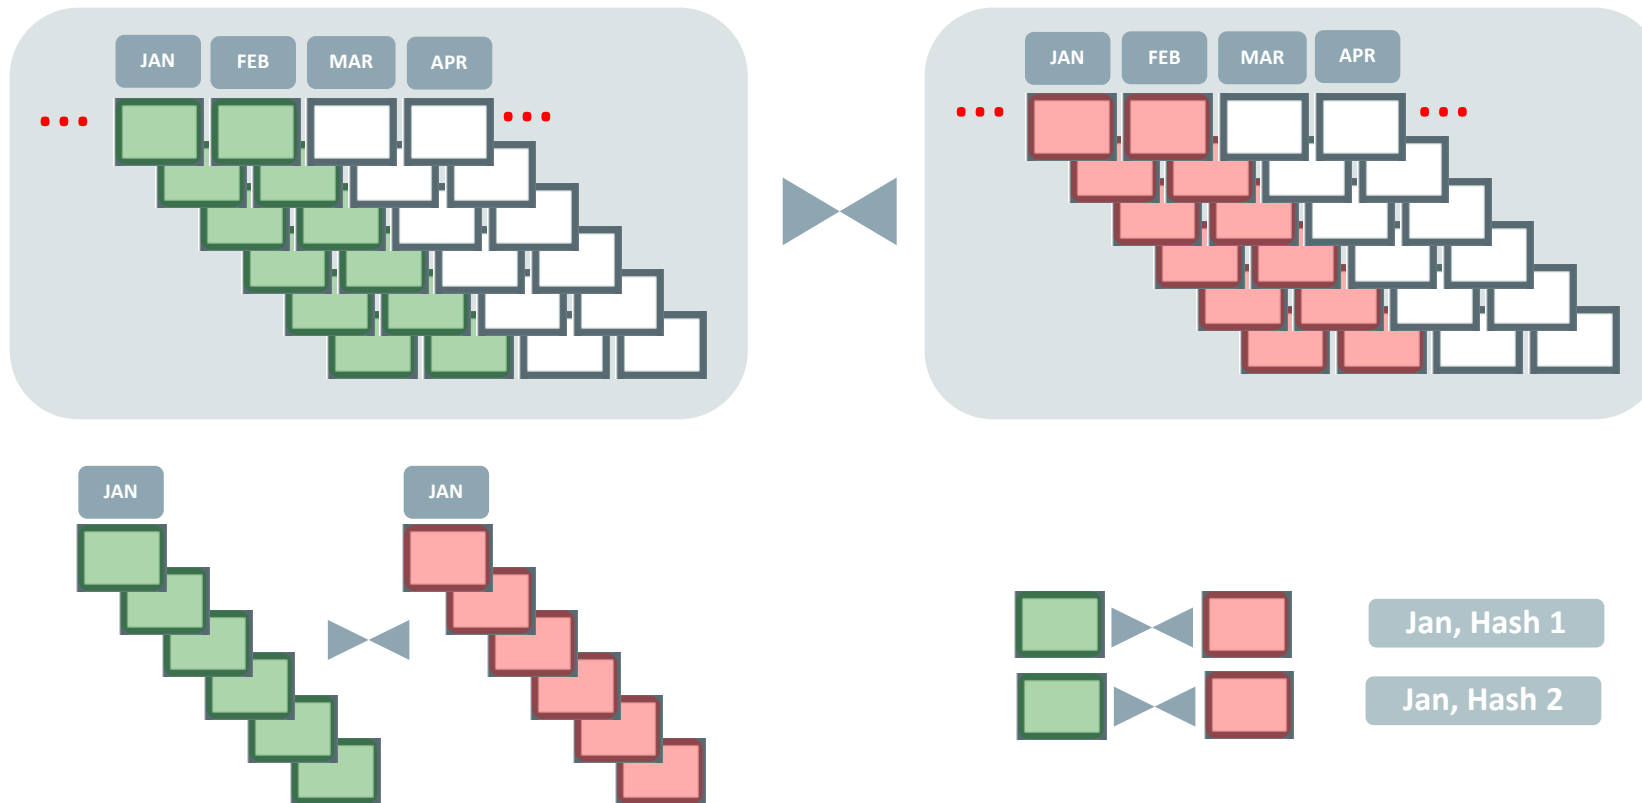
22 rows selected.
```

	Pstart	Pstop
1		16
1		16

101' AND
01'))

Partition-wise Joins

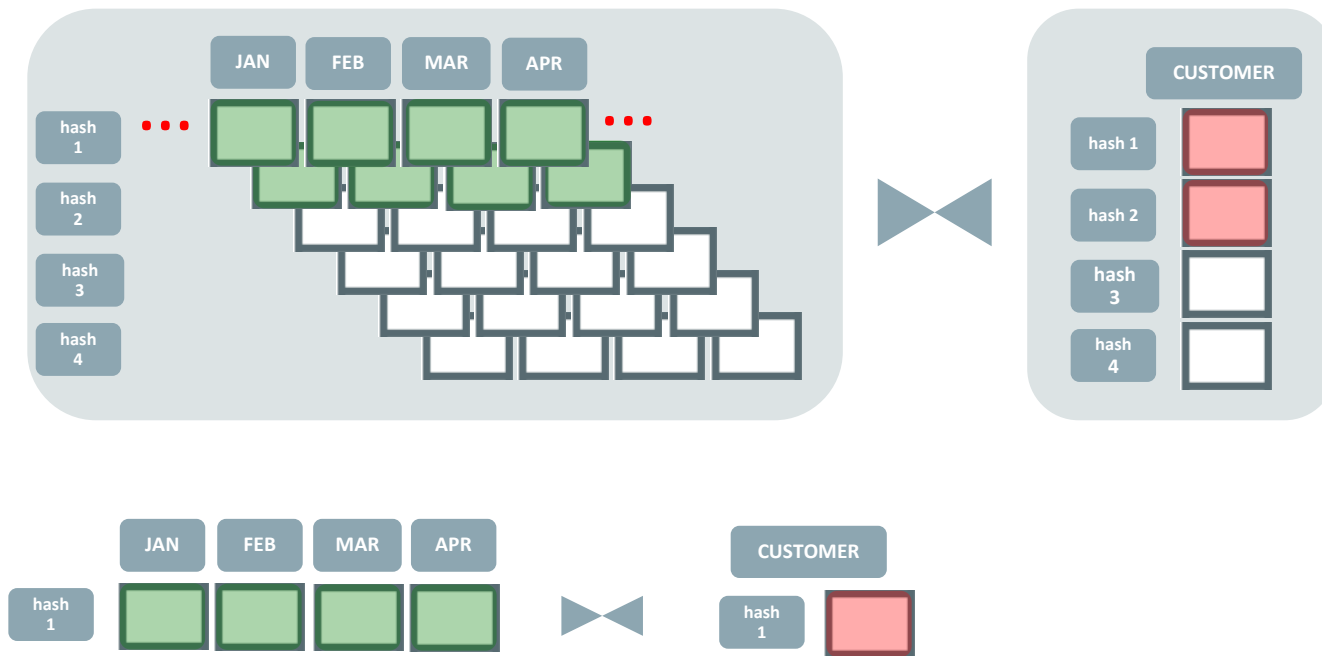
Partition pruning and PWJ's "at work"



- Large join is divided into multiple smaller joins, executed in parallel
 - # of partitions to join must be a multiple of DOP
 - Both tables must be partitioned the same way on the join column

Partition-wise Joins

Partition pruning and PWJ's “at work”



- Large join is divided into multiple smaller joins, executed in parallel
 - # of partitions to join must be a multiple of DOP
 - Both tables must be partitioned the same way on the join column

Partition Purging and Loading

- Remove and add data as metadata only operations
 - Exchange the metadata of partitions
- Exchange standalone table w/ arbitrary single partition
 - Data load: standalone table contains new data to being loaded
 - Data purge: partition containing data is exchanged with empty table
- Drop partition alternative for purge
 - Data is gone forever



Partitioning Maintenance



Partition Maintenance

Fundamental Concepts for Success

- While performance seems to be the most visible one, don't forget about the rest, e.g.
 - Partitioning must address all business-relevant areas of Performance, Manageability, and Availability
- Partition autonomy is crucial
 - Fundamental requirement for any partition maintenance operations
 - Acknowledge partitions as metadata in the data dictionary

Partition Maintenance

Fundamental Concepts for Success

- Provide full partition autonomy
 - Use local indexes whenever possible
 - Enable partition all table-level operations for partitions, e.g. TRUNCATE, MOVE, COMPRESS
- Make partitions visible and usable for database administration
 - Partition naming for ease of use
- Maintenance operations must be partition-aware
 - Also true for indexes
- Maintenance operations must not interfere with online usage of a partitioned table

Aspects of Data Management

Addressable with Partition Maintenance Operations

- Fast population of data
 - EXCHANGE
 - Per-partition direct path load
- Fast removal of data
 - DROP, TRUNCATE, EXCHANGE
- Fast reorganization of data
 - MOVE, SPLIT, MERGE

Partition Maintenance

Table Partition Maintenance Operations

```
ALTER TABLE ADD PARTITION(S)
ALTER TABLE DROP PARTITION(S)
ALTER TABLE EXCHANGE PARTITION
ALTER TABLE MODIFY PARTITION
    [PARALLEL] [ONLINE]
ALTER TABLE MOVE PARTITION [PARALLEL] [ONLINE]
ALTER TABLE RENAME PARTITION
ALTER TABLE SPLIT PARTITION [PARALLEL] [ONLINE]
ALTER TABLE MERGE PARTITION(S) [PARALLEL]
ALTER TABLE COALESCE PARTITION [PARALLEL]
ALTER TABLE ANALYZE PARTITION
ALTER TABLE TRUNCATE PARTITION(S)
Export/Import [by partition]
Transportable tablespace [by partition]
```

Index Maintenance Operations

```
ALTER INDEX MODIFY PARTITION
ALTER INDEX DROP PARTITION(S)
ALTER INDEX REBUILD PARTITION
ALTER INDEX RENAME PARTITION
ALTER INDEX RENAME
ALTER INDEX SPLIT PARTITION
ALTER INDEX ANALYZE PARTITION
```

All partitions remain available all the time

- DML Lock on impacted partitions
- Move partition online no lock at all



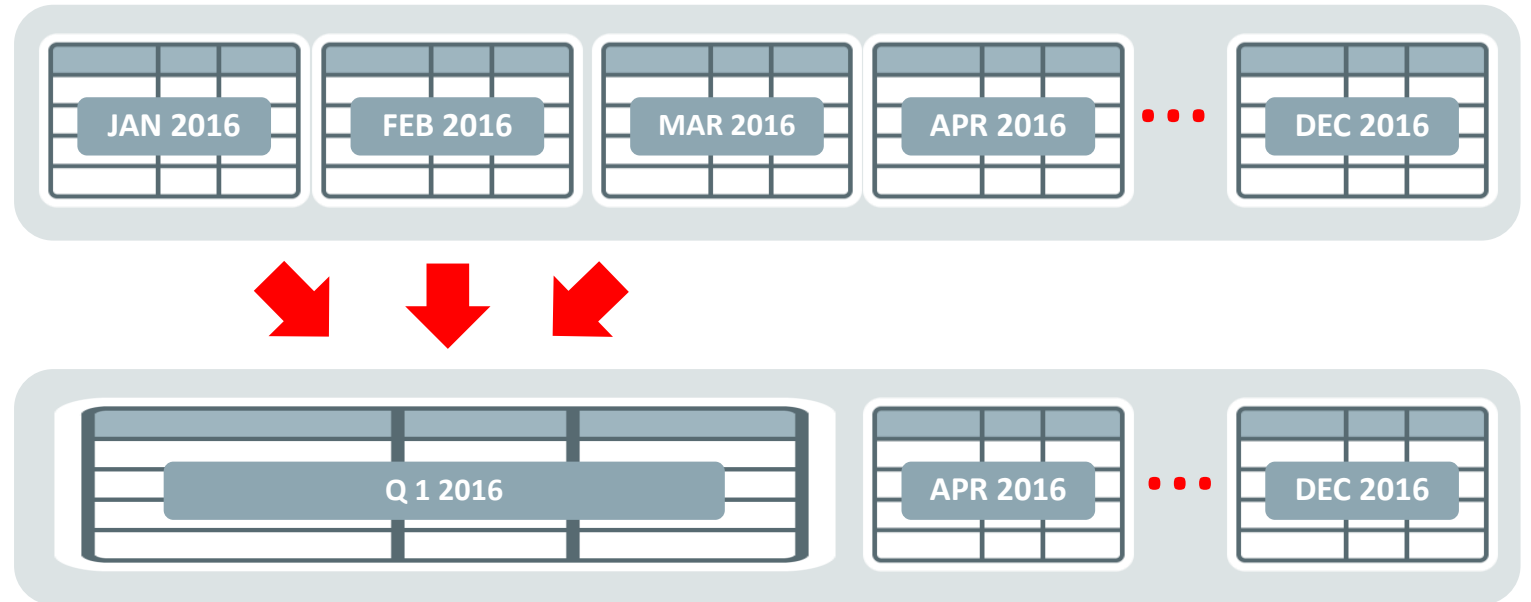
Partition Maintenance on Multiple Partitions

Introduced in Oracle 12c Release 1 (12.1)

Enhanced Partition Maintenance Operations

Operate on multiple partitions

- Partition Maintenance on multiple partitions in a single operation
- Full parallelism
- Transparent maintenance of local and global indexes



```
ALTER TABLE orders
MERGE PARTITIONS Jan2016, Feb2016, Mar2016
INTO PARTITION Q1_2016 COMPRESS FOR ARCHIVE HIGH;
```

Enhanced Partition Maintenance Operations

Operate on multiple partitions

- Specify multiple partitions in order

```
SQL > alter table pt merge partitions for (5), for (15), for (25) into  
      partition p30;
```

Table altered.

- Specify a range of partitions

```
SQL > alter table pt merge partitions part10 to part30 into partition part30;  
Table altered.
```

```
SQL > alter table pt split partition p30 into  
2   (partition p10 values less than (10),  
3   partition p20 values less than (20),  
4   partition p30);
```

Table altered.

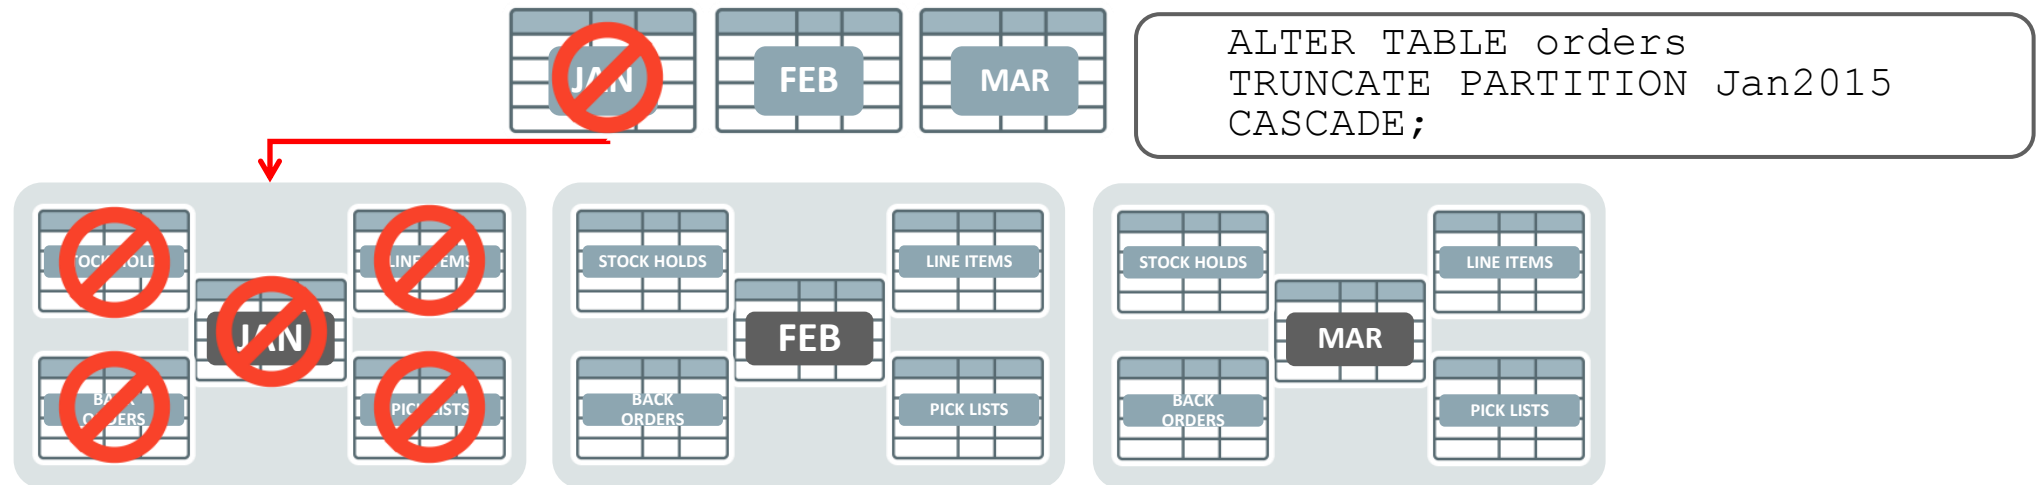
- Works for all PMOPS
 - Supports optimizations like fast split

Cascading Truncate and Exchange for Reference Partitioning

Introduced in Oracle 12c Release 1 (12.1)

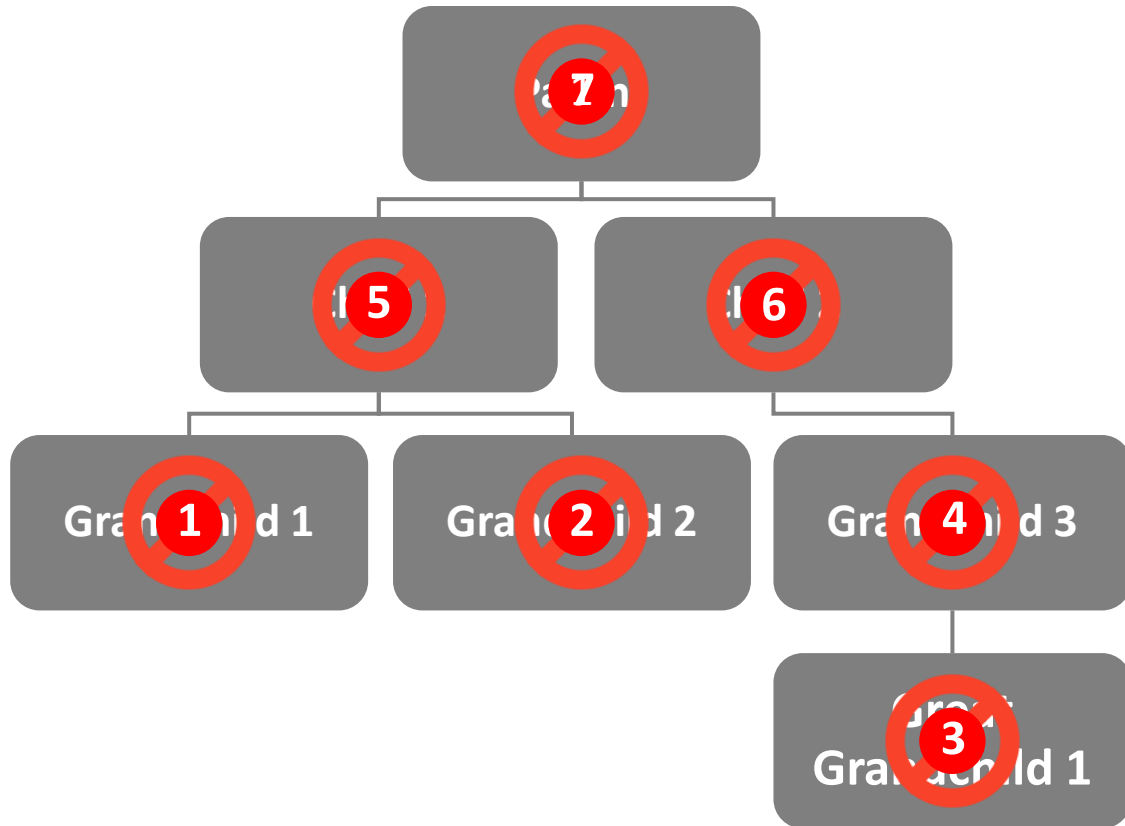
Advanced Partitioning Maintenance

Cascading TRUNCATE and EXCHANGE PARTITION

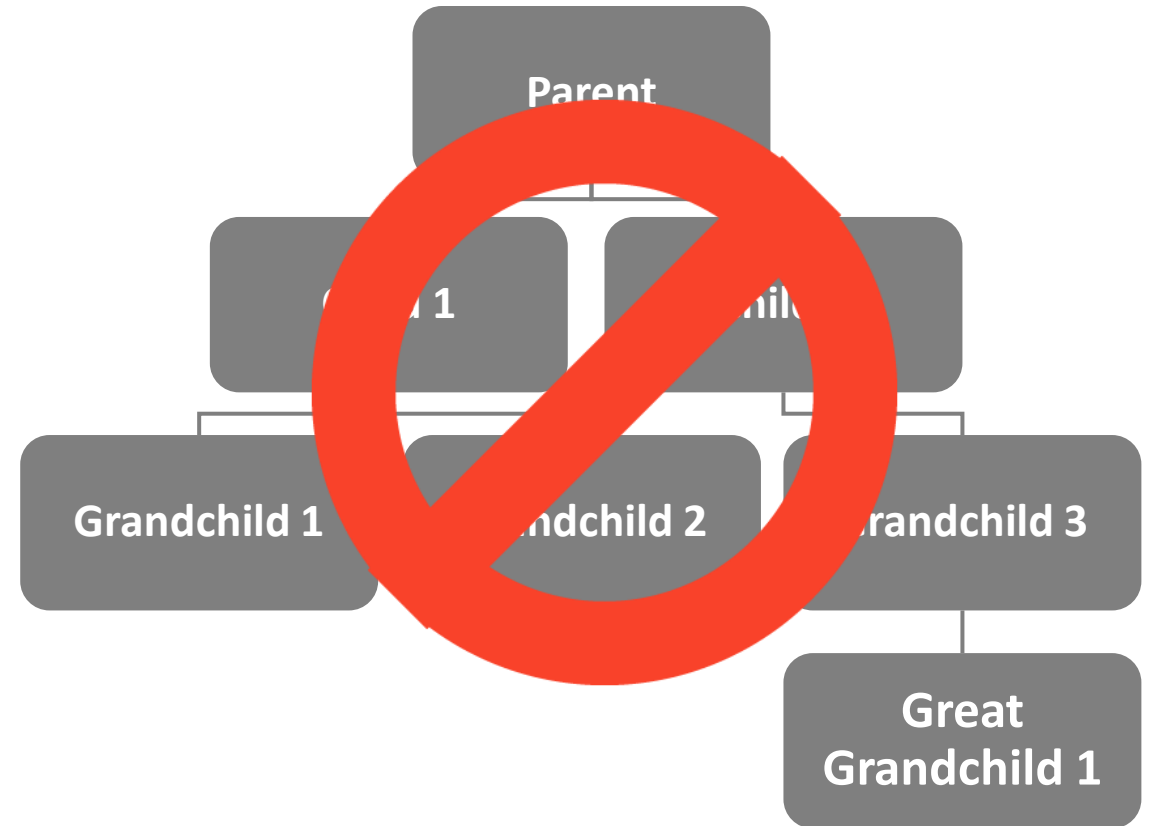


- Cascading TRUNCATE and EXCHANGE for improved business continuity
- Single atomic transaction preserves data integrity
- Simplified and less error prone code development

Cascading TRUNCATE PARTITION



- Proper bottom-up processing required
- Seven individual truncate operations



- One truncate operation

Cascading TRUNCATE PARTITION

```
SQL> create table intRef_p (pkcol number not null, col2 varchar2(200),  
2                          constraint pk_intref primary key (pkcol))  
3  partition by range (pkcol) interval (10)  
4  (partition p1 values less than (10));
```

Table created.

```
SQL>  
SQL> create table intRef_c1 (pkcol number not null, col2 varchar2(200), fkcol number not null,  
2                          constraint pk_c1 primary key (pkcol),  
3                          constraint fk_c1 foreign key (fkcol) references intRef_p(pkcol) ON DELETE CASCADE)  
4  partition by reference (fk_c1);
```

Table created.

Cascading TRUNCATE PARTITION

```
SQL> create table intRef_p (pkcol number(2)
                           constraint pkc1 primary key)
partition by range (pkcol) into
4 (partition p1 values less than
```

Table created.

```
SQL>
SQL> create table intRef_c1 (pkcol number(2)
                           constraint pkc1 primary key)
partition by reference (fk_c1);
```

Table created.

```
SQL> select * from intRef_p;
```

PKCOL	COL2
333	data for truncate - p
999	data for truncate - p

```
SQL> select * from intRef_c1;
```

PKCOL	COL2	FKCOL
1333	data for truncate - c1	333
1999	data for truncate - c1	999

```
SQL> alter table intRef_p truncate partition for (999) cascade update indexes;
```

Table truncated.

```
SQL> select * from intRef_p;
```

PKCOL	COL2
333	data for truncate - p

```
SQL> select * from intRef_c1;
```

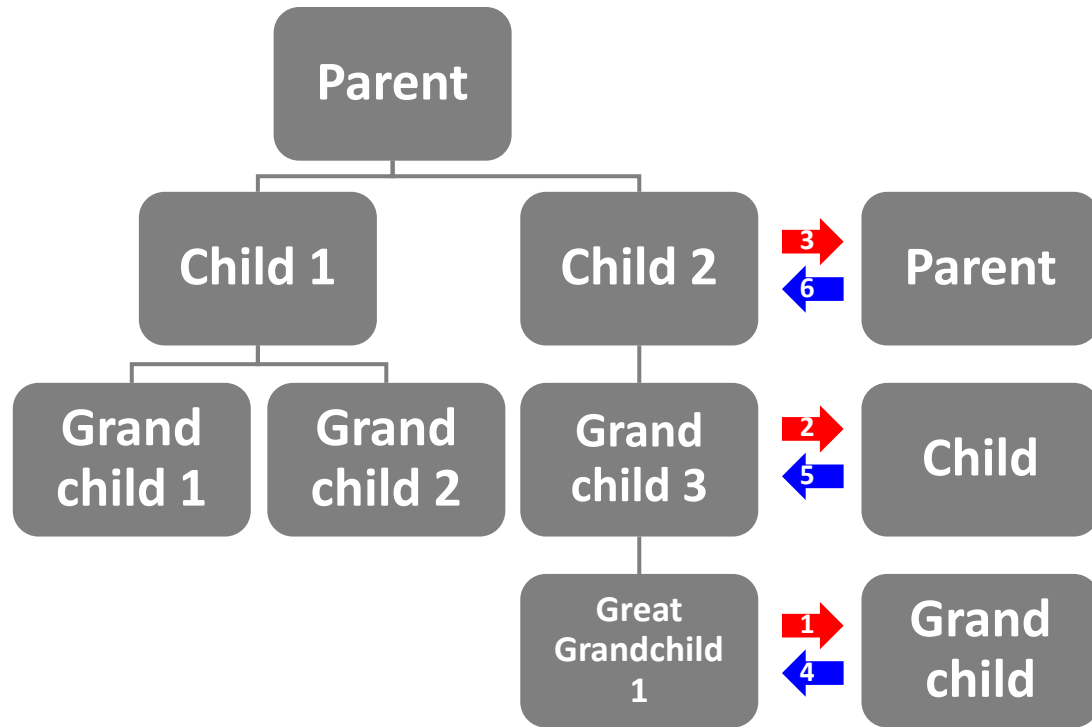
PKCOL	COL2	FKCOL
1333	data for truncate - c1	333

Cascading TRUNCATE PARTITION

- CASCADE applies for whole reference tree
 - Single atomic transaction, all or nothing
 - Bushy, deep, does not matter
 - Can be specified on any level of a reference-partitioned table
- ON DELETE CASCADE for all foreign keys required
- Cascading TRUNCATE available for non-partitioned tables as well
 - Dependency tree for non-partitioned tables can be interrupted with disabled foreign key constraints
- Reference-partitioned hierarchy must match for target and table to-be-exchanged
- For bushy trees with multiple children on the same level, each child on a given level must reference to a different key in the parent table
 - Required to unambiguously pair tables in the hierarchy tree

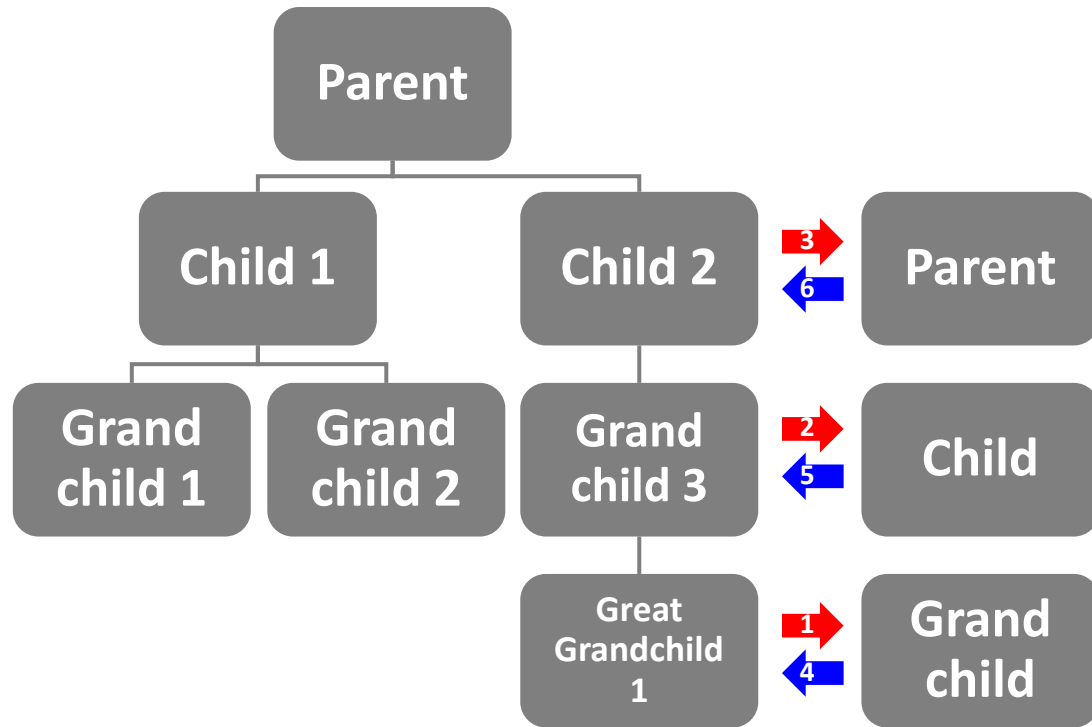


Cascading EXCHANGE PARTITION

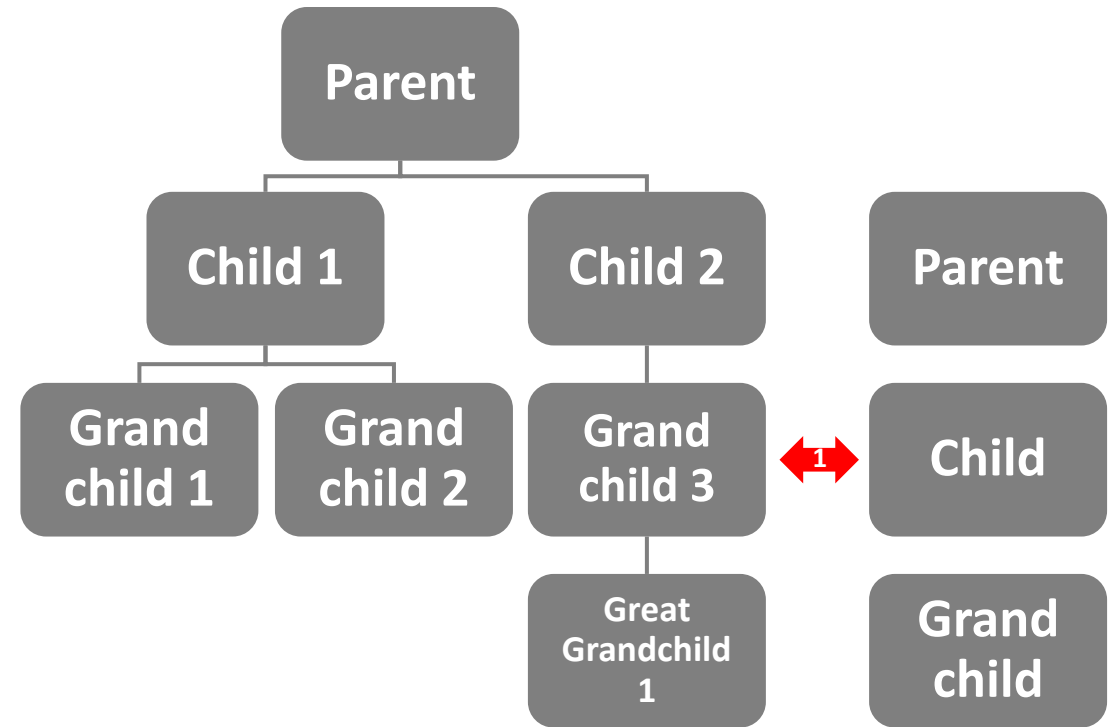


- Exchange (clear) out of target bottom-up
- Exchange (populate) into target top-down

Cascading EXCHANGE PARTITION



- Exchange (clear) out of target bottom-up
- Exchange (populate) into target top-down



- Exchange complete hierarchy tree
- One exchange operation

Cascading EXCHANGE PARTITION

```
SQL> create table intRef_p (pkcol number not null, col2 varchar2(200),
 2                          constraint pk_intref primary key (pkcol))
 3 partition by range (pkcol) interval (10)
 4 (partition p1 values less than (10));

SQL> create table intRef_c1 (pkcol number not null, col2 varchar2(200), fkcol number not null,
 2                          constraint pk_c1 primary key (pkcol),
 3                          constraint fk_c1 foreign key (fkcol) references intRef_p(pkcol) ON DELETE CASCADE)
 4 partition by reference (fk_c1);

SQL> create table intRef_gc1 (col1 number not null, col2 varchar2(200), fkcol number not null,
 2                          constraint fk_gc1 foreign key (fkcol) references intRef_c1(pkcol) ON DELETE CASCADE)
 3 partition by reference (fk_gc1);
```



Cascading EXCHANGE PARTITION

```
SQL> REM create some PK-FK equivalent table construct for exchange
SQL> create table XintRef_p (pkcol number not null, col2 varchar2(200),
  2                          constraint xpk_intref primary key (pkcol));

SQL> create table XintRef_c1 (pkcol number not null, col2 varchar2(200), fkcol number not null,
  2                          constraint xpk_c1 primary key (pkcol),
  3                          constraint xfk_c1 foreign key (fkcol) references XintRef_p(pkcol) ON DELETE CASCADE);

SQL> create table XintRef_gc1 (col1 number not null, col2 varchar2(200), fkcol number not null,
  2                          constraint xfk_gc1 foreign key (fkcol) references XintRef_c1(pkcol) ON DELETE CASCADE);
```

Cascading EXCHANGE PARTITION

```
SQL> select * from intRef_p;
```

PKCOL	COL2
333	p333 - data BEFORE exchange - p
999	p999 - data BEFORE exchange - p

```
SQL> select * from intRef_c1;
```

PKCOL	COL2	FKCOL
1333	p333 - data BEFORE exchange - c1	333
1999	p999 - data BEFORE exchange - c1	999

```
SQL> select * from intRef_gc1;
```

COL1	COL2	FKCOL
1333	p333 - data BEFORE exchange - gc1	1333
1999	p999 - data BEFORE exchange - gc1	1999

```
SQL> select * from XintRef_p;
```

PKCOL	COL2
333	p333 - data AFTER exchange - p

```
SQL> select * from XintRef_c1;
```

PKCOL	COL2	FKCOL
1333	p333 - data AFTER exchange - c1	333

```
SQL> select * from XintRef_gc1;
```

COL1	COL2	FKCOL
1333	p333 - data AFTER exchange - gc1	1333

Cascading EXCHANGE PARTITION

```
SQL> alter table intRef_p exchange partition for (333) with table XintRef_p cascade update indexes;  
Table altered.
```

Cascading EXCHANGE PARTITION

```
SQL> select * from intRef_p;
```

PKCOL	COL2
333	p333 - data AFTER exchange - p
999	p999 - data BEFORE exchange - p

```
SQL> select * from intRef_c1;
```

PKCOL	COL2	FKCOL
1333	p333 - data AFTER exchange - c1	333
1999	p999 - data BEFORE exchange - c1	999

```
SQL> select * from intRef_gc1;
```

COL1	COL2	FKCOL
1333	p333 - data AFTER exchange - gc1	1333
1999	p999 - data BEFORE exchange - gc1	1999

```
SQL> select * from XintRef_p;
```

PKCOL	COL2
333	p333 - data BEFORE exchange - p

```
SQL> select * from XintRef_c1;
```

PKCOL	COL2	FKCOL
1333	p333 - data BEFORE exchange - c1	333

```
SQL> select * from XintRef_gc1;
```

COL1	COL2	FKCOL
1333	p333 - data BEFORE exchange - gc1	1333

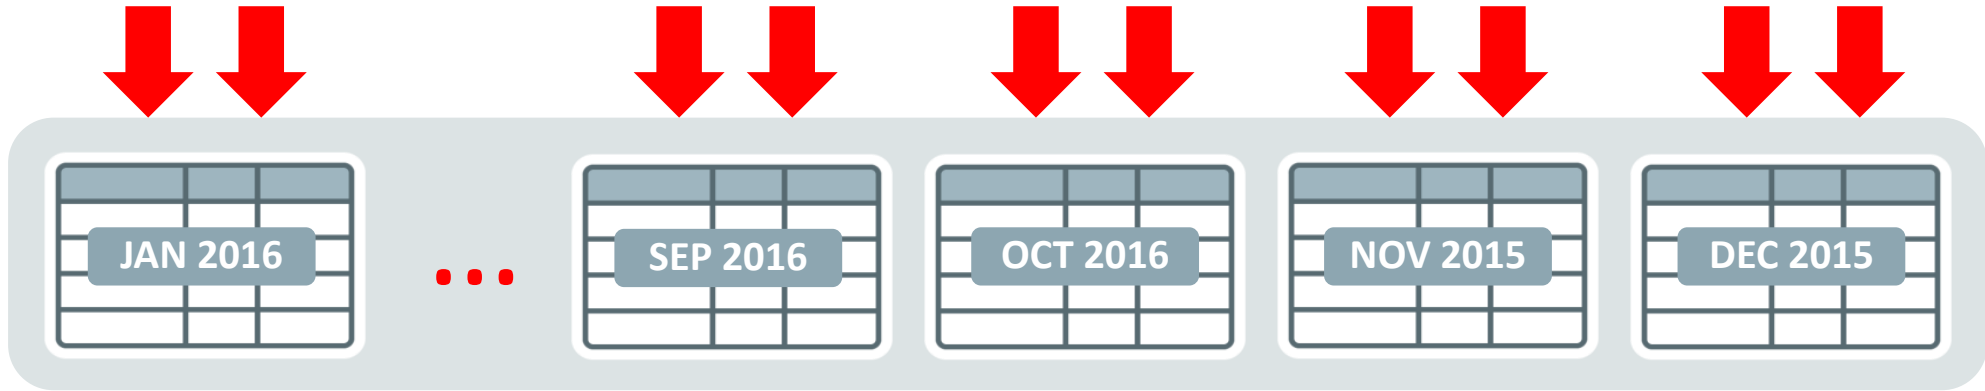


Online Move Partition

Introduced in Oracle 12c Release 1 (12.1)

Enhanced Partition Maintenance Operations

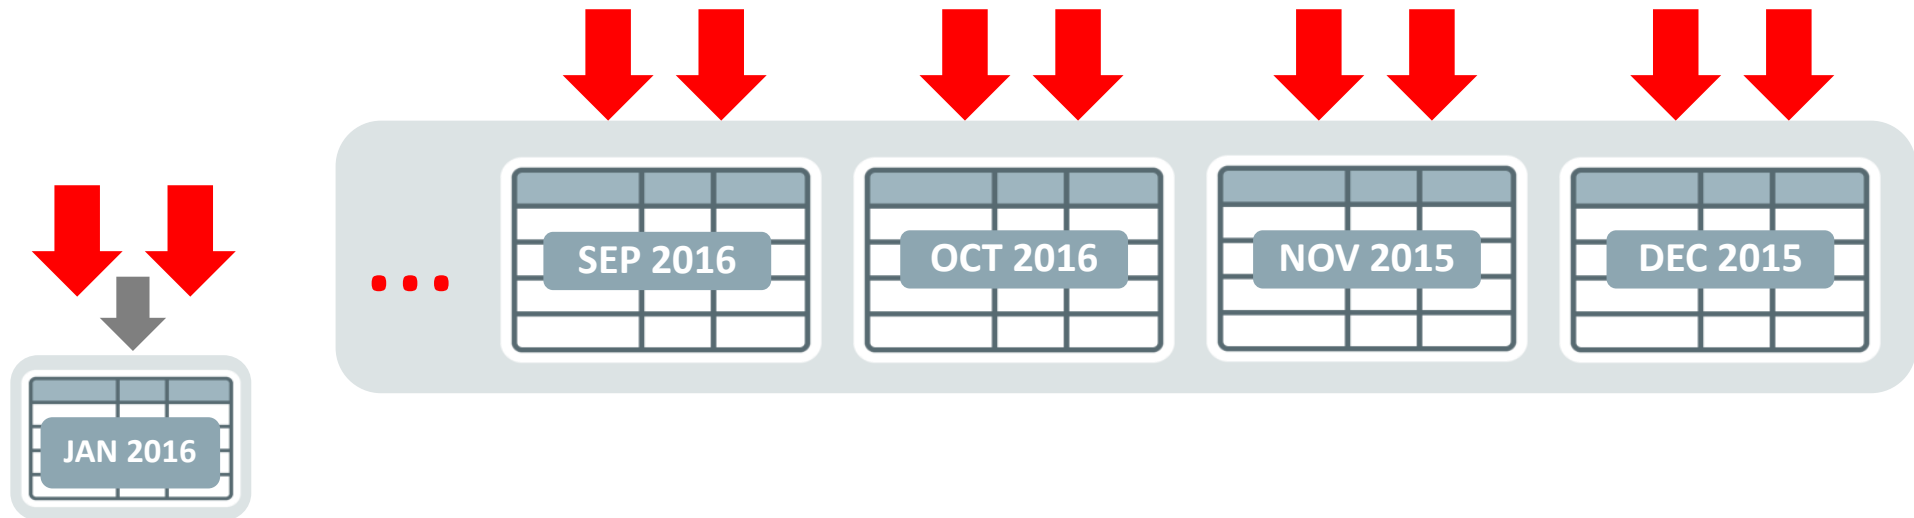
Online Partition Move



- Transparent MOVE PARTITION ONLINE operation
- Concurrent DML and Query
- Index maintenance for local and global indexes

Enhanced Partition Maintenance Operations

Online Partition Move



- Transparent MOVE PARTITION ONLINE operation
- Concurrent DML and Query
- Index maintenance for local and global indexes

Enhanced Partition Maintenance Operations

Online Partition Move – Best Practices

- Minimize concurrent DML operations if possible
 - Require additional disk space and resources for journaling
 - Journal will be applied recursively after initial bulk move
 - The larger the journal, the longer the runtime
- Concurrent DML has impact on compression efficiency
 - Best compression ratio with initial bulk move

Asynchronous Global Index Maintenance

Introduced in Oracle 12c Release 1 (12.1)

Enhanced Partition Maintenance Operations

Asynchronous Global Index Maintenance

- Usable global indexes after DROP and TRUNCATE PARTITION without index maintenance
 - Affected partitions are known internally and filtered out at data access time
- DROP and TRUNCATE become fast, metadata-only operations
 - Significant speedup and reduced initial resource consumption
- Delayed Global index maintenance
 - Deferred maintenance through ALTER INDEX REBUILD | COALESCE
 - Automatic cleanup using a scheduled job

Enhanced Partition Maintenance Operations

Asynchronous Global Index Maintenance

Before

```
SQL> select count(*) from pt partition for (9999);
```

COUNT(*)
25341440

Elapsed: 00:00:01.00

```
SQL> select index_name, status, orphaned_entries from user_indexes;
```

INDEX_NAME	STATUS	ORPHANED_ENTRIES
I1_PT	VALID	NO

Elapsed: 00:00:01.04

```
SQL>
```

```
SQL> alter table pt drop partition for (9999) update indexes;
```

Table altered.

Elapsed: 00:02:04.52

```
SQL>
```

```
SQL> select index_name, status, orphaned_entries from user_indexes;
```

INDEX_NAME	STATUS	ORPHANED_ENTRIES
I1_PT	VALID	NO

Elapsed: 00:00:00.10

After

```
SQL> select count(*) from pt partition for (9999);
```

COUNT(*)
25341440

Elapsed: 00:00:00.98

```
SQL> select index_name, status, orphaned_entries from user_indexes;
```

INDEX_NAME	STATUS	ORPHANED_ENTRIES
I1_PT	VALID	NO

Elapsed: 00:00:00.33

```
SQL>
```

```
SQL> alter table pt drop partition for (9999) update indexes;
```

Table altered.

Elapsed: 00:00:00.04

```
SQL>
```

```
SQL> select index_name, status, orphaned_entries from user_indexes;
```

INDEX_NAME	STATUS	ORPHANED_ENTRIES
I1_PT	VALID	YES

Elapsed: 00:00:00.05



Statistics Management for Partitioning

Statistics Gathering

- You must gather Optimizer statistics
 - Using dynamic sampling is not an adequate solution
 - Statistics on global and partition level recommended
 - Subpartition level optional
- Run all queries against empty tables to populate column usage
 - This helps identify which columns automatically get histograms created on them
- Optimizer statistics should be gathered after the data has been loaded but before any indexes are created
 - Oracle will automatically gather statistics for indexes as they are being created

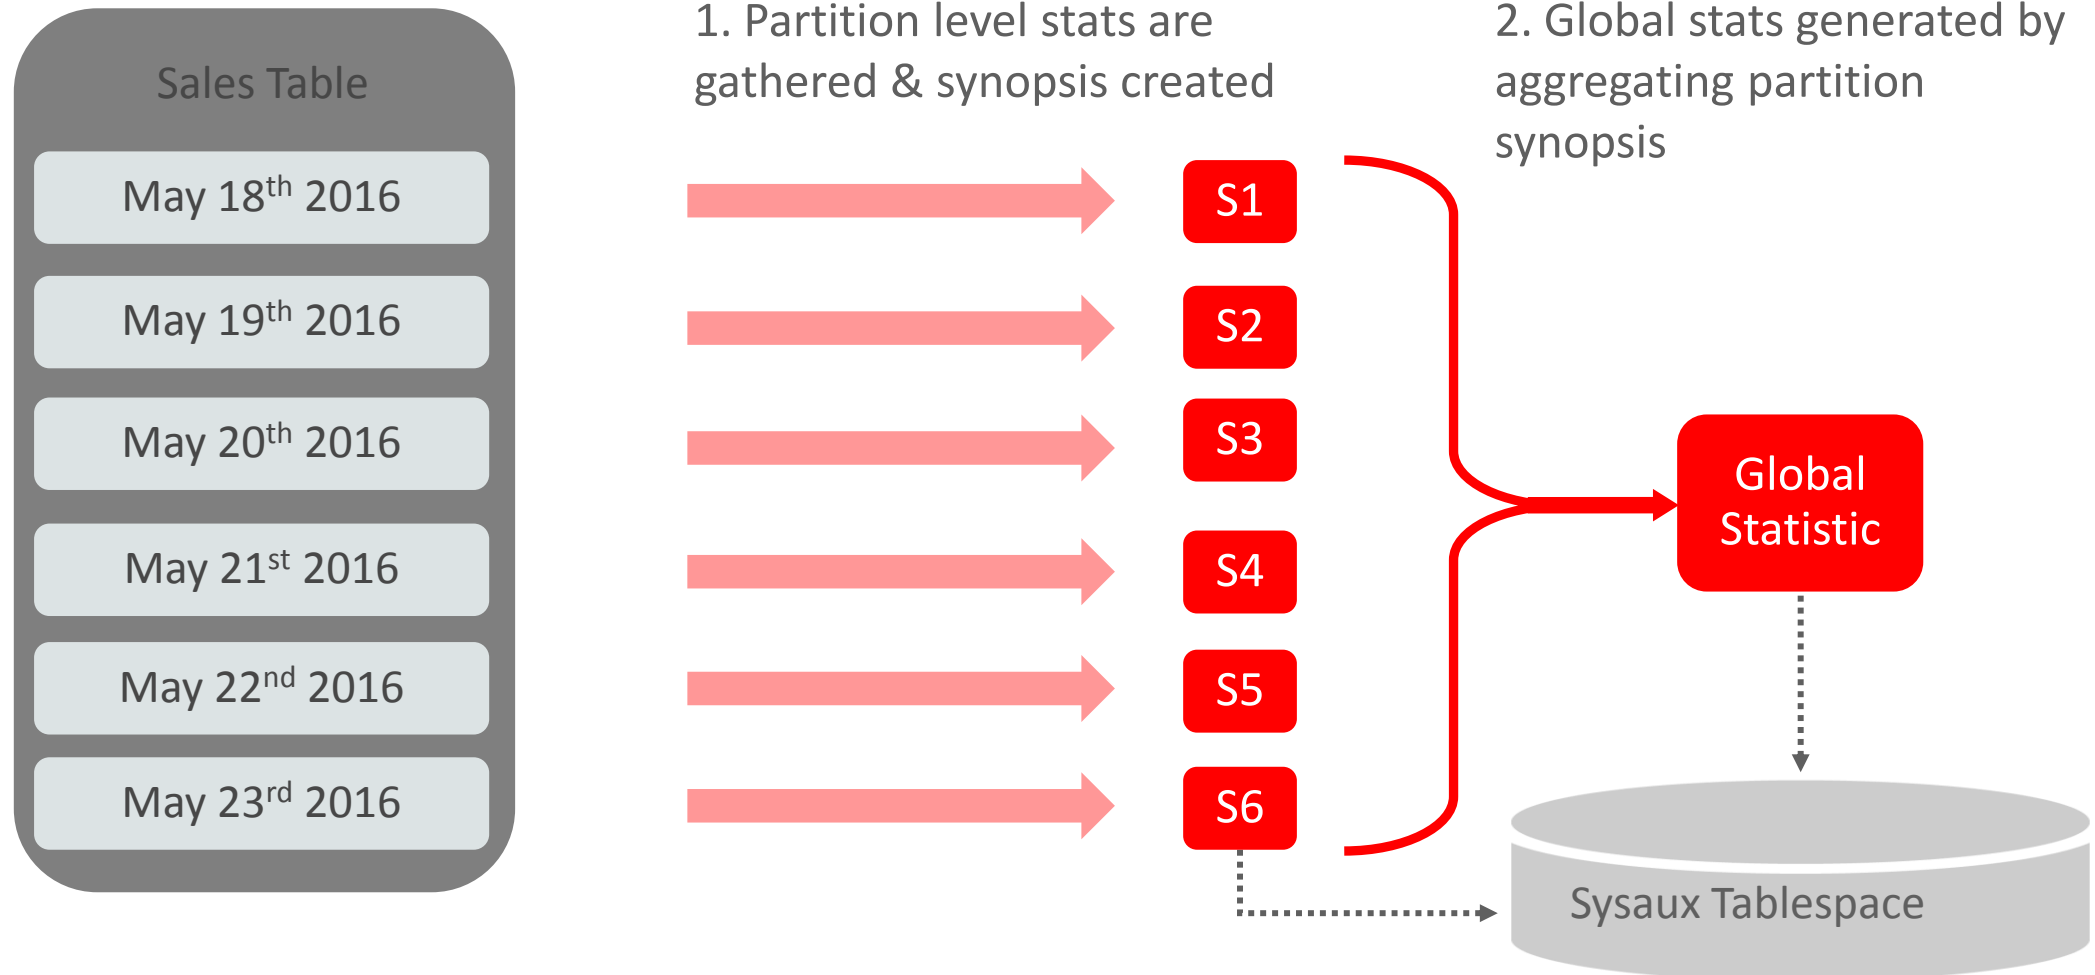
Statistics Gathering

- By default DBMS_STATS gathers the following stats for each table
 - global (table level), partition level, sub-partition level
- Optimizer uses global stats if query touches two or more partitions
- Optimizer uses partition stats if queries do partition elimination and only one partition is necessary to answer the query
 - If queries touch two or more partitions the optimizer will use a combination of global and partition level statistics
- Optimizer uses sub-partition level statistics only if your queries do partition elimination and one sub-partition is necessary to answer query

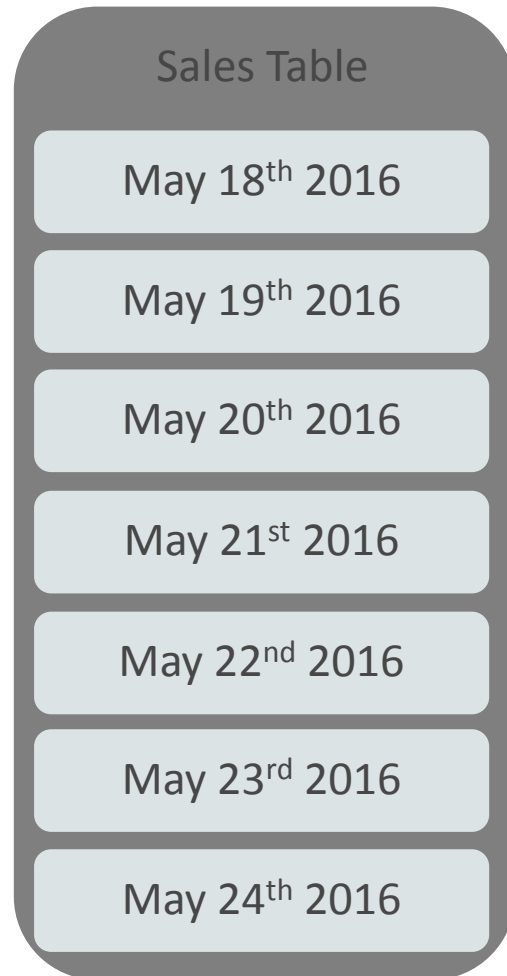
Efficient Statistics Management

- Use `AUTO_SAMPLE_SIZE`
 - The only setting that enables new efficient statistics collection
 - Hash based algorithm, scanning the whole table
 - Speed of sampling, accuracy of compute
- Enable incremental global statistics collection
 - Avoids scan of all partitions after changing single partitions
 - Prior to 11.1, scan of all partitions necessary for global stats
 - Managed on per table level
 - Static setting
 - Create synopsis for non-partitioned table to being exchanged (Oracle Database 12c)

Incremental Global Statistics



Incremental Global Statistics Cont'd

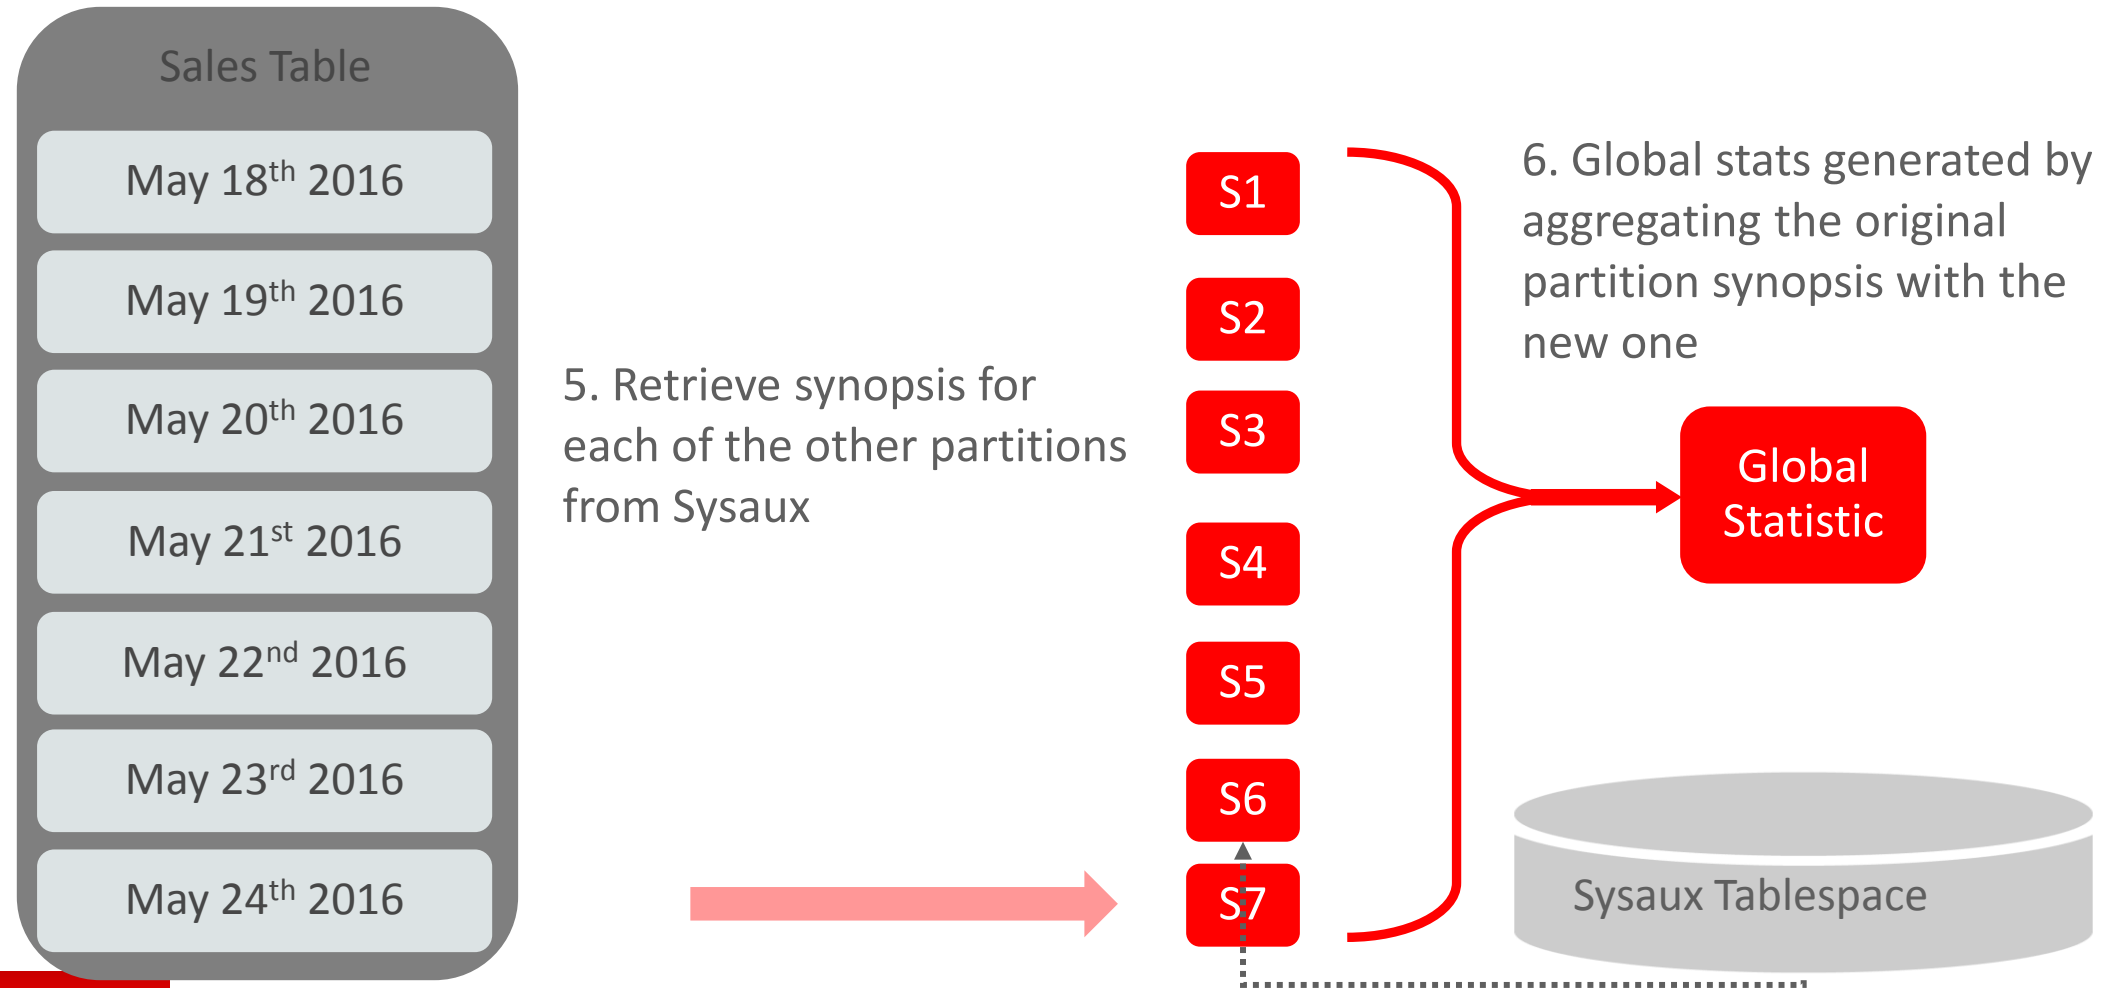


3. A new partition is added to the table and data is loaded

4. Gather partition statistics for new partition



Incremental Global Statistics Cont'd



Step necessary to gather accurate statistics

- Turn on incremental feature for the table

```
EXEC DBMS_STATS.SET_TABLE_PREFS('SH','SALES','INCREMENTAL','TRUE');
```

- After load gather table statistics using GATHER_TABLE_STATS

- No need to specify parameters

```
EXEC DBMS_STATS.GATHER_TABLE_STATS('SH','SALES');
```

- The command will collect statistics for partitions and update the global statistics based on the partition level statistics and synopsis
- Possible to set incremental to true for all tables
 - Only works for already existing tables

```
EXEC DBMS_STATS.SET_GLOBAL_PREFS('INCREMENTAL','TRUE');
```

Attribute Clustering and Zone Maps

Introduced in Oracle 12c Release 1 (12.1.0.2)

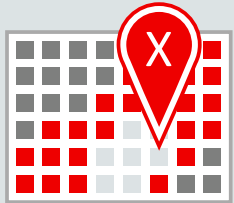


Zone Maps with Attribute Clustering



Attribute Clustering

Orders data so that columns values are stored together on disk



Zone maps

Stores min/max of specified columns per zone

Used to filter un-needed data during query execution

- Combined Benefits
- Improved query performance and concurrency
 - Reduced physical data access
 - Significant IO reduction for highly selective operations
- Optimized space utilization
 - Less need for indexes
 - Improved compression ratios through data clustering
- Full application transparency
 - Any application will benefit

Attribute Clustering

Concepts and Benefits

- Orders data so that it is in close proximity based on selected columns values: “attributes”
- Attributes can be from a single table or multiple tables
 - e.g. from fact and dimension tables
- Significant IO pruning when used with zone maps
- Reduced block IO for table lookups in index range scans
- Queries that sort and aggregate can benefit from pre-ordered data
- Enable improved compression ratios
 - Ordered data is likely to compress more than unordered data



Attribute Clustering for Zone Maps

Ordered rows

```
ALTER TABLE sales  
ADD CLUSTERING BY  
LINER ORDER (category);  
  
ALTER TABLE sales MOVE;
```

Category	Country
BOYS	AR
BOYS	JP
BOYS	SA
BOYS	US
GIRLS	AR
GIRLS	JP
GIRLS	SA
GIRLS	US
MEN	AR
MEN	JP
MEN	SA
MEN	US
WOMEN	AR
WOMEN	JP
WOMEN	SA
WOMEN	US

- Ordered rows containing category values BOYS, GIRLS and MEN.
- *Zone maps* catalogue regions of rows, or *zones*, that contain particular column value ranges.
- By default, each zone is up to 1024 blocks.
- For example, we only need to scan this zone if we are searching for category “GIRLS”. We can skip all other zones.

Attribute Clustering

Basics

- Two types of attribute clustering
 - LINEAR ORDER BY
 - Classical ordering
 - INTERLEAVED ORDER BY
 - Multi-dimensional ordering
- Simple attribute clustering on a single table
- Join attribute clustering
 - Cluster on attributes derived through join of multiple tables
 - Up to four tables
 - Non-duplicating join (PK or UK on joined table is required)

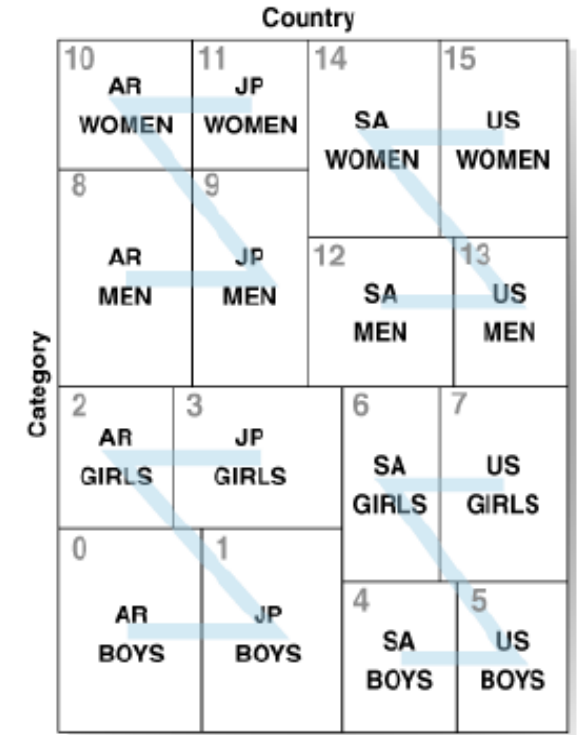
Attribute Clustering

Example

- CLUSTERING BY LINEAR ORDER (category, country)
- CLUSTERING BY INTERLEAVED ORDER (category, country)

Category	Country
BOYS	AR
BOYS	JP
BOYS	SA
BOYS	US
GIRLS	AR
GIRLS	JP
GIRLS	SA
GIRLS	US
MEN	AR
MEN	JP
MEN	SA
MEN	US
WOMEN	AR
WOMEN	JP
WOMEN	SA
WOMEN	US

- **LINEAR ORDER**



INTERLEAVED ORDER



Attribute Clustering

Basics

- Clustering directive specified at table level
 - ALTER TABLE ... ADD CLUSTERING ...
- Directive applies to new data and data movement
- Direct path operations
 - INSERT APPEND, MOVE, SPLIT, MERGE
 - Does not apply to conventional DML
- Can be enabled and disabled on demand
 - Hints and/or specific syntax

Zone Maps

Concepts and Basics

- Stores minimum and maximum of specified columns
 - Information stored per zone
 - [Sub]Partition-level rollup information for partitioned tables for multi-dimensional partition pruning
- Analogous to a coarse index structure
 - Much more compact than an index
 - Zone maps filter out what you don't need, indexes find what you do need
- Significant performance benefits with complete application transparency
 - IO reduction for table scans with predicates on the table itself or even a joined table using join zone maps (a.k.a. "hierarchical zone map")
- Benefits are most significant with ordered data
 - Used in combination with attribute clustering or data that is naturally ordered

Zone Maps

Basics

- Independent access structure built for a table
 - Implemented using a type of materialized view
 - For partitioned and non-partitioned tables
- One zone map per table
 - Zone map on partitioned table includes aggregate entry per [sub]partition
- Used transparently
 - No need to change or hint queries
- Implicit or explicit creation and column selection
 - Through Attribute Clustering: CREATE TABLE ... CLUSTERING
 - CREATE MATERIALIZED ZONEMAP ... AS SELECT ...

Attribute Clustering With Zone Maps

- CLUSTERING BY LINEAR ORDER (category, country)
- Zone map benefits are most significant with ordered data

Category	Country
BOYS	AR
BOYS	JP
BOYS	SA
BOYS	US
GIRLS	AR
GIRLS	JP
GIRLS	SA
GIRLS	US
MEN	AR
MEN	JP
MEN	SA
MEN	US
WOMEN	AR
WOMEN	JP
WOMEN	SA
WOMEN	US

Pruning with:

```
SELECT ..  
FROM table  
WHERE category =  
  'BOYS';
```

```
SELECT ..  
FROM table  
WHERE category =  
  'BOYS';  
AND country = 'US';
```

- LINEAR ORDER

Attribute Clustering With Zone Maps

- CLUSTERING BY INTERLEAVED ORDER (category, country)
- Zone map benefits are most significant with ordered data

		Country			
Category	10	11	14	15	
	AR WOMEN	JP WOMEN	SA WOMEN	US WOMEN	
	8	9	12	13	
	AR MEN	JP MEN	SA MEN	US MEN	
	2	3	6	7	
	AR GIRLS	JP GIRLS	SA GIRLS	US GIRLS	
	0	1	4	5	
	AR BOYS	JP BOYS	SA BOYS	US BOYS	

- INTERLEAVED ORDER

Pruning with:

```
SELECT ..  
FROM table  
WHERE category =  
  'BOYS';
```

```
SELECT ..  
FROM table  
AND country = 'US';
```

```
SELECT ..  
FROM table  
WHERE category =  
  'BOYS'  
AND country = 'US';
```

Zone Maps

Staleness

- DML and partition operations can cause zone maps to become fully or partially stale
 - Direct path insert does not make zone maps stale
- Single table 'local' zone maps
 - Update and insert marks impacted zones as stale (and any aggregated partition entry)
 - No impact on zone maps for delete
- Joined zone map
 - DML on fact table equivalent behavior to single table zone map
 - DML on dimension table makes dependent zone maps fully stale

Zone Maps

Refresh

- Incremental and full refresh, as required by DML
 - Zone map refresh does require a materialized view log
 - Only stale zones are scanned to refresh the MV
 - For joined zone map
 - DML on fact table: incremental refresh
 - DML on dimension table: full refresh
- Zone map maintenance through
 - DBMS_MVIEW.REFRESH()
 - ALTER MATERIALIZED ZONEMAP <xx> REBUILD;



Example – Dimension Hierarchies

ORDERS

id	product_id	location_id	amount
1	3	23	2.00
2	88	55	43.75
3	31	99	33.55
4	33	62	23.12
5	21	11	38.00
6	33	21	5.00
7	44	71	10.99

Note: a zone typically contains many more rows than show here.
This is for illustrative purposes only.

LOCATIONS

location_id	State	county
23	California	Inyo
102	New Mexico	Union
55	California	Kern
1	Ohio	Lake
62	California	Kings

```
CREATE TABLE orders ( ... )  
CLUSTERING orders  
JOIN locations ON (orders.location_id = locations.location_id)  
BY INTERLEAVED ORDER (locations.state, locations.county)  
WITH MATERIALIZED ZONEMAP ...
```

Example – Dimension Hierarchies

ORDERS

id	product_id	location_id	amount
1	3	23	2.00
2	88	55	43.75
3	31	99	33.55
4	33	62	23.12
5	21	11	38.00
6	33	21	5.00
7	44	71	10.99

Note: a zone typically contains many more rows than show here.
This is for illustrative purposes only.

Scan
Zone



LOCATIONS

location_id	State	county
23	California	Inyo
102	New Mexico	Union
55	California	Kern
1	Ohio	Lake
62	California	Kings

```
SELECT SUM(amount)
FROM orders
JOIN locations ON (orders.location.id = locations.location.id)
WHERE state = 'California';
```



Example – Dimension Hierarchies

ORDERS

id	product_id	location_id	amount
1	3	23	2.00
2	88	55	43.75
3	31	99	33.55
4	33	62	23.12
5	21	11	38.00
6	33	21	5.00
7	44	71	10.99

Note: a zone typically contains many more rows than show here.
This is for illustrative purposes only.

Scan
Zone



LOCATIONS

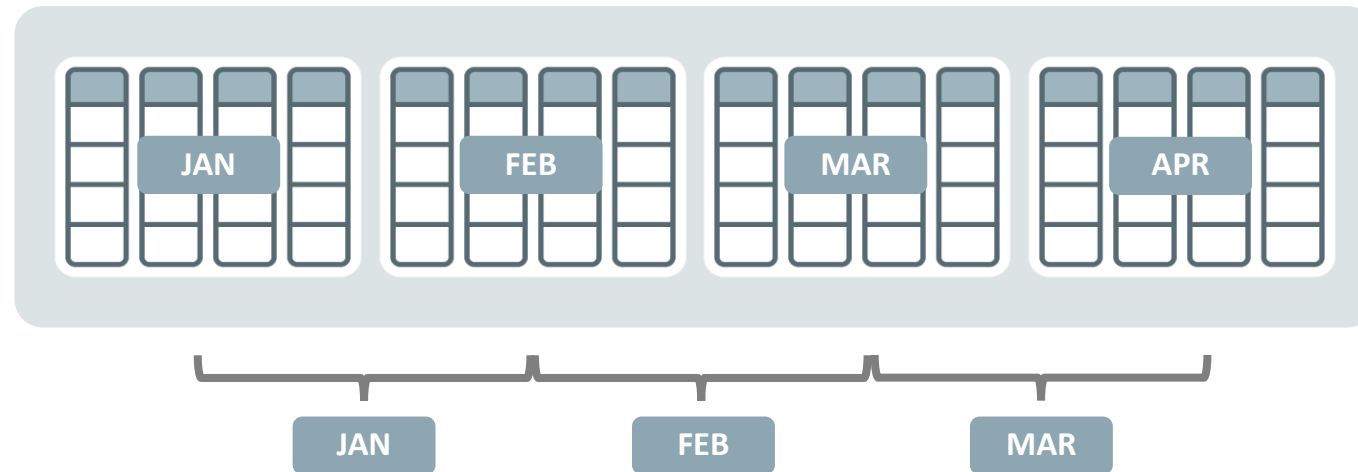
location_id	State	county
23	California	Inyo
102	New Mexico	Union
55	California	Kern
1	Ohio	Lake
62	California	Kings

```
SELECT SUM(amount)
FROM orders
JOIN locations ON (orders.location.id = locations.location.id)
WHERE state = 'California'
AND county = 'Kern';
```



Zone Maps and Partitioning

Partition Key:
ORDER_DATE



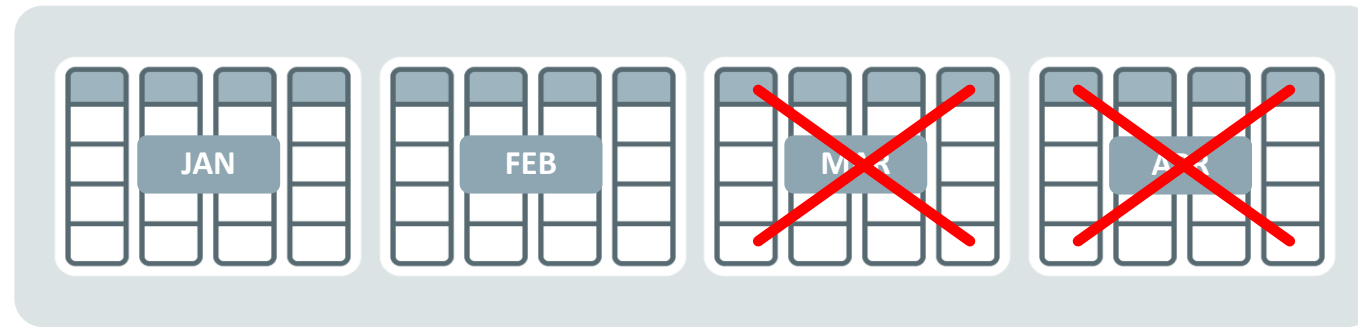
Zone map column
SHIP_DATE
correlates with
partition key
ORDER_DATE

Zone map:
SHIP_DATE

- Zone maps can prune partitions for columns that are not included in the partition (or subpartition) key

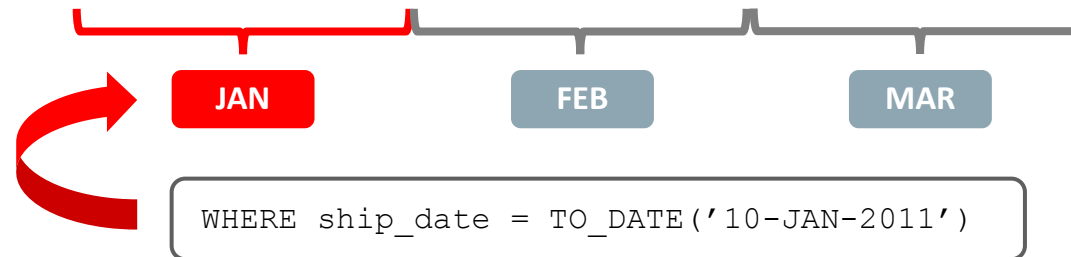
Zone Maps and Partitioning

Partition Key:
ORDER_DATE



MAR and APR partitions
are pruned

Zone map:
SHIP_DATE



```
WHERE ship_date = TO_DATE('10-JAN-2011')
```

- Zone maps can prune partitions for columns that are not included in the partition (or subpartition) key

Zone Maps and Storage Indexes

- Attribute clustering and zone maps work transparently with Exadata storage indexes
 - The benefits of Exadata storage indexes continue to be fully exploited
- In addition, zone maps (when used with attribute clustering)
 - Enable additional and significant IO optimization
 - Provide an alternative to indexes, especially on large tables
 - Join and fact-dimension queries, including dimension hierarchy searches
 - Particularly relevant in star and snowflake schemas
 - Are able to prune entire partitions and sub-partitions
 - Are effective for both direct and conventional path reads
 - Include optimizations for joins and index range scans
 - Part of the physical database design: explicitly created and controlled by the DBA



Partitioning tips and tricks

Partitioning Tips and Tricks

- Think about partitioning strategy
- Physical and logical attributes
- Eliminate hot spots
- Smart partial exchange
- Exchange with PK and UK

Think about your partitioning strategy

Choosing your Partitioning Strategy

- Think about
 - your data
 - your usage
- What do you expect from Partitioning?
 - Query performance benefits
 - Load (or purge) performance benefits
 - Data management benefits

Choosing your Partitioning Strategy

Logical shape of the data

- How is data inserted into your system?
- How is data maintained in your system?
- How is data accessed in your system?

Choosing your Partitioning Strategy

Logical shape of the data

- How is data inserted into your system?
 - Time, location, tenant, business user, ...
 - Ranges, unrelated list of values, “just lots of them”, ...
- How is data maintained in your system?
- How is data accessed in your system?

Choosing your Partitioning Strategy

Logical shape of the data

- How is data inserted into your system?
 - Time, location, tenant, business user, ...
 - Ranges, unrelated list of values, “just lots of them”, ...
- How is data maintained in your system?
 - Moving window of active data, legal requirements, data “forever”, ...
 - Don't know yet
- How is data accessed in your system?

Choosing your Partitioning Strategy

Logical shape of the data

- How is data inserted into your system?
 - Time, location, tenant, business user, ...
 - Ranges, unrelated list of values, “just lots of them”, ...
- How is data maintained in your system?
 - Moving window of active data, legal requirements, data “forever”, ...
 - Don't know yet
- How is data accessed in your system?
 - Always full, with common FILTER predicates, always index access, ...
 - Don't know yet

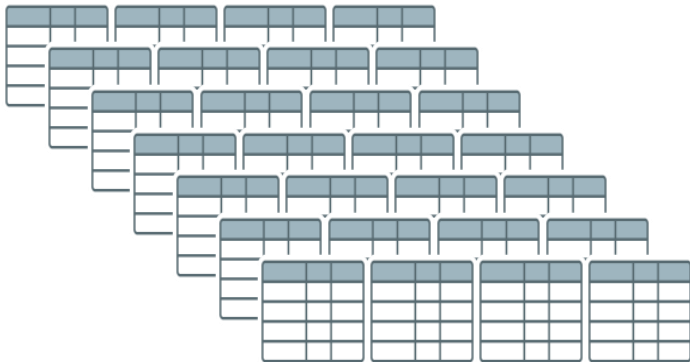
Choosing your Partitioning Strategy

Performance improvements

- Query speedup
 - Partition elimination
 - Partition-wise joins
- DML speedup
 - Alleviation of contention points
- Data maintenance
 - DDL instead of DML

Choosing your Partitioning Strategy

Data Access – Full Table Access

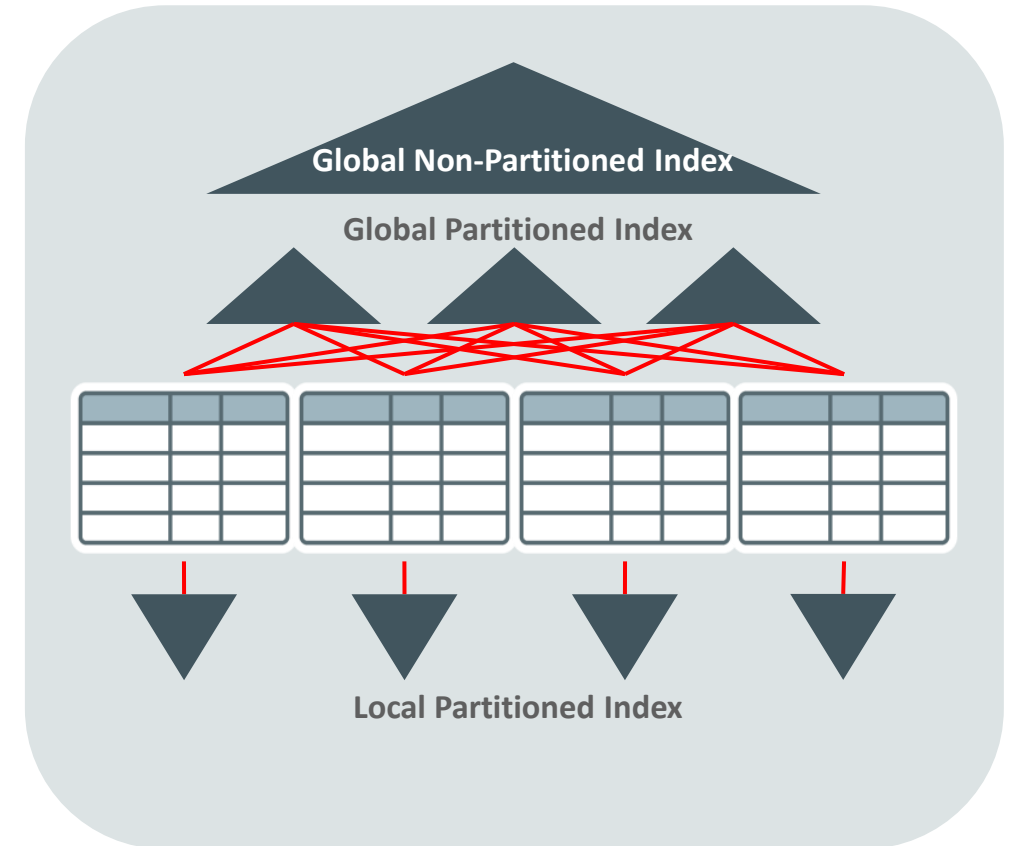


- I/O savings are linear to the number of pruned partitions
 - One of 10: ten times less IO
 - One of 100: hundred times less IO
- Runtime improvements depend on
 - Relative contribution of IO versus CPU work
 - Potential impact on subsequent operations

Choosing your Partitioning Strategy

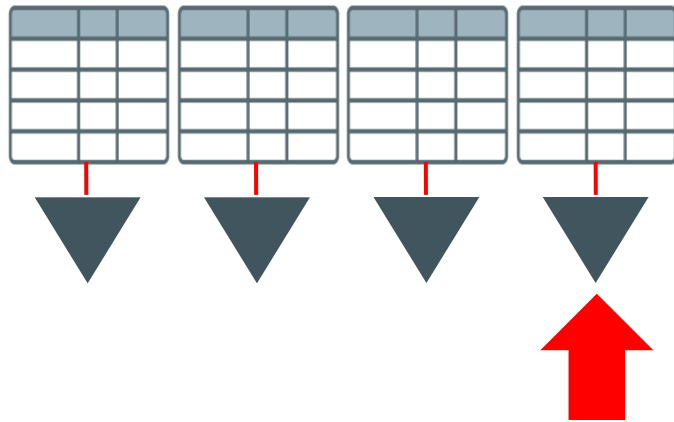
Indexing of partitioned tables

- GLOBAL index points to rows in any partition
 - Index can be partitioned or not
- LOCAL index is partitioned same as table
 - Index partitioning key can be different from index key

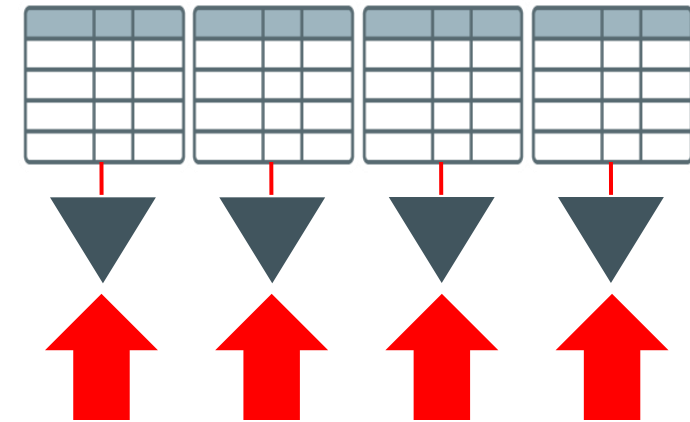


Choosing your Partitioning Strategy

Data Access – local index and global partitioned index



- Partitioned index access with single partition pruning



- Partitioned index access without any partition pruning

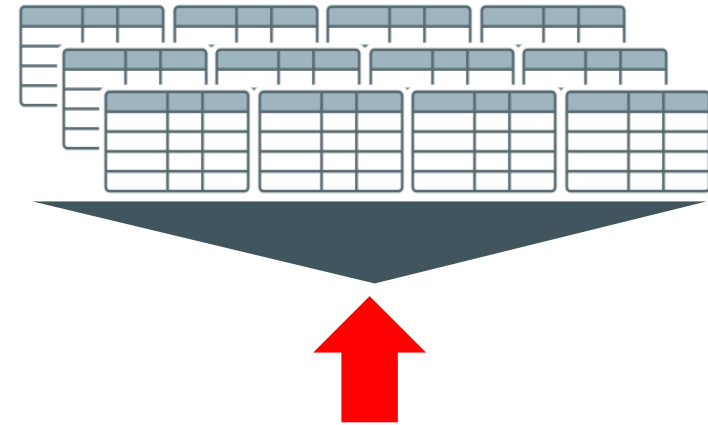
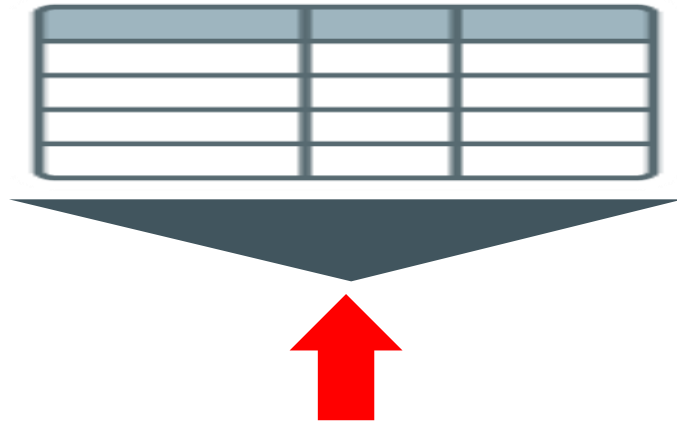
Local and Global Partitioned Indexes

Data Access

- Number of index probes identical to number of accessed partitions
 - No partition pruning leads to a probe into all index partitions
- Not optimally suited for OLTP environments
 - No guarantee to always have partition pruning
 - Exception: global hash partitioned indexes for DML contention alleviation
 - Most commonly small number of partitions
- Pruning on global partitioned indexes based on the index prefix
 - Index prefix identical to leading keys of index

Choosing your Partitioning Strategy

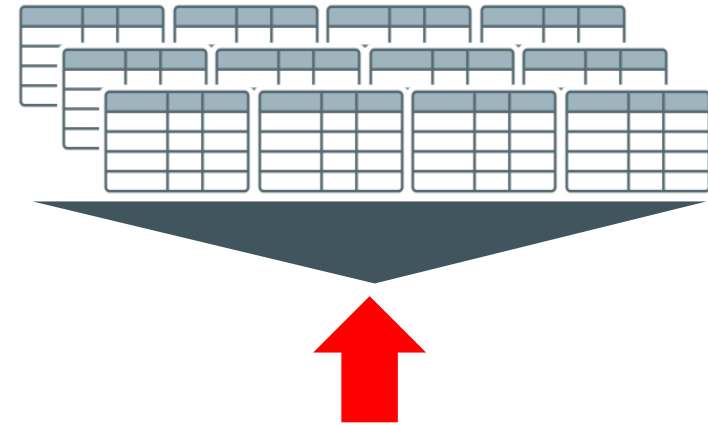
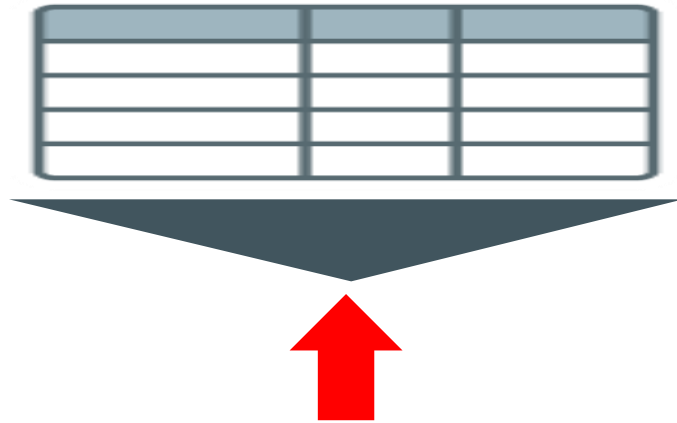
Global nonpartitioned index



- Can you see the difference?

Choosing your Partitioning Strategy

Global nonpartitioned index



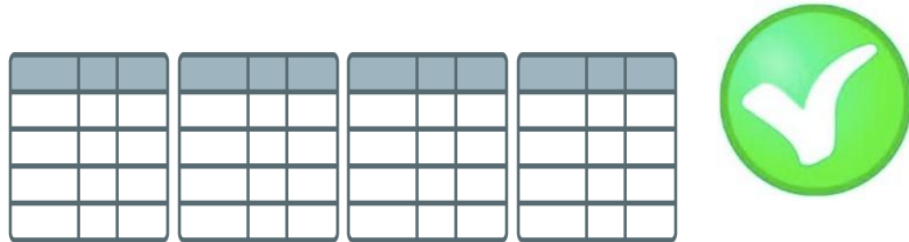
- Can you see the difference?
- There is more or less none*

* Some differences for index size, due to large rowid

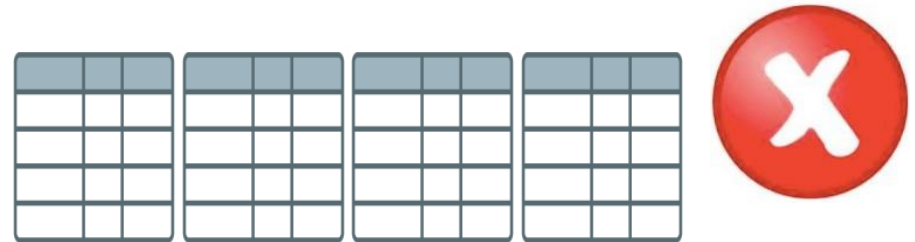
Global Indexes

Data Access

- No pruning for non-partitioned indexes
 - You always probe into a single index segment
- Global partitioned index prefix identical to leading keys of index
 - Pruning on index prefix, not partition key column(s)
- Most common in OLTP environments



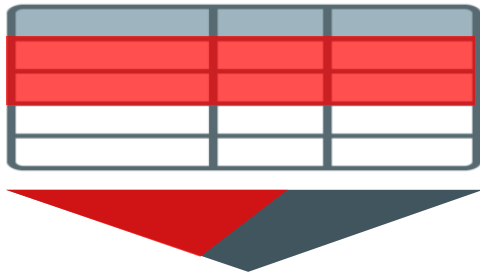
`PARTITION BY (col1), idx(col1)`



`PARTITION BY (col1), idx(col2)`

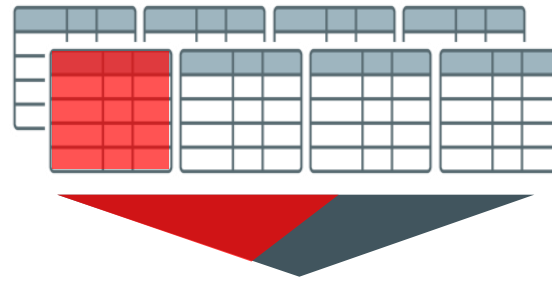
Choosing your Partitioning Strategy

Data Maintenance



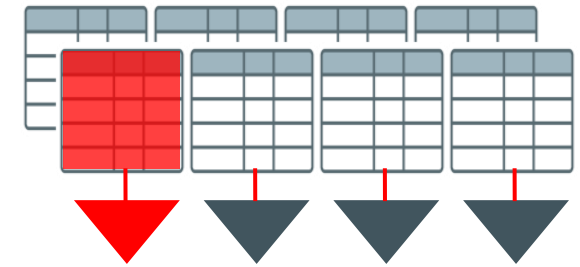
```
DELETE FROM ...  
WHERE ...
```

- Records get deleted
 - Index maintenance
 - Undo and redo



```
ALTER TABLE ... DROP PARTITION ...
```

- Partition gets dropped
 - Fast global index maintenance (12c)
 - Minimal undo



- Partition gets dropped
 - Local index gets dropped
 - Minimal undo

Local Indexes

Data Maintenance

- Incremental index creation possible
 - Initial unusable creation, rebuild of individual partitions
- Fast index maintenance for all partition maintenance operations that only touch one partition
 - Exchange, drop, truncate
- Partition maintenance that touches more than one partition require index maintenance
 - Merge, split creates new data segments
 - New index segments are created as well

Global Indexes

Data Maintenance

- Incremental index creation is hard, if not impossible
- “Fast” index maintenance for drop and truncate beginning with Oracle Database 12c
 - Fast actually means delayed index maintenance
- Partition maintenance except drop and truncate requires index maintenance
 - Conventional index maintenance equivalent to the DML operations that would represent the PMOP

How many partitions?

How many partitions?

It depends

Data Volume and Number of Partitions

- Imagine a 100TB table ...
 - With one million partitions, each partition is 100MB in size
- Imagine a 10TB table ...
 - With one million partitions, each partition is 10MB in size
- Imagine a 1TB table ...
 - With one million partitions, each partition is 1MB in size

Data Volume and Number of Partitions

- Imagine a 100TB table ...
 - With one million partitions, each partition is 100MB in size
- Imagine a 10TB table ...
 - With one million partitions, each partition is 10MB in size
- Imagine a 1TB table ...
 - With one million partitions, each partition is 1MB in size
- **How long does it take your system to read 1MB??**
 - Exadata full table scan rate is 25GB/sec ... (disk only, full rack X5-2)

Data Volume and Number of Partitions

- More is not always better
 - Every partition represents metadata in the dictionary
 - Every partition increases the metadata footprint in the SGA
- Find your personal balance between the number of partitions and its average size
 - There is nothing wrong about single-digit GB sizes for a segment on “normal systems”
 - Consider more partitions \geq 5GB segment size

Choosing your Partitioning Strategy

Customer Usage Patterns

- Range (Interval) still the most prevalent partitioning strategy
 - Almost always some time dependency
- List more and more common
 - Interestingly often based on time as well
 - Often as subpartitioning strategy
- Hash not only used for performance (PWJ, DML contention)
 - No control over data placement, but some understanding of it
 - Do not forget the power of two rule

Choosing your Partitioning Strategy

Extended Partitioning Strategies

- Interval Partitioning fastest growing new partitioning strategy
 - Manageability extension to Range Partitioning
- Reference Partitioning
 - Leverage PK/FK constraints for your data model
- Interval-Reference Partitioning (new in Oracle Database 12c)
- Virtual column based Partitioning
 - Derived attributes without little to no application change
- Any variant of the above

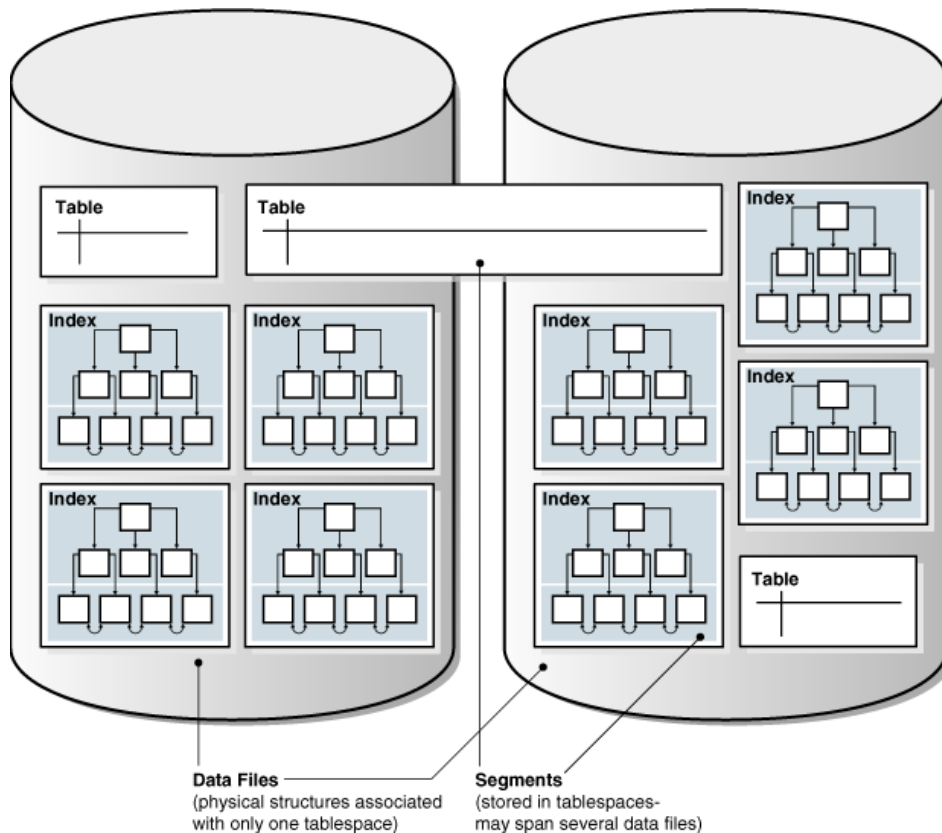
Physical and logical attributes

Physical and Logical Attributes

- Logical attributes
 - Partitioning setup
 - Indexing and index maintenance
 - Read only (in conjunction with tablespace separation)
- Physical attributes
 - Data placement
 - Segment properties in general

Nonpartitioned Tables

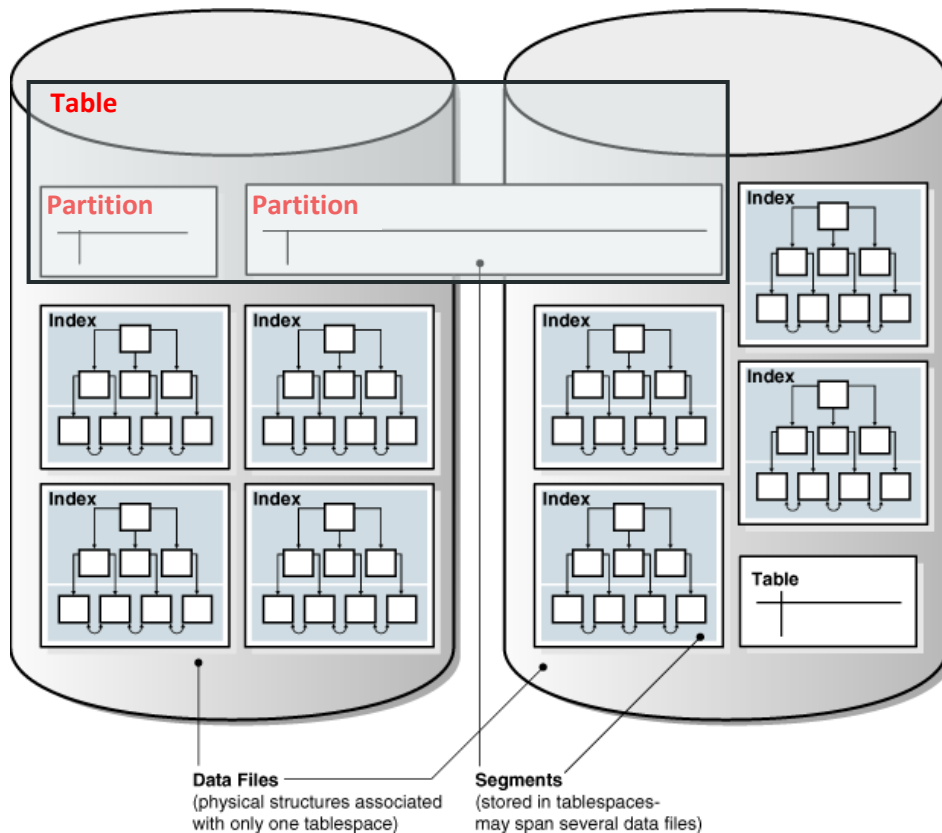
Physical and Logical Attributes



- Logical table properties
 - Columns and data types
 - Constraints
 - Indexes, ...
- Physical table properties
 - Table equivalent to segment
 - Tablespace
 - Compression, [Logging | nologging], ...
 - In-memory
 - Properties managed and changed on segment level

Partitioned Tables

Physical and Logical Attributes



- Logical **table** properties
 - Columns and data types
 - Constraints
 - **Partial** Indexes, ...
 - **Physical property directives**
- Physical **[sub]partition** properties
 - **[Sub]partition** equivalent to segment
 - Tablespace
 - Compression, [Logging | nologging], ...
 - In-memory
 - Properties managed and changed on segment level

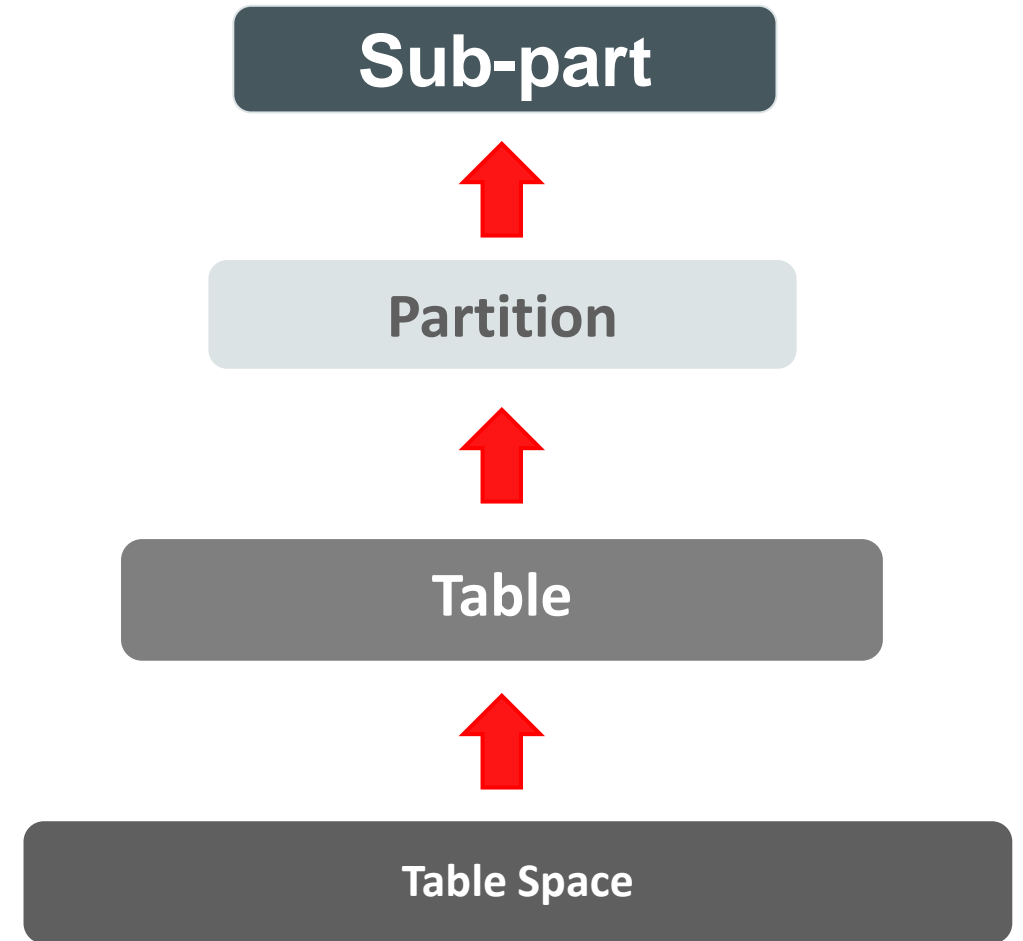
Partitioned Tables

Physical and Logical Attributes

- Table is metadata-only and directive for future partitions
 - No physical segments on table level
 - Physical attributes become directive for new partitions, if specified
- Single-level partitioned table
 - Partitions are equivalent to segments
 - Physical attributes are managed and changed on partition level
- Composite-level partitioned tables
 - Partitions are metadata only and directive for future subpartitions
 - Subpartitions are equivalent to segments

Data Placement with Partitioned Tables

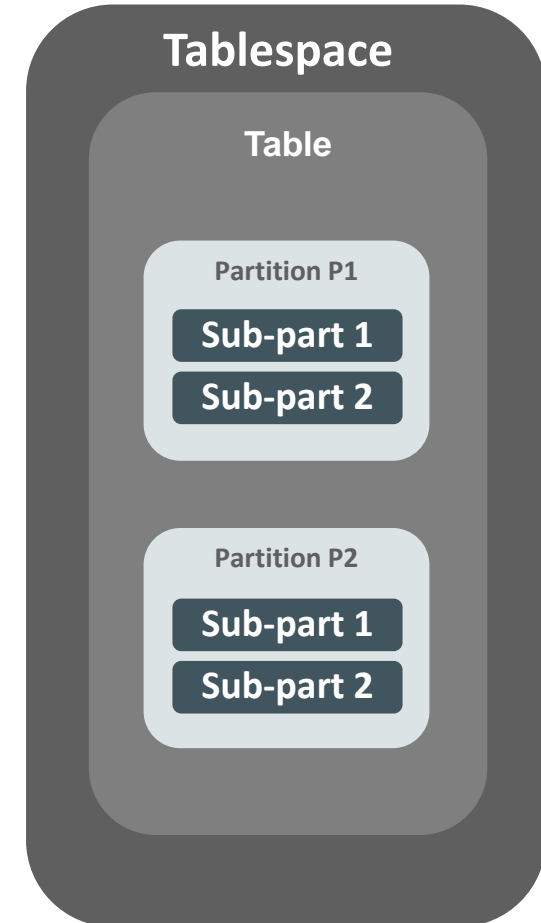
- Each partition or sub-partition is a separate object
- Specify storage attributes at each individual level
 - As placement policy for lower levels
 - For each individual [sub]partition
- If storage attributes are not specified standard hierarchical inheritance kicks in



Data Placement with Partitioned Tables

Special Case Interval Partitioning

- Interval Partitioning” pre-creates” all partitions
 - All 1 million [sub]partitions exist logically
- Physical storage is (almost) determined as well
- Partition placement
 - Inherited from table level
 - STORE IN () clause for round-robin partition placement
- Subpartition placement
 - Usage of subpartition template
 - Needs bug fix #8304261 (included in 11.2.0.3)
 - STORE IN clause currently is currently a no-op



Data Placement with Partitioned Tables

Subpartition template

- Introduced in Oracle Database 9 Release 2
 - Allows predefinition of subpartitions for **future** partitions
 - Stored as metadata in the data dictionary

```
CREATE TABLE stripe_regional_sales
  (deptno number, item_no varchar2(20),
   txn_date date, txn_amount number, state varchar2(2))
PARTITION BY RANGE (txn_date)
SUBPARTITION BY LIST (state)
SUBPARTITION TEMPLATE
  (SUBPARTITION northwest VALUES ('OR', 'WA') TABLESPACE tbs_1
   SUBPARTITION southwest VALUES ('AZ', 'UT', 'NM') TABLESPACE tbs_2
   SUBPARTITION northeast VALUES ('NY', 'VM', 'NJ') TABLESPACE tbs_3
   SUBPARTITION southeast VALUES ('FL', 'GA') TABLESPACE tbs_4
   SUBPARTITION midwest VALUES ('SD', 'WI') TABLESPACE tbs_5
   SUBPARTITION south VALUES ('AL', 'AK') TABLESPACE tbs_6
   SUBPARTITION south VALUES (DEFAULT) TABLESPACE tbs_7
  )
(PARTITION q1_2016 VALUES LESS THAN ( TO_DATE('01-APR-2016', 'DD-MON-YYYY')),
 (PARTITION q2_2016 VALUES LESS THAN ( TO_DATE('01-JUL-2016', 'DD-MON-YYYY')),
 (PARTITION q3_2016 VALUES LESS THAN ( TO_DATE('01-OCT-2016', 'DD-MON-YYYY')),
 (PARTITION q4_2016 VALUES LESS THAN ( TO_DATE('01-JAN-2015', 'DD-MON-YYYY')),
 );
```

Subpartition
definition for all
future partitions

Subpartition applied
to every partition

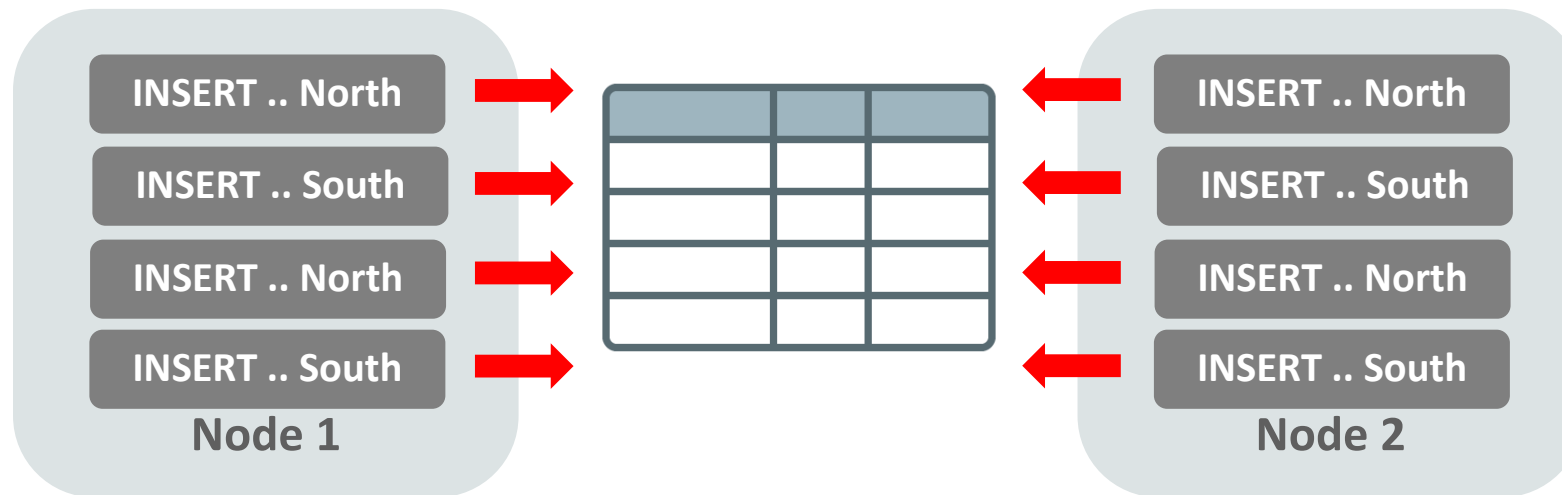


Using partitioning to eliminate hot spots

Using Partitioning to eliminate Hot Spots

Nonpartitioned table

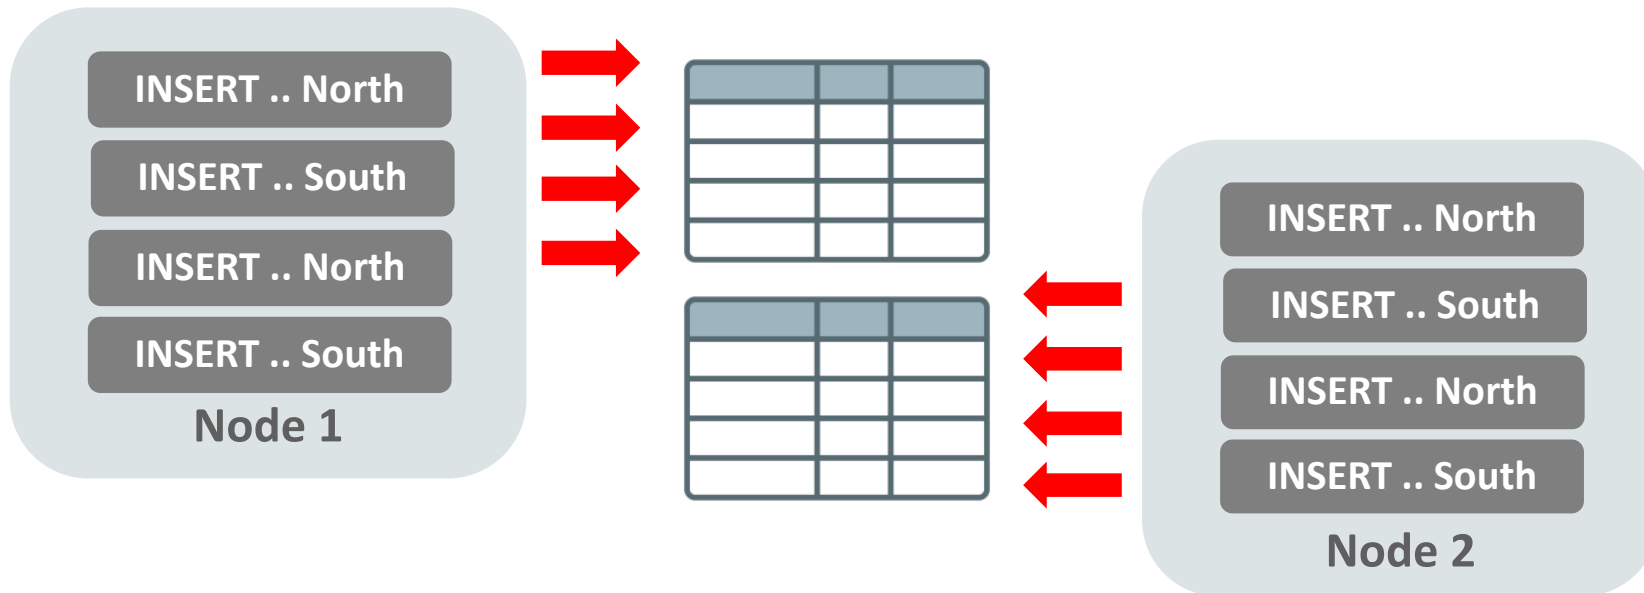
- On RAC, high DML workload causes high cache fusion traffic
 - Oracle calls this block pinging



Using Partitioning to eliminate Hot Spots

HASH partitioned table

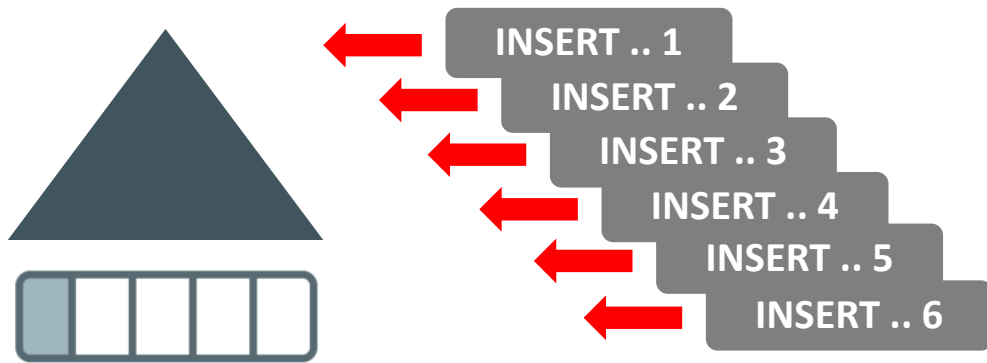
- On RAC, high DML workload causes high cache fusion traffic
 - Oracle calls this block pinging
- HASH (or LIST) partitioned table can alleviate this situation
 - Caveat: Normally needs some kind of “application partitioning” or “application RAC awareness”



Using Partitioning to eliminate Hot Spots

HASH partitioned index

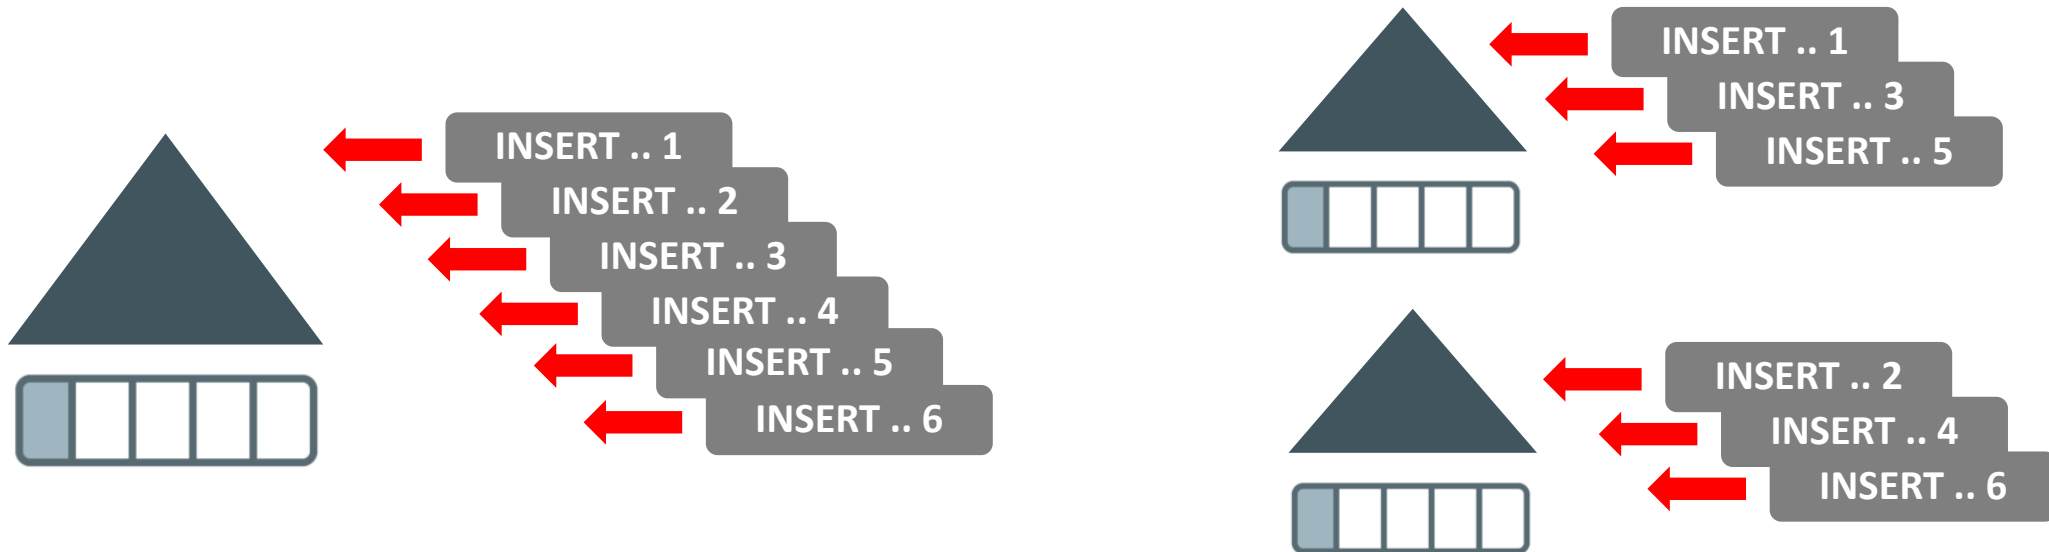
- High DML workload can create **hot spots** (contention) on index blocks
 - E.g. artificial (right hand growing) primary key index



Using Partitioning to eliminate Hot Spots

HASH partitioned index

- High DML workload can create **hot spots** (contention) on index blocks
 - E.g. artificial (right hand growing) primary key index
- With HASH partitioned index you get **warm spots**



Hot Spot Elimination – Use Case

Challenge

- Retail application using object-relational mapping
- Only “common” database functionality is used
- Every single row needs to be updated in a single transaction
- No bulk imports possible at all!
- Thousands of small SQL-Statements issued
- Sudden heavy peaks in user access
 - e.g. Cyber Monday, Christmas trade, special offers, ..
- **Experienced sporadic contention**

Hot Spot Elimination – Use Case

Performance without any application code change

Results from PoC (SKU data load)

- Reference system: 120 SKU's per second
- Exadata Machine (single node load)
 - 2,500 SKU's per second (20x faster)
- Exadata Machine X3-2 (two node load & without partitioning)
 - “only” 1,900 SKU's per second (slower than single node load !!!)
- Exadata Machine X3-2 (two node load & with proper partitioning)
 - 4,800 SKU's per second (40x faster)
- **Proper partitioning enables linear scaling**



Hot Spot Elimination – Use Case

How to (Alternative A, Hash Partitioning on store ID)

- HASH Partitioning creates <n> entry points into the table

```
CREATE TABLE <table_name> (  
  ID          NUMBER(10) NOT NULL,  
  Cn          ...)  
PARTITION BY HASH(ID) PARTITIONS <n>  
TABLESPACE <tablespace_name> STORAGE ( ... );  
  
CREATE UNIQUE INDEX <index_name> ON <table_name>  
(ID) LOCAL TABLESPACE <tablespace_name> STORAGE ( ... );  
  
INSERT INTO <table_name> (ID, ...)  
SELECT SEQ_ID.nextval, ... ;
```

Hot Spot Elimination – Use Case

How to (Alternative B, List Partitioning on instance #)

- Sequence SEQ_ID forces ID to be unique in each partition!
- List Partitioning completely separates the entry points per instance

```
CREATE TABLE <table_name> (  
  ID          NUMBER(10) NOT NULL,  
  Cn          ...          ...,  
  INSTANCE_NUMBER NUMBER(1) DEFAULT sys_context('USERENV','INSTANCE') NOT NULL)  
PARTITION BY LIST (INSTANCE_NUMBER)  
( PARTITION P1 VALUES(1),  
  PARTITION P2 VALUES(2),  
  ...  
  PARTITION Pn VALUES(n))  
TABLESPACE <tablespace_name> STORAGE ( ... );  
  
CREATE UNIQUE INDEX <index_name> ON <table_name>  
(ID, INSTANCE_NUMBER) LOCAL TABLESPACE <tablespace_name> STORAGE ( ... );  
  
INSERT INTO <table_name> (ID, ...) SELECT SEQ_ID.nextval, ... ;
```

Hot Spot Elimination – Use Case

How to (Enhanced alternative B, Hash Partitioning on instance #)

- Sequence SEQ_ID forces ID to be unique in each partition!

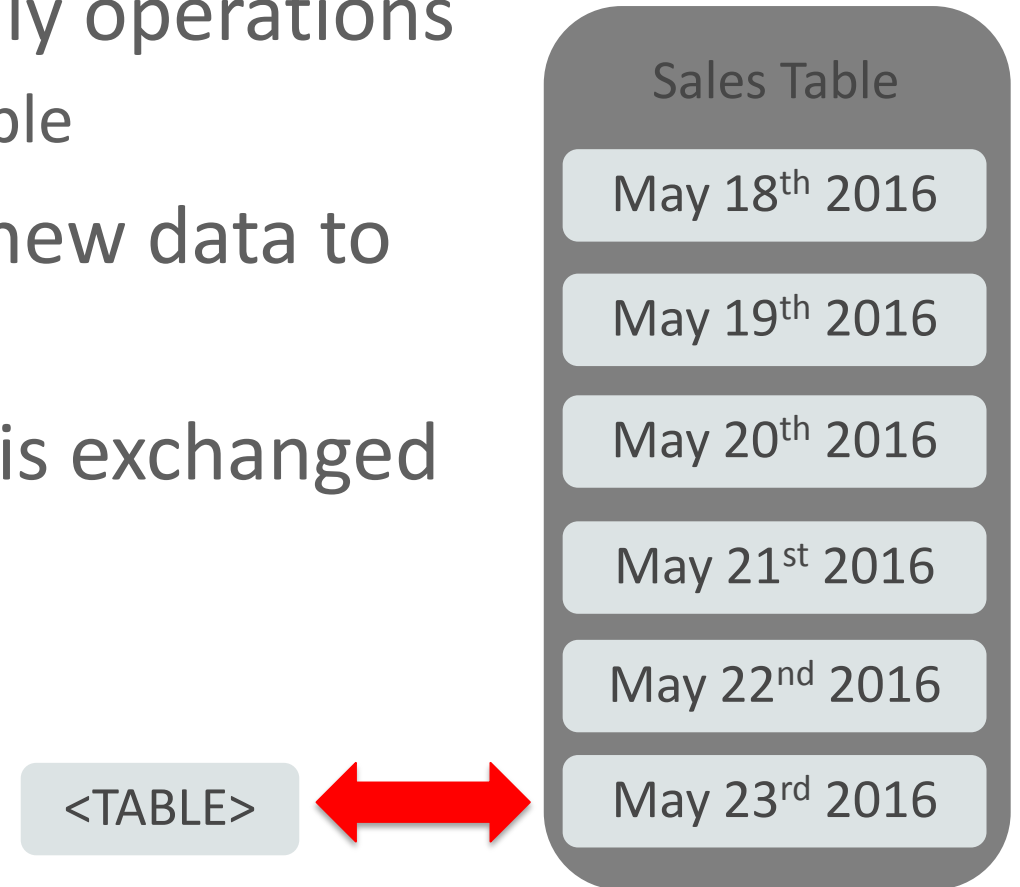
```
CREATE TABLE <table_name> (  
  ID          NUMBER(10) NOT NULL,  
  Cn          ...          ...,  
  INSTANCE_NUMBER NUMBER(1) DEFAULT sys_context('USERENV','INSTANCE') NOT NULL)  
PARTITION BY LIST (INSTANCE_NUMBER)  
SUBPARTITION BY HASH (ID) SUBPARTITIONS <m>  
( PARTITION P1 VALUES(1),  
  PARTITION P2 VALUES(2),  
  ...  
  PARTITION Pn VALUES(n))  
TABLESPACE <tablespace_name> STORAGE ( ... );  
CREATE UNIQUE INDEX <index_name> ON <table_name>  
(ID, INSTANCE_NUMBER) LOCAL TABLESPACE <tablespace_name> STORAGE ( ... );  
INSERT INTO <table_name> (ID, ...) SELECT SEQ_ID.nextval, ... ;
```

Smart partial partition exchange

“Filtered partition maintenance”

Partition Exchange for Loading and Purging

- Remove and add data as metadata only operations
 - Exchange the metadata of partition and table
- Data load: standalone table contains new data to being loaded
- Data purge: partition containing data is exchanged with empty table
- Drop partition alternative for purge
 - Data is gone forever



Smart Partial Partition Exchange

- Sounds easy but ...
- What to do if partition boundaries are not 100% aligned?
 - “Partial Purging”

Smart Partial Partition Exchange

Partial Purging

- Lock partition to being purged

```
LOCK TABLE ... PARTITION ...
```

Sales Table

May 18th 2016

May 19th 2016

May 20th 2016

May 21st 2016

May 22nd 2016

May 23rd 2016

Smart Partial Partition Exchange

Partial Purging

- Lock partition to being purged

```
LOCK TABLE ... PARTITION ...
```

- Create table containing remaining data set

```
CREATE TABLE ... AS SELECT WHERE ...
```

“REST”

Sales Table

May 18th 2016

May 19th 2016

May 20th 2016

May 21st 2016

May 22nd 2016

May 23rd 2016

Smart Partial Partition Exchange

Partial Purging

- Lock partition to being purged

```
LOCK TABLE ... PARTITION ...
```

- Create table containing remaining data set

```
CREATE TABLE ... AS SELECT WHERE ...
```

- Create necessary indexes, if any

“REST”

Sales Table

May 18th 2016

May 19th 2016

May 20th 2016

May 21st 2016

May 22nd 2016

May 23rd 2016

Smart Partial Partition Exchange

Partial Purging

- Lock partition to being purged

```
LOCK TABLE ... PARTITION ...
```

- Create table containing remaining data set

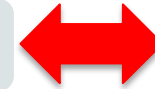
```
CREATE TABLE ... AS SELECT WHERE ...
```

- Create necessary indexes, if any

- Exchange partition

```
ALTER TABLE ... EXCHANGE PARTITION ...
```

May 18th 2016



Sales Table

“REST”

May 19th 2016

May 20th 2016

May 21st 2016

May 22nd 2016

May 23rd 2016

Smart Partial Partition Exchange

Partial Purging

- Lock partition to being purged

```
LOCK TABLE ... PARTITION ...
```

- Create table containing remaining data set

```
CREATE TABLE ... AS SELECT WHERE ...
```

- Create necessary indexes, if any

- Exchange partition

```
ALTER TABLE ... EXCHANGE PARTITION ...
```

May 18th 2016



Sales Table

“REST”

May 19th 2016

May 20th 2016

May 21st 2016

May 22nd 2016

May 23rd 2016

Exchange in the presence of unique and primary key constraints

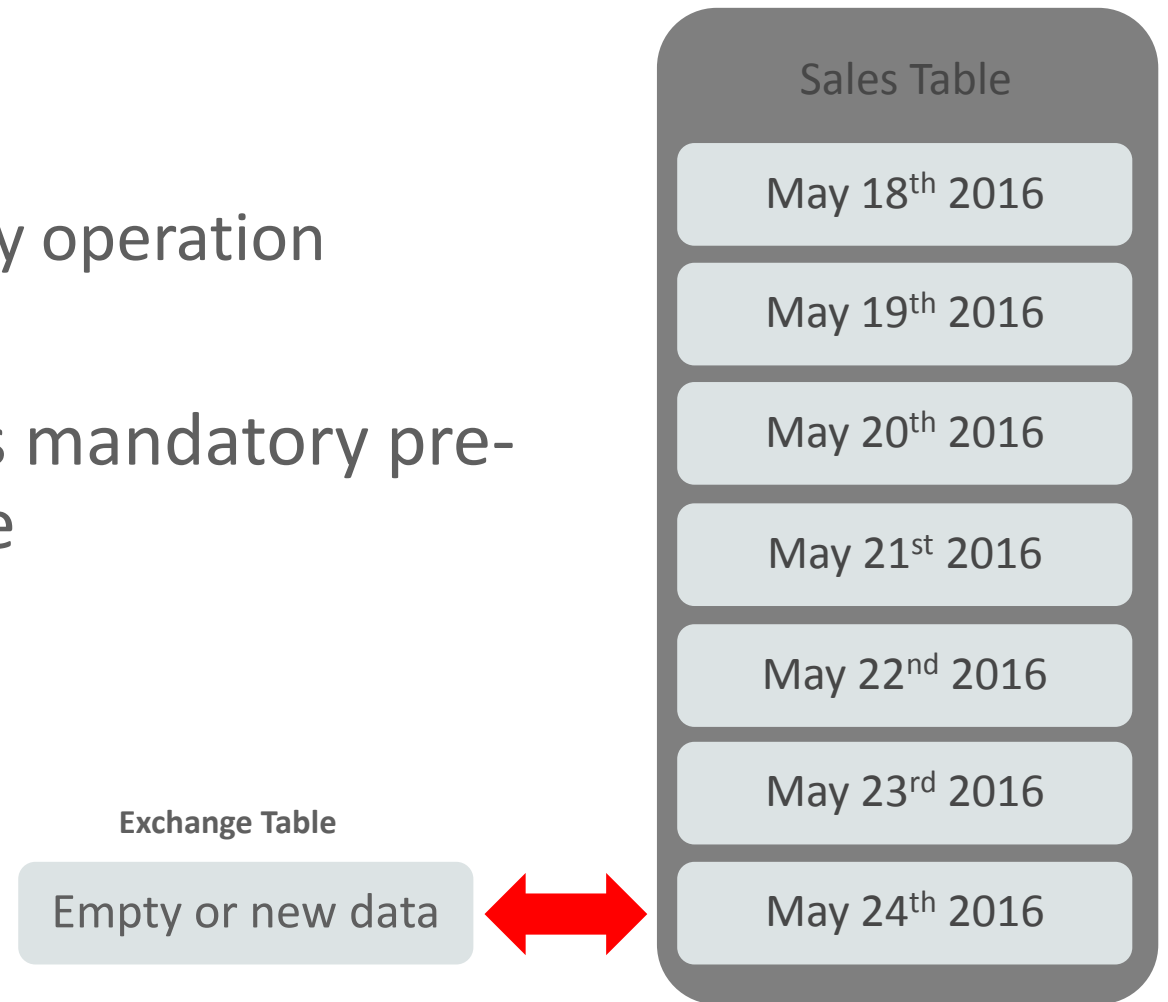
Unique Constraints/Primary Keys

- Unique constraints are enforced with unique indexes
 - Primary key constraint adds NOT NULL to column
 - Table can have only one primary key (“unique identifier”)
- Partitioned tables offer two types of indexes
 - Local indexes
 - Global index, both partitioned and non-partitioned

Partition Exchange

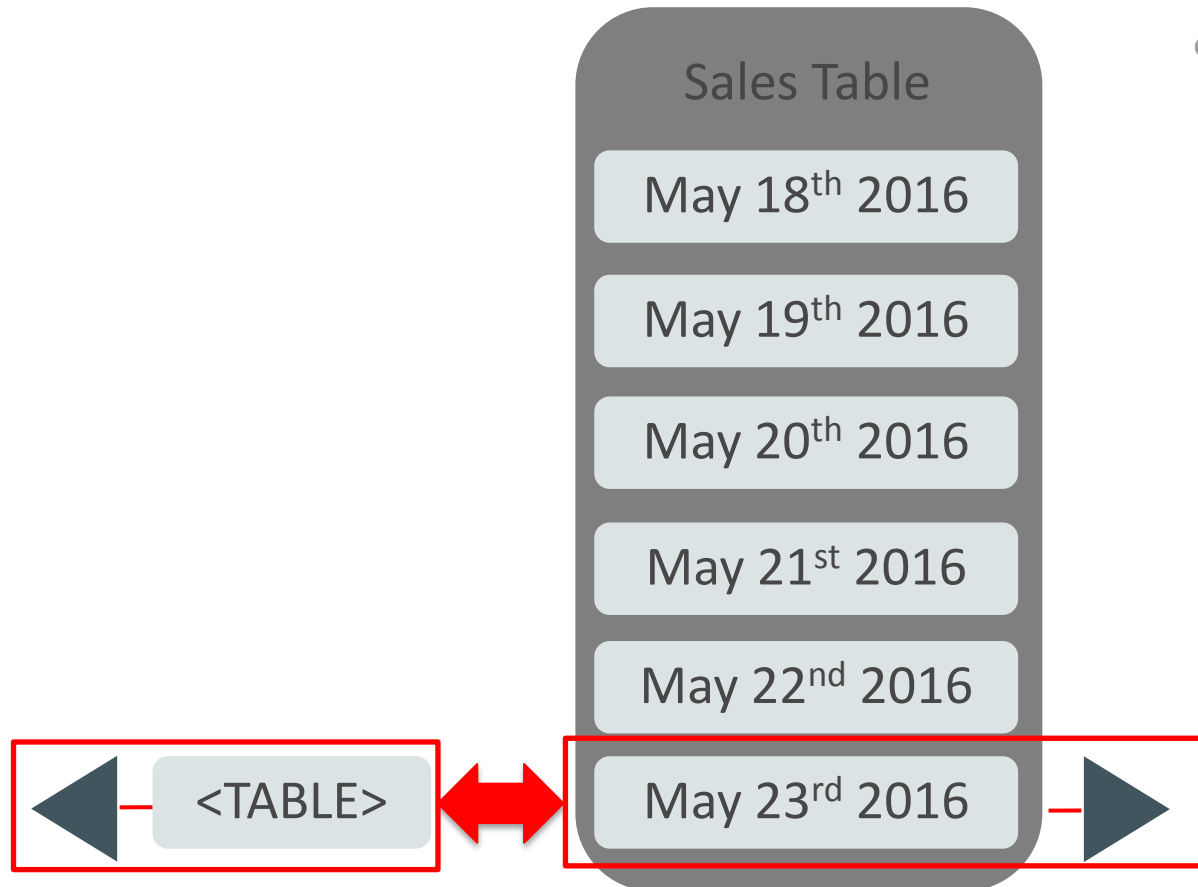
A.k.a Partition Loading and Purging

- Remove and add data as metadata-only operation
 - Exchange the metadata of partitions
- Same logical shape for both tables is mandatory pre-requirement for successful exchange
 - Same number and data type of columns
 - Note that column name does not matter
 - Same constraints
 - Same number and type of indexes



Partition Exchange

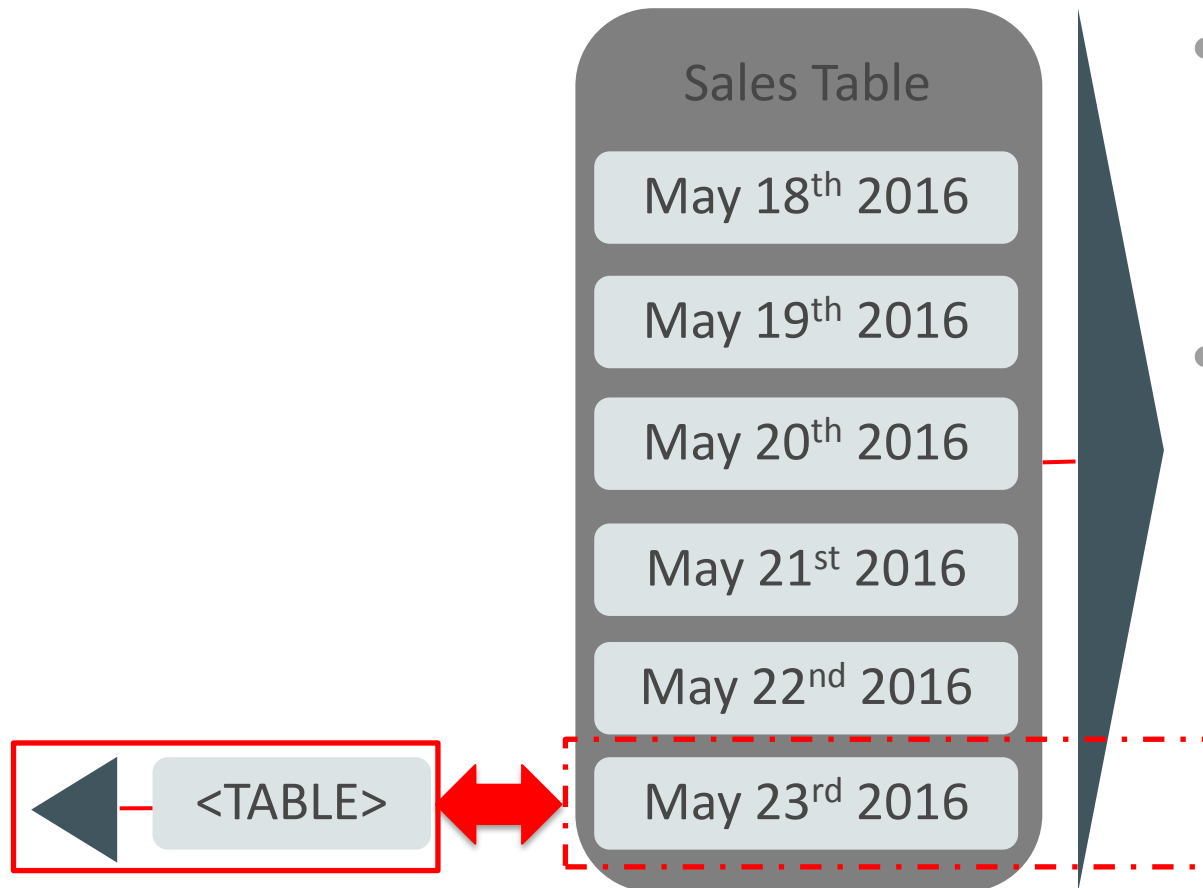
Local Indexes



- Any index on the exchange table is equivalent to a local partitioned index

Partition Exchange

Local Indexes



- Any index on the exchange table is equivalent to a local partitioned index
- What do I do when the PK index on the partitioned table needs global index enforcement?
 - Remember the requirement of logical equivalence ...

Partition Exchange and PK/Unique Constraint

The Dilemma

- Global indexes only exist for a partitioned table
 - But I need the index for the exchange table for uniqueness ...

Partition Exchange and PK/Unique Constraint

Not Really a Dilemma

- Global indexes only exist for a partitioned table
 - But I need the index for the exchange table for uniqueness ...
- Not generically true
 - Unique index only needed for **enabled** constraints
 - Enforcement for new or modified data through index probe

Partition Exchange and PK/Unique Constraint

Not Really a Dilemma

- Global indexes only exist for a partitioned table
 - But I need the index for the exchange table for uniqueness ...
- Not generically true
 - Unique index only needed for **enabled** constraints
 - Enforcement for new or modified data through index probe
 - **Disabled** constraint prevents data insertion

```
SQL> alter table tt add(constraint x unique (col1) disable validate);  
  
Table altered.  
  
SQL> insert into tt values(1,2);  
insert into tt values(1,2);  
*  
ERROR at line 1:  
ORA-25128: No insert/update/delete on table with constraint (SCOTT.X)  
disabled and validated
```



Partition Exchange and PK/Unique Constraint

The solution

- The partitioned target table
 - PK or unique constraint that is enforced by global index (partitioned or non-partitioned)
- The standalone table to be exchanged (“exchange table”)
 - Equivalent disabled validated constraint
 - No index for enforcement, no exchange problem

Partition Exchange and PK/Unique Constraint

A simple example

```
SQL > CREATE TABLE tx_simple
2      (
3          TRANSACTION_KEY          NUMBER,
4          INQUIRY_TIMESTAMP        TIMESTAMP(6),
5          RUN_DATE                 DATE
6      )
7      PARTITION BY RANGE (RUN_DATE)
8      (
9          PARTITION TRANSACTION_201605 VALUES LESS THAN (TO_DATE('20160601', 'yyyymmdd')),
10         PARTITION TRANSACTION_201606 VALUES LESS THAN (TO_DATE('20160701', 'yyyymmdd')),
11         PARTITION TRANSACTION_201607 VALUES LESS THAN (TO_DATE('20160801', 'yyyymmdd')),
12         PARTITION TRANSACTION_201608 VALUES LESS THAN (TO_DATE('20160901', 'yyyymmdd')),
13         PARTITION TRANSACTION_201609 VALUES LESS THAN (TO_DATE('20161001', 'yyyymmdd')),
14         PARTITION TRANSACTION_201610 VALUES LESS THAN (TO_DATE('20161101', 'yyyymmdd')),
15         PARTITION TRANSACTION_MAX VALUES LESS THAN (MAXVALUE)
16     )
17  /
```

Table created.

Partition Exchange and PK/Unique Constraint

A simple example

```
SQL > CREATE TABLE tx_simple
 2  (
 3      TRANSACTION KEY          NUMBER,
 4      INQUIRY_TIMESTAMP        TIMESTAMP(6),
 5      RUN_DATE                 DATE
 6  )
 7  PARTITION BY RANGE (RUN_DATE)
 8  (
 9      PARTITION TRANSACTION_201605 VALUES LESS THAN (TO_DATE('20160601', 'yyyymmdd')),
10      PARTITION TRANSACTION_201606 VALUES LESS THAN (TO_DATE('20160701', 'yyyymmdd')),
11      PARTITION TRANSACTION_201607 VALUES LESS THAN (TO_DATE('20160801', 'yyyymmdd')),
12      PARTITION TRANSACTION_201608 VALUES LESS THAN (TO_DATE('20160901', 'yyyymmdd')),
13      PARTITION TRANSACTION_201609 VALUES LESS THAN (TO_DATE('20161001', 'yyyymmdd')),
14      PARTITION TRANSACTION_201610 VALUES LESS THAN (TO_DATE('20161101', 'yyyymmdd')),
15      PARTITION TRANSACTION_MAX VALUES LESS THAN (TO_DATE('20170101', 'yyyymmdd'))
16  )
17  /

Table created.

SQL > INSERT into tx_simple (
 2      select object_id, LAST_DDL_TIME,
 3      add months(TO_DATE('20160501', 'yyyymmdd'), mod(OBJECT_ID,
12))
 4      from DBA_OBJECTS
 5      where object_id is not null)
 6  /

73657 rows created.
```



Partition Exchange and PK/Unique Constraint

A simple example

```
SQL > CREATE TABLE tx_simple
2   (
3     TRANSACTION_KEY      NUMBER,
4     INQUIRY_TIMESTAMP    TIMESTAMP(6),
5     RUN_DATE              DATE
6   )
7   PARTITION BY RANGE (RUN_DATE)
8   (
9     PARTITION TRANSACTION_201605 VALUES LESS THAN (TO_DATE('20160601', 'yyyymmdd')),
10    PARTITION TRANSACTION_201606 VALUES LESS THAN (TO_DATE('20160701', 'yyyymmdd')),
11    PARTITION TRANSACTION_201607 VALUES LESS THAN (TO_DATE('20160801', 'yyyymmdd')),
12    PARTITION TRANSACTION_201608 VALUES LESS THAN (TO_DATE('20160901', 'yyyymmdd')),
13    PARTITION TRANSACTION_201609 VALUES LESS THAN (TO_DATE('20161001', 'yyyymmdd')),
14    PARTITION TRANSACTION_201610 VALUES LESS THAN (TO_DATE('20161101', 'yyyymmdd')),
15    PARTITION TRANSACTION_MAX VALUES LESS THAN (MAXVALUE)
16  )
```

```
SQL > INSERT into tx_simple (
2     select object_id, LAST_DDL_TIME,
3     add months(TO_DATE('20160501', 'yyyymmdd'), mod(OBJECT_ID,
12))
```

```
SQL > CREATE UNIQUE INDEX tx_simple_PK ON tx_simple (TRANSACTION_KEY) nologging
2   GLOBAL PARTITION BY RANGE (TRANSACTION_KEY) (
3     PARTITION P_Max VALUES LESS THAN (MAXVALUE)
4   )
5   /
```

Index created.

```
SQL > ALTER TABLE tx_simple ADD ( CONSTRAINT tx_simple_PK PRIMARY KEY (TRANSACTION_KEY)
2   USING INDEX nologging);
```

Table altered.



Partition Exchange and PK/Unique Constraint

A simple example, cont.

```
SQL > create table DAILY_ETL_table  
2      as  
3      select * from tx_simple partition (TRANSACTION_201607);
```

Table created.

```
SQL > alter table daily_etl_table add ( constraint pk_etl primary key (transaction_key) disable validate);
```

Table altered.

```
SQL > alter table tx_simple  
2      exchange partition TRANSACTION_201607  
3      with table daily_ETL_table  
4      including indexes  
5      --excluding indexes  
6      WITHOUT VALIDATION  
7      UPDATE GLOBAL INDEXES  
8      /
```

Table altered.

ORACLE®